

DATA MANAGEMENT FOR DYNAMIC MULTIMEDIA ANALYTICS AND RETRIEVAL

Inauguraldissertation

zur

Erlangung der Würde eines Doktors der Philosophie

vorgelegt der

Philosophisch-Naturwissenschaftlichen Fakultät

der Universität Basel

von

Ralph Marc Philipp Gasser

Basel, 2023

Originaldokument gespeichert auf dem Dokumentenserver
der Universität Basel

edoc.unibas.ch



Dieses Werk ist unter dem Vertrag "Nicht kommerziell - Keine Bearbeitungen 4.0 International (CC BY-NC-ND 4.0)" lizenziert. Die vollständige Lizenz kann unter creativecommons.org/licenses/by-nc-nd/4.0/ eingesehen werden.

Genehmigt von der Philosophisch-Naturwissenschaftlichen Fakultät
auf Antrag von

Prof. Dr. Heiko Schuldt
Erstbetreuer

Prof. Dr. Christian Tschudin
Zweitbetreuer

Prof. Dr. Björn Þór Jónsson
Externer Experte

Basel, den 15.11.2022

Prof. Dr. Marcel Mayor
Dekan



Namensnennung – Nicht kommerziell – Keine Bearbeitungen 4.0 International
(CC BY-NC-ND 4.0)

Dies ist eine allgemeinverständliche Zusammenfassung der Lizenz (die diese nicht ersetzt).

Sie dürfen:

Teilen – das Material in jedwedem Format oder Medium vervielfältigen und weiterverbreiten

Der Lizenzgeber kann diese Freiheiten nicht widerrufen solange Sie sich an die Lizenzbedingungen halten.

Unter den folgenden Bedingungen:

Namensnennung – Sie müssen angemessene Urheber- und Rechteangaben machen, einen Link zur Lizenz beifügen und angeben, ob Änderungen vorgenommen wurden. Diese Angaben dürfen in jeder angemessenen Art und Weise gemacht werden, allerdings nicht so, dass der Eindruck entsteht, der Lizenzgeber unterstützte gerade Sie oder Ihre Nutzung besonders.

Nicht kommerziell – Sie dürfen das Material nicht für kommerzielle Zwecke nutzen.

Keine Bearbeitungen – Wenn Sie das Material remixen, verändern oder darauf anderweitig direkt aufbauen, dürfen Sie die bearbeitete Fassung des Materials nicht verbreiten.

Keine weiteren Einschränkungen – Sie dürfen keine zusätzlichen Klauseln oder technische Verfahren einsetzen, die anderen rechtlich irgendetwas untersagen, was die Lizenz erlaubt.

Hinweise: Sie müssen sich nicht an diese Lizenz halten hinsichtlich solcher Teile des Materials, die gemeinfrei sind, oder soweit Ihre Nutzungshandlungen durch Ausnahmen und Schranken des Urheberrechts gedeckt sind.

Es werden keine Garantien gegeben und auch keine Gewähr geleistet. Die Lizenz verschafft Ihnen möglicherweise nicht alle Erlaubnisse, die Sie für die jeweilige Nutzung brauchen. Es können beispielsweise andere Rechte wie Persönlichkeits- und Datenschutzrechte zu beachten sein, die Ihre Nutzung des Materials entsprechend beschränken.

Quelle: <https://creativecommons.org/licenses/by-nc-nd/4.0/deed.de>, 18. Januar 2023

Look again at that dot. That's here. That's home. That's us. On it everyone you love, everyone you know, everyone you ever heard of, every human being who ever was, lived out their lives. The aggregate of our joy and suffering, thousands of confident religions, ideologies, and economic doctrines, every hunter and forager, every hero and coward, every creator and destroyer of civilization, every king and peasant, every young couple in love, every mother and father, hopeful child, inventor and explorer, every teacher of morals, every corrupt politician, every "superstar," every "supreme leader," every saint and sinner in the history of our species lived there-on a mote of dust suspended in a sunbeam. [...]

— Carl Sagan, Pale Blue Dot: A Vision of the Human Future in Space

Zusammenfassung

Multimediale Daten in ihren verschiedensten Ausprägungen bergen eine einmaliige Herausforderung bei ihrer Speicherung und Verwaltung, besonders wenn die Durchsuchbarkeit und Analyse von grossen Beständen ermöglicht werden soll. Die inhärente Unstrukturiertheit der Daten selbst und der "Fluch der Dimensionalität", welcher den Umgang mit den daraus abgeleiteten Repräsentationen erschwert, bringen eine Vielzahl von Problemen mit sich, welchen auf unterschiedlichsten Ebenen begegnet werden kann. Dies ist Ausgangslage für ein Forschungsgebiet, welches sich mit der effizienten und effektiven Suche in multimedialen Datensammlungen beschäftigt und aus dem über die Jahre eine Vielzahl von spezialisierten Systemen und Anwendungen entstanden sind.

Fortschritte im Bereich der interaktiven, multimedialen Suche sowie der multimedialen Analytik haben allerdings gezeigt, dass eine Vielzahl der Annahmen, welche für die klassische Multimedia-Suche getroffen werden, für solche interaktiven Anwendungsbereiche nicht oder nur beschränkt gelten. Einerseits kann eine Vielzahl der Anforderungen nicht durch einfache Ähnlichkeitssuche ausgedrückt werden und benötigen deshalb ein Abfragemodell, welches die notwendige Flexibilität bringt. Andererseits ist die oft getroffene Annahme, dass Datenbestände sich zur Abfragezeit nicht verändern, in der Praxis nicht haltbar, insbesondere da Analytik das explizite Ziel verfolgt, neue Informationen und Wissen zu generieren und zu speichern. In Konsequenz ist es angesichts der zahlreichen möglichen Anfragearten inadäquat anzunehmen, dass ein/e Benutzer/-in dem System genau vorgeben kann, wie diese ausgeführt werden sollen.

Ausgehend von diesen Lücken und motiviert von der Tatsache, dass sehr ähnliche Herausforderungen bereits im Kontext klassischer Datenbankanwendungen für strukturierte Daten angegangen und gelöst wurden, präsentiert diese Dissertation drei Beiträge, um den zuvor erwähnten Problemen zu begegnen. Wir präsentieren ein Abfragemodell, welches die Idee von distanzbasierter Suche verallgemeinert und formalisieren die Verbindung solcher Abfragen zu hoch-dimensionalen Indexstrukturen. Dies wird komplementiert durch ein Kostenmodell, welches den oft impliziten Kompromiss zwischen Ausführungs geschwindigkeit und Qualität der Resultate bei der Verwendung dieser Indexstrukturen transparent macht. Zu guter Letzt beschreiben wir ein Modell, welches die persistente und transaktionale Anlage und Pflege solcher Indexstrukturen ermöglicht.

Alle erwähnten Beiträge sind im öffentlich zugänglichen und frei lizenzierten Datenbanksystem Cottontail DB implementiert, auf dessen Basis wir eine Evaluation präsentieren, welche die Effektivität der Modelle demonstriert. Wir schliessen mit einer Diskussion über mögliche Richtungen für zukünftige Entwicklungen mit Blick auf eine weitere Konvergenz zwischen der Forschung im Bereich der Datenbanksysteme und der (interaktiven) Suchsysteme für multimediale Bestände.

Abstract

Multimedia data in its various manifestations poses a unique challenge from a data storage and data management perspective, especially if search, analysis and analytics in large data corpora is considered. The inherently unstructured nature of the data itself and the curse of dimensionality that afflicts the representations we typically work with in its stead are cause for a broad range of issues that require sophisticated solutions at different levels. This has given rise to a huge corpus of research that puts focus on techniques that allow for effective and efficient multimedia search and exploration. Many of these contributions have led to an array of purpose-built, multimedia search systems.

However, recent progress in multimedia analytics and interactive multimedia retrieval, has demonstrated that several of the assumptions usually made for such multimedia search workloads do not hold once a session has a human user in the loop. Firstly, many of the required query operations cannot be expressed by mere similarity search and since the concrete requirement cannot always be anticipated, one needs a flexible and adaptable data management and query framework. Secondly, the widespread notion of staticity of data collections does not hold if one considers analytics workloads, whose purpose is to produce and store new insights and information. And finally, it is impossible even for an expert user to specify exactly how a data management system should produce and arrive at the desired outcomes of the potentially many different queries.

Guided by these shortcomings and motivated by the fact that similar questions have once been answered for structured data in classical database research, this Thesis presents three contributions that seek to mitigate the aforementioned issues. We present a query model that generalises the notion of proximity-based query operations and formalises the connection between those queries and high-dimensional indexing. We complement this by a cost-model that makes the often implicit trade-off between query execution speed and results quality transparent to the system and the user. And we describe a model for the transactional and durable maintenance of high-dimensional index structures.

All contributions are implemented in the open-source multimedia database system Cottontail DB, on top of which we present an evaluation that demonstrates the effectiveness of the proposed models. We conclude by discussing avenues for future research in the quest for converging the fields of databases on the one hand and (interactive) multimedia retrieval and analytics on the other.

Acknowledgements

Finishing a project such as the one described in this thesis is never the sole achievement of one person alone. And while I may be the author of the text it contains and the mind behind the concepts and systems it describes, there is a great many people that have somehow enabled the successful completion of this dissertation. It is those sung and unsung heroes that I shall dedicate these acknowledgements to.

First and foremost, I would like to thank Prof. Dr. Heiko Schuldt, who offered me this position and who convinced me five years ago, that doing a PhD is a unique opportunity not to be missed. The past five years indeed were time very well spent on interesting projects with interesting people and I am grateful for the experience and the guidance you have given me along this journey. I also wish to thank Prof. Dr. Björn Þór Jónsson, who volunteered as my external examiner. Reading this Thesis must have been a daunting task indeed and I very much hope that the experience was at least somewhat pleasant.

I would also like to take the opportunity to give a big shout-out to our DBIS research group with all its amazing members both past and present. I really enjoyed our many lunches, discussions about politics, progress reports, and of course the mandatory gossip on the second floor. A special thanks goes to Silvan Heller, Loris Sauter, Marco Vogt, Alexander Stiemer, Florian Spiess and Mahnaz Parian-Sherb as well as our former members Luca Rossetto and Ivan Giangreco who I do not only consider valued colleagues but also dear friends.

The process of writing and finishing a Thesis like this can be taxing and at times downright frustrating. In this situation, it is invaluable to have colleagues who share the same predicament and with whom one can talk about the progress, the setbacks, and just the general state of mind. Therefore, a loud, emphatic “thank you” to both Marco and Silvan for establishing and holding on to our “Thesis Reading Circle”, and for the many hours of discussions, and the invaluable feedback you have provided.

Team DRES (Loris and Luca) also deserves a special mention: It is not often that one finds people whom working with is as productive and pleasant as it is with you two and I really would like to thank you for this. We have achieved a lot together and I am certain that there are still many interesting ideas to be realised, side-projects to be started, papers to be written, and discussions to be had.

I would also like to acknowledge and thank my dear friend Andrea Schäfer for reading my entire Thesis and providing me with feedback on the English language. I guess without you, reading this damn thing would be even more painful than it likely already is anyway. I also acknowledge Luca Rossetto for occasionally providing feedback on bits and pieces.

Of course, I also want to thank my friends and family – all those of you who allowed and enabled me to wind down from the at times stressful experience of working on this project. A heartfelt “thank you” and, once again, a big shout-out to all of you – also for the patience you have shown over the past few years.

There is one person who deserves a special mention here – my wonderful and beloved partner, Jonas. We have been together for over thirteen years, but the past five years probably were among the most challenging. Balancing two occupations, the purchase and renovation of a house, the well-being of a herd of rabbits, and life in general certainly is no easy task and it would not have been possible without your unwavering support. Thank you for being there for me – in good and in bad times, as they say – and for having my back no matter what adventure I embark on.

And last but not least, thanks a lot to everyone who has not been mentioned explicitly yet is implicitly meant: The students and colleagues whom I have had the honour of working with, the administrative staff who makes sure that everything runs smoothly and all the generations of researchers whose efforts have enabled this small contribution that I am trying to make with this Thesis. To everyone who has supported me directly or indirectly, knowingly or unknowingly, willingly or unwillingly – thank you very much!

*We learned about gratitude and
humility - that so many people
had a hand in our success, from
the teachers who inspired us to the
janitors who kept our school clean
and we were taught to value
everyone's contribution and treat
everyone with respect.*

— Michelle Obama, 2012

Contents

Zusammenfassung	ix
Abstract	xi
Acknowledgements	xiii
List of Figures	xxi
List of Tables	xxvii
List of Acronyms	xxix
List of Symbols	xxxiii

I Introduction 1

1 Introduction	3
1.1 Working with Multimedia Data	4
1.1.1 Multimedia Analytics	4
1.1.2 Interactive Multimedia Retrieval	6
1.1.3 Multimedia Data Management	7
1.2 Research Gap and Objective	8
1.3 Contribution and Outline	10

2 Applications and Use Cases 13

2.1 Use Case 1: Interactive Multimedia Retrieval	13
2.2 Use Case 2: Analysis of Social Media Streams	18
2.3 Use Case 3: Signal Analysis in Medical Imaging	20
2.4 Deriving Requirements	23
2.4.1 Traditional Database Capabilities	23
2.4.2 Multimedia Query Support	24
2.4.3 Multimedia Data is Big Data	24
2.4.4 Modelling Retrieval Quality	25
2.4.5 Complex Data Types as First-class Citizens	25
2.4.6 Functions as First-class Citizens	26

2.5 Mapping Requirements to Contributions	27
II Foundations	29
3 On the Design of a Database Management System	33
3.1 Data Management and Data Models	34
3.2 The Relational Data Model	35
3.2.1 Keys and Normal Forms	38
3.2.2 Relational Algebra	41
3.2.2.1 Set Operations	41
3.2.2.2 Cartesian Product	42
3.2.2.3 Rename	42
3.2.2.4 Selection	42
3.2.2.5 Projection	43
3.2.2.6 Natural Join	43
3.2.2.7 Expressing Queries	43
3.2.3 Extensions	45
3.2.3.1 Relations vs. Bags vs. Sequences	45
3.2.3.2 Extended Projection	46
3.2.3.3 Aggregation	47
3.2.3.4 Sorting and Ranking	47
3.2.3.5 Recursive Data Access	48
3.3 Database Management System (DBMS)	49
3.3.1 Query Interface	49
3.3.2 Query Parser	51
3.3.3 Query Planner & Optimiser	52
3.3.4 Execution Engine	56
3.3.4.1 Query Parallelism	58
3.3.4.2 Optimising Execution Speed	58
3.3.5 Storage Manager	59
3.3.5.1 Data Organisation and Disk Manager	59
3.3.5.2 Indexes	60
3.3.5.3 Buffer Pool Manager	62
3.3.6 Common Components	63
4 On Multimedia Analysis and Retrieval	65
4.1 Multimedia Data and Multimedia Collections	66
4.1.1 A Data Model for Multimedia Collections	68

4.1.2	Descriptive Metadata	70
4.1.3	Features	71
4.2	Multimedia Retrieval	73
4.2.1	Similarity Search and the Vector Space Model	75
4.2.1.1	Nearest Neighbour Search (NNS)	77
4.2.1.2	Range Search	77
4.2.1.3	Reverse Nearest Neighbour Search (RNNS)	78
4.2.2	NNS and High-Dimensional Index Structures	78
4.2.2.1	Vector Approximation File (VAF)	80
4.2.2.2	Locality Sensitive Hashing (LSH)	82
4.2.2.3	Product Quantisation (PQ)	83
4.2.3	Beyond Similarity Search	84
5	Multimedia Data Management	87
5.1	Early Multimedia Retrieval Systems	88
5.2	Libraries for Multimedia Retrieval	89
5.3	Vector Database Engines	90
5.3.1	Milvus	91
5.4	Databases for Data Science	92
5.5	vitrivr, ADAM and ADAM _{pro}	92
III Dynamic Multimedia Data Management		95
6	Modelling a Database for Dynamic Multimedia Data	97
6.1	Generalised Proximity-Based Operations	98
6.1.1	Revisiting Distance Computation	99
6.1.2	Extending the Relational Model	103
6.1.2.1	Extended Projection and DFCs	103
6.1.2.2	Ranked Relations	105
6.1.2.3	s,k-selection	107
6.1.2.4	Combining $\pi_P, \omega_O, \lambda_k$	108
6.1.2.5	Algebraic Properties of $\pi_P, \omega_O, \lambda_k$	109
6.1.3	DFCs and Query Optimisation	112
6.1.4	DFCs and High-Dimensional Index Structures	114
6.2	Cost Model for Multimedia Databases	118
6.2.1	Cost Policies	121
6.2.2	Estimating Cost of Approximation	122
6.3	High-Dimensional Index Maintenance	124

6.3.1	Storage Model	125
6.3.2	Write Model of an Index	126
6.3.3	Failure Model and State of an Index	127
6.3.4	Concurrent Index Rebuilding	128
6.3.4.1	Failure Cases and Limitations	132
6.3.5	Case Studies	133
6.3.5.1	Vector Approximation File (VAF)	133
6.3.5.2	Product Quantisation (PQ)	135
7	Cottontail DB	137
7.1	Data Model and Nomenclature	137
7.2	Architecture	139
7.2.1	Query Implementation and Execution	139
7.2.1.1	Execution Model	140
7.2.2	Query Planner	143
7.2.2.1	Optimisation Algorithm	143
7.2.2.2	Plan Caching	145
7.2.2.3	Logical Optimisation	145
7.2.2.4	Physical Optimisation	146
7.2.3	Query Parsing and Binding	147
7.2.3.1	Binding Context	149
7.2.4	Functions and Function Generators	149
7.2.5	Storage	150
7.2.5.1	Storage Engine: MapDB	151
7.2.5.2	Storage Engine: Xodus	152
7.2.5.3	Storage Engine: HARE	153
7.2.6	Transactional Guarantees	154
IV	Discussion	155
8	Evaluation	157
8.1	Evaluation Setup	158
8.1.1	Metrics	159
8.2	Simulated Multimedia Analytics Workloads	163
8.2.1	Experiment 1: Influence of Parallelism & Use of Indexes . .	164
8.2.2	Experiment 2: Influence of Cost Policy	169
8.2.3	Benchmark 3: Influence of Optimisation	173
8.3	Large-Scale Similarity Search	175

8.3.1	Cottontail DB	175
8.3.2	Milvus	180
8.3.2.1	Qualitative Assessment	183
8.4	High-Dimensional Index Maintenance	184
8.5	Summary, Related Work and Discussion	190
8.5.1	Generalised Proximity-Based Queries	190
8.5.2	Query Planning and Cost Model	193
8.5.3	High-Dimensional Index Maintenance	193
8.5.4	A Remark on Benchmarking Multimedia Analytics	195
9	Conclusion and Outlook	197
9.1	Motivating Scenarios	198
9.2	Future Work	199
9.2.1	Query Model	200
9.2.2	Multimedia Indexing	200
9.2.3	Storage Model	201
9.2.4	Cost Model and Query Planning	202
9.2.5	Execution Model	203
9.2.6	Implementation and Vision	203
Appendix		205
A	Evaluation Supplements	207
B	Additional Results	211
	Bibliography	215

List of Figures

1.1	Exploration-search axis of multimedia analytics [ZW14].	5
2.1	The idealised multimedia retrieval problem: A user expresses an information need as a query and the multimedia retrieval system returns matching items from a collection as a result.	14
2.2	A simplified model of knowledge formation from an analogue input in humans (adapted from [JLX21]) and computers.	15
2.3	The interactive multimedia retrieval problem: Even when searching for a known item, the search effort typically constitutes a continuous back and forth between system and user involving different types of workload as the user tries to find the item of interest through query (re-)formulation, browsing and examination.	17
2.4	The architecture of a fictitious social media analytics platform (based on [ZCW ⁺ 15; CSW ⁺ 19; YPM ⁺ 19; BPS ⁺ 20]).	19
2.5	Simplified illustration of the physical processes enabling MRI. . . .	20
2.6	Simplified block diagram of a MRF system (adapted from [BMCFM ⁺ 19]), reconstruction of brain scan kindly provided by Manuel Hürbin). Storage and indexing typically take place offline.	22
3.1	Different levels of abstraction for data modelling.	34
3.2	The general architecture of a DBMS and its individual components. .	49
3.3	The query from Example 3.5 represented as a tree of relational operators also called a logical query execution plan. By convention, the tree is executed and read from left to right (bottom to top), and information flows from the leaves to the root.	51
3.4	Cost-based query planning in a DBMS based on Example 3.5. A parsed query is transformed to a series of equivalent, logical and physical plans with known costs. At the end, the most cost-effective plan is selected. Red boxes indicate adjustments to the operator tree.	54
3.5	Illustration of the iterator model for query listed in Example 3.5. Every operator calls the upstream operator with <code>next()</code> to request the next tuple for processing.	57
3.6	Visualisation of an example B^+ -tree with $b = 4$. Green boxes indicate the keys held by a node. Red boxes indicate pointers to children. .	61

4.1	An extended ERM of the multimedia data model used in the <i>vitrivr</i> project. The model centres around the notion of media objects and media segments which are described by metadata and features.	69
4.2	An extended ERM of the multimedia data model used for the purpose of this Thesis. It has been derived from <i>vitrivr</i> 's data model and foregoes the explicit segment entity.	70
4.3	Flow of information in similarity search at ingest and query time.	75
4.4	Illustration of the different similarity search algorithms in $\mathbb{F} \subset \mathbb{R}^2$ using the Euclidean distance (L^2). Matches are highlighted in red.	79
4.5	Simplified illustration of basic VAF indexing scheme in for $\mathbb{F} \subset \mathbb{R}^2$. A vector is assigned to a cell in an overlay grid, which gives rise to a signature that is stored in a linear list with the vector or its tuple identifier.	81
4.6	A simplified illustration of the basic LSH indexing scheme. A vector is hashed to a bucket by a locality preserving hash function h and stored in a hash table.	82
4.7	A simplified illustration of the basic PQ indexing scheme. A vector is divided into smaller subvectors and each subvector is assigned to the nearest centroid in the codebook of the respective subspace, which produces the signature.	83
6.1	Function hierarchy of all $f \in \mathcal{F}$. As we go down the hierarchy, the DBMS gains knowledge about a function's structure.	113
6.2	Quality of results produced by a PQ index in a NNS query in terms of recall and normalised DCG at different levels observed and averaged over a 1000 queries using different query vectors. The mean value remains more or less constant while minimum and maximum exhibit a large spread for smaller levels.	124
6.3	Illustration of a change C to the data that is processed by an index \mathcal{I} 's write and failure model and leads to potential state changes. An index knows the three states CLEAN, DIRTY, and STALE.	129
6.4	Illustration of asynchronous index rebuilding, which takes place in two separate phases BUILD and REPLACE. Information from transactions that take place concurrently to the BUILD phase are LEAK-ed through a side-channel upon commit of the respective transaction.	130

6.5	Illustration of the VAF index's write model (insert). If a vector lies within the bounds spanned by the minimum and maximum, the signature can be appended. If not, the closest cell is assigned, leading to a poor estimate of bounds during scanning.	134
7.1	Architecture diagram of Cottontail DB's main components. The directed arrows indicate the path a query takes within the system (dotted for a potential SQL interface). The dashed double arrows indicate interactions between components.	140
7.2	Operator pipeline for a proximity-based query that scans the entity <i>paintings</i> , calculates the L2 distance between a <i>feature</i> column and a query, orders by the obtained distance, and limits the number of outputs (kNN). Pipeline breaking operators are indicated by the red dot.	142
7.3	Operator pipeline that performs kNN on the entity <i>paintings</i> . The <i>scan</i> and <i>distance function</i> steps are partitioned and merged before <i>order</i> , <i>limit</i> , and <i>projection</i> is applied, which allows for parallel execution.	143
7.4	Decomposition of conjunctive predicates. The filter on the column <i>death</i> may be pushed down at a later stage during physical optimisation if an index were to be available.	146
7.5	Illustration of deferring FETCH operations. Over multiple passes, the FETCH of the columns <i>id</i> , <i>name</i> , and <i>death</i> is pushed up until after the FILTER has been applied. Columns <i>name</i> and <i>death</i> are eventually eliminated because they are not required.	147
8.1	Latency in seconds (x-axis) on different entities using different access methods and levels of parallelisation (y-axis) for the analytics workloads during the first experiment. The individual queries Q1a-Q1e are highlighted in different colours. The use of the VAF and PQ index was enforced using query hints.	165
8.2	Execution plans produced by Cottontail DB for the mean (Q1b), range (Q1c), and NNS (Q1d) query workloads during the first experiment. The use of the VAF and PQ index was enforced using query hints.	166

8.3	Quality of results (x-axis) in terms of recall (mint) and DCG (red) on different entities using different access methods (y-axis) for the range (Q1c) and NNS (Q1d) query workloads during the first experiment. The use of the VAF and PQ index was enforced using query hints.	168
8.4	Score and rank of different query plan options for the mean query (Q1b). The preferred option for $w_Q = 0.0$ involves the scan of a PQ index. As w_Q goes up, the cost of impaired quality increases and a full entity scan (P_2) takes over, since none of the available indexes can satisfy the query. The scores for P_2 and P_3 do not depend on the quality parameter and incur almost the same cost, with P_2 being the less expensive variant.	169
8.5	Score and rank of different query plan options for the range query (Q1c). The preferred option for $w_Q = 0.0$ involves the scan of a PQ index. As w_Q goes up, the cost of impaired quality increases and a full entity scan takes over, since none of the available indexes can satisfy the query. The scores for P_4 through P_6 do not depend on the quality parameter and incur almost the same cost, with P_4 being the least expensive.	171
8.6	Score and rank of different query plan options for the NNS query (Q1d). The preferred option for $w_Q = 0.0$ involves the scan of a PQ index. As w_Q goes up, the cost of impaired quality increases and a VAF index scan (P_3) takes over, which can also satisfy NNS queries. Since VAF does not incur quality costs, it remains the favoured option.	172
8.7	Latency in seconds on different collections (y-axis) for the query workloads Q1a through Q1e (x-axis) with (mint) and without (red) query optimisation.	174
8.8	Cottontail DB's latency in seconds for different workloads (x-axis) on different shards of the Deep 1B dataset using different execution strategies (y-axis).	176
8.9	Execution plans produced by Cottontail DB for different query workloads prior to intra-query parallelisation. The use of the VAF and PQ index was enforced using query hints.	178
8.10	Milvus's latency in seconds for different workloads (x-axis) on different shards of the Deep 1B dataset using different access methods (y-axis), with (red) and without (mint) time used for loading the data from disk.	181

8.11 Index adaptiveness measurements for VAF index for two different values of c_{ins} and c_{del} , showing the size of the collection (top left), the number of accumulated insert and delete operations (top right), speed-up for NNS using the index (bottom left), and quality of the NNS query with respect to ground truth (bottom right) over time (x-axis).	185
8.12 Index adaptiveness measurements for PQ index for two different values of c_{ins} and c_{del} , showing the size of the collection (top left), number of accumulated insert and delete operations (top right), speed-up for NNS using the index (bottom left), and quality of NNS with respect to the ground truth (bottom right) over time (x-axis).	186
8.13 Index adaptiveness measurements for the VAF and PQ index, showing the size of the collection (top left), the number of accumulated insert and delete operations (top right), speed-up for NNS using the index (bottom left), and quality of the NNS query with respect to the ground truth (bottom right) over time (x-axis). Data and queries exhibit a random noise and index rebuild is started after 30 min (grey, dashed line).	188

List of Tables

3.1	The relational operators proposed by Codd et al. [Cod70; GUW09]. . .	41
4.1	List of distance functions often used in similarity search.	76
4.2	Non-exhaustive list of high-dimensional index structures employed in similarity search and retrieval.	80
6.1	Proximity-based queries supported by the extended relational algebra.	109
6.2	Relational operators and their algebraic properties.	110
7.1	Data types supported by Cottontail DB. Types in the numeric, vector, and complex class allow for domain specific arithmetics.	138
7.2	Main types of physical operators implemented by Cottontail DB alongside with their arity, their correspondence to relational operators and whether they require materialisation.	141
7.3	Primitive operations offered by the storage engines' Tx abstractions to interact with DBOs including the argument and return types ($\text{arg} \rightarrow \text{ret}$).	151
8.1	List of the data collections used for this evaluation.	159

List of Acronyms

1NF	First Normal Form
2NF	Second Normal Form
3NF	Third Normal Form
ACID	Atomicity, Consistency, Isolation, Durability
ANNS	Approximate Nearest Neighbour Search
API	Application Programming Interface
ARIES	Algorithm for Recovery and Isolation Exploiting Semantics
AST	Abstract Syntax Tree
AVX	Advanced Vector Extensions
BASE	Basically Available, Soft state, Eventually consistent
BAT	Binary Association Table
BCNF	Boyce–Codd Normal Form
CK	Candidate Key
CLI	Call Level Interface
CNN	Convolutional Neural Network
CP	Cluster Pruning
CPU	Central Processing Unit
CRUD	Create, Read, Update, Delete
CWA	Closed-world Assumption
DBMS	Database Management System
DBO	Database Object
DCG	Discounted Cumulative Gain
DDL	Data Definition Language
DFC	Distance Function Class
DML	Data Manipulation Language
DQL	Data Query Language
ERM	Entity Relationship Model

FAISS	Facebook AI Similarity Search
FIFO	First In First Out
FK	Foreign Key
FNS	Farthest Neighbour Search
FPGA	Field-programmable Gate Array
GPU	Graphics Processing Unit
HNSW	Hierarchical Navigable Small World
HOG	Histogram of Oriented Gradients
IO	Input / Output
JDBC	Java Database Connectivity
JVM	Java Virtual Machine
kFN	k-Farthest Neighbour Search
KLT	Karhunen-Loëve Transform
kNN	k-Nearest Neighbour Search
LIRe	Lucene Image Retrieval
LRU	Least Recently Used
LSC	Lifelog Search Challenge
LSH	Locality Sensitive Hashing
MAP	Mean Average Precision
MIPS	Maximum Inner Product Search
MRF	Magnetic Resonance Fingerprinting
MRI	Magnetic Resonance Imaging
MVCC	Multi-Version Concurrency Control
NDCG	Normalised Discounted Cummulative Gain
NGT	Neighbourhood Graph and Tree
NMR	Nucleic Magnetic Resonance
NNS	Nearest Neighbour Search
NoSQL	Not only SQL
NUMA	Non-Uniform Memory Access

NVM	Non-Volatile Memory
ODBC	Open Database Connectivity
OLAP	Online Analytical Processing
OLTP	Online Transaction Processing
PCP	Pitch Class Profiles
PK	Primary Key
POSIX	Portable Operating System Interface
PQ	Product Quantisation
RAID	Redundant Array of Independent Disks
RAM	Random Access Memory
RDF	Resource Description Framework
RNNS	Reverse Nearest Neighbour Search
S2PL	Strict 2-Phase Locking
SI	Snapshot Isolation
SIFT	Scale Invariant Feature Transformation
SIMD	Single Instruction Multiple Data
SQL	Structured Query Language
SSD	Solid State Drive
SSE	Streaming SIMD Extensions
SURF	Speeded-up Robust Features
TPC	Transaction Processing Performance Council
UDF	User Defined Function
VAF	Vector Approximation File
VBS	Video Browser Showdown
YCSB	Yahoo! Cloud Serving Benchmark

List of Symbols

Basics

\mathbb{N}	Set of natural numbers
\mathbb{R}	Set of real numbers
\mathbb{C}	Set of all complex numbers
\wedge	Logical conjunction
\vee	Logical, inclusive disjunction
\neg	Logical negation

Relational Algebra

\mathcal{R}	A relation. Sometimes with subscript, e.g., $\mathcal{R}_{\text{name}}$ to specify name or index.
\mathcal{R}^O	A ranked, relation that exhibits a partial ordering of elements induced by O . Sometimes used with subscript, e.g., $\mathcal{R}^O_{\text{name}}$ to specify name or index.
\mathcal{D}	A data domain of a relation. Sometimes with subscript, e.g., $\mathcal{D}_{\text{name}}$, to specify name or index.
\mathcal{A}	An attribute of a relation, i.e., a tuple $(\text{name}, \mathcal{D})$. Sometimes with subscript to specify name or index.
\mathcal{A}^*	An attribute that acts as a primary key.
$\overline{\mathcal{A}}$	An attribute that acts as a foreign key.
t	A tuple $t \in \mathcal{R}$ of attribute values $a_i \in \mathcal{D}_i$. Sometimes with subscript to specify the index of the tuple in \mathcal{R} .
\mathbb{D}	The set system of all data domains \mathcal{D} supported by a DBMS.
SCH	The schema of a relation \mathcal{R} , i.e., all attributes $\text{SCH}(\mathcal{R}) = (\mathcal{A}_1, \mathcal{A}_2, \dots, \mathcal{A}_N)$ that make up \mathcal{R} .
σ	The unary selection operator used to filter a relation \mathcal{R} . Usually printed with a subscript to describe the Boolean predicate S , e.g. σ_S .

π	The unary projection operator used to restrict the attributes \mathcal{A} of a relation \mathcal{R} . Usually used with subscript to list the projection attributes P , e.g., $\pi_{\mathcal{A}_1, \mathcal{A}_2}$.
ρ	The unary rename operator used to rename the attributes \mathcal{A} of a relation \mathcal{R} . Usually used with subscript to list the rename operations, e.g., $\rho_{\mathcal{A}_1 \rightarrow \mathcal{A}_2}$.
ω	The binary sort operator used to construct a ranked relation \mathcal{R}^O . Usually printed with a subscript to list order attributes O , e.g., $\omega_{\mathcal{A}_1 \uparrow, \mathcal{A}_2 \downarrow}$.
λ	The unary k-selection operator restricts \mathcal{R} to the first k tuples. Usually printed with a subscript to define k , e.g., λ_k .
γ	The unary group operator aggregates tuples in \mathcal{R} . Usually printed with a subscript to list the group attributes G , e.g., γ_G .

Multimedia Retrieval

\mathbb{O}	The media collection domain $\mathbb{O} = \{o_1, o_2, \dots, o_N\}$
\mathbb{F}	The feature domain $\mathbb{F} = \{f_1, f_2, \dots, f_M\}$
t	Feature transformation function $t: \mathbb{O} \rightarrow \mathbb{F}$
\mathfrak{d}	Dissimilarity or distance function $\mathfrak{d}: \mathbb{F} \times \mathbb{F} \rightarrow \mathbb{R}$
s	Similarity function $s: \mathbb{F} \times \mathbb{F} \rightarrow [0, 1]$
c	A correspondence function $c: \mathbb{R} \rightarrow [0, 1]$

System Model

$\hat{\mathfrak{d}}$	A DFC $\hat{\mathfrak{d}}: \mathcal{D}_f \times \mathcal{D}_f \times \mathcal{D}_1 \dots \times \mathcal{D}_N \rightarrow \mathbb{R}$, which is a generalisation of a distance or dissimilarity function.
\mathcal{I}	An index \mathcal{I}_I on relation \mathcal{R} indexing attributes $I \subset \text{SCH}(\mathcal{R})$
OP	A (relational) database operation, typically acting on a relation \mathcal{R} .
\mathcal{P}	A query execution plan, i.e., a concatenation of N operators $\text{OP}_1 \circ \text{OP}_2, \dots \circ \text{OP}_N$.

\mathcal{T}	A transaction, which is used to execute multiple queries. Subscript is typically used to indicate a point in time, e.g., \mathcal{T}_1 .
\mathcal{S}	A consistent snapshot of the database and all objects it contains (e.g., relations and indexes). Subscript is typically used to indicate a point in time, e.g., \mathcal{S}_1 .
OP_c	A change operation $\text{OP}_c(\mathcal{R}, \cdot)$, which can either be an insert, update or delete.
C	A change on relation \mathcal{R} with $C = (\text{OP}_c, \mathcal{R}, t, t')$.
WRITEM	A write model function that determines if a change C can be propagated to an index \mathcal{I} .
FAILM	A failure model function that adjusts the state of an index \mathcal{I} for a change C .
REBUILD	A rebuild function that can rebuild the index \mathcal{I} based on primary data in \mathcal{R} .

PART I

Introduction

1

At first I was afraid, I was petrified. [...]

— Gloria Gaynor, 1978

Introduction

The term *multimedia* describes the combination of different forms of content – sometimes called *medium* or *modality* – into a single, sensory experience for the purpose of observation and interpretation. Those content formats include but are not limited to images and videos (visual), music, sound effects and speech (aural) or textual information. However, exotic content forms such signals produced by sensors can also be seen as modalities, even though experience by a consumer may depend on pre-processing by specialised hard- and software.

Nowadays, people encounter digital media and multimedia on a daily basis when watching videos on Netflix or YouTube, when listening to music on Spotify or when browsing a private photo collection on their computer. Multimedia content makes up a large part of today's Internet and constitutes a major driving force behind its growth, as both its volume and variety increases at an ever increasing pace. Important contributing factors are social media platforms, where users act both as consumers and producers of digital content (so-called "prosumers" [RJ10; RDJ12]). Current estimates suggest that there are roughly 4.66 billion active Internet users worldwide, of which 4.2 billion can be considered active social media users.¹ Facebook alone contributed to 144 thousand uploaded images per minute in 2020. And many more of these platforms, such as Instagram, Twitter or TikTok, serve millions of users with mixed, self-made content involving text, images, videos, or a combination thereof. A study found that by 2025, we are bound to produce an estimated annual amount of 175 ZB (i.e., 175×10^{21} B) worth of data,² a large part of which can be considered multimedia, which must be managed, manipulated, queried, explored and analysed in an efficient and effective manner.

¹ See <https://www.statista.com/>, "Social media usage worldwide"; accessed January 2021.

² See <https://www.statista.com/>, "Big Data"; accessed January 2021.

1.1 Working with Multimedia Data

At a very high level multimedia data collections consist of individual multimedia items, such as video, image, or audio files. Each item, in turn, comprise *content* and *metadata* [BdB⁺07]. Unlike traditional data collections that contain mainly (semi-)structured information such as text or numbers, the content of the multimedia item itself is unstructured at a data level, which is why *feature representations* or *descriptors* that reflect a media item's content in some way and that can be handled by data processing systems are required [BdB⁺07; ZW14]. Traditionally, such feature representations have often been real-valued vectors. However, in theory, any mathematical object that can be processed by a computer can act as such.

Multimedia analysis – which has its roots in computer vision and pattern recognition and started in the early 1960s – deals with the automated, computer-aided analysis of media data, i.e., the extraction and processing of feature representations. In the early days of computer vision, for example, much effort went into the engineering of feature representations that captured certain aspects of an image's content, such as the colour distribution, texture or relevant keypoints [Low99; BTG06]. Once such features have been obtained, they could be used to perform various tasks such as classification, clustering or statistical analysis. With the advent of deep learning, the extraction of features could largely be automated through deep neural network architectures such as the Convolutional Neural Network (CNN) and sometimes even be integrated with the downstream analysis [GBC16].

Obviously, such analysis is not restricted to the visual domain and can be applied to other types of media such as speech, music, video, or 3D models with specific applications including speech recognition, audio fingerprinting in music (re-)identification, movement detection in videos or classification of 3D models, all of which fall into the broader category of multimedia analysis.

1.1.1 Multimedia Analytics

Multimedia analytics aims at generating new knowledge and insights from large-scale multimedia collections by combining techniques from multimedia analysis and visual analytics [ZW14; SHS⁺17]. While multimedia analysis deals with the different media types and how meaningful representations and models can be derived from them, visual analytics deals with the user's interaction with the data and the models themselves [CTW⁺10; KKE⁺10]. Zahálka and Worring

[ZW14] describe an abstract multimedia analytics model consisting of four tiers with the data analyst (the user [SHS⁺17]) at its center:

Objectives are the high-level goals that guide the analyst and involve aims such as answering a specific question or confirming a hypothesis ([SHS⁺17] also refers to this as the task).

Procedure is a formalisation of the activities that take the analyst from the beginning of a task to the fulfilment of their objective. For complex tasks, this is a multi-step process that involves different intents at different stages.

Intent is the aim of the analyst at a given point within the overarching analytics task, e.g., comparing individual items. It typically involves the application of different methods.

Methods comprise low-level techniques for data analysis (e.g., aggregation), data visualisation (e.g., as a chart or grid) and interaction with those visualisations (e.g., zoom-in).

Simply put, multimedia analytics can be seen as a control loop between the multimedia data collection, a data analysis component, data visualisation and a user. Analysis is used to generate models from data, which in turn enable visualisations that are then examined and refined by the user and their input. This is an iterative process that generates new knowledge and may in and of itself lead to adjustments to the models and to new information being stored.

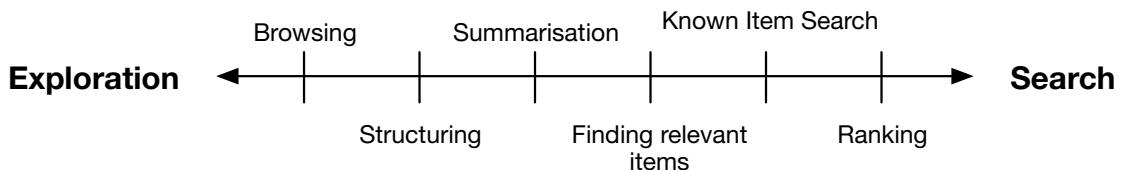


Figure 1.1 Exploration-search axis of multimedia analytics [ZW14].

The exploration-search axis proposed in [ZW14] and depicted in Figure 1.1 is used to characterise the different types of activities carried out by the analyst as they pursue their high-level objective. Exploration, on the one end, is performed when the user knows very little about the collection and its content. Search, on the other end, is applicable when the analyst knows exactly what they are looking for. During a multimedia analytics session, a user's activities often move between the two ends of the spectrum until the objective has been fulfilled. It is important to note, that all the activities rely on different methods that come with distinct requirements on data transformation and processing.

1.1.2 Interactive Multimedia Retrieval

Traditionally, multimedia retrieval has been regarded as a special niche within the multimedia analysis domain. It constitutes a dedicated field of research that deals with searching and finding items of interest within large bodies of multimedia data, based on information that may not necessarily be available in a structured manner, such as, motives depicted in an image. Even though search and retrieval may sound like the main function offered by a traditional database system, it is a very different task for multimedia than it is for structured data [BdB⁰⁷]. On the one hand, given the structure of, e.g., a relational database and languages like SQL, a user can specify exactly what elements from the database should be selected using predicates that lead to a binary decision whether items in a collection match or not. For example, when considering a product database that contains price information for individual items on sale, one can formulate a query that selects all items above a specific price threshold.

Retrieving multimedia data, on the other hand, comes with indirections due to the unstructured nature of the content, the feature representations used as a proxy for it and the *semantic gap* [SWS⁰⁰; BdB⁰⁷] associated with their generation. A very popular model to work with feature representation in multimedia retrieval involves calculation of (dis-)similarity scores from the features and sorting and ranking of items based on this score. This is commonly referred to as the *vector space model* of multimedia retrieval and similarity search [ZAD⁰⁶]. It is inherently non-binary in nature, i.e., an item matches a query to some degree and may be among the results even though the score indicates that it is a bad match. Over the years, many different combinations of features and ranking models have been proposed to facilitate content-based retrieval of different media types, such as images, audio, or video [HXL¹¹; DA13; MK18] as have been different types of query formulation, that complement traditional, text-based input. These models include methods such as *Query-by-Example* [KCH95], *Query-by-Sketch* [SMM99], *Query-by-Humming* [GLC⁹⁵] or *Query-by-Sculpting* [BGS²⁰].

It is worth noting, that when considering concrete, “real-world” multimedia retrieval system implementations today, the lines between traditional multimedia retrieval and multimedia analytics quickly start to blur. This is because traditional multimedia retrieval operates under the assumption that a specific retrieval model produces the correct result in the top ranks (e.g., the top 5) and that an item can thus be retrieved by a single query. In practice, this is often not the case due to limitations of the model itself (e.g., semantic and pragmatic gap [ZW14]), imprecisions in query formulation or simply because an information

need may be general rather than specific (e.g., finding a specific, known item versus finding items that match a general description). Therefore, finding an item of interest often requires multiple queries and other types of interaction with the data collection. This is sometimes referred to interactive multimedia retrieval and it usually includes a user in the loop.

Therefore, in addition to the extraction of appropriate features and the conception of effective ranking algorithms, interactive multimedia retrieval systems also concern themselves with aspects such as query (re-)formulation and refinement, results presentation and efficient exploration [LKM⁺19]. Moreover, such systems do not simply operate on feature representations anymore but combine similarity, fulltext and Boolean search [RGL⁺20]. Consequently, one could argue that interactive multimedia retrieval systems perform a specific type of multimedia analytics task, which is that of finding unknown items that satisfy a specific information need [SHS⁺17]. This makes many of the arguments made about data processing and data management requirements for multimedia analytics applicable to interactive multimedia retrieval as well.

1.1.3 Multimedia Data Management

As multimedia collections become large enough for the relevant units of information to no longer fit into main memory (RAM), multimedia retrieval and analytics quickly become issues of scalable data management [JWZ⁺16; PYC⁺18; VHC⁺19]. Therefore, it is no longer just a question of effective analysis and visualisation but also of strategies to store and index the original data as well as derived descriptors in a way that allows for efficient and effective processing in the face of interactive users, changing query workloads and changing data [SWS⁺00]. This data management aspect brings about many challenges that bear similarity with issues that have been addressed in database research.

However, as of today, state of the art in multimedia analytics “[...] has up to now [...] not explicitly considered the issue of data management, despite aiming for large-scale analytics.” [JWZ⁺16] (p. 296), a statement that can also be made about multimedia retrieval. In contrast, the database research community has identified the support for unstructured and multimedia data as well as the integration of information retrieval and database functionality as important research directions in both the *Lowell* and the *Claremont* reports on the state of database research in 2005 and 2008, respectively [AAB⁰⁵; AAB⁰⁸]. Despite these initiatives, however, the two fields have remained largely disjointed areas of research until today, with only a few isolated contributions that straddle the two.

1.2 Research Gap and Objective

Despite some theoretical work towards multimedia data management [MS96; AN97], the conception of concrete multimedia database systems [GS16; YLF⁺20; WYG⁺21] as well as multimedia retrieval systems that refactor data management into distinct components [CHS⁺95; RGT⁺16; GRS19a], there is still a gaping disconnect between multimedia retrieval, analytics and database research. While [Gia18] proposes important contributions towards a unified data-, query- and execution-model required for effective search in multimedia collections, scalability aspects and the need for near real-time query performance, especially in the face of changing data, are not systematically considered. Contrarily, the proposed models – despite being seminal for data management in multimedia retrieval applications – postulate assumptions, that have considerable impact on the practical applicability of systems implementing them.

The starting point for the research described in this thesis is therefore the current state-of-the-art for data management in multimedia retrieval and analytics, as briefly touched upon in the previous sections. Starting from and inspired by the models and solutions proposed in [GS16; Gia18] and motivated by the “Ten Research Questions for Scalable Multimedia Analytics” [JWZ⁺16], this thesis challenges three basic assumptions currently employed and operated upon in multimedia data management and explores the ramifications of doing so, with the higher-level goal of bridging certain gaps between research conducted in multimedia retrieval, analysis and analytics on the one hand, and classical data management and databases on the other. These assumptions are:

Assumption 1: Nearest Neighbour Search The metric space model employed in multimedia retrieval [ZAD⁺06] relies on a notion of similarity search that is usually expressed as finding the k nearest neighbouring feature vectors $f_{i \in [1, k]} \in \mathbb{F}$ to a query vector $q \in \mathbb{R}^d$ in a collection $\mathbb{F} \subset \mathbb{R}^d$ given a certain distance function. A common strategy to accommodate this type of search is to add this operation as a new query primitive in addition to whatever an existing data model might already bring [GdR⁺09; GS16; YLF⁺20], without systematically considering the interplay between this primitive and other database operations.

Assumption 2: Staticity of Data Collections Multimedia retrieval systems today often make a distinction between an *offline* phase, during which media items are analysed, features are generated, and derived data is ingested into a data management system, and an *online* phase, during which queries

take place. Usually, no changes to the data collection are allowed during the online phase. This model is, for example, advertised by [FSA⁺95; KCG⁺03; Gia18; Ros18] and to the best of our knowledge, most existing multimedia retrieval and analytics systems implement this either explicitly or implicitly, with a few exceptions, such as the system presented in [WYG⁺21].

Assumption 3: User Driven Query Execution Database systems usually evaluate and select the execution plan for an incoming query during a step that is referred to as *query planning*. The underlying assumption is that the database has all the information required to determine the most effective execution path in terms of cost parameters, such as required IO, CPU, and memory usage. In contrast, this is often not the case for multimedia retrieval systems and instead, it falls to the user to make decisions as to how a query should be executed, e.g., in terms of indexes that should be employed.

On the issue of Assumption 1, one can say that while the described model may be very concise and simple to implement, it merely allows for the ranking of results and is therefore only able to accommodate the search-end of the exploration-search axis, assuming that features are, in fact, real-valued vectors. However, the model cannot be generalised to other operations that involve distance calculation and becomes too limited for tasks such as browsing, structuring and summarisation, thus delegating the required data processing to upper-tier system components. Referring to [JWZ⁺16], it would however be desirable to offer such primitives at a database level. Furthermore, the metric space model with its Nearest Neighbour Search (NNS) paradigm may be too limited for certain use cases and a broader range of operations therefore be preferable.

It is also worth noting, that Assumption 2 and 3 both go against well established design principles usually found in modern database systems [Pet19; Ams14]. While it may be convenient from a perspective of system design, to assume a data collection to be static – because we do not have to deal with transaction isolation – this is almost never the case and such a mode of operation is utterly limiting when considering data that is subject to incremental change, e.g., when performing online analysis or when working with applications that offer CRUD support. A similar argument can be made for the manual selection of query execution paths. Even though it may be simplifying the process of query planning, such a design decision falsely assumes, that a user is always a technical expert and that all the potential query workloads are known in advance. Consequently, it limits the optimisation options available to the system.

On a more general point, it cannot be assumed that all the query workloads a multimedia database management system must handle can be anticipated at build time, especially if a session involves a user in the loop. This was emphasised by Smeulders et al. in 2000 already, who stated that “*the interacting user brings about many new challenges for the response time of the system. Content-based image retrieval is only scalable to large data sets when the database is able to anticipate what interactive queries will be made. A frequent assumption is that the image set, the features, and the similarity function are known in advance. In a truly interactive session, the assumptions are no longer valid. A change from static to dynamic indexing is required.*” [SWS⁺00] (p. 1369). This statement is still of high, practical relevance and summarises what we have expressed in the three assumptions. It can thus be regarded as yet another argument to consider multimedia data management as an important and interesting research challenge.

1.3 Contribution and Outline

In this thesis we address the research gaps identified in Section 1.2 and thereby try to bridge the disparity between the fields of databases and multimedia retrieval systems. The contribution of this thesis can be summarized as follows:

- We propose a *generalised model for proximity-based search* as an extension to the relational data model and explore the implications of such a model on aspects, such as, optimisation of query execution time and query planning.
- We propose a *cost model that factors in quality* of generated results in addition to common performance metrics, and, based on that model, derive mechanisms for the user to express their preference for either accuracy or speed at different levels of the system.
- We examine the impact of challenging the *data staticity assumption* on index structures commonly used for similarity search and describe an *high-dimensional index maintenance* model by which a multimedia database system can cope with and make sense of changing index structures.
- We introduce the multimedia database Cottontail DB [GRH⁺20] that implements the aforementioned models and present an evaluation thereof.

This thesis is presented in four parts, each of which consists of several chapters that incrementally build on one another.

Part I provides an introduction and problem statement (Chapter 1) and describes motivating use cases and the requirements they bring (Chapter 2).

Part II gives a brief summary of the theoretical foundation and state of the art in database (Chapter 3) and multimedia analysis and retrieval (Chapter 4) research and tries to combine the two aspects on a systems level in a brief survey of multimedia database systems (Chapter 5). In addition to establishing a common foundation and framework for the remainder of this thesis, we also use this part to survey related work.

Part III formalises the theoretical models that make up the contributions (Chapter 6) and introduces and describes our reference implementation Cotton-tail DB (Chapter 7).

Part IV presents and discusses the evaluation of the contributions (Chapter 8) and provides a conclusion (Chapter 9).

2

Science and everyday life cannot and should not be separated.

— Rosalind Franklin, 1940

Applications and Use Cases

A fascinating aspect of computer science and information technology is, that it has applications in every conceivable domain. It is this combination of purely technological aspects – the computational part – and the domain knowledge, that make the field challenging as well as interesting. Digital transformation – the act of improving businesses and processing using digital technology [Via19] – has been a driving force behind economic activity over the past decades, and while many non-technological factors define digital transformation, optimisation through the application of disruptive technologies remains the main incentive behind it.

Consequently, we do not want to detach the work presented in this thesis from its concrete applications and instead aim to motivate it based on real-world scenarios. In the sections that follow, we will therefore be introducing three use cases, for which we believe that a general-purpose multimedia database system would fulfil an important function. From these use cases, we then go on to derive a series of requirements for such a multimedia database.

2.1 Use Case 1: Interactive Multimedia Retrieval

Multimedia retrieval in a broader sense describes the act of finding items of interest in large multimedia collections, wherein “multimedia” can refer to any type of modality, such as text, images, videos or audio and any combination thereof. Consequently, a multimedia retrieval system must be able to have a user express their *information need*, which is translated to a *query* that the system can use to produce the item(s) of interest as a *result*. The process is visualised in Figure 2.1 and a formal problem definition will be provided in Chapter 4.



Figure 2.1 The idealised multimedia retrieval problem: A user expresses an information need as a query and the multimedia retrieval system returns matching items from a collection as a result.

There are many potential applications of multimedia retrieval. For example, one could consider libraries of media items produced by radio or TV stations [WOK⁺98], collections of cultural heritage data as maintained by museums, archives, and archaeology departments [Tsa07], or medical image databases containing X-rays or MRIs [MMB⁺04] that must be made searchable.

What makes multimedia retrieval a challenging problem, is the inherently unstructured nature of the data, with the consequence that the raw, (digital) information and its derivatives may not directly reflect the content of the item as a human consumer might perceive or describe it. Let us take, for example, the image of a rabbit the user in Figure 2.1 is looking for:¹ If our task were to describe that image so that we can retrieve it from a large database, we might use terms like “rabbit”, “bunny” or “grass”. We might assume, judging from the image, that the scene is taking place in a “garden” and maybe we notice that the rabbit is currently “munching” on a “leaf” or “blade of grass”. We might even know the rabbit’s breed or name, if we happened to be its owner.

All of this information is formed based on a series of steps ranging from perception and interpretation to cognition as illustrated in Figure 2.2. Each of these steps is accompanied by a loss, distortion or framing of the available information [JLX21], which makes this process highly subjective and context dependent (e.g., based on prior knowledge). Most importantly, however, all the aforementioned descriptions are not explicitly contained in the raw image data, which is merely an array of colour values that constitutes the image and usually does neither come pre-labelled nor pre-described. Consequently, in order to facilitate retrieval, a system must hold the same descriptions.

¹ Intentionally, we do not consider the aspect of reconstructing the image from memory, which makes the problem even more challenging.

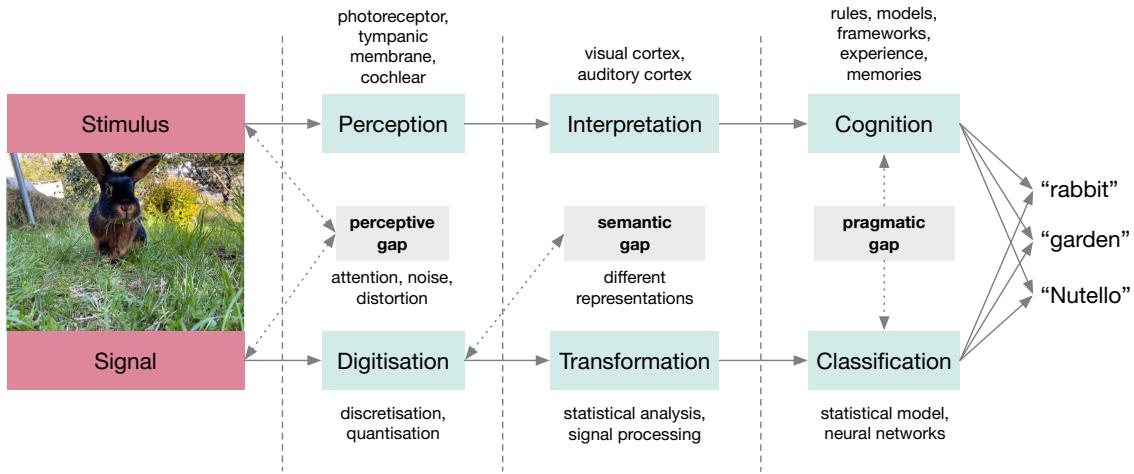


Figure 2.2 A simplified model of knowledge formation from an analogue input in humans (adapted from [JLX21]) and computers.

This impedance mismatch between the content being stored and the information required to find it is the main challenge that unstructured data brings with it, as opposed to structured data, which by its very nature consists of retrievable units of information. Bridging this mismatch is therefore one of the major obstacles to overcome in multimedia retrieval. Over the years, different solutions have emerged:

Manual Labelling The obvious solution is to manually embed the aforementioned labels and descriptions into the multimedia item so that we can use this information to find it at a later point in time. While as of today this approach is still widely used, it comes with two important disadvantages: It remains a subjective task, because as we have argued, the process of forming these labels may yield different results depending on who is assigning them. Most importantly, though, the task of manual label assignment is laborious and cannot be scaled to the large volumes and the ever-increasing velocity at which multimedia data is produced. This is related to the general challenge of working with Big Data [KUG14].

Automatic Labelling Recent advances in machine learning have made it possible to automatically label [RDG⁺16] and describe [RKH⁺21] different types of multimedia using high-level concepts. So instead of annotating the content manually, one can leave this to pre-trained models that, at a high level, have a similar mode of operation as a human operator would have, as is illustrated in Figure 2.2. While this solves the problem of scalability, training such models still remains, to some extent, a subjective task, which can lead to biases [BK17]. Furthermore, there typically also are discrepancies between

the original signal (e.g., the image) and the representations used in the process, due to the perceptive [Ros18] and semantic gap [ZAD⁺06]. In addition, there is often a mismatch between the model used for classification and the mental model employed during query formulation (pragmatic gap [ZW14]), which are typically carried out at different points in time by different people.

Content-Based Search The last approach and the one that “classical” multimedia retrieval has concerned itself with for decades is that of content-based search. Instead of using the higher-level concepts in the form of textual labels and descriptions, content-based search works with intermediate representations, called descriptors, which are derived from the data directly using signal processing or statistical analysis. These descriptors – which are often real-valued vectors [ZAD⁺06] – can then be used to establish a notion of similarity between items in a collection and a query. Instead of a list of exact matches, this type of query returns a ranked list of items that match the query well enough. This is called similarity search [BdB⁺07] and in this context, the query may refer to more than simply textual input. Given the example in Figure 2.1, one could use another image of a rabbit (Query-by-Example, [KCH95]) or try to create a sketch (Query-by-Sketch, [SMM99]). One could even come-up with more sophisticated types of queries, such as searching for poses held by people in an image or video [HAG⁺22] or using visual-text co-embeddings [RKH⁺21; SGH⁺22] to map textual descriptions to images and vice versa.

Research has shown that while individual methods falling into any of the three categories may yield acceptable or even exceptional results, there is clear indication that ultimately, a combination of the three is the most powerful solution, especially, since the exact type of information need cannot always be anticipated [RGL⁺20]. Furthermore, and partly because different techniques may perform better or worse depending on the information need, we often find that instead of the simple “issue a query, get a result” scheme illustrated in Figure 2.1, actual multimedia retrieval applications are more interactive and require a back and forth between system and user [LBB⁺22] that involves query (re-)formulation potentially using different descriptors, exploration, examination, query refinement and relevance feedback [LKM⁺19; GSJ⁺19] as illustrated in Figure 2.3. These types of interactive multimedia retrieval systems are explicitly tested in evaluation campaigns, such as the Video Browser Showdown (VBS) [Sch19] or the Lifelog Search Challenge (LSC) [GJS⁺21].

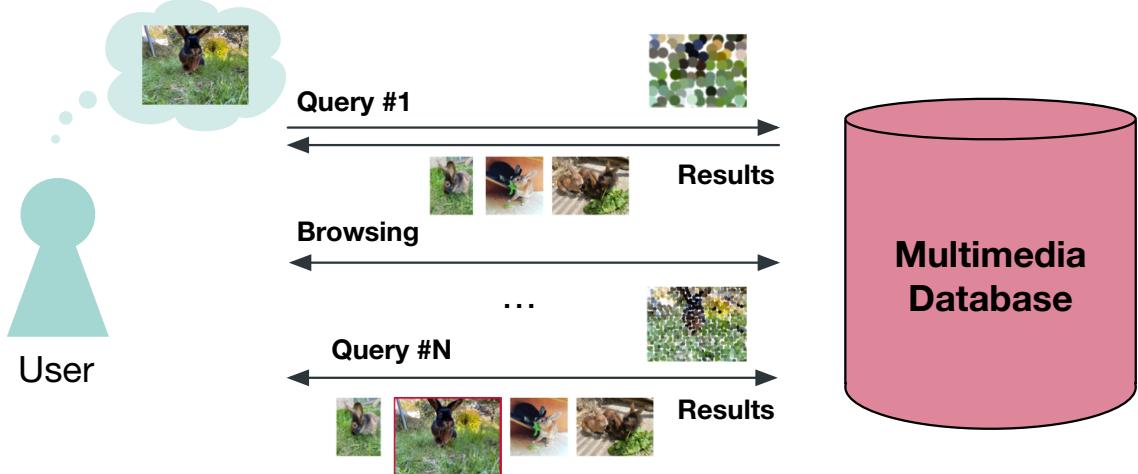


Figure 2.3 The interactive multimedia retrieval problem: Even when searching for a known item, the search effort typically constitutes a continuous back and forth between system and user involving different types of workload as the user tries to find the item of interest through query (re-)formulation, browsing and examination.

Both aspects have important implications for data management systems that must support (interactive) multimedia retrieval: Firstly, such a system must consider different data and query models when dealing with information from the aforementioned categories as well as a wide range of query workloads. While searching for labels and textual descriptions at scale can be achieved by the well-established *Boolean retrieval* model, we require the completely different *vector space retrieval* model when dealing with descriptors derived from multimedia data and, ideally, we should have the ability to combine the two [HSS⁺20]. This problem has been acknowledged by other authors before [JWZ⁺16] and that has been partially addressed, for example, by [Gia18; GS16; WYG⁺21].

Secondly, any database management system supporting interactive multimedia retrieval must be able to satisfy the different types of query workloads in a reasonable amount of time, so that the user in the loop must not wait for results too long. One can even state that retrieval speed is more important than perfect accuracy, especially in timeboxed settings, such as VBS or LSC, because (i) the retrieval process itself is inherently inaccurate due to the previously-named gaps, and, (ii) there is a user in the loop who examines results and makes decisions.

A final aspect to consider is the staticity of the data itself. In multimedia retrieval we often implicitly assume that data collections remain static while they are being queried but it must be emphasised that for most use cases, this is very likely not to be the case, either because other users input data simultaneously or because certain models make adjustments to the data themselves.

2.2 Use Case 2: Analysis of Social Media Streams

Social media platforms such as Twitter, Facebook, Instagram and TikTok have gained enormous traction over the past few decades. Current estimates suggest that there are roughly 4.66 billion active Internet users worldwide, 4.2 billion of which can be considered active social media users². This is an enormous market that can be used to place advertisements, reach customers or potential electors and to observe consumer reaction to products and services.

Consequently, social media constitutes an important communication channel for companies, politicians and other interest groups [BZ18]. The field of social media analytics aims “*to derive actionable information from social media*” [ZCL⁺10] (p. 15) for economic, political or scientific purposes. Probably one of the most prominent research challenges involves the identification and countering of “fake news” [LBB⁺18], by which we mean information that comes disguised as authentic but is in fact fabricated, misleading, or even wrong. The term became highly polarised in the 2016 U.S. presidential election [QFB⁺19] – when orchestrated campaigns on social media were targeting voters and may have tipped the scales – and it has remained important ever since [FCC⁺20].

While the problem of fake news can arguably not be solved by mere technological means, the issue has brought about many different research questions in the realm of social media analysis and analytics with potential applications in different domains. One focus lies on the direct detection of misinformation on different channels by different means, such as content, propagation patterns [ZZ20] or sources. Others try to tackle the detection of bots [DVF⁺16; Cre20]—user-accounts that are not backed by real humans and that are often used to create or spread misinformation. And then there is the topic of sentiment analysis [YCL⁺19], which can yield important clues about how (fake) news stories are received by their consumers.

While many of the aforementioned problems primarily deal with textual information, multimedia obviously plays an important role as well. Instagram and TikTok, for example, focus solely on the sharing of images and videos respectively. A very recent and preliminary analysis suggests [Ciu22], that the sharing of images and videos plays an important role in the international covering of Russia’s war on Ukraine. Solutions to potential problems could, for example, be the automatic detection of image forgeries [Far09] or the analysis of the propagation patterns of images [ZCB⁺18].

² See Statista.com, “Social media usage worldwide”; accessed January 2021.

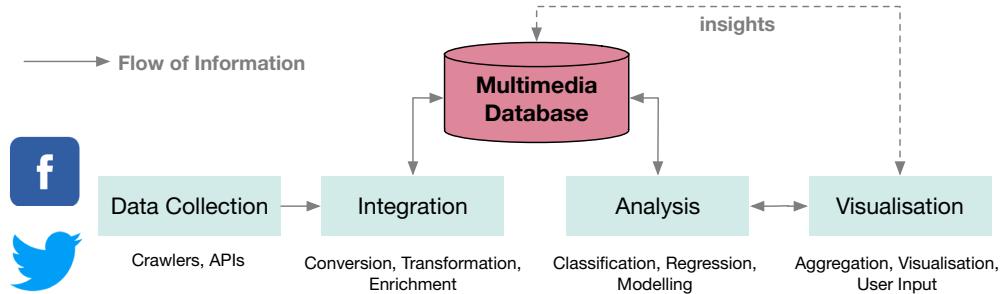


Figure 2.4 The architecture of a fictitious social media analytics platform (based on [ZCW⁺¹⁵; CSW⁺¹⁹; YPM⁺¹⁹; BPS⁺²⁰]).

If we examine what a fictitious and generalised pipeline for social media analysis and analytics would look like, we might arrive at something as illustrated in Figure 2.4. The many different channels in existence require a broad range of tools that deal with continuous data collection using crawlers or APIs provided by the respective source. Since all these potential sources are likely to use very different types of data models, some sort of integration into a common framework is necessary. This step may also be accompanied by data enrichment using external services or data existing in the system's own collection. This is then followed by what we call analysis, which may include a wide range of machine learning and data analysis techniques that may or may not rely on existing data. A common theme for this step is that, similarly to multimedia retrieval and analytics, it probably involves generation and storage of formal descriptors of the multimedia content. Therefore, the entire pipeline is very likely to be backed by some type of persistent data store, which contains labels, descriptors, and metadata and which allows queries by both the analysis and the visualisation components built on top of it [VHC⁺¹⁹]. Overall, the architecture resembles the model proposed in [ZW14].

From the described architecture we can derive certain needs with respect to the multimedia database: Firstly, for this use case, the database must be able to deal with non-static data, since the various sources potentially generate a steady flow of new information which is likely to be contributing to future inquiries. Secondly, and similarly to the multimedia retrieval problem, the pipeline is likely to deal with structured, semi-structured and unstructured information that must be stored and processed jointly, and the multimedia database must be able to cope with this variety. And finally, with social media being a very classical field for applying multimedia analytics techniques [PYC⁺¹⁸; JWZ⁺¹⁶], the different tools built upon the multimedia database are likely to rely on a range of query workloads, which the database should be able to accommodate.

2.3 Use Case 3: Signal Analysis in Medical Imaging

Multimedia retrieval for medical applications has been an important area of research for many decades now [MU17; MMB⁺04] and many of the arguments that we have made in Section 2.1 can be repeated for this specialised sub-domain.

However, in this section, we focus on a very specific application, in which we do not consider media in the classical sense but in a broader sense of raw signals stemming from a Magnetic Resonance Imaging (MRI) device. One can also refer to this as “science data” [SBZ⁺13]. MRI is a non-invasive, non-ionizing imaging technique that has gained huge importance in various fields of modern medicine. MRI is enabled by Nucleic Magnetic Resonance (NMR), which was first described by Purcell and Bloch in 1946 [Blo46; PTP46].

The process behind it is roughly illustrated in Figure 2.5. In simple terms, an MRI machine operates by applying a very strong, external, magnetic field – referred to as B_0 – which forces nuclei with a non-zero spin, e.g., an ^1H nucleus (proton), to align to and precess around it. That static field is disrupted by applying a second radio-frequency pulse that interferes with the nucleus’ alignment. As this pulse is switched off, the protons start to re-align with B_0 over time in what is called relaxation. This process is characterised by the relaxation times T_1 (spin-lattice relaxation) and T_2 (spin-spin relaxation), which are specific to a type of tissue, and it is accompanied by a measurable, oscillating magnetic field – the free induction decay – that induces a current in the receiver coil. This recorded signal can be used to reconstruct magnetic properties of the material. To optimise the signal to noise ratio, the procedure is repeated multiple times. Furthermore, spatial resolution is attained by repetition while scanning the body or a specific part thereof. Both aspects contribute to a direct trade-off between acquisition time and image quality.

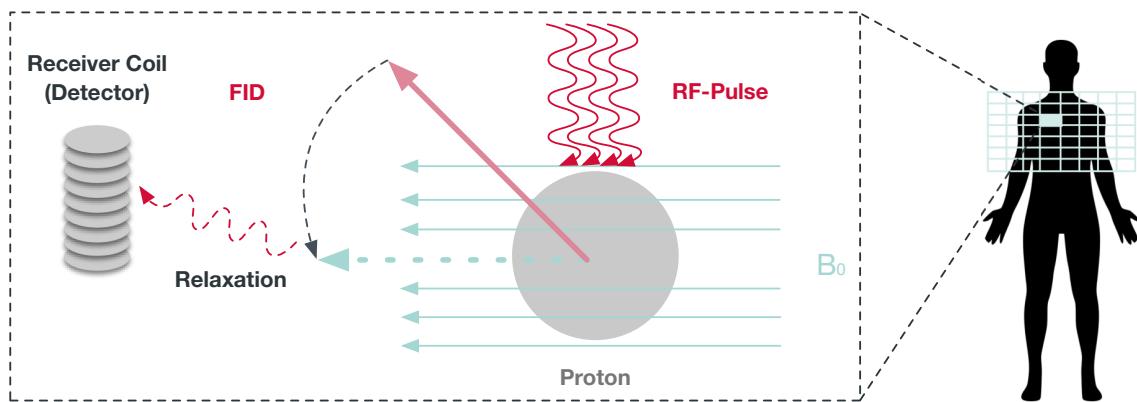


Figure 2.5 Simplified illustration of the physical processes enabling MRI.

MRI has the ability to generate fine-grained contrasts in different types of soft tissue. However, it comes with two critical disadvantages: Firstly, there is an inherently long acquisition time that can range from several minutes to over an hour per scan, especially when one tries to achieve quantification of magnetic properties of soft tissue. This is an issue because the patient is required to lie still for that time, which they are often unwilling or unable to do. Secondly, and due to this lack of fast, quantitative tools, MRI scans have become mostly qualitative in nature. Anatomical areas are often referred to as either “hyperintense” or “hypointense” with respect to their immediate surroundings, but the difference between such areas or even the absolute values themselves are not measured directly.

Magnetic Resonance Fingerprinting (MRF) is a technique [MGS⁺13] that offers a solution to this problem and allows for the quantitative mapping of multiple properties simultaneously while lowering acquisition time. The method is described in [BMCFM⁺19] and illustrated in Figure 2.6. In summary, MRF is based on the assumption that different types of tissue produce a distinct signal evolution – called *fingerprint* – when exposed to a specific acquisition scheme. This means that during data acquisition, measurements take place under varying (pseudo-random), external magnetic field and RF-pulse configurations and signals are recorded for the different configurations resulting in the fingerprint. That fingerprint can then be matched against a database (called dictionary) of pre-calculated, discretised fingerprints that have been simulated using Bloch’s theorem of magnetic resonance [Blo46]. Through that matching step, values for B_0 , T_1 and T_2 (and other parameters) can be obtained and subsequently used to reconstruct an image. Spatial resolution in MRF – similarly as in MRI – is attained by scanning, and the described process is therefore repeated.

Most MRF setups implement the pattern matching step as an exhaustive search that retrieves the simulated entry from the dictionary that maximises the inner product with the obtained signal [BMCFM⁺19]. This is called Maximum Inner Product Search (MIPS) and it involves evaluating the inner product between a complex signal vector and all the vectors in the database. This problem is not unlike similarity search used in multimedia retrieval, with the difference that the vectors are complex and not real-valued, that the inner product is maximised instead of minimised and that there is interest in a single result per query only. This step constitutes one bottleneck of the entire process, since many queries must be performed in sequence. Several attempts to accelerate it have been made over the years [MPM⁺14; CSM⁺15; CZR18].

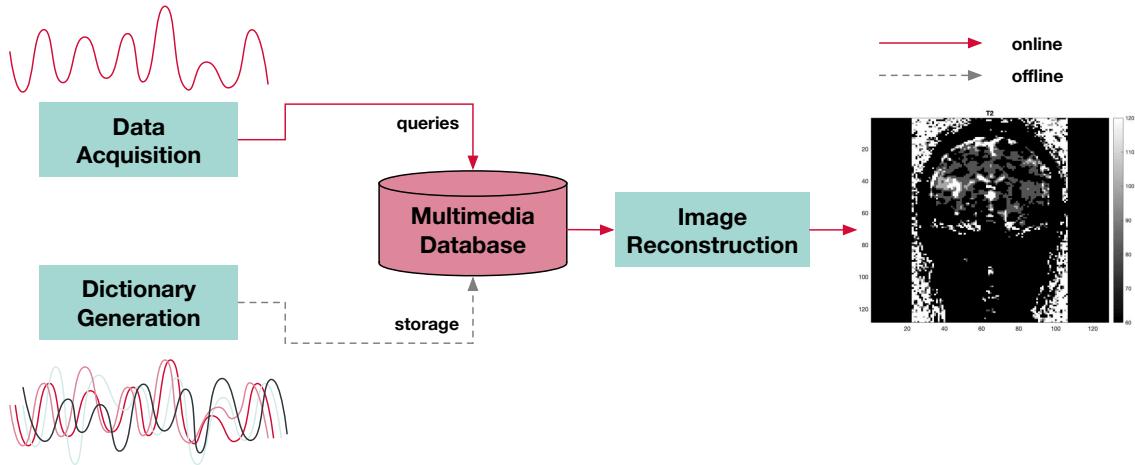


Figure 2.6 Simplified block diagram of a MRF system (adapted from [BMCFM⁺19], reconstruction of brain scan kindly provided by Manuel Hürbin). Storage and indexing typically take place offline.

Once again, we assume that this pattern matching and image reconstruction step should be supported by a dedicated multimedia database, which is a reasonable assumption once dictionaries become so large that they no longer fit into main memory. And again, we can try to describe the requirements imposed on such a database:

Firstly, and similarly to the multimedia retrieval problem, the database requires a dedicated data- and query model to store complex vectors and perform MIPS. This model differs from the classical, metric-space model applied in multimedia retrieval because it does not restrict itself to real-valued vectors. While considering complex vectors has been discussed in previous work already [RM97], it remains a niche for specialised applications. At the same time, the database must also be able to store and access very classical data that fits into the Boolean retrieval model such as the individual parameters B_0 , T_1 or T_2 .

Secondly, and also similar to interactive multimedia retrieval, query time is an important aspect and should be minimised as much as possible, especially in a clinical setup. However, in contrast to multimedia retrieval, this optimisation cannot be traded against retrieval quality, because the MRF problem does not have a human in the loop who might correct an error in an individual query. Instead, the process relies on the lookup of the correct signal vector for every single query. Preliminary work on applying multimedia retrieval techniques to speeding up MRF pattern matching has shown that sacrificing accuracy leads to sub-optimal image reconstruction results [Hür20; Zih21] and thus high-dimensional indexing methods known from multimedia retrieval may not be applicable.

2.4 Deriving Requirements

In Sections 2.1, 2.2 and 2.3 we have discussed three applications that, (i) deal with large amounts of unstructured (multimedia) data that must be stored, managed and queried, (ii) use some mathematical descriptors of that unstructured (multimedia) data to enable analysis, comparison and queries, (iii) and that combine these descriptors with data such as labels, descriptions and other information, that must be accessed and sometimes queried as well.

We acknowledge, that while the details may vary, the basic needs outlined in the aforementioned use cases are overlapping to such an extent that they would benefit from a dedicated multimedia database management system. In fact, the argument for such a system has been made by different authors before [AN97; SWS⁺00; JWZ⁺16; JGA⁺19; Kha21] and steps towards such a system have been taken [Gia18; WYG⁺21]. However, all concrete approaches somehow narrow the scope in a way that restricts their generalisability.

In this section, we assume that a general-purpose multimedia database is desirable, and based on the arguments presented thus far, we formalise the requirements imposed on such a system. We do not claim that we address all of the listed requirements in this thesis. We merely attempt to raise and formalise them, similarly to and based on [AN97; JWZ⁺16; Kha21].

2.4.1 Traditional Database Capabilities

Regardless of application, any multimedia database is very likely to rely on traditional DBMS capabilities at some point. Be it when data changes or is accessed concurrently and guarantees regarding consistency must be provided or when executing Boolean queries on some aspect of the available metadata. This has been pointed out by [AN97; Kha21] and we can therefore reiterate, that a multimedia database must possess the same abilities that any traditional DBMS would with respect to data modelling, data management, execution of queries and the enforcement of guarantees such as ACID [HR83] or BASE [Pri08].

Requirement 2.1 A Multimedia Database is a DBMS

At its core, any general-purpose multimedia database is a classical Database Management System and should provide the same functionality in terms of data modelling, data management, query execution and guarantees with regards to data integrity, concurrency control, data persistence, and recovery.

2.4.2 Multimedia Query Support

Multimedia queries are fundamentally different in that, as opposed to Boolean queries, they do not produce exact results but instead rely on scores and ranking to quantify similarity between a query and an object in a database. A multimedia database system must therefore have the ability to support such queries at all levels from query formulation (i.e., query language) to query execution [AN97]. Furthermore, the combination of similarity search with Boolean queries should be possible and well-defined in terms of outcome.

It is worth noting, that by “Boolean query” we merely refer to the ability to produce exact matches based on Boolean predicates without making an assumption about the underlying data model. If a multimedia database system chooses to implement the graph model, then traditional graph traversal queries fall under the category “Boolean query” as well.

Requirement 2.2 Multimedia Query Support

A multimedia database must be able to support multimedia queries based on similarity and scores in addition to Boolean queries, and it must allow the combination of the two search paradigms in a unified framework.

2.4.3 Multimedia Data is Big Data

At its core, dealing with multimedia data equates to dealing with Big Data, in that a system deals with large *volumes* of information, that may change at a high *velocity* and exhibit large *variety*, ranging from highly structured metadata attributes to completely unstructured, raw data.³ A multimedia database can, of course, favour one of the three aspects over the others (e.g., assume that velocity at which data changes is zero) but in general, all three aspects must be considered jointly, which results in interesting engineering and research problems, e.g., for high-dimensional indexing [HJB19] or distribution of data and computation.

Requirement 2.3 Multimedia Data is Big Data

A multimedia database must be able to deal with large volumes of heterogeneous data that may be subject to change at a high velocity within a single, well-defined framework.

³ This is known as the “3 Vs of Big Data”. Over time, more of these V’s have been described in literature, for example in [KUG14].

2.4.4 Modelling Retrieval Quality

Classical databases produce results that match or do not match a given query, i.e., there is no notion of uncertainty or imprecision in the results. In contrast, the models used in multimedia retrieval are inherently inaccurate to some extent, in that the “best matching” entries are obtained, meaning that even bad results may be produced for the lack of better options. This imprecision or inaccuracy is influenced by many factors, not the least of which are the design and choice of the multimedia descriptors themselves.

While the inherent (in-)accuracy of any given descriptor is impossible to predict and thus out of scope for the database that holds it, the notion of results quality still opens a relevant avenue, since some of that quality can be deliberately sacrificed in order to gain speed, which is often done in high-dimensional indexing. Since some use cases are very robust to such decisions while others are not, this aspect must be made explicit at a system level.

Requirement 2.4 Model for Retrieval Quality

A multimedia database must employ an internal model of quality of results. That model must (i) enable a user to specify the desired quality of the result, (ii) allow the database system to derive the expected quality for a query execution path, and, (iii) take both aspects into account when planning query execution.

2.4.5 Complex Data Types as First-class Citizens

Since multimedia analysis and analytics workloads rely on mathematical objects such as vectors, complex numbers, or matrices, we argue, that the data model of a multimedia database must be extended to consider those types of data as *first-class citizens* in addition to well-established data types such as numbers or strings. This means, that the structure as well as the mathematical properties of those types must be considered for storage, query planning and execution.

This is not a new idea, per se, and to some extent self-evident since new use cases require additions to existing data models. Therefore, the idea of extensible data models and type systems was already discussed as early as 1988 [LKD⁺88] and it was also argued by [Gia18] that to support multimedia retrieval workloads in a relational database, one must extend the relational model to support those vectors as basic data types. However, it has been pointed out by Michael Stonebraker that the mere ability to model the data, e.g., in the scope of the

relational model, is not enough to provide satisfactory performance [SBZ⁺13]. He states: “*In summary, RDBMSs seem inappropriate for the vast majority of science applications because they have the wrong data model, the wrong query language, lack important features, and do not provide scalability in an open-source code base*” [SBZ⁺13] (p. 56). Similarly, [WLK⁺10] demonstrated that the mere ability to extend a type system with user-defined types through (what they call) loose-coupling, is not enough to attain satisfactory performance for queries involving spatial data. The same observation was made when adding information retrieval support [WLL⁺15] and we add that this argument can be extended to the handling of multimedia data and associated data types, which are potentially even more complex.

Requirement 2.5 Complex Data Types as First-class Citizens

A multimedia database must treat composite data types (e.g., vectors, complex numbers, matrices, tensors) as first-class citizens of the data-, storage- and execution-model in addition to primitive types such as numbers or strings.

2.4.6 Functions as First-class Citizens

All applications considered so far involve the execution of functions that operate on complex data atoms, e.g., to obtain the similarity between a query and items in the database. For a multimedia database to be able to operate effectively, those functions must be considered first-class citizens as well, in that the multimedia database must have the ability to reason about them and make decisions regarding their execution. Hence, they should not be regarded as mere black-boxes, as, for example, a User Defined Function (UDF). We argue, that similarly to user-defined types, there is a trade-off between loose and tight integration of such functions into the DBMS.

This has also been laid out in the 2020 *Seattle Report on Database Research*, where the combination of “*relational algebra and linear algebra in a richer query paradigm*” [AAA⁺20] (p. 52) is listed as a possible research focus.

Requirement 2.6 Functions as First-class Citizens

A multimedia database must treat mathematical functions as first-class citizens of the execution model and must have the ability to reason about those functions’ properties during query planning and execution.

2.5 Mapping Requirements to Contributions

In reference to the contributions listed in Section 1.3, we can make the following mapping between those contributions and the requirements listed in Section 2.4:

Generalised Model for Similarity Search addresses Requirement 2.2 on a theoretical level in that it combines the different query paradigms in a consistent yet versatile framework. The proposed model and especially its implementation also address Requirement 2.5 and Requirement 2.6 to some extent, with one focus lying on the optimisation of function execution within a query as well as the connection between index structures and the execution of such functions. However, there remain many aspects to be explored.

Adaptive Index Management provides an answer to the velocity aspect of Requirement 2.3, because managing indexes in the face of changing data is a consequence of changing data. As we have argued, these changes can take place fast and they take place concurrently to other query operations, which the multimedia database must handle elastically. The proposed solution is a general-purpose approach and can be complemented with more specialised optimisations at an index level.

Cost Model for Quality is a direct consequence of Requirement 2.3 because, as we will demonstrate, responding to changes in the data can lead to the deterioration of indexing quality. Hence, the multimedia database is forced to choose one over the other. Since some use cases may tolerate such deterioration while others do not, the contribution also directly addresses Requirement 2.4.

PART II

Foundations

Preface to Part II

Since the contributions of this thesis rely on decades of research in database systems as well as multimedia retrieval and analysis, we use this part to provide an overview over the relevant fundamentals. Furthermore, we use it to survey related work and to establish the common terminology, which we think is necessary to understand the contributions of this Thesis.

Chapter 3 discusses the theory surrounding Database Management System both at a data model and systems level. Most of the fundamentals are guided by the books *DATABASE SYSTEMS The Complete Book* [GUW09] and *Database Internals* [Pet19]. Furthermore, some information was adapted from a video lecture kindly provided by Andy Pavlo, Associate Professor at Carnegie Mellon University. That lecture is publicly available on YouTube.⁴

Chapter 4 shifts the focus to the analysis and retrieval of multimedia data – including but not limited to images, audio, and video. In addition to the basic techniques, we also discuss concrete system implementations. The two major sources used in this chapter are the books *Multimedia Retrieval* [BdB⁺07] and *Similarity Search – The Metric Space Approach* [ZAD⁺06].

Chapter 5 tries to combine the two previous chapters and briefly surveys what efforts have been made to converge the two domains, which are, to this day, still largely separate areas of research. This is where the transition to the next part is staged.

All the aforementioned textbooks and resources provide an excellent overview of their respective fields and this is as a good time as any to thank the authors for their invaluable work. While we do indicate when specific ideas were taken directly from those sources, some information that we considered to be “general knowledge” is not always cited explicitly.

⁴ See <https://www.youtube.com/c/CMUDatabaseGroup/>; accessed February 2022.

3

*The key, the whole key and
nothing but the key, so help me
Codd!*

— Mnemonic

On the Design of a Database Management System

Database Management Systems (DBMSs), or simply “databases”, power everything from small and simple websites to large data warehouses that serve millions of users simultaneously. These systems play a crucial role in banking, e-commerce, science, entertainment and practically every aspect of our socio-economic lives, because more and more we rely on the processing and the insights generated from the analysis of data [Dha13] and DBMS form the backbone of many (if not most) information and data processing systems.

The first commercial DBMS were introduced in the 1960s [GUW09] and they have evolved ever since to adapt to a wide range of requirements, leading to a long list of open-source and commercial systems such as PostgreSQL, Microsoft SQL or MongoDB. Even though many different flavours of DBMS have emerged over the years, at their core, they still serve the same purpose:

Management DBMS enable their users to organise data corpora that can range from a few megabytes to hundreds of terabytes according to a data model. For example, data can be structured into documents [HR16], graphs [AG08], or tables [Cod70] that in turn can be organised into collections or schemata.

Definition DBMS provide users with the ability to alter the organisation of their data within the constraints of the data model using a Data Definition Language (DDL).

Manipulation DBMS provide users with the ability to modify the data within the constraints of the data model using a Data Manipulation Language (DML). Modifications may include adding, removing, or changing entries.

Querying DBMS provide users with the ability to query the data using a Data Query Language (DQL). Such queries can be used to answer specific “questions” about the data and to generate the aforementioned insights.

Guarantees DBMS usually provide guarantees, such as assuring durability upon failure or providing access control for concurrent operations. A well-known set of guarantees offered by many modern DBMS is known by its acronym ACID [HR83], which was later complemented by the more relaxed BASE guarantees offered by many NoSQL systems [Pri08].

3.1 Data Management and Data Models

A data model is a formal framework that describes any type of data or information. It usually involves a formal description of the data’s *structure*, the *operations* that can be performed and the *constraints* that should be imposed on the data [GUW09]. The purpose of any data model is to formalise how data governed by it can be accessed, modified, and queried, and any given DBMS usually adopts a specific type of data model (or multiple models, for that matter).

In the context of database systems and data management, it has become common practice to distinguish between different levels of abstraction for data models, as indicated by Figure 3.1. At the top, there is the *conceptual data model*, which is often closely tied to some real-world object, fact, or process. For example, in an online shop, one can think in terms of customers that place orders, products that are being sold, and invoices that must be sent out. In the Entity Relationship Model (ERM) [Che76] – a popular framework used to describe conceptual data models – those “real things” would be modelled as *entities* that come with *attributes* that describe them and that have *relationships* among them as required by the concrete application.

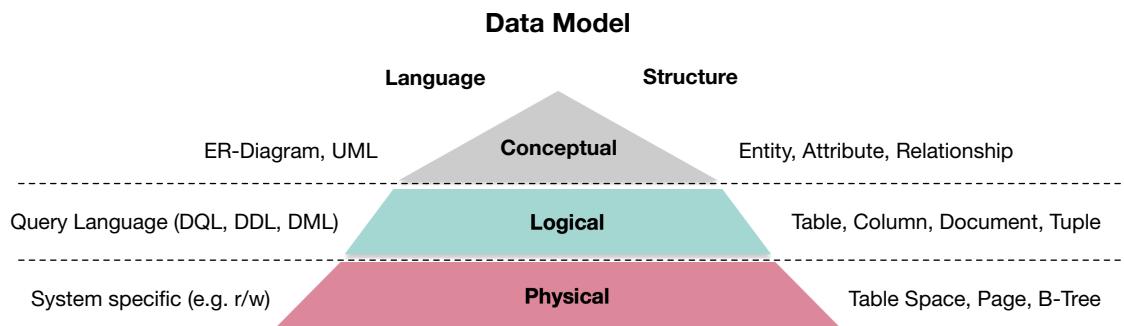


Figure 3.1 Different levels of abstraction for data modelling.

The *logical data model* is more closely tied to the type of database being used. Different logical data models have been conceived over the years, including, but not limited to, models centred around documents [HR16], graphs [AG08], key-value pairs, and tables [Cod70], each with their own (often domain-specific) advantages and disadvantages. At this level, one can use DDL to describe and modify the data organisation and DQL or DML to query and modify the data itself. As an example, one could store each entity described before in a dedicated table wherein each attribute occupies a different column. This is known as the relational data model [Cod70], and an important query language for the relational model is the Structured Query Language (SQL), which became an international standard in 1987 under ISO 9075 and has evolved ever since [Cha12].

At the very bottom, we usually find the *physical data model*, which is specific to the database implementation and describes low-level aspects such data access in terms of *read* and *write* operations to data structures such as *table spaces*, *data pages*, or *trees*. An important argument in favour of separating the logical from the physical data model, even though they are both somewhat specific to a DBMS implementation, is that the end-users should not concern themselves with how data is organised and accessed at the lowest level but should instead describe their intent in terms of the information they wish to retrieve. Mapping user-intent to operations then becomes a task of the DBMS.

3.2 The Relational Data Model

In June 1970, E. F. Codd published his pivotal research paper *Relational Model of Data for Large Shared Data Banks* [Cod70], where he describes a logical data model, which he himself refers to as “relational”. This model has become the fundament on which many modern database management systems have been built, with examples dating back to the 1970s [ABC⁺76] and prominent, contemporary examples including systems such as Maria DB, PostgreSQL, or Oracle.

The relational model is structured around *relations*, which are a mathematical construct but can be visualised as two-dimensional tables. Such a table consists of columns – which are called *attributes* – and rows – which are called *tuples* and hold *attribute values*. Semantically, a relation can be seen as a knowledge-base about some fact – such as, the paintings held by a museum – under a Closed-world Assumption (CWA) [Rei81]. That is, the relation contains all the information available about the fact and can thus be used do derive conclusive answers or results given a stated question or query.

In order to formalise the structure of and the operations that can be executed on the data represented by a relation, one can turn to Definition 3.1 [Cod70].

Definition 3.1 Relation

Let \mathcal{D}_j be sets we call *data domains* with $j \in [1, N]$ and $N \in \mathbb{N}_{>0}$. A *relation* \mathcal{R} constructed over these data domains is a set of M *tuples* $t_i = (a_{i,1}, a_{i,2} \dots a_{i,N})$ with $i \in [1, M]$ and $M \in \mathbb{N}_{>0}$ such that the first attribute value of any tuple comes from \mathcal{D}_1 , the second from \mathcal{D}_2 , and so forth. That is, values $a_{i,1} \in \mathcal{D}_1, a_{i,2} \in \mathcal{D}_2 \dots a_{i,N} \in \mathcal{D}_N$. A relation can thus be seen as a subset of the Cartesian product of all data domains over which it was constructed, i.e., $\mathcal{R} \subset \mathcal{D}_1 \times \mathcal{D}_2 \dots \times \mathcal{D}_N$.

The number of data domains N is referred to as its *degree* and we call such a relation *N-ary*. The number of tuples M is called the *cardinality* of \mathcal{R} with $M = |\mathcal{R}|$. It must be noted, that the relational model does not dictate what the data domains \mathcal{D} are (apart from the constraint that its values must be atomic). However, in a relational DBMS, they usually correspond to the data types supported by the database and the programming environment it was written in, for example, integer and floating-point numbers or text. In this Thesis, we therefore sometimes use *data domain* synonymously with *data type*.¹

In addition to Definition 3.1, we will use the Definitions 3.2 to 3.5.

Definition 3.2 Attributes of a Relation

Let \mathcal{R} be a N -ary relation over the data domains $\mathcal{D}_i, i \in [1, N]$ with corresponding names or indexes l_i . We call the combination of a data domain with a human readable label *attribute*, that is, $\mathcal{A}_i = (\mathcal{D}_i, l_i)$. To simplify notation, we sometimes use the label of an attribute in the subscript instead of the index.

Definition 3.3 Schema of a Relation

Let \mathcal{R} be a N -ary relation over the attributes $\mathcal{A}_i, i \in [1, N]$. We call the list of all attributes over which \mathcal{R} was constructed the *heading* or *schema* $SCH(\mathcal{R})$ of \mathcal{R} .

$$SCH(\mathcal{R}) = (\mathcal{A}: \text{if } \mathcal{A} \text{ is an attribute of } \mathcal{R})$$

¹ Strictly speaking, the data domains of a given relation \mathcal{R} are merely subsets of the sets that represent the respective data type which, in turn, are subsets of even more basic sets such as \mathbb{N} or \mathbb{R} for Int and Float respectively. The relationship between types and data domains is the subject of a more categorical approach to data models and is nicely laid out in [Spi09].

Definition 3.4 Accessing Attribute Values

Let \mathcal{R} be a N -ary relation over attributes \mathcal{A}_i , $i \in [1, N]$ and let further $t \in \mathcal{R}$. To address the attribute value $a_i \in t$ that belongs to attribute $\mathcal{A}_i \in \text{SCH}(\mathcal{R})$, we use the *attribute value accessor* $t[\mathcal{A}_i]$.

$$a_i = t[\mathcal{A}_i] = \{a_j \in t : i = j\}$$

Definition 3.5 Supported Data Domains

For a given DBMS, we call \mathbb{D} the set system of data domains or data types supported by the system.

$$\mathbb{D} = \{\mathcal{D} : \text{if } \mathcal{D} \text{ is supported by DBMS}\}$$

In its original form, the relational model assumes the following properties to be true for a relation \mathcal{R} and its attributes [Cod70]:

Ordering of Tuples Tuples $t \in \mathcal{R}$ are inherently unordered, and two relations are considered equal if they contain the same tuples, regardless of order.

Ordering of Attributes Attribute values a_i within the tuple $t \in \mathcal{R}$ always occur in the order of the attributes $\mathcal{A}_i \in \text{SCH}(\mathcal{R})$. This order can evolve over time but remains constant in a momentary snapshot of \mathcal{R} . It follows from the definition that $|t| = |\text{SCH}(\mathcal{R})| \forall t \in \mathcal{R}$.

Duplicates Since relations are sets, they do not allow for duplicates, i.e., every tuple $t \in \mathcal{R}$ must be unique in terms of their attribute values.

Given the idea of a relation, the sum of all data managed by a DBMS can logically be regarded as a collection of different relations \mathcal{R}_k , $k \in \mathbb{N}_{\geq 0}$ of assorted degrees N_k over data domains $\mathcal{D} \in \mathbb{D}_{\text{dbms}}$ (i.e., a collection of tables). The schema of a database can then be seen as the set system of all $\text{SCH}(\mathcal{R}_k)$. As [Cod70] points out, relations are subject to change over time. These changes can take place on the level of any relation's structure, i.e., $\text{SCH}(\mathcal{R}_k)$, the relations \mathcal{R}_k itself, e.g., by tuples being added to (insert) or removed from (delete) a relation or the level of a tuple $t \in \mathcal{R}_k$, e.g., by altering one or multiple attribute values.

Example 3.1 features an example relation painting visualised as a table. Each entry in the table represents a painting and the related attribute values.

Example 3.1 Table Representation of a Relation $\mathcal{R}_{\text{painting}}$

The following table lists the schema and extent of a ternary relation ($N = 5$) $\mathcal{R}_{\text{painting}}$. The attributes $\mathcal{A}_{\text{title}}, \mathcal{A}_{\text{year}}, \mathcal{A}_{\text{artist}}, \mathcal{A}_{\text{birth}}, \mathcal{A}_{\text{death}}$ correspond to the table's columns. The tuples t_i are “valid” combinations of attributes under a CWA and constitute the rows.

$\mathcal{R}_{\text{painting}}$	$\mathcal{A}_{\text{title}}$	$\mathcal{A}_{\text{year}}$	$\mathcal{A}_{\text{artist}}$	$\mathcal{A}_{\text{birth}}$	$\mathcal{A}_{\text{death}}$
t_1	Mona Lisa	1506	Leonardo da Vinci	1452	1660
t_2	The Starry Night	1890	Vincent van Gogh	1853	1890
t_3	Las Meninas	1656	Diego Velázquez	1599	1660
t_4	Irises	1890	Vincent van Gogh	1853	1890

3.2.1 Keys and Normal Forms

The notion of a relation provides us with the basic tools for data modelling. Codd proposed a range of constraints [Cod70] to guarantee proper data definition using the concept of functional dependencies between attributes and different types of keys that can be derived from those dependencies by applying Armstrong's axioms [Arm74].

Definition 3.6 Functional Dependency between Attributes

Let $X, Y \subset \text{SCH}(\mathcal{R})$ be sets of attributes that belong to a relation \mathcal{R} . We say that the attributes in Y are functionally dependent on (determined by) the attributes in X , if the values of all attributes in Y are implied by the choice of all values for attributes in X for all instances of \mathcal{R} . In this case, we write $X \rightarrow Y$.

Definition 3.7 Superkey of a Relation

A superkey S is a set of attributes $S \subseteq \text{SCH}(\mathcal{R})$ that uniquely identify any tuple $t \in \mathcal{R}$. Hence, all the non-key attributes $Y = \{a : a \in \text{SCH}(\mathcal{R}) \wedge a \notin S\}$ are functionally dependent on S and thus $S \rightarrow Y$.

Definition 3.8 Candidate Key of a Relation

A Candidate Key (CK) K is a set of attributes $K \subseteq \text{SCH}(\mathcal{R})$ such that K is a superkey that is minimal (or irreducible). Consequently, given $K \rightarrow Y$, no attribute can be removed from K without changing Y .

Using $\mathcal{R}_{\text{painting}}$ from Example 3.1, we can see that a painting's artist and the year of its conception are functionally determined by its title and therefore $\{\mathcal{A}_{\text{title}}\} \rightarrow \{\mathcal{A}_{\text{artist}}, \mathcal{A}_{\text{year}}\}$. Furthermore, the date of birth and death are implied by the artist, and thus, $\{\mathcal{A}_{\text{artist}}\} \rightarrow \{\mathcal{A}_{\text{birth}}, \mathcal{A}_{\text{death}}\}$. Given these two dependencies, there is a superkey $\{\mathcal{A}_{\text{title}}, \mathcal{A}_{\text{artist}}\}$ that determines all the other attributes. However, a CK must also be minimal and $\mathcal{A}_{\text{artist}}$ is implied by $\mathcal{A}_{\text{title}}$ (transitivity). Hence, the CK for $\mathcal{R}_{\text{painting}}$ must be $\{\mathcal{A}_{\text{title}}\}$.

Definition 3.9 Primary Key

A Primary Key (PK) of relation \mathcal{R} is a particular choice of CK (in case a relation has multiple). By definition, a PK cannot contain empty, that is, NULL values. In this thesis, we indicate a PK attribute with a star.

Definition 3.10 Foreign Key

A Foreign Key (FK) F of relation \mathcal{R} is a set of attributes $F \subseteq \text{SCH}(\mathcal{R})$ that reference a PK of a relation \mathcal{R}' . FKs can be used to model relationships between relations. In this thesis, we indicate a FK attribute with an overline.

In addition to Definitions 3.6 to 3.10, Codd proposed a series of *normal forms* [Cod70; Cod72; Cod74] that build onto one another and that are an indication of the quality of a data model based on the functional dependency of the attributes on a CK. The basic idea is to avoid redundancy by normalisation, i.e., by putting data that belongs together into dedicated relations and modelling relationships between them using FKs. The major forms proposed by Codd are the following:

First Normal Form (1NF) [Cod70] requires that all attributes in a relation are atomic, i.e., do not contain multiple values.

Given $\mathcal{R}_{\text{painting}}$ in Example 3.1, this means, for instance, that $\mathcal{A}_{\text{artist}}$ cannot hold multiple entries for a single tuple, in case one would want to model multiple artists that contributed to the same piece of artwork. Since it does not, $\mathcal{R}_{\text{painting}}$ is already in 1NF.

Second Normal Form (2NF) [Cod72] requires that \mathcal{R} is in 1NF and that every attribute that is not part of a CK is determined by the entirety of all CKs.

Given $\mathcal{R}_{\text{painting}}$ in Example 3.1, we can see that it is also in 2NF since all attributes only depend on $\mathcal{A}_{\text{title}}$.

Third Normal Form (3NF) [Cod72] requires that \mathcal{R} is in 2NF and that every attribute that is not part of a CK is non-transitively determined by every CK.

Given $\mathcal{R}_{\text{paintings}}$ in Example 3.1, we must acknowledge that it is not in 3NF since, as we have argued, $\{\mathcal{A}_{\text{title}}\} \rightarrow \{\mathcal{A}_{\text{birth}}, \mathcal{A}_{\text{death}}\}$ is a transitive dependency. To resolve this, we can split the relation into two relations as done in Example 3.2. The attribute $\overline{\mathcal{A}}_{\text{artist}} \in \mathcal{R}_{\text{painting}}$ now becomes a FK.

Boyce–Codd Normal Form (BCNF) [Cod74] is a stronger version of 3NF formulated by R. F. Boyce and E. F. Codd to eliminate certain anomalies found in 3NF. BCNF requires that for every functional dependency $X \rightarrow Y$ found in a relation \mathcal{R} , X must be a superkey of \mathcal{R} . Consequently, every attribute in \mathcal{R} can only depend on that key.²

Given $\mathcal{R}_{\text{paintings}}$ in Example 3.2, this means that $\mathcal{R}_{\text{paintings}}$ cannot hold any attributes that are not directly implied by $\mathcal{A}^*_{\text{title}}$ (which there are not).

Additional normal forms up to 6NF have been defined. For the sake of brevity, we will omit those since they are not relevant to the discussion ahead.

Example 3.2 Relations with Primary and Foreign Keys

The following tables list the schema and extent of $\mathcal{R}_{\text{painting}}$ and $\mathcal{R}_{\text{artist}}$, which are in BCNF. The attributes $\mathcal{A}^*_{\text{title}}$ and $\mathcal{A}^*_{\text{artist}}$ act as the respective PKs and $\overline{\mathcal{A}}_{\text{artist}} \in \mathcal{R}_{\text{painting}}$ is a FK that references $\mathcal{R}_{\text{artist}}$.

$\mathcal{R}_{\text{painting}}$	$\mathcal{A}^*_{\text{title}}$	$\overline{\mathcal{A}}_{\text{artist}}$	$\mathcal{A}_{\text{year}}$
t_1	Mona Lisa	Leonardo da Vinci	1506
t_2	The Starry Night	Vincent van Gogh	1889
t_3	Las Meninas	Diego Velázquez	1665
t_4	Irises	Vincent van Gogh	1890

$\mathcal{R}_{\text{artist}}$	$\mathcal{A}^*_{\text{artist}}$	$\mathcal{A}_{\text{birth}}$	$\mathcal{A}_{\text{death}}$
t_1	Leonardo da Vinci	1452	1519
t_2	Vincent van Gogh	1853	1890
t_3	Diego Velázquez	1599	1660

Note that we assume artists and paintings to be uniquely identified by their names and titles respectively, which is a simplification.

² This has led to the mnemonic “*The key, the whole key, and nothing but the key, So help me Codd.*” in reference to a similarly structured oath often used in courts of law.

Table 3.1 The relational operators proposed by Codd et al. [Cod70; GUW09].

Name	Symbol	Arity	Description
Union	\cup	2	Set union of two input relations.
Intersection	\cap	2	Set intersection of two input relations.
Difference	\setminus	2	Set difference of two input relations.
Cartesian Product	\times	2	Concatenates each tuple from the left with every tuple of the right input relation.
Rename	$\rho_{\mathcal{A}_A \rightarrow \mathcal{A}_B}$	1	Renames attribute \mathcal{A}_A in input relation to \mathcal{A}_B .
Selection	σ_S	1	Eliminates tuples (i.e., rows) from the input relation that do not match predicate S .
Projection	π_P	1	Eliminates attributes (i.e., columns) from the input relation that are not included in P .
Natural join	\bowtie	2	Pairs each tuple from the left with every tuple of the right input relation if their shared attributes match and concatenates them.

3.2.2 Relational Algebra

Having introduced the aspect of data representation and constraints, we can now move to that of operations that can be performed on the data upon querying. For that purpose, [Cod70] proposed the idea of a *relational algebra*, which follows a simple yet powerful idea: All query operations performed on relations are expressed by a set of *relational operators* that take one or multiple relations as input, and output a new relation as expressed by Equation (3.1).

$$\text{OP} : \mathcal{R}_1, \dots, \mathcal{R}_n \rightarrow \mathcal{R}_O \quad (3.1)$$

Those relational operators can then be composed to express a query of arbitrary complexity as indicated by Equation (3.2).

$$\text{QUERY} = \text{OP}_1 \circ \text{OP}_2, \dots, \circ \text{OP}_m \quad (3.2)$$

Codd proposed a minimal set of relational operators listed in Table 3.1 and explained in the following sections. The notation and operators may differ slightly depending on the source. We mainly use [GUW09] as our reference, with a few adjustments of our own for the sake of internal consistency.

3.2.2.1 Set Operations

Since relations are in essence sets of tuples, all basic operations known from set theory can be applied, namely *union*, *intersection*, and *difference*, with the only constraint that the two input relations $\mathcal{R}_L, \mathcal{R}_R$ must be *union compatible* and thus exhibit the same attributes, i.e., $\text{SCH}(\mathcal{R}_L) = \text{SCH}(\mathcal{R}_R)$.

The set union $\mathcal{R}_L \cup \mathcal{R}_R$ generates a new relation of all tuples contained in either \mathcal{R}_L or \mathcal{R}_R , as expressed by Equation (3.3). Due to relations being sets, duplicates resulting from a union operation are implicitly eliminated.

$$\mathcal{R}_L \cup \mathcal{R}_R = \{t : t \in \mathcal{R}_L \vee t \in \mathcal{R}_R\} \quad (3.3)$$

The intersection $\mathcal{R}_L \cap \mathcal{R}_R$ generates a new relation of all tuples contained in \mathcal{R}_L and \mathcal{R}_R , as expressed by Equation (3.4).

$$\mathcal{R}_L \cap \mathcal{R}_R = \{t : t \in \mathcal{R}_L \wedge t \in \mathcal{R}_R\} \quad (3.4)$$

The difference $\mathcal{R}_L \setminus \mathcal{R}_R$ generates a new relation of all tuples contained in \mathcal{R}_L and not in \mathcal{R}_R , as expressed by Equation (3.5).

$$\mathcal{R}_L \setminus \mathcal{R}_R = \{t : t \in \mathcal{R}_L \wedge t \notin \mathcal{R}_R\} \quad (3.5)$$

These set operations simply combine two input relations without changing their structure, i.e., $\text{SCH}(\mathcal{R}_L) = \text{SCH}(\mathcal{R}_R) = \text{SCH}(\mathcal{R}_O)$.

3.2.2.2 Cartesian Product

The binary, *Cartesian product* or *cross product* $\mathcal{R}_L \times \mathcal{R}_R$ of two input relations $\mathcal{R}_L, \mathcal{R}_R$ concatenates every tuple $t_L \in \mathcal{R}_L$ with every tuple $t_R \in \mathcal{R}_R$ to form a new output tuple, as expressed by Equation (3.6).

$$\mathcal{R}_L \times \mathcal{R}_R = \{(t_L, t_R) : t_L \in \mathcal{R}_L \wedge t_R \in \mathcal{R}_R\} \quad (3.6)$$

The result is a relation that contains all the attributes of \mathcal{R}_L and \mathcal{R}_R , i.e., $\text{SCH}(\mathcal{R}_L \times \mathcal{R}_R) = \text{SCH}(\mathcal{R}_L) \cup \text{SCH}(\mathcal{R}_R)$ and every possible permutation of tuples from the input relations.

3.2.2.3 Rename

The unary *rename* operator $\rho_{\mathcal{A}_A \rightarrow \mathcal{A}_B}(\mathcal{R})$ renames the attribute \mathcal{A}_A to \mathcal{A}_B without changing any value. This can be useful to eliminate collisions before applying the Cartesian product, or to enable a natural join on differently named attributes.

3.2.2.4 Selection

The unary *selection* operator $\sigma_S(\mathcal{R})$ applied on an input relation \mathcal{R} creates an output relation that contains a subset of tuples $t \in \mathcal{R}$, such that only tuples that match the predicate S are retained as expressed by Equation (3.7).

$$\sigma_S(\mathcal{R}) = \{t \in \mathcal{R} : S(t)\} \subset \mathcal{R} \quad (3.7)$$

The predicate S can be any conditional statement consisting of individual atoms that involve attributes of \mathcal{R} or any constant value and comparison operators $=, \neq, >, <, \geq, \leq$. Individual atoms can also be combined by logical operators \wedge, \vee or \neg . Examples could be $\mathcal{A}_1 \geq 2$, $\mathcal{A}_2 = \mathcal{A}_3$ or $\mathcal{A}_1 \geq 2 \wedge \mathcal{A}_2 = \mathcal{A}_3$ to express that an attribute should be greater than a constant, equal to another attribute, or combine the two with $A_1, A_2, A_3 \in \text{SCH}(\mathcal{R})$.

3.2.2.5 Projection

The unary *projection* operator $\pi_P(\mathcal{R})$ with $P \subset \text{SCH}(\mathcal{R})$ applied on an input relation \mathcal{R} creates an output relation that only contains the attributes listed in P , i.e., $\text{SCH}(\pi_P(\mathcal{R})) = P$, as expressed by Equation (3.8).

$$\pi_P(\mathcal{R}) = \{t[P] : t \in \mathcal{R}\} \text{ with } t[P] = \{t[\mathcal{A}] : \mathcal{A} \in P\} \quad (3.8)$$

All the tuples in \mathcal{R} are retained, however, resulting duplicates are removed.

3.2.2.6 Natural Join

The binary, *natural join* operator $\mathcal{R}_L \bowtie \mathcal{R}_R$ on two input relations $\mathcal{R}_L, \mathcal{R}_R$ concatenates every tuple $t_L \in \mathcal{R}_L$ with every tuple $t_R \in \mathcal{R}_R$ to form a new output tuple, if the attribute values of t_L and t_R are the same for the shared attributes $\xi = \{\mathcal{A} : \mathcal{A} \in \text{SCH}(\mathcal{R}_L) \wedge \mathcal{A} \in \text{SCH}(\mathcal{R}_R)\}$ as expressed by Equation (3.9)

$$\mathcal{R}_L \bowtie \mathcal{R}_R = \{t_L \cup t_R : t_L \in \mathcal{R}_L \wedge t_R \in \mathcal{R}_R \wedge t_L[\xi] = t_R[\xi]\} \quad (3.9)$$

The result is a relation that contains all the attributes of \mathcal{R}_L and \mathcal{R}_R , i.e., $\text{SCH}(\mathcal{R}_L \bowtie \mathcal{R}_R) = \text{SCH}(\mathcal{R}_L) \cup \text{SCH}(\mathcal{R}_R)$. Shared attributes are only retained once.

3.2.2.7 Expressing Queries

The following Example 3.3 illustrates how relational operators can be combined to form complex queries. Since expressing queries in such a way is quite inconvenient, in practice, queries are usually formulated in a human-readable query language, which is then translated to the relational operators. A famous example of such a language is SQL [Cha12] for relational databases.

Example 3.3 Searching for Paintings Using Relational Algebra

The following tables list the schema and extent of $\mathcal{R}_{\text{painting}}$ and $\mathcal{R}_{\text{artist}}$.

$\mathcal{R}_{\text{painting}}$	$\mathcal{A}^*_{\text{title}}$	$\mathcal{A}_{\text{artist}}$	$\mathcal{A}_{\text{year}}$
t_1	Mona Lisa	Leonardo da Vinci	1506
t_2	The Starry Night	Vincent van Gogh	1889
t_3	Las Meninas	Diego Velázquez	1665

$\mathcal{R}_{\text{artist}}$	$\mathcal{A}^*_{\text{artist}}$	$\mathcal{A}_{\text{birth}}$	$\mathcal{A}_{\text{death}}$
t_1	Leonardo da Vinci	1452	1519
t_2	Vincent van Gogh	1853	1890
t_3	Diego Velázquez	1599	1660

Using relational algebra, the query “return the names of all paintings that were painted by an artist who died after 1800” can be expressed by joining $\mathcal{R}_{\text{painting}}$ and $\mathcal{R}_{\text{artist}}$, followed by a selection and projection:

$$\mathcal{R}_{\text{result}} = \pi_{\mathcal{A}^*_{\text{title}}}(\sigma_{\mathcal{A}_{\text{death}} > 1800}(\mathcal{R}_{\text{painting}} \bowtie \mathcal{R}_{\text{artist}}))$$

This query produces the relation $\mathcal{R}_{\text{result}}$ that contains $t_2[\mathcal{A}^*_{\text{title}}]$ of $\mathcal{R}_{\text{painting}}$.

3.2.3 Extensions

While the relational model and its algebra forms the foundation of many modern DBMS, the model as originally proposed by Codd has often turned out to be too limited to accommodate certain functionalities as required, e.g., by the ANSI SQL standard [Lib03; Xopb]. For example, many applications require storage of duplicate data, and therefore the notion of a relation – which is a set of tuples and thus does not allow for duplicates – is inadequate. Another example could be the support for sorting or aggregation (`ORDER BY` and `GROUP BY` in SQL), which also is not covered by the original algebra.

Over the years, this has led to a growing list of proposals for extensions, some of which have been adopted while a majority has remained theoretical in nature. In the following sections, we will introduce a few examples. While at a first glance, these adaptions seem to be an elegant way of extending the expressiveness of the relational data model, there are important implications: Most importantly, some operators may not behave in the way specified in Section 3.2.2, or we may require additional or different operators to accommodate all the functionality needed. This adds complexity to the DBMS, especially when considering aspects such as query planning.

3.2.3.1 Relations vs. Bags vs. Sequences

The motivation for considering other mathematical structures than sets as an algebraic foundation for a data model are twofold: Firstly, relations are unable to provide functionality that may be desirable in a DBMS, such as duplicate entries or explicit ordering of tuples. Secondly, some of the mathematical convenience of the relational model, e.g., prohibiting duplicates, may be inefficient to actually implement [GUW09]. It is therefore not surprising that ANSI SQL [Xopb] formally operates on bags rather than sets, which allow for duplicate entries [GUW09; Cha12]. If one would want to express explicit ordering of tuples, they would even have to move to sequences, in which every tuple occupies a specific position.

We will not elaborate on all the consequences such a transition may have, and refer to [GUW09], which discusses this issue in great detail. However, for illustrative purposes, we still provide the Example 3.4 inspired by [GUW09]. It demonstrates that even minor changes to the properties of the purely relational data model must be taken into account when implementing systems, e.g., during algebraic query optimisation.

Example 3.4 Bag vs. Set Semantics

The following tables list the schema and extent of \mathcal{R}_{p1} and \mathcal{R}_{p2} , which are union compatible and have one tuple in common (t_1).

\mathcal{R}_{p1}	$\mathcal{A}^*_{\text{title}}$	$\overline{\mathcal{A}}_{\text{artist}}$	$\mathcal{A}_{\text{year}}$
t_1	Mona Lisa	Leonardo da Vinci	1506
t_2	The Starry Night	Vincent van Gogh	1889
t_3	Las Meninas	Diego Velázquez	1665

\mathcal{R}_{p2}	$\mathcal{A}^*_{\text{title}}$	$\overline{\mathcal{A}}_{\text{artist}}$	$\mathcal{A}_{\text{year}}$
t_1	Mona Lisa	Leonardo da Vinci	1506
t_2	The Birth of Venus	Sandro Botticelli	1485
t_3	The Night Watch	Rembrandt Harmenszoon van Rijn	1642

If now we consider the union operation, i.e., $\mathcal{R}_p = \mathcal{R}_{p1} \cup \mathcal{R}_{p2}$, we will see, that for ordinary relations, $|\mathcal{R}_p| = 5$, whereas for bags, $|\mathcal{R}_p| = 6$, since the common tuple will appear twice.

Now let us further consider the set difference and the distributive law with union, i.e., $(\mathcal{R}_{p1} \cup \mathcal{R}_{p2}) \setminus \mathcal{R}_{p1} = (\mathcal{R}_{p1} \setminus \mathcal{R}_{p1}) \cup (\mathcal{R}_{p2} \setminus \mathcal{R}_{p1})$. We can easily see, that this identity holds for sets but not for bags, since $|(\mathcal{R}_{p1} \setminus \mathcal{R}_{p1}) \cup (\mathcal{R}_{p2} \setminus \mathcal{R}_{p1})| = 2$ whereas $|(\mathcal{R}_{p1} \cup \mathcal{R}_{p2}) \setminus \mathcal{R}_{p1}| = 3$, since one of the duplicates will be retained.

3.2.3.2 Extended Projection

The extended project as described by [GUW09] is an addition to the projection operator π_P specified in the relational data model and necessitated by the ANSI SQL standard [Xopb]. It involves a more general definition of P , that can now contain any of the following elements for a relation \mathcal{R} :³ (i) any attribute $\mathcal{A} \in \text{SCH}(\mathcal{R})$, (ii) any literal value $\mathcal{L} \in \mathcal{D}_{\mathcal{L}}$, with $\mathcal{D}_{\mathcal{L}} \in \mathbb{D}$, or, (iii) any N-ary function expression $f: \mathcal{G}_1 \times \dots \times \mathcal{G}_A \rightarrow \mathcal{A}_{\text{out}}$ with \mathcal{G}_j being either an attribute $\mathcal{A}_i \in \text{SCH}(\mathcal{R})$ or a literal $\mathcal{L}_j \in \mathcal{D}_{\mathcal{L}}$.

Therefore, in addition to projecting on existing attributes, the extended projection can be used to project onto literal values and arbitrary function expressions, thus appending new attributes to a relation. This is a very powerful extension, because it enables the generation of new values based on existing ones. Many DBMS also allow users to define their own User Defined Function (UDF).

³ We intentionally deviate from [GUW09] in that we do not limit ourselves to algebraic expressions and special functions.

3.2.3.3 Aggregation

Aggregation is the act of summarising certain attribute values based on the membership of a tuple in a specific category or group, which is established based on other attribute values. An example could be a customer who would like to see the maximum price of all products per brand. The ANSI SQL standard [Xopb] defines the GROUP BY operator for this purpose.

[GUW09] formalises the grouping operator γ_G , which can be used to aggregate on attributes in \mathcal{R} . G is a list of elements, which can either be: (i) any attribute $A \in \text{SCH}(\mathcal{R})$, the attribute(s) that the operator aggregates by, or, (ii) any aggregation or set function $\text{AGR}: \mathcal{A} \rightarrow \mathcal{A}_{\text{out}}$.

The output relation $\gamma_G(\mathcal{R})$ is then constructed as follows: First, the tuples in the input relation \mathcal{R} are grouped into N groups $\mathcal{G}_1, \dots, \mathcal{G}_N \subset \mathcal{R}$ based on the values in the specified grouping attribute(s). Subsequently, the aggregation function is executed for all tuples in each group, giving rise to an output relation that contains N tuples – one per group – each bearing the grouping attributes as well as the output of the set functions. The ANSI SQL standard [Xopb] defines a list of supported set functions, including but not limited to functions such as COUNT, MIN, MAX, SUM and MEAN.

3.2.3.4 Sorting and Ranking

Sorting is the act of arranging items in a specific order based on their attribute values. Consider, for example, a customer who wants to look at prices of products from lowest to highest. The ANSI SQL standard [Xopb] defines the ORDER BY operator for this purpose.

However, formally, the act of sorting is difficult in the context of “pure” relational algebra, because neither sets nor bags exhibit any ordering. [GUW09] proposes the ω_T operator to make this operation explicit. Given a relation \mathcal{R} , T is simply a list of attributes to sort on, i.e., $T \subset \text{SCH}(\mathcal{R})$ and the result $\omega_T(\mathcal{R})$ is a sorted list of values. The conversion between set, bag, and list takes place implicitly through application of the operator and the consequences for the algebra are not addressed.

More formal approaches to sorting were proposed, e.g., by [RDR⁺98] who introduced four additional operators and a *sequence algebra* that can deal with ordered queries. Similarly, [LCI⁺05] proposes a more systematic approach to *ranking*, which is sorting based on an external function. They introduce the idea of a *rank relation*, which carries explicit properties with respect to the ranking.

3.2.3.5 Recursive Data Access

Recursion is an important concept in computer programming, e.g., to test for reachability in graph-structured data, or to model trees. Yet, neither the relational model nor the original ANSI SQL [Xopb] standard provided support for recursion [Lib03]. It was not until ANSI SQL-99, that support for recursive queries in the form of Common Table Expressions (CTEs) was added [PBB⁺10].

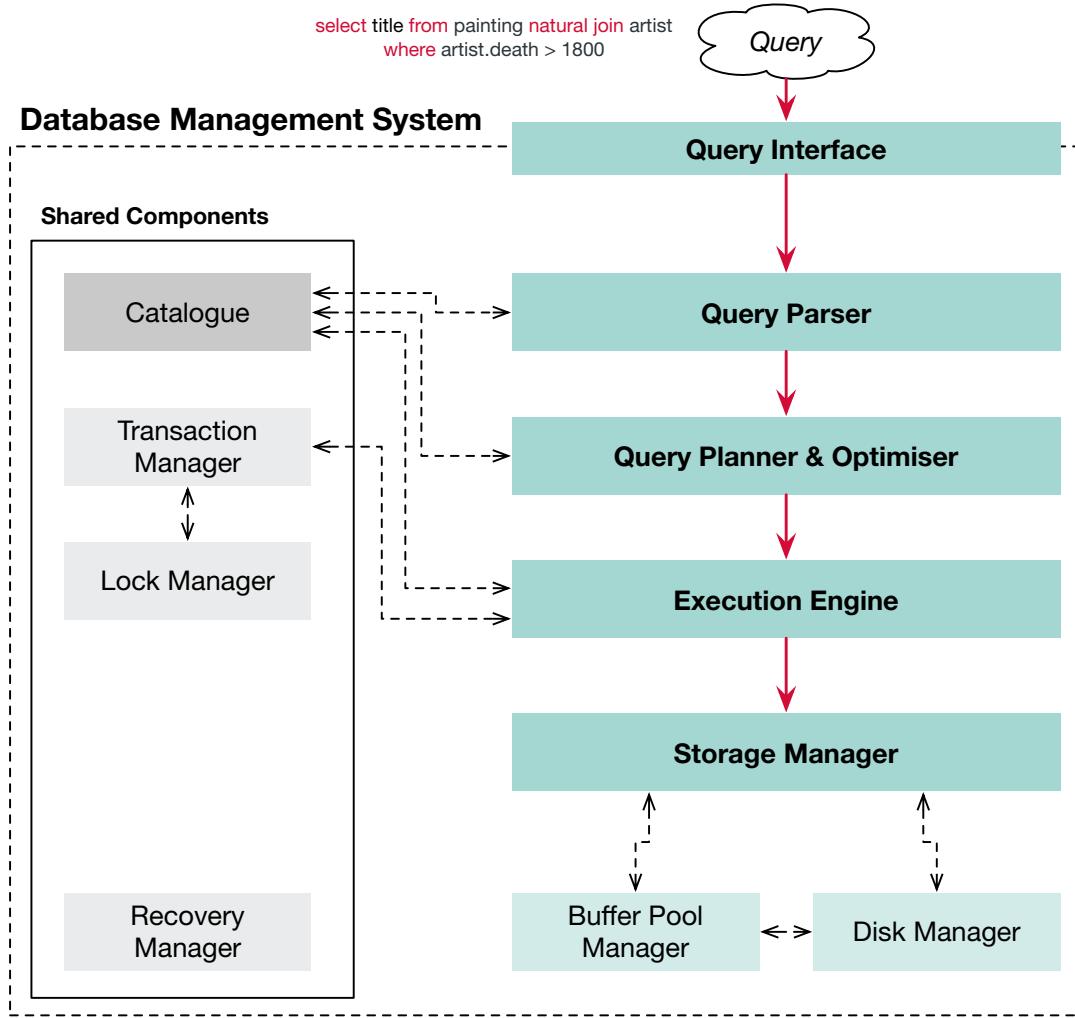


Figure 3.2 The general architecture of a DBMS and its individual components.

3.3 Database Management System (DBMS)

Irrespective of the concrete system and data model, DBMS generally exhibit a common internal structure [Pet19; HSH07], even though specifics may vary widely between implementations. This overall structure is depicted in Figure 3.2 and in this section we use the path a query⁴ takes within the system, as indicated by the red arrows, to illustrate the components involved.

3.3.1 Query Interface

The query interface is the DBMS's channel to the outside world and enables its use by other systems. The query interface takes care of client connections, accepts and forwards queries for processing, and returns results to the caller.

⁴ Henceforth, the term “query” refers to any DDL, DML, or DQL expression.

Queries accepted by the query interface are usually expressed in a specific query language – which are often designed for human users and therefore human-readable. We have already mentioned SQL [Xopb; Cha12] as the most prominent example of such a query language, which can be used to express queries using a declarative, textual syntax (see Example 3.5). In addition to SQL, there also exist other, domain-specific languages such as SPARQL [PAG09] to query Resource Description Framework (RDF) data, Cypher to query labelled property graphs [FGG⁺18] found in graph databases such as Neo4j⁵ or the document-oriented query language used by MongoDB.⁶

Example 3.5 Searching for Paintings Using SQL

Let $\mathcal{R}_{\text{painting}}$ and $\mathcal{R}_{\text{artist}}$ be the two relations from Example 3.3, the query “return the names of all paintings that were painted by an artist who died after 1800” can be expressed in SQL as follows:

```
select title from painting natural join artist
      where artist.death > 1800
```

The select clause expresses the projection π to attribute `title`, the natural join clause expresses the natural join \bowtie between `painting` and `artist` and the where clause expresses the selection σ to tuples whose `artist.death` attribute is greater than 1800.

Early attempts at a client-facing standardisation of the query interface for SQL-based systems led to the inception of the Call Level Interface (CLI) standard [Xopa], which in its early version enabled embedding of SQL commands into C or Cobol programmes. The CLI standard was later complemented by Open Database Connectivity (ODBC) and Java Database Connectivity (JDBC), which provide database connectivity to a wide range of SQL and NoSQL systems. In addition, most of the SQL (e.g., MySQL, PostgreSQL, Microsoft SQL, Oracle) and NoSQL (e.g., MongoDB, Neo4j, Redis) systems provide database specific and therefore un-standardised connectivity through a plethora of client libraries that are maintained for different platforms and programming environments by either companies and/or open-source contributors.

⁵ See <https://neo4j.com/>; accessed June 2022.

⁶ See <https://www.mongodb.com/>; accessed June 2022.



Figure 3.3 The query from Example 3.5 represented as a tree of relational operators also called a logical query execution plan. By convention, the tree is executed and read from left to right (bottom to top), and information flows from the leaves to the root.

3.3.2 Query Parser

The query parser's main task is to transform a query provided in a human-readable query language to a logical representation that can be processed by a DBMS. This often involves a conversion of a syntax to the operators specified by the data model, e.g., a tree of the relational operators that represent an SQL query as illustrated in Figure 3.3, for example. The result is what is generally called a *logical query (execution) plan*, since it outlines all the operations necessary to generate the desired results without, however, specifying concrete access methods or algorithms.

Practically, query parsing involves several steps [Gra93; GUW09]: First, the syntax of the textual query must be analysed and converted to a data structure that can be processed. Very often, Abstract Syntax Trees (ASTs) or *parse trees* based on formal grammars of the query language, and libraries such as ANTLR⁷ or JavaCC⁸ are employed in this step. Subsequently, the leaf-nodes in the AST must be mapped to Database Objects (DBOs) contained in the DBMS in a step called preprocessing. This involves lookups in the *catalogue* – a data structure that lists a DBMS' units of data organisation. For a relational database, for example, the catalogue would be used to look up named tables and columns (e.g., tables artist and painting from Example 3.5). And finally, the nodes in the AST must be converted to an internal representation that matches the data model, e.g., a tree of relational operators. This conversion step may also involve basic sanity checks for type compatibility or the existence of requested DBOs. Certain optimisations can also be applied, e.g., simplifying trivial identities such as $1 = 1$ (which is always true) or $0 > 1$ (which is always false). If we take the query from Example 3.5, the logical query plan would look as illustrated in Figure 3.3.

⁷ See <https://www.antlr.org/>; accessed July 2022.

⁸ See <https://javacc.github.io/>; accessed July 2022.

3.3.3 Query Planner & Optimiser

A query planner tries to transform the *logical* query plan produced by the parser to a *physical* plan in a way that allows for efficient and effective query execution in terms of minimal execution time or resource usage [JK84; GUW09]. The distinction between a logical and physical representation of a query is very common in DBMS and illustrated in Example 3.6. While the former determines *what* should be produced terms of some logical algebra, the latter defines the *how* in terms of operators and concrete access methods. Every logical relation, attribute, and operation has one (or multiple) physical counterparts. This is the context in which query optimisation can take place.

Example 3.6 Logical Query Plan vs. Physical Query Plan

Let $\mathcal{R}_{\text{painting}}$ and $\mathcal{R}_{\text{artist}}$ be the two relations from Example 3.3, the query “return the names of all paintings that were painted by an artist who died after 1800” can be expressed in relational algebra as follows:

$$\mathcal{R}_{\text{result}} = \pi_{\mathcal{A}_{\text{title}}}(\sigma_{\mathcal{A}_{\text{death}} > 1800}(\mathcal{R}_{\text{painting}} \bowtie \mathcal{R}_{\text{artist}}))$$

The naive, unoptimised physical implementation can be obtained by mapping every relational algebra operator to a physical operator.

Logical	Physical	Description
$\mathcal{R}_{\text{painting}}$	SCAN _{painting}	Scans every tuple in $\mathcal{R}_{\text{painting}}$.
$\mathcal{R}_{\text{artist}}$	SCAN _{artist}	Scans every tuple in $\mathcal{R}_{\text{artist}}$.
\bowtie	HASHJOIN _{artist}	Performs a hash-join on the two input relations based on the common attribute.
$\sigma_{\mathcal{A}_{\text{death}} > 1800}$	FILTER _{death > 1800}	Filters every tuple based on comparison.
$\pi_{\mathcal{A}_{\text{title}}}$	PROJECT _{title}	Projects onto attribute title.

Despite being a central task of every DBMS and a huge body of research, there are only very few comprehensive surveys and overviews of established techniques [JK84; Gra93; Cha98b]. This has been pointed out as early as 1984 and one of the reasons identified is that query optimisation is achieved “[...] by integrating a large number of techniques and strategies, ranging from logical transformations of queries to the optimization of access paths and the storage of data on the file system level.” [JK84] (p. 112). Consequently, many techniques are specific to a concrete implementation and the underlying data and execution model.

Example 3.7 Implementation of a JOIN between two relations.

Let $\mathcal{R}_{\text{painting}}$ and $\mathcal{R}_{\text{artist}}$ be the two relations from Example 3.5 with $N = |\mathcal{R}_{\text{artist}}|$, $M = |\mathcal{R}_{\text{painting}}|$ and $N < M$. The JOIN between the two relations, i.e., $\mathcal{R}_{\text{artist}} \bowtie \mathcal{R}_{\text{painting}}$ can be implemented by the following algorithms [Gra93]:

Nested Loop Join The nested loop join iterates over all elements in the smaller relation, i.e., $\mathcal{R}_{\text{artist}}$ and tries to find matching entries in the larger relation, i.e., $\mathcal{R}_{\text{painting}}$ for every element. If $\mathcal{R}_{\text{painting}}$ exhibited an index on the join column, that index could be used for lookup. Otherwise, a full scan of $\mathcal{R}_{\text{painting}}$ must be performed for every element in $\mathcal{R}_{\text{artist}}$. Therefore, computational complexity ranges between $O(N \log M)$ and $O(NM)$.

Hash Join The hash join algorithm builds up a hash-table for one side of the JOIN to speed up the lookup of join keys, hence, omitting the inner loop of the nested-loop join. If resource consumption of building and accessing the hash-table is ignored, the computational complexity becomes $O(N + M)$. However, building the hash-table is typically accompanied by a non-negligible cost, especially, if it does not fit into main memory.

Sort-Merge Join The sort-merge join can only be used if $\mathcal{R}_{\text{artist}}$ and $\mathcal{R}_{\text{painting}}$ are sorted by the join key (i.e., the name of the artist) and if the comparison operator checks for equality between the two keys. Given this pre-condition, the join can be realised in a single loop that collects the entries for distinct values of the join keys, i.e., computational complexity is $O(N + M)$, if we assume that the relations come pre-sorted.

Over the years, two high-level strategies of optimising a query have emerged: bottom-up and top-down [JK84]. Early research focused on bottom-up approaches wherein individual operators and special edge-cases were studied and optimised. This turned out to be important foundational work since at the lowest level, query optimisation can be achieved by employing efficient algorithms. Example 3.7 tries to demonstrate this using possible implementations of a JOIN between two relations.

However, as queries became more complex, bottom-up optimisation started to reach its limits, especially in terms of generalisability. If we turn to Example 3.7, we will realise that the optimal path of execution does not only depend on the choice of JOIN algorithm but also on the join order, the presence or absence of indexes that allow for more efficient access, and the estimated size of

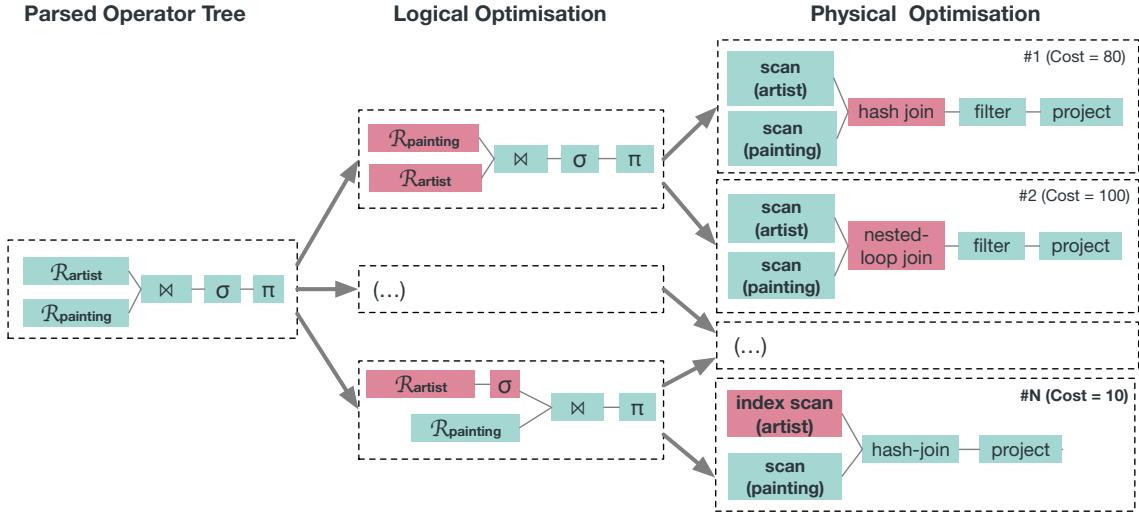


Figure 3.4 Cost-based query planning in a DBMS based on Example 3.5. A parsed query is transformed to a series of equivalent, logical and physical plans with known costs. At the end, the most cost-effective plan is selected. Red boxes indicate adjustments to the operator tree.

intermediate data structures and thus the size of the relations themselves. If now, in addition, we consider queries that involve multiple JOINS, we must acknowledge that these metrics may differ for every combination of relation, i.e., there may not be a single, optimal choice of algorithm.

As a consequence, attention has shifted towards top-down approaches that focused on holistic optimisation of the entire query plan. Early top-down approaches tried to focus on heuristic methods [JK84]. However, the System-R paper [SAC⁺79] published in 1979 established the idea of *cost-based query optimisation* as a gold standard for DBMS. A cost-based query planner enumerates physical plans that employ different strategies and selects the one that minimises the estimated cost. Those costs may involve computation (time used on a CPU), storage access (data accessed on disk) or transmission (data transferred between nodes) required by a query [JK84; GUW09]. Which of those metrics are used is highly dependent on the system: In the not-too-distant past before the invention of the Solid State Drive (SSD), access to secondary storage was considered to be the defining cost [GUW09]. However, for distributed databases, the transfer of data between nodes has become another factor to consider [FJK96; BJZ13].

Cost-based query planning relies on an internal representation of the query, e.g., the logical plan illustrated in Figure 3.3, which is transformed into equivalent alternatives over multiple iterations as illustrated in Figure 3.4. As queries become more complex (e.g., multiple joins) the solution space grows exponentially, i.e., the problem is NP-hard and must be solved heuristically. Often, the

query optimisation process is divided into three phases as described for the *Volcano Optimiser Generator* [JK84; GM93] and the *Cascades Framework* [Gra95] (which still serve as a foundation for systems today, e.g., [SAR⁺14; BCH⁺18]):

1. Transformation of the logical input plan to either simplify it or to replace parts that can be covered by a specific edge-case. This results in different, alternative logical query plans.
2. Mapping the transformed, logical query plans into sequences of elementary operations with known costs by selecting specific algorithms or access patterns. This results in a (potentially larger) set of physical query plans.
3. Estimate the total cost for the physical query plans and select the best one.

Each of the aforementioned steps involves the application of algebraic equivalence rules [GUW09] (e.g., commutativity of operations), heuristics [GUW09; Gra93; Swa89; BGJ10; TSF⁺12] (e.g., predicate pushdown) or special rules to handle certain (edge-)cases [JK84; Gra93] (e.g., replace table by index scan).

Indexes deserve a special mention in this context, since they are of huge practical importance. An index is an auxiliary data structure that organises access to the data in a relation based on one or multiple attributes to allow for more efficient data access when searching for an indexed attribute – $O(\log N)$ instead of $O(N)$, with $N = |\mathcal{R}|$ for B^+ -trees. Example 3.8 illustrates this for a simple filter operation. We will address indexes and the B^+ -tree in Section 3.3.5.2.

Example 3.8 Index Scan vs. Table Scan

Let us consider the logical and physical plan of Example 3.6. And let us further assume that there exists a B^+ -tree index on $\mathcal{A}_{\text{death}}$ in $\mathcal{R}_{\text{artist}}$. The following query plan could be an optimisation of the original plan:

Logical	Physical	Description
$\mathcal{R}_{\text{painting}}$	SCAN _{painting}	Scans every tuple in $\mathcal{R}_{\text{painting}}$.
$\sigma_{\mathcal{A}_{\text{death}} > 1800}(\mathcal{R}_{\text{artist}})$	INDEX _{artist} ^{death>1800}	Scans every tuple in the index that matches the predicate.
\bowtie	HASHJOIN _{artist}	Performs a hash-join on the two input relations based on the common attribute.
$\pi_{\mathcal{A}_{\text{title}}}$	PROJECT _{title}	Projects onto attribute title.

A critical factor for any cost-based query optimiser is the model used for cost and cardinality estimation [JK84; YHM15]. Most existing systems rely on statistical modelling based on the data stored in the database [GTK01; Ioa03] as well as the model for cost incurred by accessing system resources [MBK02]. Simple statistics involve information about the number of entries in a relation or statistical moments of columns but may also rely on more complex analysis of data distribution such as histograms. This pre-hoc modelling was later complemented by taking into account the attained execution speed, which requires post-hoc analysis of an execution plan with respect to the cost model and establishes a continuous control loop between planner and execution engine [ML86]. Recent work in the field also explores the application of machine learning at different levels of the cost-estimation and planning process [WCZ⁺13; Vu19; AAA⁺20].

3.3.4 Execution Engine

Once a query execution plan has been selected, the execution of that plan falls to the DBMS's execution engine. At a high level, every operator in the plan is typically mapped to an actual operator in a query execution pipeline that performs the necessary operations, e.g., executes a scan of a table or a specific JOIN algorithm. This is visualised in Figure 3.5. In the resulting *operator tree*, every operator functions independently, reads input, processes it and produces output that is then handed to the next operator in the tree (or to the client).

While this model is simple enough, there are a few subtle differences in how this can be implemented by a system. First, there is a distinction as to how data is being handed from operator to operator:

Materialisation Model In this model, every operator reads all tuples from the input, performs its work, and then forwards the entire output to the next operator. This requires potentially large buffers to store intermediate results.

Iterator Model This is the model described in [GM93; Gra93]. In this model, every operator reads a single tuple, processes it, and hands it to the next operator in the pipeline. Therefore, tuples are iterated one by one and handed from the leaves to the root of the operator tree.

Batch Model In this model, every operator reads a specific number of tuples from its inputs (a batch), processes these tuples, and then forwards the batched results to the next operator. Therefore, data is processed batch by batch.

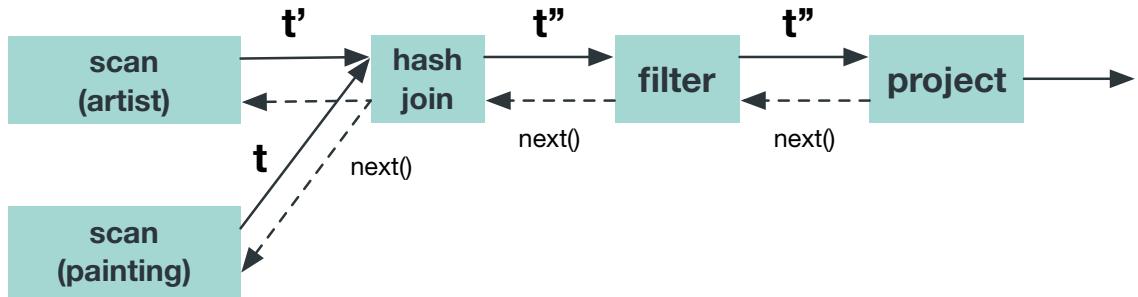


Figure 3.5 Illustration of the iterator model for query listed in Example 3.5. Every operator calls the upstream operator with `next()` to request the next tuple for processing.

It must be mentioned, that even in the iterator model, some operators still require intermediate materialisation. For example, a sort operation must be able to “see” all tuples in order to establish a total ordering. Consequently, intermediate caching is required in these cases. In addition, we can distinguish between which operator initiates data exchange:

Pull-based Operators higher-up in the operator tree request the next item from the previous operator. This is the common model for disk-based systems and leads to the slowest operator in the pipeline determining execution speed.

Push-based Operators lower in the operator tree send their results to the next operator, which has to buffer if it is not ready to process. While potentially more efficient, this requires a buffering infrastructure.

The pull-based iterator model is very easy to implement, since basically, every operator just needs to provide a `next()` method that returns the next tuple (in addition to `open()` and `close()` to signal start and end). Every operator then simply calls `next()` on its input to get the next tuple. The model is illustrated in Figure 3.5. It is probably the most common processing model found in disk-based DBMS, even though it has downsides such as a poor instruction locality and overhead due to the large number of method calls [NL14]. The materialisation model can be beneficial for small datasets and OLTP queries. However, it is not a good fit for analytical workloads on large datasets, especially, if main memory is limited. The batch model is something in between the previous two and can be beneficial when using vectorised processing and SIMD instructions such as those provided by the AVX or SSE extensions, since multiple tuples can be processed in batches. It is thus often employed in analytical databases such as MonetDB [IGN⁺12].

3.3.4.1 Query Parallelism

Since nowadays most computing platforms offer multiple CPUs or CPU cores, parallel execution of queries can provide a considerable speed-up in terms of execution time. [Gra93] distinguishes between three types of parallelism a DBMS can support:

Inter-query parallelism refers to the parallel execution of different, independent queries (potentially from different users). There are basically two approaches to achieve inter-query parallelism, namely, by mapping different queries to different system processes (e.g., PostgreSQL) or threads within the same process (e.g., MySQL), with the advantage of the latter being that all queries share the same memory address space.

Inter-operator parallelism refers to the parallel execution of different operators within a query, which are then synchronised by a dedicated operator that merges the results. For example, one could parallelise the execution of the respective order operations on two relations prior to a sort-merge join. This is a form of *intra-query parallelism*.

Intra-operator parallelism refers to the splitting up of the work performed by a single operator. This typically involves the ability to parallelise based on the data, e.g., by partitioning on it and letting different operator instances take care of parts of the workload. This is also a form of *intra-query parallelism*.

All three aspects must be handled by the query execution engine (and other component involved) and give rise to a large number of implementation details that must be considered when building a DBMS. For example, the challenges that arise from parallel query execution and the resulting concurrent data access – especially in the presence of parallel reads and writes – are covered by the *isolation* part of ACID and typically mitigated by a *transaction manager* and *concurrency control protocols*.

3.3.4.2 Optimising Execution Speed

Optimising query execution at different levels of the system has always been an important area of research. At a hardware level, we have already mentioned the use of specialised CPU instructions such as AVX or SSE [IGN⁺12; PR20; PR19; SXJ21]. Recent trends take this a step further by leveraging the GPU or specialised hardware, such as FPGAs [AAA⁺14; AAA⁺20; PLH21].

Another popular approach used to minimise the overhead incurred by the iterator model is compilation [Kvh21; FT21], e.g., implemented by HyPer [Neu11; NL14; Bgl+22]. This technique tries to pre-compile parts of the execution plan directly into machine code, i.e., merging multiple operators into one, which increases instruction locality and allows for caching of pre-compiled elements. The method can be combined with vectorised execution [Szb11; Rbm22].

3.3.5 Storage Manager

At the lowest levels of any DBMS resides the storage manager, which orchestrates access to data on primary and/or secondary storage devices. Many different approaches to storage have been implemented over the years and we will simply focus on some basic concepts and refer to [Pet19] for more information on the topic. What most DBMS have in common is that they implement a two-level memory architecture wherein data resides on disk (which is slower but durable) and in RAM (which is faster but volatile). The two aspects are handled by the two main components: the *disk manager* and the *buffer pool manager* [Pet19]. Recent years have also seen a trend towards main memory databases that keep data primarily in RAM [GS92; Fkl+17] or Non-Volatile Memory (NVM) [AP17].

3.3.5.1 Data Organisation and Disk Manager

DBMS are typically required to accommodate multiple design goals at once. They are supposed to store data, such that, (i) space occupied by the data, (ii) time required to access (i.e., read) records, and, (iii) time required to update (i.e., write) records are minimised. In order to achieve all these goals, a DBMS often combines different types of data structures, which results in a distinction between two main types of files:⁹ Larger *data files* (or primary files) and smaller *index files*. While the data files contain the data itself, index files contain metadata that allows for faster access. Most database systems access both types of files in terms of fixed-size database pages that are read and/or written in their entirety.¹⁰ It is the main task of the *disk manager* to provide these low-level primitives. The organisational structure of the (meta-)data is then superimposed onto the page structure. There are many implementation details to be considered, such as dealing with variable-length data, dealing with data that is larger than a page, checksumming, and versioning. Different techniques are being applied to address these issues.

⁹ These are not necessarily separate files in the file system.

¹⁰ The size of a page is often a fixed multiple of the block-size implemented by the file system.

In a relational database, the data files can be regarded as the physical manifestation of a relation. There are many ways to organise such data files and we refer to [Pet19] for details. At a high level, we can distinguish between index-organised, hash-organised, and heap-organised tables, which arrange the tuples in different data structures (i.e., trees, hash tables, or heaps), with advantages and disadvantages inherent to the respective choice. Recent work has given rise to novel ways of data organisation, e.g., in systems like MonetDB that organises data in fixed-size arrays called binary association tables [BKM08]

In order to be able to locate and relate individual tuples across different pages, files, and data structures, there is usually a notion of an internal *tuple identifier* TID, which uniquely identifies every record (see Definition 3.11). Conceptually, the tuple identifier connects the logical notion of a tuple to a physical world in which such a tuple must be stored, located, and accessed as a record.

Definition 3.11 Tuple Identifier of a Relation \mathcal{R}

Let \mathcal{R} be a relation. We call TID the relation's *tuple identifier* if it uniquely identifies every tuple in \mathcal{R} , i.e., has a unique value per tuple.

TID can be considered an implicit, internal CK of \mathcal{R} that is often not visible to the outside and may even change its value as records are being written. However, some systems make the tuple identifier accessible, e.g., Oracle with its ROWID and in some database implementations, an actual PK may act as TID.

3.3.5.2 Indexes

Indexes are on-disk structures that enable faster data access, especially in the presence of predicates that query an indexed attribute. In a relational database, indexes are typically built from one or multiple attributes of a relation. One can distinguish between primary and secondary indexes [Pet19]: A primary index always indexes the PK or tuple identifier and is the main mean of data access. Depending on how the data files are organised, the primary index is either an auxiliary structure or integrated into the data files, e.g., for index-organised tables or clustered indexes. Since queries may also rely on attributes other than the PK, one often creates auxiliary, secondary indexes, which are used to optimise query execution performance for certain query types by indexing selected attributes. Since every index structure must also be maintained when records are updated, there is an inherent trade-off between read and write performance.

A very popular (if not the most popular) index structure in both SQL and NoSQL systems is a B -tree [BM70] variant called the B^+ -tree [GUW09], which supports both fast point lookups as well as range search while at the same time being able to gracefully handle inserts, deletes and updates to the data. An example of a B^+ -tree is provided in Figure 3.6.

The B^+ -tree arranges the indexed attributes (the key) in the leaf nodes of a balanced search tree often alongside with the indexed data or a pointer into the data files, i.e., a tuple identifier or a set of identifiers for non-PK attributes. The leaf nodes at this bottom-most level of the tree are connected so that the keys can be traversed as an ordered, linked list to facilitate sequential access.

Every internal (or intermediate) node of the B^+ -tree can contain m children with $\lceil \frac{b}{2} \rceil \leq m \leq b$, wherein the order b is a chosen constant for the tree. Each child of a node covers a certain key interval and points to either an internal node at the next level or to a leaf node. The root node at the top is a special kind of internal node with $2 \leq m \leq b$ (unless it is the only node in the tree).

When searching for value in the B^+ -tree, the DBMS can traverse the tree from the root downwards. At every node it decides which child node to follow based on the key intervals they cover. If, for example, we were to lookup the value 87 in the B^+ -tree illustrated in Figure 3.6, the lookup algorithm would follow the left child twice (since $87 < 540$ and $87 < 248$) until it reaches the left-most leaf node. Within the leaf, the algorithm performs a sequential scan until 87 is reached. Since only a single branch of the tree is traversed, the runtime complexity of such a point lookup can be improved from $O(N)$ (sequential scan) to $O(\log N)$ with $N = |\mathcal{R}|$. The linked list at the base of the tree can furthermore be used to perform efficient range search, e.g., for $<$, \leq , $>$, \geq comparisons.

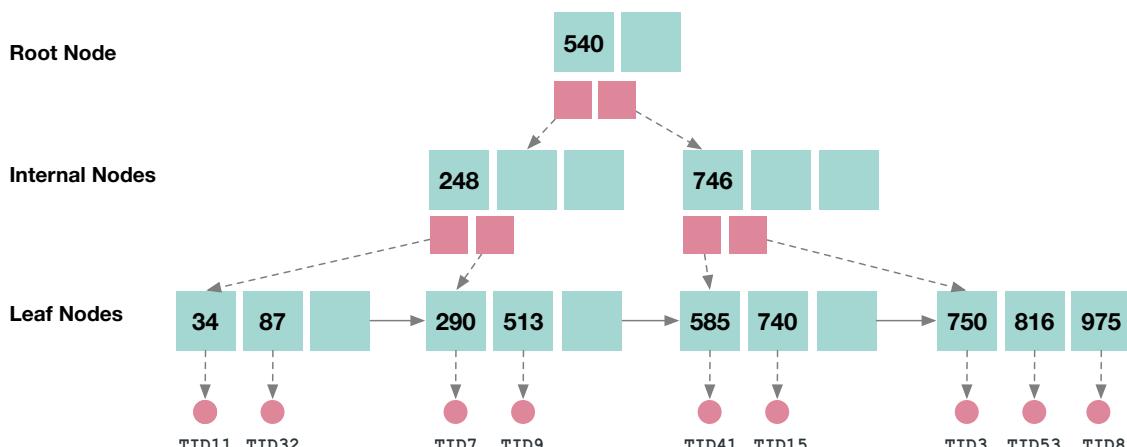


Figure 3.6 Visualisation of an example B^+ -tree with $b = 4$. Green boxes indicate the keys held by a node. Red boxes indicate pointers to children.

Logically, an index $\mathcal{I}_I^{\mathcal{R}}$ can be seen as a copy of \mathcal{R} with respect to the attribute TID as described in Definition 3.12. That is, every value of TID in \mathcal{R} is somehow contained in $\mathcal{I}_I^{\mathcal{R}}$.

Definition 3.12 Index Structure for Relation \mathcal{R}

Let \mathcal{R} be a relation with tuple identifier TID and $I = \{\mathcal{A}_1, \dots, \mathcal{A}_K\} \subset \text{SCH}(\mathcal{R})$. We call $\mathcal{I}_I^{\mathcal{R}}$ a (dense) index of relation \mathcal{R} with respect to I if $\mathcal{I}_I^{\mathcal{R}}$ contains at least TID and if $|\mathcal{I}_I^{\mathcal{R}}| = |\mathcal{R}|$ and $\exists(t_1, t_2) : t_1[\text{TID}] = t_2[\text{TID}] \forall t_1 \in \mathcal{R}, t_2 \in \mathcal{I}_I^{\mathcal{R}}$.

As a consequence, we can always use $\mathcal{I}_I^{\mathcal{R}}$ instead of \mathcal{R} and we can always reconstruct the attributes $A = \{\mathcal{A} : \mathcal{A} \in \text{SCH}(\mathcal{R}) \wedge \mathcal{A} \notin \text{SCH}(\mathcal{I}_I^{\mathcal{R}})\}$ that are contained in \mathcal{R} but not contained in $\mathcal{I}_I^{\mathcal{R}}$ by considering the equi-join on TID:¹¹

$$\mathcal{R} \equiv \mathcal{I}_I^{\mathcal{R}} \bowtie_{\text{TID}} \pi_A(\mathcal{R}) \quad (3.10)$$

However, many of the properties that go beyond Definition 3.12 depend on the type of index and the concrete implementation, e.g., which additional attributes are contained in $\mathcal{I}_I^{\mathcal{R}}$ or how data in the index is organised.

3.3.5.3 Buffer Pool Manager

The main objective of the buffer pool manager is to keep frequently accessed pages in memory for faster lookup. That is, if upper levels of the DBMS require access to a given page, they request it from the buffer pool manager, which can either serve it from memory or read it from disk through the disk manager. Meanwhile, the buffer pool manager can keep track of page access and can evict or retain pages depending on usage. In addition to caching, buffer pool managers also implement different strategies for page replacement and eviction (e.g., FIFO or LRU) and pre-fetching of pages depending on workloads.

Some DBMS, such as MonetDB [BKM08] or LMDB [Hen19], forego the implementation of a buffer pool manager and instead rely on the POSIX `mmap` system call [Sto81], which allows the mapping of files on disk into the address space of an application and leaves caching of pages to the operating system. While this simplifies the implementation because loading and eviction of pages must not be handled explicitly, research suggests that the use of `mmap` in a DBMS is not optimal, and that knowledge about the database workload (e.g., random access vs. sequential scan) allows for more fine-grained control [CLP22].

¹¹ This is merely a logical construct and not related to how data access is implemented.

3.3.6 Common Components

The following components are also part of a DBMS and are briefly described for the sake of completeness:

Catalogue The catalogue is used to persistently store and look up information used by a DBMS to organise the data it contains, i.e., information about units of data organisation such as tables, columns, collections, indexes as well as metadata such as statistics, transactions, or locks.

Transaction One or multiple queries processed by a DBMS are often wrapped in transactions, i.e., a series of operations that should be executed together and either be committed or rolled back at the end. Often, the DBMS provides guarantees with respect to the execution of transactions such as ACID [HR83] or BASE [Pri08], e.g., assuring an all-or-nothing semantic (atomicity) for operations being executed within a transaction (the A in ACID).

Transaction Managers In order to provide the aforementioned guarantees for a transaction, the query workloads must be orchestrated across different components (e.g., disk manager, buffer pool manager, execution engine, lock manager, recovery manager). This is the main task of the transaction manager in addition to keeping track of a transaction's life cycle.

Lock Manager One approach to providing transaction isolation (the I in ACID), i.e., proper protection of the data in case different transactions access it concurrently, is the implementation Strict 2-Phase Locking (S2PL). This is typically implemented by a lock manager, which keeps shared and exclusive locks on pages that are accessed and detects lock cycles.

Recovery Manager Typically, databases are required to provide some form of durability with respect to the execution of a transaction. Therefore, once a transaction commits, changes are expected to be permanent (the D in ACID). This is guaranteed by different techniques that allow for data reconciliation in case of a system failure (e.g., power outage), such as write-ahead logging. This is the main task of the recovery manager, and a famous algorithm is known as ARIES [MHL⁺92].

Obviously, some databases feature additional system components depending on concrete applications, e.g., a distribution manager for DBMS that provides distribution of data and processing across multiple nodes.

*All our knowledge begins with the
senses, proceeds to the
understanding, and ends with
reason.*

4

— Immanuel Kant, 1781

On Multimedia Analysis and Retrieval

Multimedia refers to the combination of different content forms for the purpose of transport and perception. Over the centuries, countless content forms have been conceived by mankind, natural language being one of the earliest examples. With the invention of the modern computer, more and more of these content forms were transferred to the digital world. Today, multimedia data is ubiquitous and used in different forms in every area of our socioeconomic lives. Statista.com estimates, that the amount of data created, captured, copied, and consumed in 2025 will exceed 181 ZB from an estimated 2 ZB in 2010.¹ To a large extent, that data is multimedia data that is being created, stored, analysed and shared for various purposes.

A key development in this regard was the widespread adoption of the smartphone in the early 2000s, which enabled almost every person on this planet to not only consume but also produce multimedia content turning them effectively into prosumers [RJ10; RDJ12]. Fuelled by this technological progress, the past decades have seen a staggering development in both volume and variety of multimedia data found on the Internet and in commercial and private data collections.

This extreme growth as well as the increasing velocity at which data is being produced, poses an enormous challenge for information processing. Therefore, researchers have tried to tackle the questions surrounding the efficient and effective management, analysis and retrieval of multimedia data at a large scale for many decades now, with first attempts reaching back to the early 1970s.

¹ See <https://www.statista.com/>, “Volume of data/information created, captured, copied, and consumed worldwide from 2010 to 2025”; accessed February 2022.

4.1 Multimedia Data and Multimedia Collections

The term *multimedia* refers to a combination of one or many different *content formats* or *media types*, such as but not limited to observable aural or visual information. We distinguish between these content formats based on the representation we use when working and interacting with them. Raw sound, for example, is formed by air pressure waves that are registered by our ears. In contrast, visual information involves electromagnetic waves captured by our eyes. In both cases, the brain interprets the underlying processes, thus forming the human perception of the phenomenon and the meaning it carries.

Traditionally, we often think of media as information that can somehow be directly observed by our senses. However, since a medium is merely a means to store and transport information, it may also include less apparent and sometimes even digitally native examples such as data streams stemming from sensors, which require specialised hard- and software to allow for observation.

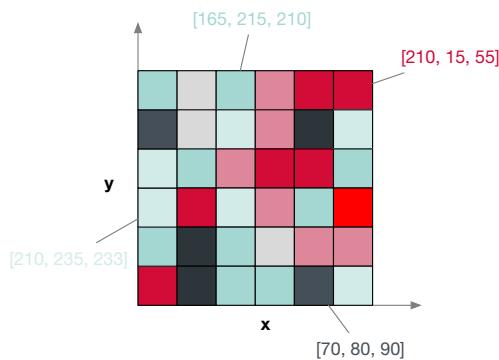
Even though the various types of media may exhibit very different characteristics in their original representation, we find some commonalities that are crucial for formalising the problem of processing and analysing their digital (data) representation (see Examples 4.1 and 4.2) in information processing systems – which is our main interest for the purpose of this Thesis.

Relationship with Time At a high level, we can roughly classify all the different media types based on their relationship with time. *Static* media types (e.g., an image) do not exhibit a temporal development, whereas the information in *dynamic* media types (e.g., an audio signal) depends on the time point that is being examined [BdB⁺07].

Unstructured Data The digital representation of a media type is usually highly unstructured in that there is no connection between the raw data and the higher-level meaning it carries. For example, there is no way to deduce the motives visible in an image just by examining the raw image data. Therefore, information processing systems rely on *derivative representations* of the original data, gained through pre-processing and data analysis.

Semantic Gap The derivative representations of the original data usually describe specific aspects thereof, and conversion to/from such representations is therefore often accompanied by a loss of information. This is a problem often referred to as the semantic gap [SWS⁰⁰; BdB⁰⁷].

Example 4.1 Digital Representation of Visual Information



The visual information in a flat image is stored as a two-dimensional array of *pixels*. The information in every pixel is typically formed by a sensor, in an array of sensors, that captures the electromagnetic signal. Each pixel holds colour values, usually one per colour channel. For example, with the RGB colour model, every pixel holds three values, one for the red, green and blue channel. The number of pixels per dimension determines the *resolution* of the image. Typically, we use a fixed number of bits per colour – the *colour depth* – which determines the number colours that can be distinguished.

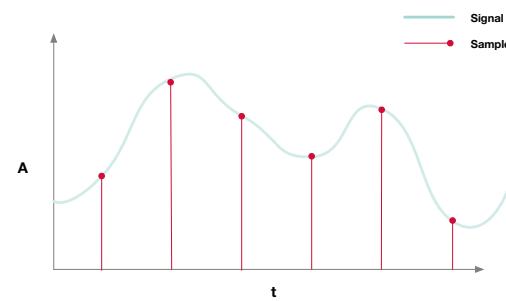
If we take, for example, a coloured image of 1000×1000 pixels, we must encode 3×10^6 individual colour values. Using 8bit per colour, we end up storing 24×10^6 bit, which amounts to 3 MB worth of uncompressed image data. Images coming from modern cameras, exhibit resolutions much higher than that.

Example 4.2 Digital Representation of Aural Information

The air pressure waves that form sound can be recorded and translated to an electrical signal by microphones. When digitising this information, the temporal development of this signal amplitude is *samped* at a fixed rate. Each sample point consists of a value that quantises the amplitude's energy at a given point in time

to a number on a given range. An audio stream is then a sequence of these values. The quality of the process is determined by the *sample rate*, i.e., the number of samples per time unit, and the *bit depth*, which determines quantisation of the amplitude power.

If we take a 1 s audio snippet, sampled at 44 kHz, we end up storing 44 000 sample points. Using a bit depth of 16 bit per sample, this amounts to 704 000 bit or 88 kB of uncompressed audio data. In modern audio systems, we typically record multiple audio channels independently, leading to even more data.



4.1.1 A Data Model for Multimedia Collections

In Chapter 3, we have introduced the notion of a data model and its importance for working with any type of data. Despite being agnostic to the concrete conceptual data model in principle, the work presented in this Thesis still relies on a formal understanding of media data that enables organisation and management thereof. Due to its tight integration into the *vitrivr* project [RGT⁺16; GRS19a; HSS⁺20], there is a particular model that we would like to introduce. Under this data model – which is illustrated in Figure 4.1 – a (multi-)media *collection* consists of multiple (typically a large number of) *media objects*. The media object describes the data at the level of individual documents – e.g., a video or image file – and comprises basic attributes such as its identifier, its media type, or its path. The path forms the link to the raw media data in the file system, which we assume to be the most suited form of storage for this type of information. That raw data must always be available for processing and presentation.

To account for the temporal development found in dynamic media types, *vitrivr*'s data model has a notion of a *media segment*, which represents a clearly defined temporal slice of the object and stands in a one-to-many relationship with it. However, the model does not make any assumption as to how segments are formed. For static media types, such as images, there is a trivial one-to-one relationship between an object and a segment. For more complex types of media, such as audio or video, media segmentation is a dedicated area of research [KC01] and can be approached in many different ways, e.g., at a fixed rate, based on changes in the content [Foo00; TYB16] or using self-adaptive deep neural networks [SML19]. Obviously, different strategies for segmentation yield different degrees of summarisation of information over time. In addition, the proposed distinction also allows for application specific media types and segmentation strategies, such as image sequences – which were introduced for LSC 2020 [HAPG⁺20] and group images per day for the purpose of lifelog analysis.

To model the different derivative representations that are used to describe a media object and its segments, the model foresees *features*, which again stand in a one-to-many relationship with the segments. That is, every segment can have an arbitrary number of such features that describe different aspects of its content. In practice, different types of features are obtained and stored in dedicated entities. Additionally, *descriptive metadata* allows for a description of the media both at the document and segment level by means of simple key-value pairs that can contain technical metadata, descriptions, labels, or other types of structured information.

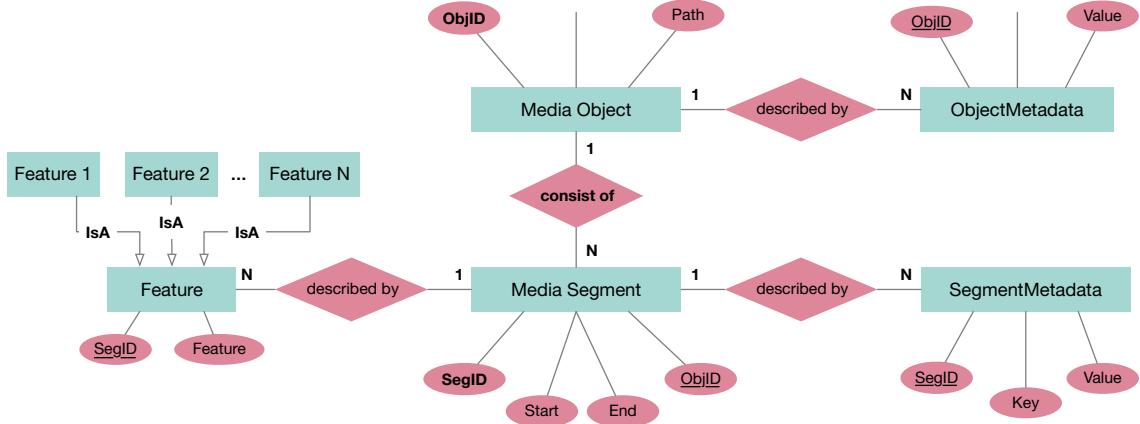


Figure 4.1 An extended ERM of the multimedia data model used in the *vitrivr* project. The model centres around the notion of media objects and media segments which are described by metadata and features.

In this Thesis, we consider a slightly altered version of the aforementioned data model, which is more in line with [BdB⁺07; ZW14] and is illustrated in Figure 4.2. This model also considers media objects to be an abstraction of individual files. However, this model foregoes the indirection introduced by the media segment and treats *features*, *annotations* and *descriptions* as different types of a more general *metadata* type that describes the media object directly. To model the temporal aspect in dynamic media types, the metadata itself exhibits information about the part of the object that is being described, denoted by an optional *start* and an *end* marker. In practice, this could be a frame number or a timestamp.

While the difference between the two models may seem marginal, we argue, that the latter offers several advantages over the former: Firstly, it eliminates a level of indirection introduced by the media segment and thus simplifies handling of static media types, which are simply described by different types of metadata that lack a start and end marker. Secondly, it does not assume segmentation to be statically defined and instead makes this a property of the metadata itself. This is reasonable, because the optimal segmentation strategy may depend on a particular feature, especially when multiple media types are involved, e.g., aural and visual information in a video. And finally, the model could be extended to also support spatial information as a basis for segmentation, e.g., in images, which would enable description of spatio-temporal aspects of any media type. However, such ideas are beyond the scope of this Thesis.

We will use the proposed conceptual data model from Figure 4.2 throughout this Thesis and whenever we refer to terms like collection, media object, feature, or descriptive metadata, we refer to the concepts introduced in this model.

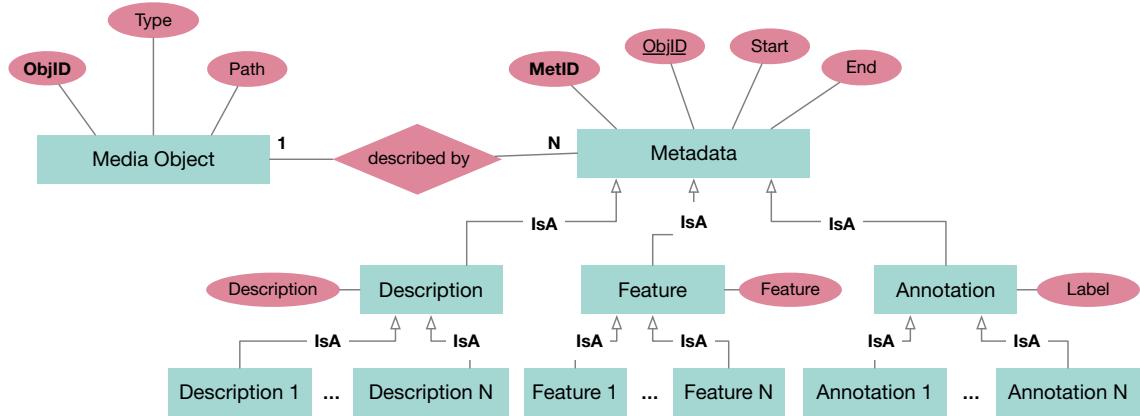


Figure 4.2 An extended ERM of the multimedia data model used for the purpose of this Thesis. It has been derived from *vitrivr*'s data model and foregoes the explicit segment entity.

4.1.2 Descriptive Metadata

Descriptive metadata in the sense of the proposed data model consists of textual descriptions (e.g., title, summary), technical information (e.g., duration, location, frame rate) and annotations (e.g., category labels). Such information has always played an important role in multimedia retrieval and analysis, because it provides (indirect) access to the media's content and can be leveraged in a structured way, e.g., in database queries or for data organisation.

Over the years, many different metadata standards have emerged. ID3² tags, for example, allow for organisation of large music libraries and classification of songs based on information about artists or albums. Similarly, the Dublin Core Metadata Element Set (DCMES or just Dublin Core)³ comprises 15 basic properties that can be used to describe any type of media, e.g., videos or images. And last but not least, EXIF⁴ can be used to describe images and also includes technical metadata, such as the camera model, the exposure time or the f-number. These well-established standards are complemented by a plethora of domain specific standardised and non-standardised metadata frameworks.

However, while useful and important for the problem of analysis and retrieval, this type of metadata comes with important disadvantages. Firstly, and notwithstanding recent developments in machine learning, e.g., in image captioning [HSS⁺19], such information is traditionally assigned manually in a laborious and time-consuming annotation process. This process is not at all able to keep up with the ever-increasing velocity at which new content is created.

² See <https://id3.org/>; accessed September 2022.

³ See <https://www.dublincore.org/>; accessed September 2022.

⁴ See <https://www.cipa.jp/>; accessed September 2022.

Secondly, descriptions and labels – especially if not standardised – are often subjective due to language, expertise, and personal experience and may therefore differ depending on the person assigning them. This is closely related to the problem of the perceptive and interpretative gap [Ros18] and highly relevant for any search process. And finally, it is often challenging to describe the content of media in a textual manner, especially if temporal development is involved.

Regardless of these challenges, however, experience shows that annotations – be they assigned manually or automatically – play an important role, especially when it comes to information retrieval. A contributing factor here is that of query formulation, which is very low-effort and familiar for textual input and becomes more challenging for other types of queries.

4.1.3 Features

Simply put, a *feature* is a mathematical object that describes “derived characteristics” [BdB⁺07] of a media object or a part thereof. Therefore, they are sometimes also referred to as *descriptors*. A formal definition is given in Definition 4.1.

Definition 4.1 Definition of a Feature in Multimedia Analysis

Let $\mathcal{O} = \{o_1, o_2, \dots, o_N\}$ be a collection of media objects o (i.e., files). A *feature* $f_{i,\text{start},\text{end}} \in \mathbb{F}$ describes a temporal interval $[\text{start}, \text{end}]$ with $\text{start}, \text{end} \in \mathbb{N}_0$, $\text{start} \leq \text{end}$ of the media object $o_i \in \mathcal{O}$ under the feature transformation t , with

$$t: \mathcal{O} \times \mathbb{N}_0 \times \mathbb{N}_0 \longrightarrow \mathbb{F} \quad (4.1)$$

We call \mathcal{O} the media object domain and \mathbb{F} the feature domain. The structure of \mathbb{F} is fully determined by t .

Features play a central role in many different domains, ranging from mere multimedia analysis to retrieval. [BdB⁺07] distinguishes between low-level and high-level features, which is a qualitative assessment of how far removed the derived characteristic of a particular feature is from a concept that has a semantic meaning to a human user. For example, a histogram of the colours found in an image would be considered a low-level feature, whereas a feature capturing salient key points or even objects in the same image would be classified as high-level. In practice, high-level features are often derived from low-level features, and again these conversions are subject to the semantic gap.

The process of generating features from media objects is called *feature extraction* [BdB⁺07] and there is an entire corpus of research that deals with the engineering of features that describe the relevant aspects of a given media type's content, with first attempts in computer vision and image understanding dating back to the early 1960s. Various surveys, such as [MB03; DZJ⁺12; SJ19], cover the topic of feature extraction and we simply name and describe a few selected examples to provide an (unrepresentative) overview. Obviously, every type of media has its own analysis techniques that can be used to generate features.

For images, one can roughly distinguish between low-level colour, texture, and shape features [SJ19], as well as higher-level local keypoint-based features. An example of a colour feature could be a simple colour histogram or a statistical moment that captures colour distribution, e.g., the average colour in a defined area of the image. Prominent examples for keypoint detection features include algorithms such as Scale Invariant Feature Transformation (SIFT) [Low99], Speeded-up Robust Features (SURF) [BTVG06] or Histogram of Oriented Gradients (HOG) [DT05]. All these features are employed in many different applications ranging from retrieval and similarity search to the detection of higher-level concepts, such as, objects or faces [DBS⁺11; Far16].

Similarly, there is a wide range of features applied in audio analysis and retrieval. Mel-frequency Cepstral Coefficients (MFCCs) – a representation of a short-term power spectrum on the perceptual *mel scale* – play an important role in speaker recognition, sound and music classification as well as audio segmentation [KS04]. They were shown to outperform audio descriptors outlined in the MPEG-7 standard [QL01], which is a collection of standardised feature descriptors for images, audio, and video. Other types of features include Pitch Class Profiles (PCP) for music retrieval and classification [Lee06; DBS19] or features based on rhythm, timbre, or speed.

So far, we have mostly considered features that were engineered through the careful application of signal processing and/or statistical analysis of the raw media data. However, advances in machine learning and deep neural network architectures have enabled the feature extraction process to be automated to some extent, leading to a shift from engineered to learned features [HE10; GAR⁺16]. There is also research that deals with the comparison of the two paradigms in multimedia retrieval and analysis [BGS⁺17] as well as other domains, such as chemical structure analysis [GLSJ⁺21]. Some findings indicate, that while learned features often outperform the engineered ones, best results are sometimes attained by a combination of the two [GLSJ⁺21].

Given the cornucopia of features and feature extraction techniques to choose from, there are two important aspects that ought to be considered with respect to data modelling and data management: Firstly, the resulting features often exhibit a comparatively simple mathematical structure, despite the vast range of feature extraction techniques. We show in Section 4.2 that in multimedia retrieval, one often deals with vectors in a high-dimensional, real-valued vector space [ZAD⁰⁶]. Secondly, when considering a complex enough use case, no single feature can satisfy all the different types of applications because, as we have pointed out, a specific feature usually focuses on a specific aspect of the content. Consequently, a combination often yields the best results [DKN08].

Both aspects have important implications for the requirements for a multimedia data management and processing system, especially, if such a system should be able support different use cases that are not necessarily known in advance [SWS⁰⁰]. We would also like to point out, that both data models introduced in Section 4.1.1 are flexible enough to handle these two aspects.

4.2 Multimedia Retrieval

Multimedia retrieval deals with algorithms, methods and systems that allow for content-based search and obtaining items of interest from large multimedia collections based on user-defined queries. In addition to text-based search, such queries may also involve new modes of query formulation such as *Query-by-Example* [KCH95], *Query-by-Sketch* [CWW¹⁰], *Query-by-Humming* [GLC⁹⁵], or *Query-by-Sculpting* [BGS²⁰]. Many different research domains for the different types of media have emerged over the years, including but not limited to content-based image retrieval [DA13], audio retrieval [Lu01], video retrieval [HXL¹¹], 3D model retrieval [YLZ07] and various subdomains [MK18].

Formally, the multimedia retrieval problem can be characterised as finding all media objects $o \in \mathcal{O}$ that may be relevant to a given *information need* expressed by a user as query q . This is achieved by obtaining some form of (*dis-*)similarity⁵ between q and o and ranking results based on it, which is why the method is also referred to as *similarity search*. In contrast to Boolean search with its strict match/miss-semantic, the results produced by similarity search are non-binary in the sense that an object o may match a query q only to a certain degree but still be considered part of the resultset.

⁵ Depending on the context, the likeness of an object to a query may be directly (similarity) or inversely (dissimilarity) proportional to the obtained score.

Since a direct comparison of a media object to a query is often not feasible – for reasons described in Section 4.1 – features are used as a proxy instead. This is a manifestation of the idea that derivative representations are considered in lieu of the actual object, leading to the formal Definition 4.2 of similarity search.

Definition 4.2 Similarity Search

Let $\mathcal{O} = \{o_1, o_2, \dots, o_N\}$ be a media collection of media objects o and let further \mathbb{F} be a feature domain of \mathcal{O} under the feature transformation t . Furthermore, we assume a query $q \in \mathbb{F}$ and the existence of an inverse t' with $t'(f) = o$.^a Similarity search then is the optimisation problem of finding the object $o \in \mathcal{O}$ that given a query q and a similarity function $s: \mathbb{F} \times \mathbb{F} \rightarrow [0, 1]$ maximises the following expression.

$$o_r = t'(\operatorname{argmax}_{f \in \mathbb{F}} s(f, q)) \quad (4.2)$$

We call the resulting item o_r most similar with respect to q , which is indicated by the similarity score $s(f, q)$. In practice, we may be interested in the top k results rather than just a single item.

^a While this may seem like a preposterous assumption at first, this is a given since we explicitly store the relationship between o and f , according to the data model in Section 4.1.

In simple terms, a multimedia retrieval system operates upon features $f \in \mathbb{F}$ derived from the original media objects by means of a feature transformation. This transformation is applied twice: Once when indexing a media object in the collection – often a step that is thought of as taking place offline – and once when deriving the same feature from the query input provided by the user, leading to q . Subsequently, the system generates the similarity measure $s(q, f)$ which quantifies the similarity between q and f from low (0.0) to almost identical (1.0) and selects the top k entries that maximise this score. The flow of information that results from this approach is illustrated in Figure 4.3.

Oftentimes, the function s may not express similarity but rather dissimilarity between objects, in which case we can simply consider the dual problem. Furthermore, the score may not always be normalised to $[0, 1]$, which makes it difficult to combine scores of multiple features, e.g., in *score-based late fusion* [DM10; Ros18]. However, a normalised similarity score can always be derived by applying an additional correspondence function $c: \mathbb{R}_{\geq 0} \rightarrow [0, 1]$. Both aspects are of high practical relevance, but they do not alter the basic problem.

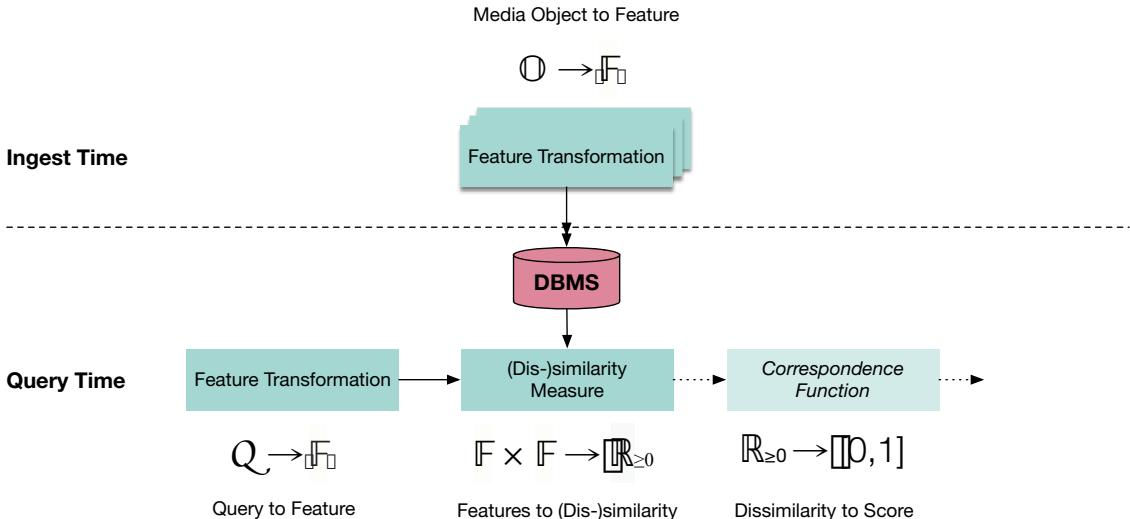


Figure 4.3 Flow of information in similarity search at ingest and query time.

4.2.1 Similarity Search and the Vector Space Model

The *vector space model* of similarity search [SWY75] assumes that the domain of the feature vectors \mathbb{F} is a subset of a mathematical vector space – that is, a set whose elements may be added and multiplied by a scalar. In practice, we often consider that vector space to be real-valued and high-dimensional, i.e., $\mathbb{F} \subset \mathbb{R}^d$ wherein d is the dimension and an inherent property of the feature and the underlying transformation t . Consequently, every feature $f \in \mathbb{R}^d$ is simply a point in that high-dimensional vector space.

Under this premise, we consider the dissimilarity function $\delta(f, q)$ to be a function that calculates the distance between a feature and a query $f, q \in \mathbb{F}$, as defined by Equation (4.3). Typically, the farther two vectors lie apart under the distance measure, the more dissimilar they are, which is why we also call this proximity-based search. A list of important distance functions used in similarity search is given in Table 4.1.

$$\delta: \mathbb{F} \times \mathbb{F} \rightarrow \mathbb{R} \quad (4.3)$$

In broad terms, one can distinguish between *continuous* and *discrete* distance functions [ZAD⁰⁶]. Continuous functions typically exhibit a very large, potentially infinite co-domain whereas discrete functions map to a limited, often pre-defined set of values. Furthermore, some distance functions induce a topology on the underlying vector space – often denoted as *metric space* (\mathbb{F}, δ) . We call these *metric distances*, and they satisfy the *non-negativity*, *symmetry*, *identity of indiscernibles*, and *triangle inequality*.

$$\begin{aligned}
\mathbf{d}(f, q) &\geq 0 && \text{(non-negativity)} \\
\mathbf{d}(f, q) &= \mathbf{d}(q, f) && \text{(symmetry)} \\
\mathbf{d}(f, q) = 0 &\iff f = q && \text{(identity of indiscernibles)} \\
\mathbf{d}(f, q) &\leq \mathbf{d}(f, g) + \mathbf{d}(g, q) && \text{(triangle inequality)} \\
\forall f, g, q \in \mathbb{F} \wedge (\mathbb{F}, \mathbf{d})
\end{aligned}$$

The constraints imposed on a metric space give rise to mathematical properties that can be exploited for efficient execution of similarity search and specialised indexing techniques [ZAD⁺06]. However, despite their mathematical convenience, some applications involve distances that do not exhibit all the aforementioned properties of a metric, e.g., quasi-metrics (lack symmetry), semi-metrics (lack triangle inequality), meta-metrics (lack identity) [ZAD⁺06], or even necessitate the use of non-metric functions [SB11].

The choice of (dis-)similarity measure mainly depends on the application. All the Minkowski distances (i.e., L^1 , L^2 , L^p) can be used for any type of real-valued vector and in practice, it often makes little difference which one is employed [Ros18]. Of all the Minkowski distances, the Euclidean distance (L^2) is the most common. The Chi-squared distance is well suited, if the feature vectors represent histograms [PW10]. The Hamming and Levenshtein [Lev65] distances belong to the class of *edit distances* and can be used to compare strings.

In the following sections, we will explore different variants of the similarity search problem, which are all variants of Definition 4.2 and illustrated in Figure 4.4 and Example 4.3.

Name	Symbol	$\mathbf{d}(f, q)$	Domain	Co-Domain	Metric
Manhattan	L^1	$\sum_{i=1}^d f_i - q_i $	\mathbb{R}^d	\mathbb{R}	metric
Euclidean	L^2	$\sqrt{\sum_{i=1}^d (f_i - q_i)^2}$	\mathbb{R}^d	\mathbb{R}	metric
Minkowski	L^p	$\sqrt[p]{\sum_{i=1}^d (f_i - q_i)^p}$	\mathbb{R}^d	\mathbb{R}	metric if $p \in \mathbb{N}_{\geq 1}$
Chi-Square	χ^2	$\sum_{i=1}^d \frac{(f_i - q_i)^2}{f_i + q_i}$	\mathbb{R}^d	\mathbb{R}	metric
Kullback-Leibler	-	$\sum_{i=1}^d q_i \log(\frac{q_i}{f_i})$	\mathbb{R}^d	\mathbb{R}	non-metric
Cosine	-	$1 - \frac{\sum_{i=1}^d f_i q_i}{\sqrt{\sum_{i=1}^d f_i^2} \sqrt{\sum_{i=1}^d q_i^2}}$	\mathbb{R}^d	$[-1, 1]$	semi-metric
Jaccard	-	$1 - \frac{ A \cap B }{ A \cup B }$	\mathbb{F}	$[0, 1]$	metric
Levenshtein [Lev65]	-	-	\mathbb{S}	\mathbb{N}	metric
Hamming	-	-	\mathbb{S}	\mathbb{N}	quasi-metric
LCS [Hir77]	-	-	\mathbb{S}	\mathbb{N}	non-metric

Table 4.1 List of distance functions often used in similarity search.

4.2.1.1 Nearest Neighbour Search (NNS)

NNS is probably one of the most prominent types of queries used in multimedia retrieval and certainly the one employed most often in the context of the *vitrivr* [RGT⁺16; GRS19a] stack. Given the query object q and an arbitrary limit $k \in \mathbb{N}$, the task is to find the k features $f \in \mathbb{F}$ that are closest to q as illustrated in Figure 4.4(a). This is also referred to as k-Nearest Neighbour Search (kNN). Formally, the resultset R is given by Equation (4.4).

$$\{R \subset \mathbb{F}: \forall f_r \in R, f \in \mathbb{F} \setminus R, d(q, f_r) \leq d(q, f) \wedge |R| = k\} \quad (4.4)$$

Algorithmically, the distance between all $f \in \mathbb{F}$ and q must be evaluated exhaustively and the top k results must be retained by ranking and ordering the results, e.g., in a bound priority queue. This is sometimes referred to as the *brute-force* approach to NNS. For $\mathbb{F} \subset \mathbb{R}^d$, its runtime complexity is $O(dN)$ and thus depends linearly on the cardinality $N = |\mathbb{F}|$ and the dimension d .

A variant of kNN is k-Farthest Neighbour Search (kFN) wherein instead of the nearest, the farthest neighbours are obtained. This can be useful, for example, to obtain negative examples to train recommender systems [PSS⁺15]. However, the underlying algorithm remains the same.

4.2.1.2 Range Search

Range search – or ϵ NN – is useful for finding objects that lie within a given range, typically expressed by a value $\epsilon \in \mathbb{R}$. Given the query object q , the task is to find all features $f \in \mathbb{F}$ that lie within ϵ of q as illustrated in Figure 4.4(b). Formally, the resultset R is given by Equation (4.5).

$$\{R \subset \mathbb{F}: \forall f_r \in R, d(q, f_r) \leq \epsilon\} \quad (4.5)$$

Algorithmically, this search strategy is similar to NNS in that all features must be evaluated exhaustively and in that its complexity for $\mathbb{F} \subset \mathbb{R}^d$ is $O(dN)$. However, in contrast to NNS, there is no need for intermediate sorting, since the selection is based on a simple numerical comparison to the constant ϵ .

Variants of ϵ NN include instances, where f must be larger than ϵ or where f must fall into an interval $[\epsilon_l, \epsilon_u]$. However, these variants are equivalent to the base case since only the selection predicate changes.

4.2.1.3 Reverse Nearest Neighbour Search (RNNS)

In RNNS [KM00] we try to find all $f \in \mathbb{F}$ that given a limit $k \in \mathbb{N}$ consider the query $q \in \mathbb{F}$ to be part of their kNN resultset as illustrated in Figure 4.4(c). Formally, the resultset R is given by Equation (4.6)

$$\{R \subset \mathbb{F} : \forall f_r \in R, q \in \text{kNN}(f_r) \wedge f \in \mathbb{F} \setminus R : q \in \text{kNN}(f)\} \quad (4.6)$$

Algorithmically, this type of search is more complex than kNN or ϵ NN, because basically, a NNS problem must be solved for every element $f \in \mathbb{F}$. Therefore, the runtime complexity is $O(dN^2)$ and scales with the square of $N = |\mathbb{F}|$.

Example 4.3 Comparing different types of similarity search algorithms

For this example, we consider $f \in \mathbb{F} \subset \mathbb{R}^2$ to be points of interest (e.g., museums, hotels, parks) on a two-dimensional city map. Furthermore, we consider $q \in \mathbb{F}$ to be another position on the same map, e.g., given by our current position or a selected point of interest. The different types of similarity queries can now be thought of as follows:

Nearest Neighbour Search is the equivalent of finding the k points of interest that are closest to q (e.g., the three museums closest to my location).

Range Search is the equivalent of finding all points of interest that are within ϵ distance from q (e.g., all museums within 3 km from my location).

Reverse Nearest Neighbour Search is the equivalent of finding all points of interest that consider q to be among their closest points of interest (e.g., finding all hotels that have a specific museum in the set of their three closest sites that must be seen when visiting the city).

4.2.2 NNS and High-Dimensional Index Structures

As we have argued in Section 4.2, even the simple NNS problem in \mathbb{R}^d exhibits a non-negligible runtime complexity of $O(dN)$ for a linear brute-force scan – thus being proportional to the cardinality $N = |\mathbb{F}|$ and the dimension d . Experience from database research shows that linear scans are too slow in terms of time required for query execution if N becomes large enough, especially when interactive query workloads are involved. This necessitates the use of index structures that allow for sub-linear data access.

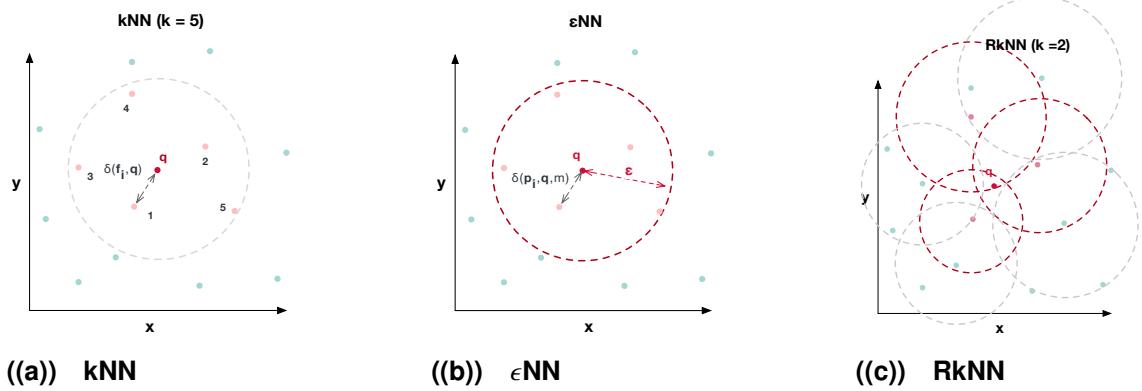


Figure 4.4 Illustration of the different similarity search algorithms in $F \subset \mathbb{R}^2$ using the Euclidean distance (L^2). Matches are highlighted in red.

The additional dependency on d has a major impact on the execution performance of any search algorithm, due to what is known as the *curse of dimensionality* [BGR⁺99; ZAD⁺06]. The impact of this “curse” is threefold: Firstly, the execution time of a brute-force algorithm is directly proportional to the dimension. With a tendency towards more complex and thus higher-dimensional vectors, there is a negative effect on runtime. Secondly, the distances between a query and the closest and the farthest point tend to become indistinguishable as dimensionality increases [BGR⁺99]. And finally, research has shown that traditional approaches, e.g., based on partitioning or classical tree structures, degenerate for higher dimensions [IM98; WSB98] with the consequence that they yield only marginal or no improvement over executing a linear scan. In fact, it was formally proven in [SR06] that the performance of almost all known index structures at that time deteriorates to that of a linear scan as d grows.

Nevertheless, the issue of index structures that allow for more efficient NNS was taken on by various researchers and is still an active area of research with results of high practical relevance. At a high-level, indexing algorithms can be classified along three dimensions: Firstly, based on whether they provide any guarantees with respect to the quality of the results, wherein quality is often measured in terms of precision and recall of a result produced by an index as compared to the result produced by a linear scan [EZP21]. Since a majority of established techniques sacrifice quality for speed [SJH⁺05], this is also known as Approximate Nearest Neighbour Search (ANNS). Secondly, based on their *organisation structure* [SR06], e.g., graph and inverted index [SZ03] or multi-index [BL14]. And finally, by the method that is employed, e.g., *quantisation*, *hashing*, *partitioning*, or *clustering*. An overview of different algorithms is provided in Table 4.2.

Name	Guarantees	Structure	Method
KDB-Tree [Rob81]	exact	tree	space partitioning
R^* -Tree [BKS ⁺ 90]	exact	tree	data partitioning
VAF & VA+ [WSB98; FTA ⁺ 00]	exact	list	scalar quantisation
LSH [IM98; WZs ⁺ 18]	approximate, probabilistic	inverted index	hashing
PQ [JDS10]	approximate	list / inverted index	vector quantisation
CP & eCP [CPR ⁺ 07; GJA10]	approximate	inverted index	clustering
NV-tree [LÁJ ⁺ 09]	approximate	tree	scalar quantisation
NGT [Iwa16]	approximate	graph / tree	proximity graph
MRNG [FXW ⁺ 17]	approximate	graph	proximity graph
HNSW [MY18]	approximate	graph	proximity graph
SPANN [CZW ⁺ 21]	approximate	graph	proximity graph

Table 4.2 Non-exhaustive list of high-dimensional index structures employed in similarity search and retrieval.

Early examples of high-dimensional indexes mainly relied on space or data partitioning, i.e., they assigned features to buckets or data files based on some partitioning of either the vector space [Ben75; Rob81; FB74] or the data [Gut84; BKS⁺90; CPZ97]. While these indexes were able to produce exact results, the speed-up they could provide quickly diminished as dimensionality increased due to the aforementioned dimensionality curse [SR06], which led to the inception of approximate technique such as LSH [IM98] or PQ [JDS10].

Recent research in high-dimensional indexing focuses on the use of graphs [SOV⁺21], for example, proximity [ZWN22] or Hierarchical Navigable Small World (HNSW) graphs [MY18; CZW⁺21], to facilitate fast ANNS. Another focus lies on the ability to execute NNS indexing and search on a GPU [JDJ19; ZTL20] or specialised hardware [LCM⁺22], dynamic indexing, i.e., indexes that allow for online processing of changes to the data [ÓJA⁺11; ZWN22] and disk-based indexing [JSDS⁺19]. In the next few sections, we will highlight a few, selected examples of high-dimensional index structures.

4.2.2.1 Vector Approximation File (VAF)

Vector Approximation File (VAF) were proposed by Weber et al [WSB98] to address the problem of the dimensionality curse. The basic idea of a VAF is the division of the data space into 2^b cells, as illustrated in Figure 4.5, where b denotes a fixed number of bits. Typically, one chooses a specific number of bits per dimension b_d and uniformly partitions the vector space along each dimension, i.e., $b = b_d d$. A feature's position in the resulting lattice is concatenated into a compact *signature* of size b and the VAF constitutes a list of those signatures.

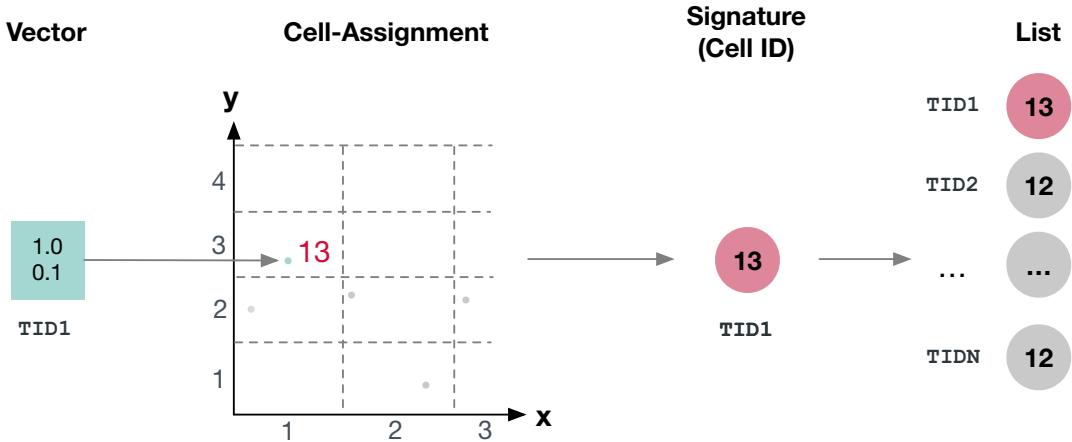


Figure 4.5 Simplified illustration of basic VAF indexing scheme in for $\mathbb{F} \subset \mathbb{R}^2$. A vector is assigned to a cell in an overlay grid, which gives rise to a signature that is stored in a linear list with the vector or its tuple identifier.

NNS is then realised as a linear scan of the file. Speed-up over a brute-force scan of the original vectors is achieved in two ways: Firstly, the signatures are more compact than the original representation (b bits per dimension), i.e., it reduces IO cost through scalar quantisation. Secondly, the signatures can be used to obtain approximate bounds that are computationally less expensive than the actual distance calculation. These bounds can be used to filter out a majority of the points in the file, further reducing IO and CPU costs. Experimental results presented in [WSB98] demonstrate that more than 90% of all features can be skipped and thus a distance must only be obtained for the remaining 10%. However, VAF remain linear in nature and [EJP21] determine analytically that for VAF, a degeneration is still to be expected for high dimensions. Nevertheless, a huge advantage of the VAF index is, that it does neither sacrifice precision nor recall to obtain speed-up.

The VA+ files described by [FTA⁺00] are an extension to the VAF and introduce a few optimisations to the original algorithm, which implicitly assumes independence of the individual dimensions within a feature (which is often not a given). They suggest a prior decorrelation of the individual dimensions by applying a Karhunen-Loève Transform (KLT). Furthermore, they propose a non-uniform assignment of bits per dimension, based on data distribution, as opposed to a uniform distribution employed in the original paper. This leads to an independent quantiser per transformed dimension. Results suggest that these adjustments lead to a better estimation of lower and upper bounds, which improves the VA file's filtering capabilities and thus further reduces IO cost.

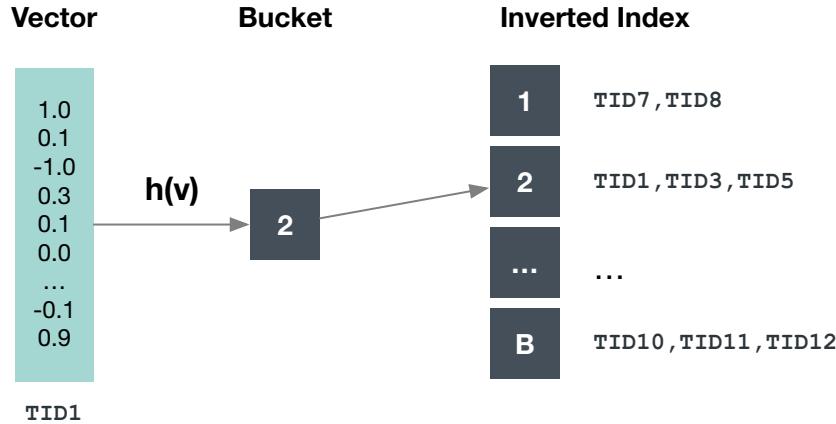


Figure 4.6 A simplified illustration of the basic LSH indexing scheme. A vector is hashed to a bucket by a locality preserving hash function h and stored in a hash table.

4.2.2.2 Locality Sensitive Hashing (LSH)

LSH refers to a family of algorithms [EJP21; WZs⁺18] that all share the same basic idea of *locality preserving hashing* as first proposed by [IM98]. The idea is centred around a hash function $h : \mathbb{F} \rightarrow B$ that maps any two points $f, q \in \mathbb{F}$ to a bucket $b \in B$ such that Equations 4.7 and 4.8 hold, wherein $(\mathbb{F}, \mathfrak{d})$ must constitute a metric space, $P_1, P_2 \in [0, 1], P_1 > P_2$ denote probabilities, $R > 0$ denotes an arbitrary threshold and $c > 1$ denotes an approximation factor.

$$\mathfrak{d}(p, q) \leq R \Rightarrow h(f) = h(q) \text{ with a probability of at least } P_1 \quad (4.7)$$

$$\mathfrak{d}(p, q) \geq cR \Rightarrow h(f) = h(q) \text{ with a probability of at most } P_2 \quad (4.8)$$

In simple terms, a locality preserving hash function has a high probability for collision (i.e., mapping to the same bucket) if two features $f, q \in \mathbb{F}$ lie close to one another, i.e., $\mathfrak{d}(q, f) \leq R$, but exhibits a low probability for collision if $\mathfrak{d}(q, f) \geq cR$. A feature $f \in \mathbb{F}$ can then be hashed to a bucket and stored in a hash table with other vectors that lie close to it as illustrated in Figure 4.6.

Speed-up for NNS can be achieved by limiting distance calculation to the vectors in the bucket a query q falls into. This is referred to as non-exhaustive search. However, due to the probabilistic nature of the hash function, this type of search strategy may exhibit errors due to vectors that were hashed into the “wrong” bucket. More refined strategies can be employed, including hash-code ranking as well as single- and multi-table lookup [WZs⁺18] to further minimise that error. Nevertheless, LSH remains an approximate search strategy. A survey of different LSH algorithms can be found in [WZs⁺18].

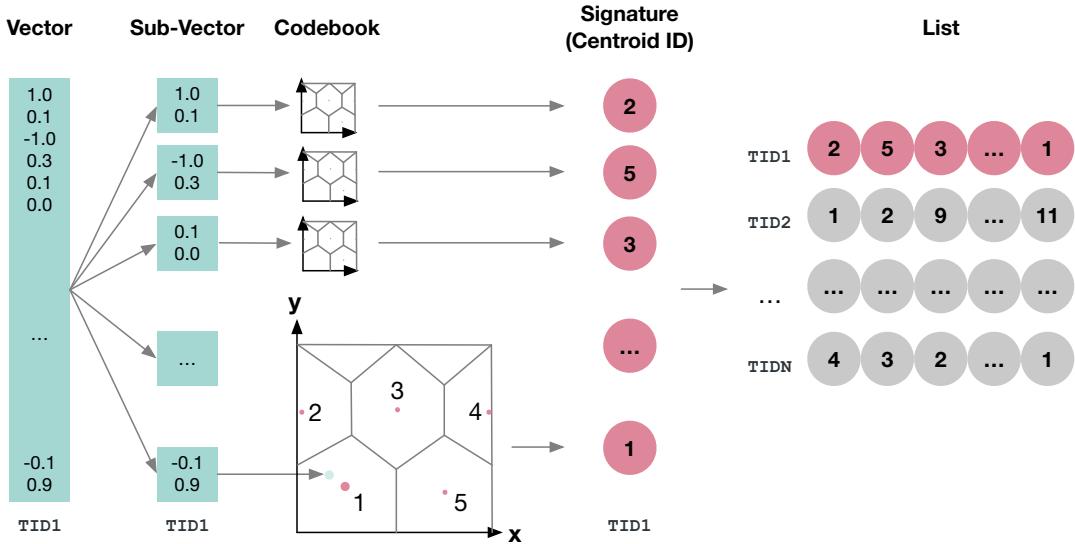


Figure 4.7 A simplified illustration of the basic PQ indexing scheme. A vector is divided into smaller subvectors and each subvector is assigned to the nearest centroid in the codebook of the respective subspace, which produces the signature.

4.2.2.3 Product Quantisation (PQ)

The idea for PQ was proposed by [JDS10] and it describes a technique for quantising high-dimensional vectors. The main idea involves decomposition of the original vector space \mathbb{R}^d into smaller subspaces \mathbb{R}^s with $s < d$ and separate quantisation of each subspace. The quantisation is achieved by learning a codebook for every subspace through k-means clustering on a representative sample of the data. Upon indexing, all subvectors of each feature vector $f \in \mathbb{F}$ are assigned to the nearest centroid in the codebook for the respective subspace. Concatenation of the centroid indexes gives rise to a compact signature that can be stored in an index, as illustrated in Figure 4.7.

Speed-up of NNS is achieved in two ways: Firstly, the resulting signature is a highly compressed representation of the original vector, which reduces IO by a large margin as is demonstrated in Example 4.4. Secondly, the signature can be used to approximate the distance between a query q and a feature f using a pre-calculated lookup table and inexpensive operations, which significantly reduces CPU costs. However, this speed-up is bought by a distortion of the approximated distance, which may lead to errors if features lie very close to one another and makes PQ a suboptimal fit if exact distance values are required. [JDS10] also discusses several optimisations, e.g., encoding the difference between feature vector and centroid – the residual – and to use an additional, coarse quantiser to assign vectors to entries in an inverted file index (IVFPQ).

Example 4.4 Size of PQ-Signature vs. original feature

Let us assume $\mathbb{F} \subset \mathbb{R}^{256}$, i.e., a 256-dimensional vector space ($d = 256$). Assuming a float data type, we require 4 B per vector component, i.e., 1024 B to store a single vector.

Using PQ, this space could now be divided into 16 subspaces with a dimensionality of 16 each, i.e., $s = 16 < d$. If now we assume, that we create a codebook consisting of 128 clusters for each subspace, then every subvector will be assigned to one of 128 centroids, which will result in 16 numbers (one per subspace) between 0 and 127. Each of these numbers can be stored in a single byte, meaning, we only need 16 B per vector to store an index entry.

4.2.3 Beyond Similarity Search

As we have argued in Section 2.1, the query workloads in real-world multimedia retrieval systems are often far more diverse than what we have described so far. In addition to one-shot queries, users find themselves searching, refining, browsing and exploring datasets with the aim of finding a particular item or a class of items [LKM⁺19; RGL⁺20]. This is exemplified by interactive retrieval competitions such as VBS [Sch19; LBS⁺18] or LSC [GJS⁺21], where participants are required to find items of interest in large, standardised data collections [BRS⁺19; RSB21] within a limited amount of time.

The many participants to the aforementioned campaigns have demonstrated interesting techniques that complement or even replace mere similarity search. While the 2022 instalment of VBS was won by a system combining search using the CLIP model [RKH⁺21] and sophisticated browsing [HSJ⁺22], the team behind [KVM⁺20] has won the 2020 iteration of VBS with a system that relies on Self-organizing Maps (SOMs) [Koh90] to generate clustered visualisations of the high-dimensional data collection initialised by search input. Similarly, the system that won in 2016 [BHM16] used hierarchical graphs and visually sorted image maps to facilitate exploration without providing any classical search functionality. Generally, it was found that efficient browsing and exploration is as important for effective retrieval as an effective similarity search model [LKM⁺19]. A similar comparison of systems at the LSC 2018 has shown, that a combination of functionality that complements similarity search is very common, including functions such as Boolean search on biometric data, facet filters, and visual clustering [GSJ⁺19]. Hence, as argued before, multimedia retrieval goes way beyond mere similarity search but often relies on it to some extent.

The trend away from mere search seems to have gained traction over the past few years: The Exquisitor system [RPK⁺19], for example, which participated to both LSC [KDJ⁺21] and VBS [KSJ⁺22], mainly uses relevance-feedback to train and adapt a semantic classifier (interactive learning). The ViRMA system [DJ22], also present at both LSC and VBS, provides sophisticated, multi-dimensional exploration in virtual reality which is enabled by the M^3 data model [DJ21]. That very same model was also used by the PhotoCube system [SWD⁺21], however, using a traditional desktop UI. Nevertheless, to some extent that transition also seems to be fueled by the novel means of interaction offered by virtual reality and 3D visualisation [SGH⁺21a; SGH⁺21b; SGH⁺22].

Information is the oil of the 21st century, and analytics is the combustion engine.

5

— Peter Sondergaard, 2011

Multimedia Data Management

The insight that special requirements must be addressed to accommodate multimedia data in traditional DBMS can be traced back to the 1990s [MS96; AN97] with many theoretical contributions but very little in terms of concrete (open-source) implementations. However, from a database research perspective, the topic received renewed attention with the publication of the *Lowell Database Research Self-Assessment* in 2005 [AAB⁺05] and the *Claremont Report on Database Research* in 2008 [AAB⁺08], both of which identified support for unstructured and multimedia data as well as the integration of information retrieval and database functionality as important research directions to consider.

In contrast, multimedia retrieval has traditionally relied on highly optimised data structures for storage and retrieval that were typically tightly coupled with the retrieval system itself in a monolithic architecture, offering little to no support for advanced data management functionality. It was not until the *Ten Research Questions for Scalable Multimedia Analytics* [JWZ⁺16] that the issue of scalable data management for multimedia analytics and a convergence with database research was raised in the face of growing data collections [BRS⁺19; RSB21]. Interestingly, many of the points raised by Jónsson et al. are related to aspects listed in the 2014 *The Beckman Report on Database Research* [AAA⁺14], namely the focus on “end-to-end processing systems” and “data understanding”, demonstrating, that the two fields do have much in common.

However, while the amalgamation of text retrieval with classical database functionality became a topic of interest, multimedia retrieval remained a largely separated field of research from that of database research with the exception of a few isolated contributions that tried to converge specific aspects such as scoring and ranking [LCI⁺05; ZHC⁺06], extensions to relational algebra based on fuzzy logic [MT99], or the conception of a multimedia query language [BBZ12].

5.1 Early Multimedia Retrieval Systems

As an alternative to a monolithic architecture, some of the early multimedia retrieval systems relied on and potentially combined existing database systems and built the special requirements for multimedia retrieval on top in either a single software component (2-tier architecture) or some type of middleware that then could be used by other systems (3-tier architecture). Examples that use the 2-tier approach include systems such as the video retrieval system OVID [OT93], the image retrieval system MARS [RHM97], or systems such as QBIC [FSA⁺95] or MUVIS [KCG⁺03]. Unfortunately, neither of these systems are publicly available and therefore could not be examined in detail.

A prominent example is IBM's QBIC system [FSA⁺95], which allowed for (visual) similarity search on image and video data. The architecture of the QBIC system indicates the existence of a dedicated database layer used for storing features. However, there is little detail as to what type of database was used and what tasks were executed at a database level. Judging by the description, however, it seems that the queries were most likely processed by the "matching engine" and the database itself was only used for storage. It is also noteworthy that QBIC strictly distinguished between "database population" and "database query" as two different phases that took place independently of one another. A separation rather atypical for traditional database applications.

MUVIS [KCG⁺03] is another 2-tier multimedia retrieval system, which, in contrast to the aforementioned systems, combined multiple databases for storing image, video, and audio information separately without, however, providing an explicit reason for doing so. Again, there is little detail as to what tasks fell to these database layers and which tasks were carried out by the browsing layer built on top of it. Interestingly, we can again observe the distinction between "indexing" and "retrieval" as two separate and independent phases in the system's life cycle.

The *Garlic Approach* [CHS⁺95] proposed by IBM is an example of a 3-tier architecture. The Garlic middleware combines and integrates different, dedicated types of datastores to support heterogenous data under an extensible and modular query and runtime system. This system can then be used by other systems via standardised APIs. It was later used to integrate QBIC [CHN⁺95], which allowed for the combined search in visual and non-visual data. In a manner of speaking, Garlic can be seen as an early attempt at a polystore approach for multimedia data, like, for example, *Polypheny-DB* [VSS18].

The 2-tier architecture involving classical databases for information and multimedia retrieval is still very common today, with a tendency towards the use of NoSQL systems [MSL⁺14; dSP⁺17]. Furthermore, there has been some work on retrofitting established DBMS for multimedia retrieval, either by using existing extension endpoints (e.g., user-defined data types or User Defined Functions (UDFs)), by writing extensions, or altering core system components [GdR⁺09; WLK⁺10; GAKS14; FC13; WLL⁺15; YLF⁺20].

5.2 Libraries for Multimedia Retrieval

Instead of tightly integrating multimedia retrieval techniques with systems that use them, the past two decades bore witness to a trend towards open-source libraries that could be (re-)used. The standardised ANNS benchmarking campaign known as *ANN-Benchmarks* [ABF17],¹ for example, lists and evaluates 25 individual libraries, such as, *Spotify Annoy*,² ScaNN [GSL⁺20] by Google Research, *Hnswlib* [MY18], Facebook’s FAISS [JDJ19], or Yahoo’s NGT [Iwa16].

A noteworthy early attempt at providing such a library for content-based image retrieval was Lucene Image Retrieval (LIRe) [LC08] – an open-source Java library based on the Apache Lucene engine, which is typically used for fulltext search. LIRe was later used and extended to create LLvRE [dLG16], a similar project but for video retrieval. While both LIRe and LLvRE were an important step in the right direction, they unfortunately tightly integrated feature extraction and storage and were thus of limited use when trying to combine them with features that did not come pre-bundled. Nonetheless, these projects demonstrated Apache Lucene’s capability of executing NNS. Consequently, Lucene based projects such as ElasticSearch³ or OpenSearch⁴ still provide NNS and ANNS through plugins and extensions such as Elastiknn.⁵

An important step towards an engine for scalable NNS was the open-source library Facebook AI Similarity Search (FAISS) [JDJ19]. FAISS brings together different methods for ANNS – for example, PQ [JDS10], HNSW [MY18], and LSH [IM98] based index structures – and provides SIMD and GPU support both for indexing and querying. While FAISS is not a multimedia database system per-se, it provides a formidable foundation for building one.

¹ See <http://ann-benchmarks.com/>; accessed July 2022.

² See <https://github.com/spotify/annoy>; accessed July 2022.

³ See <https://www.elastic.co/>; accessed July 2022.

⁴ See <https://opensearch.org/>; accessed July 2022.

⁵ See <https://github.com/erikbern/ann-benchmarks/>; accessed July 2022.

However, even though all these projects are important and useful, they typically require integration into a system and do not provide an “out-of-the-box” solution that can be used via standardised interfaces such as SQL.

5.3 Vector Database Engines

The past few years have seen the conception of novel vector databases as either dedicated systems or extensions to existing ones. We refrain from presenting all these systems in detail and instead briefly name and describe a few, with the exception of one system that receives a more detailed discussion. Most of these systems have in common that they have been built as cloud-native applications with distribution in mind and that they are open-source. For every system, we reference the respective GitHub repository as well as the current version and the time of the last update to illustrate the activity in this particular domain.

Weaviate (1.14.1, July 2022)⁶ is a vector search engine that allows for the storage of vectors and attributes. Both the vector and attributes can be queried using the *GraphQL* query language and it relies on HNSW [MY18] for ANNS. Weaviate claims full CRUD support, i.e., changes to entries are possible, which however limits the types of index structures used (currently, a custom HNSW index is employed). The recent, major update to version 1.14.0 has added support for the Euclidean distance and the dot product.

Vald (1.5.5, July 2022)⁷ is a vector search engine that relies on NGT for ANNS. It allows for similarity search as well as lookups by identifier using a gRPC based query API. A unique selling point is the *asynchronous auto indexing* feature, which refreshes outdated indexes in the background in a non-blocking fashion, which allows for changes to the data.

Qdrant (0.8.5, July 2022)⁸ is a similarity search engine that can be used to store high-dimensional vectors with additional attributes – called payload. It can be characterised as a document database. Qdrant supports similarity search queries as well as filtering for payload attributes via either a gRPC or RESTful interface. Both payload attributes as well as vectors can be indexed and HNSW [MY18] is used for the latter. A query planner determines query execution in case multiple index structures are employed.

⁶ See <https://github.com/semi-technologies/weaviate/>; accessed July 2022.

⁷ See <https://github.com/vdaas/vald/>; accessed July 2022.

⁸ See <https://github.com/qdrant/qdrant/>; accessed July 2022.

5.3.1 Milvus

Milvus (2.0.2, April 2022) [WYG⁺21] is a vector database system with its most recent 2.0 major release dating back to January 2022. Milvus is advertised as a “vector database built for scalable similarity search” and it relies on many of the libraries mentioned in Section 5.2. According to information on the official blog⁹, it was built as a cloud-native application with high scalability and availability in mind. It offers support for functionality that goes beyond mere NNS, such as, Boolean filtering as well as hybrid queries involving both similarity search and Boolean predicates on scalar data types as well as index support for both types of queries. In addition, Milvus is also capable of processing online changes to data (i.e., inserts and deletes).

Data in Milvus is organised into named collections, which can consist of multiple attributes and are distributed among multiple partitions. The logical data model and type system is similar to that of Cottontail DB, however, in addition to scalar types (e.g., int, float, double or string) it only allows for float or binary vectors. Milvus offers client libraries for Go, Java (and Kotlin), Python, and Node, which offers a user-friendly facade over the gRPC-based query interface and allows for data definition, data management, and querying.

Milvus provides fast NNS through a component named *Knowwhere*, which provides a standardised operations interface that abstracts away the underlying state-of-the-art libraries and techniques such as FAISS [JDJ19], Hnswlib [MY18] and Spotify’s Annoy¹⁰ and which offers support for SIMD and GPU execution as well as various index structures. Dynamic data support for indexes is enabled by an eventual consistency model, wherein data is being (re-)indexed on a partition by partition basis, i.e., if changes occur, only the affected partitions must be rebuilt. This mechanism is combined with a MVCC-based snapshot isolation, which also allows for querying the state of the database at a specific point in time. Deletes are filtered on the fly during query time based on *tombstone* markers in the database, a concept known from Apache Cassandra.

It is important to mention that while data held by Milvus resides on disk, a collection must be loaded into main memory in its entirety in order to be able to query it. Therefore, for Milvus to be able to execute NNS or Boolean filtering, all the data (or at least the indexes), must fit into the combined main memory of the available query nodes.

⁹ See <https://milvus.io/blog/>; accessed July 2022.

¹⁰ See <https://github.com/spotify/annoy/>; accessed in June 2022.

5.4 Databases for Data Science

With the start of the NoSQL movement some researchers began to address the shortcomings of the relational data model in the context of scientific applications and proposed different architectures, such as array databases. Early examples include SimDB [SAA⁺10] or SciDB [SBP⁺11; SBZ⁺13]. The topic of science-oriented databases has been listed as highly relevant in the 2020 database research self-assessment report, due to the continued trend towards data driven applications, data science, and data analytics [AAA⁺20]. In addition to challenges in the domain of data governance, one focus that was identified is the convergence between linear algebra and relational algebra in an efficient and effective framework [LGG⁺18], an effort that could also greatly benefit multimedia retrieval and analytics applications [BDO95]. Another area of interest is the use of specialised hardware, such as FPGAs or GPUs, to accelerate query evaluation.

5.5 *vitrivr*, ADAM and ADAM_{pro}

To the best of our knowledge, the open-source *vitrivr* video retrieval [RGT⁺16] and later multimedia retrieval [GRS19a] system was one of the first systems to advertise a 3-tier architecture that included a specialised multimedia database in addition to a dedicated feature extraction and score fusion engine as well as a dedicated user interface. The first-generation database layer came in the form of ADAM [GAKS14], which can be seen as an early attempt at bridging the gap between databases and multimedia retrieval. At its core, ADAM was an extension to PostgreSQL – called *dbADAM* – that added support for NNS and VAF-based indexes [WSB98]. In that respect it is comparable to projects like *pgvector*,¹¹ *pgANN*,¹² and PASE [YLF⁺20], which however use more recent index structures such as PQ [JDS10] and HNSW [MY18]. The *dbADAM* component was complemented by *wsADAM*, which could orchestrate query execution over multiple PostgreSQL instances using MapReduce [DG08]. ADAM’s successor ADAM_{pro} [GS16] tried to address performance limitations with a multi-model (polyglot) database system approach that allowed for query execution on different storage engines and distribution of data and computation using *Apache Spark*, a large-scale data analytics platform. ADAM_{pro} still relied on PostgreSQL for storage of relational data but used *Apache Parquet* to store and access features.

¹¹ See <https://github.com/pgvector/pgvector/>; accessed July 2022.

¹² See <https://github.com/netrasy/pgANN/>; accessed July 2022.

The work on ADAM_{pro} was complemented by a theoretical model [Gia18]: Giangreco et al. demonstrated that for a database system to be able to support similarity search given the relational model as described in Section 3.2, one can extend the set system of allowed data domains \mathbb{D} by $\mathcal{D} \subset \mathbb{R}^{\text{dim}}, \text{dim} \in \mathbb{N}_{>1}$ and postulate the existence of a relational similarity operator $\tau_{\delta(\cdot,\cdot),a,q}(\mathcal{R})$ that can be used to express similarity search. According to the definition, such an operation introduces an implicit attribute in the underlying relation \mathcal{R} , which in turn induces an ascending ordering of the tuples. Two variants of τ were proposed, namely $\tau_{\delta(\cdot,\cdot),a,q}^{kNN}$ (NNS) and $\tau_{\delta(\cdot,\cdot),a,q}^{\epsilon NN}$ (range search), which select tuples by their cardinality k or a maximum cut-off distance ϵ respectively.

While postulating two new relational operators is well within the tradition of how relational algebra has been extended over the years (see Section 3.2.3), it comes with certain limitations. In its proposed form, τ addresses several functions at once: It (i) specifies the distance function that should be evaluated, (ii) generates an implicit distance attribute on the underlying relation \mathcal{R} , (iii) imposes an ascending ordering and limits the cardinality of \mathcal{R} based on a predicate or a specified limit.

Even though it may be straightforward to implement, the amalgamation of all this functionality into a single operation is very specifically tailored to the use case of NNS and range search and only of limited value when considering more general distance-based query operations. If one were to, for example, obtain the k farthest neighbours rather than the k nearest neighbours, as necessary when doing MIPS or obtaining negative examples, they would have to either change the distance function, define new operators, or extend the definition of τ .

Another important issue with the definition of τ in its proposed form is that despite the two operations $\tau_{\delta(\cdot,\cdot),a,q}^{kNN}$ and $\tau_{\delta(\cdot,\cdot),a,q}^{\epsilon NN}$ serving a very similar purpose and sharing the core definition, they behave very differently with respect to other operations. Generally, $\tau_{\delta(\cdot,\cdot),a,q}^{kNN}(\mathcal{R})$ does not commute with any selection σ due to the inherent limiting of the cardinality to a constant value, hence:

$$\sigma(\tau_{\delta(\cdot,\cdot),a,q}^{kNN}(\mathcal{R})) \neq \tau_{\delta(\cdot,\cdot),a,q}^{kNN}(\sigma(\mathcal{R})) \quad (5.1)$$

The left side of Equation (5.1) filters the results of a kNN-search on \mathcal{R} , thus returning $n \leq k$ results, wherein $n = k$ only if σ matches all tuples. The right side of Equation (5.1) performs a kNN-search on a pre-filtered relation \mathcal{R} , also returning $n \leq k$ entries. However, n will only be smaller than k if σ selects fewer than k tuples.

PART III

Dynamic Multimedia Data Management

6

*All the world will be your enemy,
Prince with a Thousand Enemies,
and whenever they catch you, they
will kill you. [...]*

— Richard Adams, 1972

Modelling a Database for Dynamic Multimedia Data

We have argued in Chapter 2 that multimedia retrieval, analysis, and analytics workloads (and workloads that exhibit similar characteristics) pose specific requirements on a database system. In this chapter, we focus on four aspects and derive a formal model to accommodate them on a systems level:

Support for Complex Data Types Multimedia workloads process features that are often real-valued, high-dimensional vectors but can be any type of mathematical object (e.g., complex numbers). A multimedia DBMS must support these objects as first-class citizens (see Requirement 2.5).

Support for Distance Computation The notion of proximity between features plays a crucial role in all types of multimedia workloads, which involves the evaluation of functions. A multimedia DBMS must be able to execute these functions efficiently, using the available data structures and execution methods (see Requirement 2.6).

Dynamic Data Data is usually subject to constant change and a multimedia DBMS must be able to propagate these changes to all the relevant data structures while maintaining an internal consistency model (see Requirement 2.3 and Requirement 2.1).

Quality vs. Time Multimedia queries often trade retrieval quality against speed or vice versa, e.g., when using high-dimensional indexes. A multimedia DBMS must make these decisions explicit and allow the user to express their preference at different levels of the system (see Requirement 2.4).

To address these requirements, we propose and formalise a set of functionalities, which will be specified in the next sections. Firstly, we propose a notion of *generalised, proximity-based queries* as an extension to the relational model. Secondly, we describe a *cost model* that takes the quality of results into account. And finally, we introduce the formal requirements to enable high-dimensional indexes to cope with dynamic data and derive a model for maintaining such indexes. All these contributions can be seen as extensions to the functionality provided by a traditional DBMS. To establish a frame of reference, we assume the relational model and ACID properties for the model DBMS. However, the introduced concepts should generalise to other data and consistency models.

6.1 Generalised Proximity-Based Operations

Starting with the vector space model for similarity search [ZAD⁺06] presented in Chapter 4, we propose to extend the notion of distance-based similarity search to that of a more general *proximity-based query* following Definition 6.1.

Definition 6.1 Proximity-Based Query

Let \mathcal{R} be a relation. Any database query that relies on the evaluation of a distance function $\mathfrak{d} : \mathcal{D}_f \times \mathcal{D}_f \rightarrow \mathbb{R}_{\geq 0}$ that calculates the distance between an attribute $\mathcal{A}_f \in \text{SCH}(\mathcal{R})$ and some query $q \in \mathcal{D}_f$, is called a *proximity-based query*. We call q the *query* and \mathcal{A}_f the *probing argument*, which both share the same *query data domain* \mathcal{D}_f . It is to be noted, that $\mathcal{D}_f \equiv \mathbb{F}$.

Definition 6.1 simply requires a notion of proximity between an attribute and a query to be obtained by evaluating a distance function. The definition does not make any assumptions as to what data domains \mathcal{D}_f the query and probing attribute belong to, nor how the distance is being used within the query.

It is obvious, that similarity search falls into the category of a proximity-based query, wherein the distance is used to rank a relation and subsequently limit its cardinality. In addition to this common application, the definition of proximity-based queries also includes operations such as evaluating a distance function followed by some filtering (e.g., range search) or aggregation based on the computed value (e.g., clustering). That is, we are not limited to mere NNS, since we consider obtaining the distance as one step, which can be freely combined with all other operations supported by the underlying algebra.

6.1.1 Revisiting Distance Computation

Since the choice of the distance function \mathfrak{d} is a crucial part of any proximity-based query, it is worth revisiting its definition. Again, starting with the metric space model [ZAD⁺06], we identify the following (implicit) constraints with respect to the distance function:

1. The domain (i.e., the input) of the distance function \mathfrak{d} is assumed to be $\mathbb{R}^{\text{dim}} \times \mathbb{R}^{\text{dim}}$, that is, the distance function is assumed to be a binary function and arguments are restricted to real-valued vectors.
2. The codomain (i.e., the output) of the distance function \mathfrak{d} is assumed to be $\mathbb{R}_{\geq 0}$, hence, the generated distance value is a positive, real number.
3. $(\mathcal{D}_f, \mathfrak{d})$ usually constitutes a metric, thus satisfying the non-negativity, identity of indiscernibles, symmetry, and subadditivity properties.

Upon closer examination and if we consider the use cases presented in Chapter 2, one must come to challenge some of these constraints. If, for example, we turn to Example 6.1, we realise that it is not reasonable to impose a general restriction of the domain of the distance function to \mathbb{R}^{dim} with $\text{dim} \in \mathbb{N}^+$

Example 6.1 Maximum Inner Product Search for MRF

In MRF (see Section 2.3), we try to obtain the signal vector $a_{i \in \mathbb{N}} \in \mathcal{D}_f \subset \mathbb{C}^{\text{dim}}$ so that it maximises the inner product to a signal (query) vector $q \in \mathbb{C}^{\text{dim}}$. In this case, the distance function \mathfrak{d} has the form $\mathfrak{d}: \mathbb{C}^{\text{dim}} \times \mathbb{C}^{\text{dim}} \rightarrow \mathbb{R}_{\geq 0}$. Hence, the domain of \mathfrak{d} is a complex vector space.

This limitation of the definition of \mathfrak{d} can be easily remediated simply by extending the set of supported data domains \mathbb{D} by \mathbb{C}^{dim} , similarly to how it was done for \mathbb{R}^{dim} as proposed by [Gia18]. Nevertheless, we have to acknowledge, that the textbook definition of a distance function is too limited for some real-world applications, and that the query and probing arguments of a distance function could be any type of value supported by the DBMS. Another, similar example could be the *Levenshtein distance* [Lev65] – often used in computer linguistics – which is a distance metric on string values. If now in addition, we consider Example 6.2, we see that restricting oneself to binary functions with $\mathbb{R}_{\geq 0}$ is also too limiting when facing certain practical scenarios.

Example 6.2 Distance Between a Vector and a Hyperplane

To find positive and negative examples $a_i \in \mathcal{D}_f \subset \mathbb{R}^{\text{dim}}$ given a linear classifier, e.g., provided by a SVM, we evaluate the distances between the attributes and a (query-)hyperplane described as $\mathbf{q}^T \mathbf{x} - b = 0$ with $\mathbf{q}, \mathbf{x} \in \mathbb{R}^{\text{dim}}$ and $b \in \mathbb{R}$. The distance function is then given by:

$$d(a_i, \mathbf{q}, b) = \frac{\|\mathbf{q}^T \mathbf{a}_i + b\|}{\|\mathbf{q}\|} \quad (6.1)$$

d now has the form $d: \mathbb{R}^{\text{dim}} \times \mathbb{R}^{\text{dim}} \times \mathbb{R} \rightarrow \mathbb{R}$. Hence, d is no longer a binary but a ternary function with arguments \mathbf{a}_i , \mathbf{q} and b . Furthermore, the distance may take negative values depending on whether a result is considered a positive or negative example.

Again, we are confronted with an example that violates the textbook definition of a distance function. While the idealised idea that distance functions used in proximity-based queries must always constitute a metric on some real-valued vector space may be very common [ZAD⁺06], we see in practice that this assumption is often violated [SB11; BS19] and that many functions used to calculate proximity between objects are not actual metrics. Even though this can be a disadvantage when considering high-dimensional index structures that exploit the geometric properties of metric spaces, it is obvious that such functions exist and must thus be considered in any generalised multimedia database application supporting proximity-based queries.

In contrast, there is good reason to assume the codomain of a distance function d to lie in \mathbb{R} . On the one hand, it is convenient both for the underlying mathematics as well as from an implementation perspective. More importantly, however, real numbers – unlike, for example, complex numbers or vectors – come with a natural, total ordering, which is required for downstream operations such as sorting, ranking or clustering.

To summarise, we have established that, on the one hand, we require a proper definition of what is considered to be a valid distance function so as to be able to efficiently plan and execute queries that use them, while on the other hand we must keep that definition open enough to accommodate the many different applications. To address the aforementioned limitations while retaining support for classical, metric distances, we propose the extension of the distance function to the more general notion of a Distance Function Class (DFC) following Definition 6.2.

Definition 6.2 Distance Function Class

A DFC $\hat{d}: \mathcal{D}_f \times \mathcal{D}_f \times \mathcal{D}_1 \dots \times \mathcal{D}_N \rightarrow \mathbb{R}$ is an N-ary but at least binary function $\hat{d}(p, q, s_1, \dots, s_n)$ that outputs a distance $d \in \mathbb{R}$ between a probing argument $p \in \mathcal{D}_f$ and a query argument $q \in \mathcal{D}_f$ using a defined number of *support arguments* $s_k \in \mathbb{N}$ from any of the data domains in \mathbb{D} .

As a convention for notation, the probing argument p of a DFC $\hat{d}(p, q, s_1, \dots, s_n)$ will always come first, followed by the query argument q , again followed by its support arguments s_k in no particular order.

Example 6.3 Examples of Distance Function Classes

A merely illustrative yet widely used example of a DFC would be the family of *Minkowski Distances* $\hat{d}_M: \mathbb{R}^{\text{dim}} \times \mathbb{R}^{\text{dim}} \times \mathbb{R} \rightarrow \mathbb{R}_{\geq 0}$. In this context, support argument m is used to convert the higher-order function to a classical distance function, such as, the Manhattan (L_1) or Euclidean (L_2) distance:

$$\hat{d}_M(\mathbf{p}, \mathbf{q}, m) = \left(\sum_{i=1}^n |p_i - q_i|^m \right)^{\frac{1}{m}} \quad (6.2)$$

$$d_{L1}(\mathbf{p}, \mathbf{q}) = \hat{d}_M(\mathbf{p}, \mathbf{q}, 1) = \sum_{i=1}^n |p_i - q_i| \quad (6.3)$$

$$d_{L2}(\mathbf{p}, \mathbf{q}) = \hat{d}_M(\mathbf{p}, \mathbf{q}, 2) = \sqrt{\sum_{i=1}^n (p_i - q_i)^2} \quad (6.4)$$

However, using the idea of a DFC, we can express many different types of much more complex distance functions in a simple, yet powerful framework. For example, the following distance function is used to obtain a dissimilarity between poses using the joint locations of an 18-keypoints skeleton (more details are presented in [HAG⁺22]):

$$\hat{d}(\mathbf{p}, \mathbf{q}, \mathbf{w}_q, \mathbf{w}_f, s) = \sum_{i=1}^n |(p_i - q_i)| \min(w_{q,i}, w_{f,i}) + 2\pi(s - \sum_{i=1}^n \min(w_{q,i}, w_{f,i})) \quad (6.5)$$

The query and support arguments \mathbf{p} and \mathbf{q} are the actual, normalised keypoints and are accompanied by two weight-vectors \mathbf{w}_q and \mathbf{w}_f , that encode the presence or absence of a joint in a scene. To normalise for missing joints, a normalisation factor based on the weight-sum s is employed.

A DFC (and any function for that matter) is uniquely defined by its *signature*, a $(N + 1)$ -tuple specifying the name of the DFC and the data domains \mathcal{D}_i of the arguments it can accept as defined in Equation (6.6).

$$\text{SIG}(\hat{\mathbf{d}}) = (\text{NAME}, \mathcal{D}_f, \mathcal{D}_f, \mathcal{D}_1, \dots, \mathcal{D}_{N-2}) \quad (6.6)$$

In the context of a concrete execution plan for a proximity-based query, we observe that the query argument effectively remains constant, while the probing argument changes as tuples are being evaluated. The support arguments may or may not be subject to change. For example, for simple NNS, there is a single, constant query vector that is compared against all the vectors in the DBMS, i.e., the probing argument. Hence, irrespective of the arity of the DFC, the function used during query execution comes with an arity of at most $N - 1$. We call this the *implementation* of $\hat{\mathbf{d}}$ as denoted by Equation (6.7).

$$\mathbf{d} = \text{IMP}(\hat{\mathbf{d}}) \ni \text{SIG}(\mathbf{d}) \subset \text{SIG}(\hat{\mathbf{d}}) \quad (6.7)$$

With this idea of a DFC and its (partial) implementations in mind, we can start to reason about its properties in the broader context of database operations in general and proximity-based queries in particular:

DFC and Proximity-Based Queries In extension to Definition 6.1, we realise that any query execution plan that involves the evaluation of a (partial) implementation of a DFC, falls into the category of a proximity-based query.

Variable and Constant Arguments In the context of a proximity-based query, every attribute of a DFC can either be regarded as *constant* or *variable*. That is, as elements in the database are iterated during query execution, they either remain the same or assume new values. This can be used for optimisation.

Purity of Domain Since the purpose of a DFC is to quantify proximity between a probing argument p and a query argument q , we can assume that $p, q \in \mathcal{D}_f$. Furthermore, we require that the query argument is always constant, whereas the probing argument is always variable.

Scalarity of Codomain A DFC and its implementations always produce a single, scalar output $d \in \mathbb{R}$ to quantify the distance. Depending on whether a function generates a distance or a similarity score, a DBMS could use special, annotated types (e.g., DISTANCE or SCORE) to allow for specialised downstream operations.

As an aside, we must address the role of dimensionality \dim in the case of data domains that are subsets of vector spaces, i.e., $\mathcal{D}_f \subset \mathbb{R}^{\dim}$ or $\mathcal{D}_f \subset \mathbb{C}^{\dim}$. One could argue that the dimensionality of such a vector space can also be seen as a parameter of the DFC. Irrespective of the merits such an argument might have, we consider the dimensionality of the vector space to be a strictly structural property of the underlying data domain \mathcal{D}_f and thus a constant in the context of a query.

6.1.2 Extending the Relational Model

Using Definition 6.2, one can now start to integrate DFCs into the relational data model. In line with [Gia18], we first assume \mathbb{ID} – the set system of data domains supported by the database – to be extended by whatever data domain \mathcal{D}_f is required, such as but not limited to \mathbb{R}^{\dim} or \mathbb{C}^{\dim} .

6.1.2.1 Extended Projection and DFCs

We use the idea of an extended projection π_P – as proposed by [GHQ95; GUW09] and introduced in Section 3.2.3 – wherein P denotes a list of projection expressions p that involve attributes $\mathcal{A}_i \in \mathcal{R}$. That is, in addition to the simple projection onto attributes \mathcal{A}_i , an extended projection may also include the evaluation of literals or functions as expressed in Definition 6.3.

Definition 6.3 Extended Projection π_P

Let \mathcal{R} be a N-ary relation over attributes $\mathcal{A} \in \text{SCH}(\mathcal{R})$ with tuples $t \in \mathcal{R}$. We call π_P the *extended projection* with K projection elements $p_i \in P$, i.e., $P = \{p_1, p_2, \dots, p_K\}$, with $p_i \in \text{SCH}(\mathcal{R}) \vee p_i \in \mathcal{D}_{\mathcal{L}} \vee p_i \in \mathcal{F}$, where \mathcal{F} denotes a set of all functions $f_i: \mathcal{D}_1 \times \mathcal{D}_2 \dots \mathcal{D}_M \rightarrow \mathcal{D}_{\text{out}}$ known to the DBMS and $\mathcal{D}_{\mathcal{L}} \in \mathbb{ID}$ denotes a domain of literal values. The extended projection is then defined as follows:

$$\pi_P(\mathcal{R}) = \{t[P] : t \in \mathcal{R}\} \text{ with } t[P] = \bigcup_{p_i \in P} \begin{cases} t[p_i] & \text{if } p_i \in \text{SCH}(\mathcal{R}) \\ p_i & \text{if } p_i \in \mathcal{D}_{\mathcal{L}} \\ p_i(t[P']) & \text{if } p_i \in \mathcal{F} \end{cases}$$

Note, that the domain of every function $f_i \in \mathcal{F}$ is again an implicit, extended projection $t[P']$, where P' contains all projection elements required as function arguments. Hence, function invocations can be nested arbitrarily.

If now we assume any DFC supported by the DBMS to be a member of \mathcal{F} , the evaluation of a DFC as part of a query is given by Definition 6.4 as illustrated in Example 6.4.

Definition 6.4 Distance Function Class in Extended Projection π_P

Let $\hat{\delta}: \mathcal{D}_f \times \mathcal{D}_f \times \mathcal{D}_1 \dots \times \mathcal{D}_{N-2} \rightarrow \mathbb{R}$ be a N-ary DFC. The *extended projection* $\pi_{\delta(p_1, p_2, \dots, p_N)}(\mathcal{R})$ with $p_i \in \text{SCH}(\mathcal{R}) \vee p_i \in \mathcal{D}_{\mathcal{L}} \vee p_i \in \mathcal{F}$ denotes the evaluation of $\hat{\delta}$ using N arguments. Applying $\pi_{\delta(\cdot)}$ on a M-ary relation \mathcal{R} introduces a new distance attribute \mathcal{A}_d , i.e., $\text{SCH}(\pi_{\delta(\cdot), \mathcal{A}_1, \dots, \mathcal{A}_M}(\mathcal{R})) = \text{SCH}(\mathcal{R}) \cup \{\mathcal{A}_d\}$.

Example 6.4 Obtaining a Distance using Extended Projection

The following table lists the schema and extent of $\mathcal{R}_{\text{painting}}$, which now contains a column of feature vectors $\mathcal{A}_{\text{feature}}$ in addition to the scalar attributes, i.e., $\mathcal{D}_{\text{feature}} \subset \mathbb{R}^3$.

$\mathcal{R}_{\text{painting}}$	$\mathcal{A}^*_{\text{title}}$	$\mathcal{A}_{\text{artist}}$	$\mathcal{A}_{\text{year}}$	$\mathcal{A}_{\text{feature}}$
t_1	Mona Lisa	Leonardo da Vinci	1506	[2.0, 7.0, -8.0]
t_2	The Starry Night	Vincent van Gogh	1889	[1.0, 0.0, 3.0]
t_3	Las Meninas	Diego Velázquez	1665	[-1.0, 3.0, 9.0]

The extended projection $\pi_P(\mathcal{R})$ with $P = (\mathcal{A}_{\text{title}}, \delta([0, 0, 0], \mathcal{A}_{\text{feature}}))$, projects onto the attribute $\mathcal{A}_{\text{title}}$ and the Manhattan distance $\delta \in \mathcal{F}$ between query argument $q = [0, 0, 0] \in \mathcal{D}_f \subset \mathbb{R}^3$ and probing argument $\mathcal{A}_{\text{feature}}$. This is the relational algebra equivalent to the following SQL query and leads to $\mathcal{R}_{\text{result}}$.

```
select title, d([0,0,0],feature) from paintings
```

$\mathcal{R}_{\text{result}}$	$\mathcal{A}^*_{\text{title}}$	\mathcal{A}_d
t_1	Mona Lisa	17.0
t_2	The Starry Night	4.0
t_3	Las Meninas	13.0

From Definition 6.4, we can follow that the evaluation of multiple DFCs in an extended projection (e.g., to support aggregate NNS [PTM⁺05]) or the combination of attribute projections with the evaluation of DFCs are also allowed.

With the proposed extended projection, the obtained distance value becomes an explicit attribute \mathcal{A}_d in the relation $\mathcal{R}' = \pi_{\delta(\cdot)}(\mathcal{R})$, which can then be used by downstream relational operators and/or be returned to any calling system by the query. This makes sense, for example, in NNS, where the distance is often converted to a similarity score by some additional correspondence function.

6.1.2.2 Ranked Relations

To support sorting by the obtained distance (or any other attribute), we use a variant of the ranked relation as proposed by [LCI⁺05] and described in Section 3.2.3. A ranked relation \mathcal{R}^O exhibits an ordering of tuples $t \in \mathcal{R}^O$ induced by O . In contrast to [LCI⁺05] and more in line with [GUW09], we assume O to be a sequence of attributes by which the relation should be sorted¹, as specified in Definitions 6.5 and 6.6 and shown in Example 6.5 .

Definition 6.5 Ranked Relation \mathcal{R}^O

Let \mathcal{R} be a N-ary relation over attributes $\mathcal{A} \in \text{SCH}(\mathcal{R})$ with M tuples $t_m \in \mathcal{R}$. We call \mathcal{R}^O a *ranked relation* with K ranking attributes $O = (o_1, o_2, \dots, o_K)$, such that $o_a = (\mathcal{A}_a, \leq_a)$ with $\mathcal{A}_a \in \text{SCH}(\mathcal{R})$ and $\leq_a \in \{\leq_\uparrow, \leq_\downarrow\}$. O induces a non-strict, total ordering of tuples t_m using the binary, lexicographic ranking relation \leq :

$$t_i \leq t_k = \bigvee_{o \in O} t_i \leq_o t_k, \text{ with } \begin{cases} t_i \leq_o t_k \leftrightarrow t_i [\mathcal{A}_o] \leq t_k [\mathcal{A}_o] & \text{if } \leq_o = \leq_\uparrow \\ t_i \leq_o t_k \leftrightarrow t_i [\mathcal{A}_o] \geq t_k [\mathcal{A}_o] & \text{if } \leq_o = \leq_\downarrow \end{cases}$$

It is to be noted, that the ranking O is an inherent property of a ranked relation \mathcal{R}^O as is, for example, its schema $\text{SCH}(\mathcal{R})$. It is retained or induced as operators are being applied to a relation. Furthermore, if a relation does not exhibit a specific ordering, then O is empty and $\mathcal{R}^O = \mathcal{R}^\emptyset = \mathcal{R}$.

Definition 6.6 Order Operation ω_O

Let \mathcal{R} be a N-ary relation over attributes $\mathcal{A} \in \text{SCH}(\mathcal{R})$ and let O be list of ranking attributes with $o[0] \in \text{SCH}(\mathcal{R}) \forall o \in O$. The *relational order operator* imposes order O on relation \mathcal{R} , overriding any previous ordering:

$$\omega_O(\mathcal{R}) = \mathcal{R}^O$$

¹ We argue that the evaluation of a *scoring function*, as proposed by [LCI⁺05], can be expressed in the scope of the extended projection and does not need to be part of the order operation.

Example 6.5 Applying Order using the Order Operation ω_O

The following table lists the schema, extent, and rank of ranked relation $\mathcal{R}_{\text{painting}}^\emptyset$, which now contains an attribute $\mathcal{A}_{\text{distance}}$ (e.g., obtained by executing a proximity-based query) and does not exhibit an explicit ordering.

$\mathcal{R}_{\text{painting}}^\emptyset$	$\mathcal{A}^*_{\text{title}}$	$\overline{\mathcal{A}}_{\text{artist}}$	$\mathcal{A}_{\text{year}}$	$\mathcal{A}_{\text{distance}}$
t_1	Mona Lisa	Leonardo da Vinci	1506	17.0
t_2	The Starry Night	Vincent van Gogh	1889	4.0
t_3	Las Meninas	Diego Velázquez	1665	13.0
t_4	Mars and Venus	Sandro Botticelli	1483	17.0

The order operation $\omega_O(\mathcal{R}_{\text{painting}})$ with $O = ((\mathcal{A}_{\text{distance}}, \leq_\downarrow), (\mathcal{A}_{\text{title}}, \leq_\uparrow))$ sorts by $\mathcal{A}_{\text{distance}}$ in descending and $\mathcal{A}_{\text{title}}$ in ascending order. This is the relational algebra equivalent of the following SQL query and leads to $\mathcal{R}_{\text{result}}$.

```
select * from paintings
order by distance desc, title asc
```

$\mathcal{R}_{\text{result}}^O$	$\mathcal{A}^*_{\text{title}}$	$\overline{\mathcal{A}}_{\text{artist}}$	$\mathcal{A}_{\text{year}}$	$\mathcal{A}_{\text{distance}}$
t_1	Mars and Venus	Sandro Botticelli	1483	17.0
t_2	Mona Lisa	Leonardo da Vinci	1506	17.0
t_3	Las Meninas	Diego Velázquez	1665	13.0
t_4	The Starry Night	Vincent van Gogh	1889	4.0

On a practical note, it must be mentioned, that until a query is evaluated, the ordering induced by O must not necessarily be materialised. To determine the position of a tuple within a relation upon materialisation, we assume the existence of a ranking function.

$$\text{RANK}_O : \mathcal{R} \rightarrow [1, M] \text{ with } M = |\mathcal{R}| \quad (6.8)$$

The ranking function assigns a position $m \in [1, M]$ to a tuple based on O , i.e., the first tuple t_1 gets position 1, second tuple t_2 position 2 etc. Unordered relations also exhibit an internal ordering of tuples – the *natural ordering* of \mathcal{R} – which is not necessarily related to any observable property. Consequently, RANK is also defined for such relations, even though positions may appear arbitrary.

6.1.2.3 s,k-selection

To be able to execute a specific type of proximity-based query – namely, kNN or kFN – we require one last extension to the relational algebra in the form of the *s,k-selection* operator $\lambda_{s,k}(\mathcal{R})$, which skips s tuples and limits the cardinality to k tuples according to Definition 6.7 and as demonstrated in Example 6.6.

Definition 6.7 s,k-selection Operation $\lambda_{s,k}$

Let \mathcal{R} be a N-ary relation over attributes $\mathcal{A} \in \text{SCH}(\mathcal{R})$ and let $s, k \in \mathbb{N}$ and $s < k$. The s,k-selection operation $\lambda_{s,k}$ selects tuples based on ranking as follows:

$$\lambda_{s,k}(\mathcal{R}^O) = \{t \in \mathcal{R}^O : s \leq \text{RANK}_O(t) \leq k\}$$

The s,k-selection operation is inherently sensitive to ranking since it selects tuples based on it. Furthermore, since oftentimes $s = 0$, we define a shortcut $\lambda_k = \lambda_{0,k}$ to express a simple selection to k entries without any skipping.

Example 6.6 Selection Based on Ranking using $\lambda_{s,k}$

The following table lists schema, extent, and ranking of ranked relation $\mathcal{R}_{\text{painting}}^\emptyset$, which does not exhibit an explicit ordering.

$\mathcal{R}_{\text{painting}}^\emptyset$	$\mathcal{A}^*_{\text{title}}$	$\overline{\mathcal{A}}_{\text{artist}}$	$\mathcal{A}_{\text{year}}$
t_1	Mars and Venus	Sandro Botticelli	1483
t_2	Mona Lisa	Leonardo da Vinci	1506
t_3	The Starry Night	Vincent van Gogh	1889
t_4	Las Meninas	Diego Velázquez	1665

$\lambda_{0,3}(\mathcal{R}_{\text{painting}})$ limits the result to the first three entries. It is the relational algebra equivalent of the following SQL query and produces $\mathcal{R}_{\text{result}}$.

```
select * from paintings limit 3
```

$\mathcal{R}_{\text{result}}^\emptyset$	$\mathcal{A}^*_{\text{title}}$	$\overline{\mathcal{A}}_{\text{artist}}$	$\mathcal{A}_{\text{year}}$
t_1	Mars and Venus	Sandro Botticelli	1483
t_2	Mona Lisa	Leonardo da Vinci	1506
t_3	The Starry Night	Vincent van Gogh	1889

6.1.2.4 Combining $\pi_P, \omega_O, \lambda_k$

The proposed operations can now be freely combined to form proximity-based queries. In Example 6.7 we demonstrate this for kNN and range search. A list of additional types of queries is provided in Table 6.1.

Example 6.7 kNN and ϵ NN using Extended Relational Algebra

The following table lists the schema, extent and rank of $\mathcal{R}_{\text{painting}}^0$, with their title $\mathcal{A}_{\text{title}}$ and a feature $\mathcal{A}_{\text{feature}}$. Let further $P = (\mathcal{A}_{\text{title}}, \mathfrak{d}([0,0,0], \mathcal{A}_{\text{feature}}))$ be a list of projections, $\mathfrak{d} \in \mathcal{F}$ be the Manhattan distance between query argument $q = [0,0,0] \in \mathcal{D}_f \subset \mathbb{R}^3$ and probing argument $\mathcal{A}_{\text{feature}}$.

$\mathcal{R}_{\text{(p)ainting}}^0$	$\mathcal{A}^*_{\text{title}}$	$\overline{\mathcal{A}}_{\text{artist}}$	$\mathcal{A}_{\text{year}}$	$\mathcal{A}_{\text{feature}}$
t_1	Mona Lisa	Leonardo da Vinci	1506	[2.0, 7.0, -8.0]
t_2	The Starry Night	Vincent van Gogh	1889	[1.0, 0.0, 3.0]
t_3	Las Meninas	Diego Velázquez	1665	[-1.0, 3.0, 9.0]

kNN-search can be expressed as $\lambda_k(\omega_{\mathcal{A}_d \uparrow}(\pi_P(\mathcal{R})))$ with $k = 1$. This is the relational algebra equivalent of the following SQL query and leads to $\mathcal{R}_{\text{result}}$.

```
select title, d(feature, [0.0,0.0,0.0]) as distance
from paintings order by distance limit 1
```

Range search (or ϵ NN) can be expressed as $\sigma_{\mathcal{A}_d > \epsilon}(\pi_P(\mathcal{R}))$ with and $\epsilon = 15.0$. This is the relational algebra equivalent of the following SQL query and also leads to $\mathcal{R}_{\text{result}}$.

```
select title, d(feature, [0.0,0.0,0.0]) as distance
from paintings where distance > 15.0 order by distance
```

$\mathcal{R}_{\text{result}}$	$\mathcal{A}^*_{\text{title}}$	$\mathcal{A}_{\text{distance}}$
t_1	Mona Lisa	17.0

Even though both queries generate the same result for the purpose of this example, it is worth pointing out that they do so by applying very different semantics – that of selection by cardinality on the one hand and that of selecting by predicate on the other.

Table 6.1 Proximity-based queries supported by the extended relational algebra.

Name	Result	Algebraic Form
NNS / kNN	$\mathcal{R}^{\mathcal{A}_d \uparrow}$	$\lambda_k(\omega_{\mathcal{A}_d \uparrow}(\pi_{\mathcal{A}_y, \mathbf{d}(\mathcal{A}_p)}(\mathcal{R}_p)))$
FNS / kFN	$\mathcal{R}^{\mathcal{A}_d \downarrow}$	$\lambda_k(\omega_{\mathcal{A}_d \downarrow}(\pi_{\mathcal{A}_y, \mathbf{d}(\mathcal{A}_p)}(\mathcal{R}_p)))$
ϵ NN / Range Search	$\mathcal{R}^{\mathcal{A}_d \uparrow}$	$\omega_{\mathcal{A}_d \uparrow}(\sigma_{\mathcal{A}_d \leq \epsilon}(\pi_{\mathcal{A}_y, \mathbf{d}(\mathcal{A}_p)}(\mathcal{R}_p)))$
NNS w. Selection	$\mathcal{R}^{\mathcal{A}_d \uparrow}$	$\lambda_k(\omega_{\mathcal{A}_d \uparrow}(\pi_{\mathcal{A}_p, \mathbf{d}(\mathcal{A}_p)}(\sigma_{\mathcal{A}_y=1889}(\mathcal{R}_p))))$
Group NNS and Aggregate NNS [PST ⁺ 04; PTM ⁺ 05]	$\mathcal{R}^{\mathcal{A}_d \uparrow}$	$\lambda_k(\omega_{\mathcal{A}_d \uparrow}(\pi_{f(\mathbf{d}_1(\mathcal{A}_p), \dots, \mathbf{d}_n(\mathcal{A}_p))}(\mathcal{R}_p)))$
Batch NNS [SHS ⁺ 08]	\mathcal{R}^\emptyset	$\lambda_k(\omega_{\mathcal{A}_d \uparrow}(\pi_{\mathbf{d}_1(\mathcal{A}_p)}(\mathcal{R}_p))) \cup \lambda_k(\omega_{\mathcal{A}_d \uparrow}(\pi_{\mathbf{d}_2(\mathcal{A}_p)}(\mathcal{R}_p)))$
Multi Order	$\mathcal{R}^{\mathcal{A}_{d1} \uparrow, \mathcal{A}_{d2} \uparrow}$	$\omega_{\mathcal{A}_{d1} \uparrow, \mathcal{A}_{d2} \uparrow}(\pi_{\mathbf{d}_1(\mathcal{A}_p), \mathbf{d}_2(\mathcal{A}_p)}(\mathcal{R}_p))$
Aggregation by Distance	\mathcal{R}^\emptyset	$\gamma_{\text{MEAN}}(\pi_{\mathbf{d}(\mathcal{A}_p)}(\mathcal{R}_p))$

6.1.2.5 Algebraic Properties of $\pi_P, \omega_O, \lambda_k$

Since a query optimiser of a multimedia DBMS must be able to manipulate the operators that were introduced before, there is a series of algebraic properties that must be taken into account. For unary operators we must consider *commutativity*, i.e., whether changing the order of application leads to the same result. For binary operators, we must consider the *commutativity of operands* and *distributivity* with respect to other operations. The identities are provided in Equation (6.9).

$$\begin{aligned}
 \text{OP}_1(\text{OP}_2(\mathcal{R})) &= \text{OP}_2(\text{OP}_1(\mathcal{R})) && \text{(commutativity of operators)} \\
 \text{OP}_1(\text{OP}_2(\mathcal{R}_L, \mathcal{R}_R)) &= \text{OP}_2(\text{OP}_1(\mathcal{R}_L), \text{OP}_1(\mathcal{R}_R)) && \text{(distributivity of operators)} \\
 \text{OP}(\mathcal{R}_L, \mathcal{R}_R) &= \text{OP}(\mathcal{R}_R, \mathcal{R}_L) && \text{(commutativity of operands)} \\
 \text{OP}(\text{OP}(\mathcal{R})) &= \text{OP}(\mathcal{R}) && \text{(idempotence of operators)}
 \end{aligned} \tag{6.9}$$

In addition, to account for ranked relations, we must distinguish between relational operators that are either *order retaining* (order of the input remains the same) or *order inducing* (order of the input is changed). Furthermore, we call operators that do not commute with order inducing operators *order sensitive*.

It follows directly from its definition, that ω_O is an order inducing operator. Furthermore, it commutes with other unary operators if they are neither *order sensitive* nor *order inducing*, namely, π , σ and ρ . Importantly, ω_O does not commute with its own kind unless the operators are identical (idempotence)!

Furthermore, σ , π , ρ , and λ are *order retaining*, since they do not alter the input relation with respect to ordering. It is a bit more complicated for the

Table 6.2 Relational operators and their algebraic properties.

Name	Symbol	Order	Com.	Com. Op.	Dist.
Union	$\mathcal{R}_L^{O_1} \cup \mathcal{R}_R^{O_2}$	induces \emptyset	-	✓	-
Cartesian Product	$\mathcal{R}_L^{O_1} \times \mathcal{R}_R^{O_2}$	induces \emptyset	-	✗	-
Intersection	$\mathcal{R}_L^{O_1} \cap \mathcal{R}_R^{O_2}$	induces O_1	-	✗	-
Difference	$\mathcal{R}_L^{O_1} \setminus \mathcal{R}_R^{O_2}$	induces O_1	-	✗	-
Natural Join	$\mathcal{R}_L^{O_1} \bowtie_P \mathcal{R}_R^{O_2}$	induces O_1	-	✗	-
Rename	$\rho_{\mathcal{A}_A \rightarrow \mathcal{A}_B}(\mathcal{R}^O)$	retains O	$\pi, \sigma, \rho, \lambda, \omega$	-	$\cup, \cap, \setminus, \times, \bowtie$
Selection	$\sigma_S(\mathcal{R}^O)$	retains O	$\pi, \sigma, \rho, \omega$	-	$\cup, \cap, \setminus, \times, \bowtie$
Extended Projection	$\pi_P(\mathcal{R}^O)$	retains O	$\pi, \sigma, \rho, \lambda, \omega$	-	\cup
Limit	$\lambda_{s,k}(\mathcal{R}^O)$	retains O	π, ρ	-	✗
Order	$\omega_O(\mathcal{R}^{O_1})$	induces O_2	π, σ, ρ	-	\cap, \setminus, \bowtie

binary operators, which do not behave consistently: On the one hand, the set difference (\setminus), the set intersection (\cap) and the natural join (\bowtie) retain the order of the left operand \mathcal{R}_L , since the operators select elements from \mathcal{R}_L based on their presence in \mathcal{R}_R . As a consequence, the operands do not commute, since it is always the left operand that defines the output ordering. The set union (\cup) and the cross product (\times), on the other hand, combine two relations \mathcal{R}_L and \mathcal{R}_R . The resulting relation does not have an explicit order that can be expressed within the proposed framework and must thus be considered unordered, i.e., the two operators are order inducing in that they establish a new natural ordering. In addition, distributivity with all the order retaining binary operations is given as expressed by Equation (6.10).

$$\omega_O(\mathcal{R}_L^{O_L} \cap \mathcal{R}_R^{O_R}) = \omega_O(\mathcal{R}_L^{O_L}) \cap \omega_O(\mathcal{R}_R^{O_R}) = (\mathcal{R}_L \cap \mathcal{R}_R)^{O_L} \quad (6.10)$$

$$\omega_O(\mathcal{R}_L^{O_L} \setminus \mathcal{R}_R^{O_R}) = \omega_O(\mathcal{R}_L^{O_L}) \setminus \omega_O(\mathcal{R}_R^{O_R}) = (\mathcal{R}_L \setminus \mathcal{R}_R)^{O_L} \quad (6.11)$$

$$\omega_O(\mathcal{R}_L^{O_L} \bowtie_P \mathcal{R}_R^{O_R}) = \omega_O(\mathcal{R}_L^{O_L}) \bowtie_P \omega_O(\mathcal{R}_R^{O_R}) = (\mathcal{R}_L \bowtie_P \mathcal{R}_R)^{O_L} \quad (6.12)$$

However, the order operator does not distribute with the operators \cup and \times .

$$\omega_O(\mathcal{R}_L^{O_L} \cup \mathcal{R}_R^{O_R}) = (\mathcal{R}_L^{O_L} \cup \mathcal{R}_R^{O_R})^O \stackrel{!}{\neq} \omega_O(\mathcal{R}_L^{O_L}) \cup \omega_O(\mathcal{R}_R^{O_R}) = (\mathcal{R}_L \cup \mathcal{R}_R)^O$$

$$\omega_O(\mathcal{R}_L^{O_L} \times \mathcal{R}_R^{O_R}) = (\mathcal{R}_L^{O_L} \times \mathcal{R}_R^{O_R})^O \stackrel{!}{\neq} \omega_O(\mathcal{R}_L^{O_L}) \times \omega_O(\mathcal{R}_R^{O_R}) = (\mathcal{R}_L \times \mathcal{R}_R)^O$$

The extended projection π_P inherits all its properties from the normal projection π if we require $f \in \mathcal{F}$ to be stateless and side-effect free, i.e., it must not influence the state of the DBMS in any other way than producing the desired output and it must not maintain an internal state itself. Irrespective of π_P , this

is a reasonable assumption. However, it is important to note that operators may rely on attributes introduced by a projection π_P and therefore commutativity is limited to cases, where no such interdependence exists.

Finally, the s,k-selection operator λ is also idempotent if $s = 0$ and it commutes with most ρ and π . However, it does neither commute with the selection operator σ nor the order operator ω .

$$\lambda(\omega(\mathcal{R})) \neq \omega(\lambda(\mathcal{R}))$$

$$\lambda(\sigma(\mathcal{R})) \neq \sigma(\lambda(\mathcal{R}))$$

Lack of commutativity of λ with ω follows directly from its definition: λ is sensitive to the ranking of \mathcal{R}^O , and ω influences that ranking. Consequently, executing λ before or after imposing a certain order will yield different results. For commutativity with σ , one can turn to Proof 6.1. A summary of all properties can be found in Table 6.2.

Proof 6.1 Non-commutativity of σ and λ

Let \mathcal{R} be a unary relation with $M = |\mathcal{R}|$ and $SCH(\mathcal{R}) = \{\mathcal{A}_R\}$. Furthermore, let $t[\mathcal{A}_R] = \text{RANK}(t) \forall t \in \mathcal{R}$, i.e., the attribute holds the rank of the tuple according to the natural ordering as a fixed value. If σ and λ_k were to commute, then the following identity would hold for any $x < M \in \mathbb{N}_{>0}$:

$$\sigma_{\mathcal{A}_R > x}(\lambda_x(\mathcal{R})) = \lambda_x(\sigma_{\mathcal{A}_R > x}(\mathcal{R}))$$

However, we can proof by contradiction that this does not hold since:

$$\sigma_{\mathcal{A}_R > x}(\lambda_x(\mathcal{R})) = \emptyset \neq \lambda_x(\sigma_{\mathcal{A}_R > x}(\mathcal{R})) = \{t_{x+1}, \dots, t_{\min(x+x, M)}\}$$

That is, whenever λ_k removes tuples that would match the predicate of σ , the order of execution influences the result and the two operators do not commute. An example for $M = 6$ and $x = 2$ is given below:

\mathcal{R}	\mathcal{A}_R^*	$\lambda_x(\mathcal{R})$	$\sigma_{\mathcal{A}_R > x}(\mathcal{R})$	$\sigma_{\mathcal{A}_R > x}(\lambda_x(\mathcal{R}))$	$\lambda_x(\sigma_{\mathcal{A}_R > x}(\mathcal{R}))$
t_1	1	✓	✗	✗	✗
t_2	2	✓	✗	✗	✗
t_3	3	✗	✓	✗	✓
t_4	4	✗	✓	✗	✓
t_5	5	✗	✓	✗	✗
t_6	6	✗	✓	✗	✗

6.1.3 DFCs and Query Optimisation

In addition to the algebraic properties of the extended relational algebra, we can also leverage the requirements imposed on the structure of a function $f \in \mathcal{F}$ to enable a DBMS to plan and optimise the execution of a proximity-based query by altering the execution plan in terms of DFC implementations that should be employed. To do so, we must make a few assumptions as to how queries are being executed:

1. We assume that a DBMS maintains the collection of functions \mathcal{F} . Since every function can be identified by its signature $SIG(f)$, \mathcal{F} can be used to obtain the correct function implementation given the specification of the query. We refer to this data structure as *function registry*, which can also be used to store additional properties and metadata.
2. We assume that the execution of a (proximity-based) query takes place in a *query context*, which remains constant as the query is being evaluated, with the exception of the tuples that are being loaded and pushed into the operator tree.
3. We require all $f \in \mathcal{F}$ to be stateless and side-effect free with respect to that query context.

Assuming these properties, we can distinguish between functions that remain *constant* within the scope of the query context – which can be converted to a literal value at parsing time – and such that do not. This very basic property allows for early evaluation and caching of function values. Since the extended projection also allows for the nesting of functions, such an optimisation can be relevant. In addition, we propose the notion of a *function hierarchy* depicted in Figure 6.1. As we move down that hierarchy, more information about the structure of a function employed in a query plan becomes available. That information can be leveraged to choose the best execution path:

General Function Every function $f \in \mathcal{F}$ is registered in the catalogue with information about its name and domain (i.e., its signature $SIG(f)$). This enables the DBMS to provide optimised implementations or replacements for composit operations.

Distance Function Class Since DFCs constitute a specific type of function, we can embed equivalences between a DFC and different types of index struc-

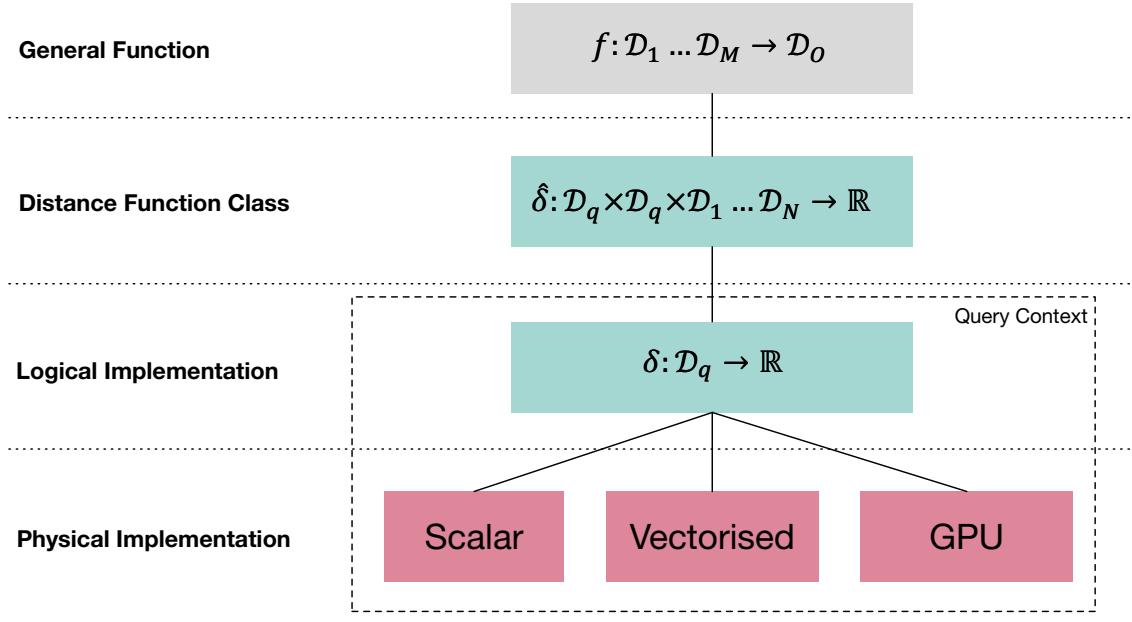


Figure 6.1 Function hierarchy of all $f \in \mathcal{F}$. As we go down the hierarchy, the DBMS gains knowledge about a function's structure.

tures directly within the function registry. Those equivalences allow a query planner to select one over the other.

Logical Implementation As an implementation of a query manifests, the structure of a DFC may change through implementation $\mathfrak{d} = \text{IMP}(\hat{\delta})$. For example, certain attributes may turn out to remain constant, which can be leveraged by the planner to optimise execution.

Physical Implementation As the physical execution plan emerges, the planner can make decisions as to what type of implementation for a DFC should be selected. For example, it may select a vectorised version, which leverages SIMD instructions of the CPU over a scalar implementation that calculates the distance in a tight loop.

Optimising implementations of DFCs and general function invocations can be achieved in multiple ways. As a first step, it is reasonable to provide and select implementations that are specific for the data type of the function arguments, in order to avoid type conversion in a tight loop over a large amount of data. This is especially true for DFCs that operate on complex datatypes, such as vectors and goes into the direction of Requirement 2.5.

Another aspect that can be optimised is the use of vectorised implementations for computations that involve vectors or matrices, i.e., exchanging a scalar implementation of a DFC by an implementation that uses SIMD instructions.

This can be beneficial both for a multimedia DBMS that employs the iterator model and even more so for a batched processing model. However, for the iterator model it is to be expected that one challenge lies in finding the threshold at which usage of such an optimised version pays off. If we take real-valued vectors as an example, it is likely that vectorised execution will not benefit low-dimensional vectors while it will greatly benefit high-dimensional vectors.

Finally, one can try to match constellations of functions and replace them with an optimised version or even inline them, similarly to HyPer [BGL⁺22]. A simple example could be the expression $a * b + c$, i.e., $\text{add}(\text{mul}(a, b), c)$, which can be replaced by a single *fused multiply-add* function $\text{fma}(a, b, c)$ to avoid loss of precision. Replacing the evaluation of a DFC with a lookup in a high-dimensional index also falls into this category, however, this transformation depends on more than just the function itself.

6.1.4 DFCs and High-Dimensional Index Structures

One identity that is particularly relevant for the execution of proximity-based queries is the relationship between the execution of a DFC and the use of a high-dimensional index. As we have discussed in Section 4.2.2, there exist many different types of indexing techniques that can be used to speed up NNS in different ways. We propose to formalise the use of high-dimensional indexes² by defining equivalence classes between an index and the relational algebra expressions it replaces. For this, we rely on the index definitions presented in Section 3.3.5. In addition to Definition 3.12, the following restrictions may apply in addition to implementation-specific restrictions for a specific high-dimensional index: (i) High-dimensional indexes are always derived from a feature attribute $\mathcal{A}_f \in \text{SCH}(\mathcal{R})$. Therefore, an index can only replace a DFC if they operate on an unaltered version of \mathcal{A}_f . Changes to \mathcal{A}_f , e.g., through application of a nested function $\text{fun}(\mathcal{A}_f)$, will render the index unusable unless the same alteration was applied when creating the index. (ii) High-dimensional indexes are sometimes trained for a specific distance function (e.g., Euclidean) or class of distance functions (e.g., Minkowski). Therefore, an index can only replace an expression if the DFC used within coincides with the original function(-class). Using these constraints, we can now formalise the replacement of a relational expression by an $\mathcal{I}_I^{\mathcal{R}}$ which, for high-dimensional indexes and depending on what expression is matched, falls into one of three classes given by Definitions 6.9, 6.10 or 6.11.

² By “index” we refer to a concrete implementation rather than a family since all techniques introduced in Section 4.2.2 can be employed in different ways.

Definition 6.8 Exact Index Replacement

Let \mathcal{R} be a relation, \mathcal{P} a relational (sub-)expression (i.e., an execution plan) involving \mathcal{R} , and $\mathcal{I}_I^{\mathcal{R}}$ an index on \mathcal{R} . We call the replacement of \mathcal{P} with $\mathcal{I}_I^{\mathcal{R}}$ an *exact index replacement* if $\mathcal{P} \equiv \mathcal{I}_I^{\mathcal{R}}$.

Definition 6.9 High-Dimensional Index Replacement of Class 1

Let \mathcal{R} be a relation, $\mathcal{A}_f \in \text{SCH}(\mathcal{R})$ an attribute of \mathcal{R} , $\hat{\mathbf{d}}$ a DFC, $q \in \mathcal{D}_f$ a query and π_P the extended projection involving execution of that DFC using the query, i.e., $\mathbf{d} = \hat{\mathbf{d}}(\mathcal{A}_f, q, \dots) \in P$. We call a replacement with a high-dimensional index $\mathcal{I}_{\mathcal{A}_f, \hat{\mathbf{d}}}^{\mathcal{R}} \equiv \pi_{\hat{\mathbf{d}}}(\mathcal{R})$ *class 1*, if:

$$\pi_P(\mathcal{R}) = \pi_{\hat{\mathbf{d}}}(\mathcal{R}) \bowtie_{\text{TID}} \pi_{P \setminus \hat{\mathbf{d}}}(\mathcal{R}) = \mathcal{I}_{\mathcal{A}_f, \hat{\mathbf{d}}}^{\mathcal{R}} \bowtie_{\text{TID}} \pi_{P \setminus \hat{\mathbf{d}}}(\mathcal{R})$$

Class 1 index replacements constitute the simplest and most versatile case. The index acts as a drop-in replacement for the distance calculation. One high-dimensional index structure that can provide this functionality is the simplest variant of the PQ index, namely, the exhaustive SDC and ADC algorithms described in [JDS10]. However, by its nature PQ does not constitute an exact replacement.

Definition 6.10 High-Dimensional Index Replacement of Class 2

Let \mathcal{R} be a relation, $\mathcal{A}_f \in \text{SCH}(\mathcal{R})$ an attribute of \mathcal{R} , $\hat{\mathbf{d}}$ a DFC, $q \in \mathcal{D}_f$ a query, and π_P the extended projection involving execution of that DFC using the query, i.e., $\mathbf{d} = \hat{\mathbf{d}}(\mathcal{A}_f, q, \dots) \in P$ and $\omega_{\mathcal{A}_{d\downarrow}}$ be an order operation involving the derived distance attribute \mathcal{A}_d . We call a replacement with a high-dimensional index $\mathcal{I}_{\mathcal{A}_f, \hat{\mathbf{d}}}^{\mathcal{R}} \equiv \omega_{\mathcal{A}_{d\downarrow}}(\pi_{\hat{\mathbf{d}}}(\mathcal{R}))$ *class 2* if:

$$\omega_{\mathcal{A}_{d\downarrow}}(\pi_P(\mathcal{R})) = \omega_{\mathcal{A}_{d\downarrow}}(\pi_{\hat{\mathbf{d}}}(\mathcal{R})) \bowtie_{\text{TID}} \pi_{P \setminus \hat{\mathbf{d}}}(\mathcal{R}) = \mathcal{I}_{\mathcal{A}_f, \hat{\mathbf{d}}}^{\mathcal{R}} \bowtie_{\text{TID}} \pi_{P \setminus \hat{\mathbf{d}}}(\mathcal{R})$$

Class 2 index replacements return a ranked relation that comes pre-sorted according to the derived distance in addition to obtaining the actual distance value. This can be convenient for certain applications, but it is also a rare trait and we do not have a concrete example of a high-dimensional index that provides this property. However, the caching of (sub-)query results could lead to such a replacement.

Definition 6.11 High-Dimensional Index Replacement of Class 3

Let \mathcal{R} be a relation, $\mathcal{A}_f \in \text{SCH}(\mathcal{R})$ an attribute of \mathcal{R} , $\hat{\mathbf{d}}$ a DFC, $q \in \mathcal{D}_f$ a query, and π_P the extended projection involving execution of that DFC using the query, i.e., $\mathbf{d} = \hat{\mathbf{d}}(\mathcal{A}_f, q, \dots) \in P$, $\omega_{\mathcal{A}_{d\uparrow}}$ an order operation involving the derived distance attribute \mathcal{A}_d , λ_k a k -selection and σ_S a selection involving a comparison of the distance attribute \mathcal{A}_d . We call a replacement with a high-dimensional index $\mathcal{I}_{\mathcal{A}_f, \hat{\mathbf{d}}}^{\mathcal{R}} \equiv \text{OP}(\omega_{\mathcal{A}_{d\uparrow}}(\pi_{\mathbf{d}}(\mathcal{R})))$ with $\text{OP} \in \{\lambda_k, \sigma_S\}$ class 3 if:

$$\text{OP}(\omega_{\mathcal{A}_{d\uparrow}}(\pi_P(\mathcal{R}))) = \text{OP}(\omega_{\mathcal{A}_{d\uparrow}}(\pi_{\hat{\mathbf{d}}}(\mathcal{R}))) \bowtie_{\text{TID}} \pi_{P \setminus \hat{\mathbf{d}}}(\mathcal{R}) = \mathcal{I}_{\mathcal{A}_f, \hat{\mathbf{d}}}^{\mathcal{R}} \bowtie_{\text{TID}} \pi_{P \setminus \hat{\mathbf{d}}}(\mathcal{R})$$

Class 3 replacements constitute an optimisation of a specific edge-case, which is that of kNN, kFN, and or ϵ NN search. It is probably the most common since it allows for most optimisation but is also the least versatile type of index replacement since it merely supports one specific type of search. An example of such an index replacement is provided in Example 6.8

Example 6.8 Index Replacement for kNN Search

The following table lists the schema, extent, and rank of $\mathcal{R}_{\text{painting}}^{\emptyset}$, with their title $\mathcal{A}_{(t)\text{title}}$, the year of their creation $\mathcal{A}_{(p)\text{ainted}}$, and some arbitrary feature vector $\mathcal{A}_{(f)\text{eature}}$. Let further $\hat{\mathbf{d}}$ be the Manhattan distance.

$\mathcal{R}_{\text{painting}}^{\emptyset}$	$\mathcal{A}^*(t)\text{itle}$	$\overline{\mathcal{A}}_{(a)\text{rtist}}$	$\mathcal{A}_{(p)\text{ainted}}$	$\mathcal{A}_{(f)\text{eature}}$
t_1	Mona Lisa	Leonardo da Vinci	1506	[2.0, 7.0, -8.0]
t_2	The Starry Night	Vincent van Gogh	1889	[1.0, 0.0, 3.0]
t_3	Las Meninas	Diego Velázquez	1665	[-1.0, 3.0, 9.0]
...
t_N	Mars and Venus	Sandro Botticelli	1483	[-3.0, 1.0, 0.0]

The result of the kNN-search can be produced by a class 3 index replacement with $\mathcal{I}_{\mathcal{A}_f, \hat{\mathbf{d}}}^{\mathcal{R}_p}$.

$$\mathcal{R}_{\text{result}}^{\mathcal{A}_d \uparrow} = \lambda_k(\omega_{\mathcal{A}_d \uparrow}(\pi_{\mathcal{A}_t, \hat{\mathbf{d}}(\mathcal{A}_f)}(\mathcal{R}_p))) = \mathcal{I}_{\mathcal{A}_f, \hat{\mathbf{d}}}^{\mathcal{R}_p} \bowtie_{\text{TID}} \pi_{\mathcal{A}_t}(\mathcal{R}_p)$$

An index structure that could be employed here is the VAF index [WSB98].

With the aforementioned definitions and restrictions, we have defined the basic relationship and rules that can be applied by a multimedia DBMS, when deciding whether or not to use a high-dimensional index structure. However, the definitions so far consider the replacement an exact transformation, i.e., the result of the original expression and the index is supposed to be identical. For approximate index structures, this is often not the case, which leads to Definition 6.12.

Definition 6.12 Approximate Index Replacement

Let \mathcal{R} be a relation, \mathcal{P} a relational (sub-)expression involving \mathcal{R} and $\mathcal{I}_I^{\mathcal{R}}$ an index on \mathcal{R} . We call the replacement of \mathcal{P} with $\mathcal{I}_I^{\mathcal{R}}$ an *approximate index replacement* if $\mathcal{P} \approx_b \mathcal{I}_I^{\mathcal{R}}$ with:

$$\mathcal{P} \approx_b \mathcal{I}_I^{\mathcal{R}} \leftrightarrow Q(\mathcal{P}, \mathcal{I}_I^{\mathcal{R}}) \geq b, b \in [0, 1]$$

Q is a *quality estimator* that derives a score between zero and one by comparing the results of \mathcal{P} and \mathcal{I} .

In a manner of speaking, \mathcal{P} acts as a ground truth and is compared to the approximate results produced by $\mathcal{I}_I^{\mathcal{R}}$. The choice of quality metric Q is implementation-specific. One could, for example, pick recall if one is primarily interested in set-based overlap rather than the ranking of individual items, which is enough for unranked query plans:

$$Q_1(\mathcal{P}, \mathcal{I}) = \frac{|\mathcal{P} \cap \mathcal{I}|}{|\mathcal{P}|}$$

An alternative could be the use of a normalised Discounted Cumulative Gain (DCG) [JK02], if the ranking of items should be taken into account as well:

$$Q_2(\mathcal{P}, \mathcal{I}) = \frac{\text{DCG}(\mathcal{P}, \mathcal{I})}{\text{iDCG}(\mathcal{P})}$$

Regardless of the choice of metric, the challenge lies in predicting its value for a query prior to its execution. Since the result of Q can only be produced in hindsight once a query has been executed, we need techniques to anticipate its value for an incoming query so that the DBMS can use this downstream to make decision about query execution.

6.2 Cost Model for Multimedia Databases

As we have explained in Chapter 3, most DBMS rely on a cost model for planning and selecting the execution and access path of a user-specified query. In fact, ever since the System R paper [SAC⁺79], cost-based query planning and optimisation have been considered a gold standard for database systems [MCS88]. Most traditional DBMS exhibit cost models that mainly rely on the cost incurred by accessing (reading / writing) database pages from or to disk [MCS88; GUW09; Pet19]. Based on what has been discussed thus far, we argue that this model must be adapted to be useful for multimedia databases so that it takes the following aspects into account when choosing a query execution plan:

Disk Access (IO) Persistent storage and access to information residing on disk is the factor that contributes the most to long query execution times and is thus a factor to minimise [SAC⁺79]. This is also true for proximity-based queries, where part of the optimisation lies in reducing access to vectors through compression and/or early filtering, e.g., [WSB98; CPR⁺07].

CPU Due to the computational complexity of proximity-based operations, especially when the evaluation of functions on high-dimensional vectors are involved, the processing time on CPU is a factor that can no longer be ignored and thus must be taken into account as well. In fact, several index structures achieve speed-up by reducing the computational complexity of the distance calculation [JDS10] or by avoiding it altogether [WSB98].

Memory Some algorithms can greatly benefit from pure in-memory processing. While memory used to be a scarce resource, some environments allow for complete in-memory processing even for very large datasets. This is exploited by systems like Milvus [WYG⁺21], that load entire data collections into memory.

Quality The quality of a result incurred, e.g., by the choice of a certain high-dimensional index structure [IM98; JDS10], is a price that can be paid to attain speed-up. This is common practice in approximate NNS and should be made transparent to the system and user for reasons explained in Requirement 2.4, especially since this trade-off is not always viable (see Section 2.3).

These four factors can be used to characterise query workloads in a purely local setup consisting of a single node. For distributed databases, the cost incurred by data and message exchange must obviously be considered as well.

While this is not within the scope of the work presented in this Thesis, the proposed model can be extended to accommodate those (and other) types of costs. The aforementioned factors lead us to Definition 6.13 and Definition 6.14 for the cost incurred by an operator OP and query execution plan \mathcal{P} .

Definition 6.13 Cost of a Relational Operator OP

Let OP be a (relational) operator. We call the 4-tuple $C_{\text{OP}} = (c_{\text{IO}}, c_{\text{CPU}}, c_{\text{MEM}}, c_{\text{Q}})$ the *cost* of OP , which captures the atomic costs in terms of disk access (IO), CPU usage (CPU), memory usage (MEM) and quality (Q).

Definition 6.14 Cost of a Query Execution Plan \mathcal{P}

Let $\mathcal{P} = \text{OP}_1 \circ \text{OP}_2, \dots \circ \text{OP}_N$ be a query execution plan consisting of N operators. The plan's total cost $C_{\mathcal{P}}$ is given by the component-wise sum of the individual cost components:

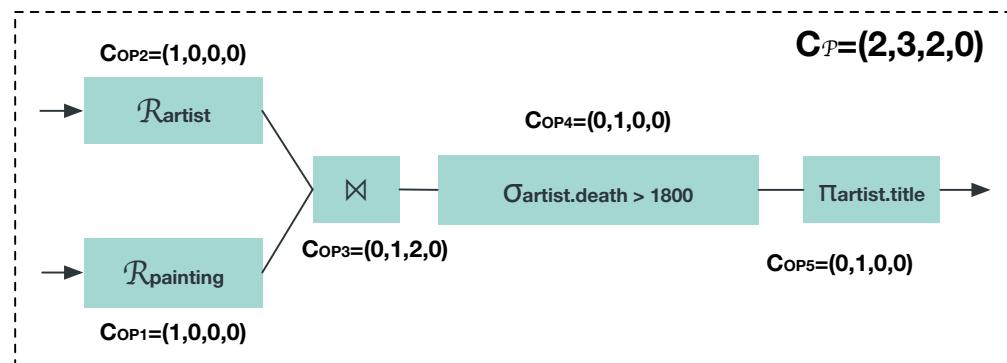
$$C_{\mathcal{P}} = \sum_{i=1}^N C_{\text{OP}_i} = (\sum_{i=1}^N c_{\text{IO},i}, \sum_{i=1}^N c_{\text{CPU},i}, \sum_{i=1}^N c_{\text{MEM},i}, \sum_{i=1}^N c_{\text{Q},i})$$

Example 6.9 Cost of a Query Execution Plan

Let $\mathcal{R}_{\text{painting}}$ and $\mathcal{R}_{\text{artist}}$ be relations from Example 3.3, the query “return the names of all paintings that were painted by an artist who died after 1800” results in the following, unoptimised execution plan:

$$\mathcal{P} = \pi_{\mathcal{A}_{\text{title}}}(\sigma_{\mathcal{A}_{\text{death}} > 1800}(\mathcal{R}_{\text{artist}} \bowtie \mathcal{R}_{\text{painting}}))$$

This query gives rise to the following operator tree with individual and total costs (the cost factors are purely illustrational and do not relate to any real values):



The role of $C_{\mathcal{P}}$ is twofold: Firstly, its absolute values can be used by the DBMS to determine hard constraints imposed for query execution. For example, if there is a fixed limit on memory that can be used by a query, an execution plan may be discarded simply based on exceeding that memory limit regardless of other factors. Similarly, a hard constraint on quality could be imposed. Secondly, it can be used to relatively compare execution plans to one another. Example 6.9 illustrates cost calculation for a simple execution plan.

The cost factors per operator can be determined in many ways and the model does not assume any specific approach. However for c_Q , we will propose an approach in Section 6.2.2. We do require, though, that the method for determining the factors must be internally consistent and should not change so as to keep costs comparable. For illustrative purposes, we outline the following, rough metrics to estimate c_{IO} , c_{CPU} , c_{MEM} : IO cost can be estimated by the number of bytes that must be read / written by an operator. This is equivalent to the estimated number of tuples read (T_r) or written (T_w) times the estimated size of a tuple s in bytes times the atomic IO cost a_r or a_w for reading or writing a byte. One can also introduce a correction factor c to distinguish between different types of accesses (e.g., sequential vs. random).

$$c_{IO} = (a_r T_r + a_w T_w)sc$$

CPU cost can be estimated by the number of memory accesses and/or the number of arithmetic operations that are executed by an operator. Here we assign an atomic CPU cost a_{CPU} to a type of operation and the total cost can then be expressed as that atomic cost times the number of tuples T being processed (T_p).

$$c_{CPU} = a_{CPU}T_p$$

Memory cost can be estimated by the size of the in-memory data structures maintained by an operator (e.g., for sorting or hash-joins), i.e., the number of tuples materialised in memory T_m times the size of a tuple s .

$$c_{MEM} = sT_m$$

Since the different cost factor stem from completely different domains and are therefore not directly comparable, and because we only compare plans relative to one another, we require a cost normalisation step, which we base on available plan candidates $\mathcal{P} \in PLANS$ and which is outlined in Definition 6.15.

Definition 6.15 Cost Normalisation

Let $\text{PLANS} = \mathcal{P}_1, \dots, \mathcal{P}_N$ be N equivalent query execution plans with associated costs $\text{COSTS} = C_{\mathcal{P}_1}, \dots, C_{\mathcal{P}_N}$. We consider the maximum cost C_{\max} for those plans to be the component-wise maximum of all $C_{\mathcal{P}} \in \text{COSTS}$. Using C_{\max} , the normalised cost $C_{\mathcal{P},i,\text{norm}}$ can be expressed as:

$$C_{\mathcal{P},i,\text{norm}} = \left(\frac{c_{\text{IO},i}}{c_{\text{IO},\max}}, \frac{c_{\text{CPU},i}}{c_{\text{CPU},\max}}, \frac{c_{\text{MEM},i}}{c_{\text{MEM},\max}}, \frac{c_{\text{Q},i}}{c_{\text{Q},\max}} \right)$$

It follows from the definition, that each component of a normalised cost falls into an interval between 0 and 1 due to the normalisation, i.e., $c_{*,\text{norm}} \in [0, 1]$.

It is important to note, that the normalised cost can only be used for relative comparison within an existing set of plans. It generally does not provide any information about the original cost value and the underlying properties they may represent. Using these normalised costs, we can now derive the *cost score* according to Definition 6.16.

Definition 6.16 Simple Cost Score

Let \mathcal{P} be a query execution plan with associated, normalised cost $C_{\mathcal{P},\text{norm}}$. We consider the *cost score* to be the sum of the atomic costs.

$$\mathcal{S} = c_{\text{IO},\text{norm}} + c_{\text{CPU},\text{norm}} + c_{\text{MEM},\text{norm}} + c_{\text{Q},\text{norm}}$$

The basic cost score \mathcal{S} falls into an interval between 0 and N wherein N denotes the number of atomic costs, i.e., $\mathcal{S} \in [0, N]$ for $N = |C_{\mathcal{P},\text{norm}}|$.

The cost score \mathcal{S} can be used to directly compare query plans in a given context and, e.g., sort them in ascending order during plan enumeration in order to select the one plan that minimises the costs. However, it assumes that all the components have the same importance, which in practice is often not the case.

6.2.1 Cost Policies

The definitions employed so far treat every component of a cost $C_{\mathcal{P},\text{norm}}$ equally with regards to the resulting cost score \mathcal{S} . This is rarely desirable since the different components influence the outcome to a different extent (e.g., in terms of query execution time or result quality). This can be expressed by employing a *cost policy* as outlined in Definition 6.17.

Definition 6.17 Cost Policy and Weighted Score

Let \mathcal{P} be a query execution plan with associated, normalised cost $C_{\mathcal{P}_{\text{norm}}}$. We call the 4-tuple \mathcal{W} a cost policy.

$$\mathcal{W} = (w_{\text{IO}}, w_{\text{CPU}}, w_{\text{MEM}}, w_{\text{Q}}), \text{ with } w_{\text{IO}}, w_{\text{CPU}}, w_{\text{MEM}}, w_{\text{Q}} \in [0, 1]$$

The weighted cost score $S_{\mathcal{W}}$ with respect to \mathcal{W} can be expressed as the linear combination of $C_{\mathcal{P}_{\text{norm}}}$ and \mathcal{W} :

$$S_{\mathcal{W}} = C_{\mathcal{P}_{\text{norm}}} \cdot \mathcal{W} = w_{\text{IO}} C_{\text{IO}, \text{norm}} + w_{\text{CPU}} C_{\text{CPU}, \text{norm}} + w_{\text{MEM}} C_{\text{MEM}, \text{norm}} + w_{\text{Q}} C_{\text{Q}, \text{norm}}$$

The cost policy provides us with a blunt but simple instrument to assign importance to individual cost factors. For example, one can prioritise execution speed by assigning more weight to the CPU and IO components or one can favour the quality of the results by having a high w_{Q} . In an actual multimedia DBMS implementation, \mathcal{W} can be influenced at different levels, which supersede one another, namely (i) as a system-wide configuration, (ii) for a user-session (connection-level) or transaction, or (iii) for an individual query. Consequently, different users should have the ability to express different preferences in terms of cost policy.

6.2.2 Estimating Cost of Approximation

In order to complete the cost model, we must be able to generate an estimate of how the quality of a result is influenced by choosing a particular execution plan over another. We propose a simple method that is directly related to the approximate index replacement problem given in Definition 6.12. It involves a post-hoc comparison of query results to estimate the quality measure at different levels using the quality function Q .

Definition 6.18 Empirical Estimation of Quality Cost

Let \mathcal{P} be a naive, unoptimised query plan and let \mathcal{P}^* be an optimised version thereof. Let further λ_k be the s, k -selection and Q be a quality estimator that derives a score by comparing the results produced by \mathcal{P} and \mathcal{P}^* . The cost estimate for quality $c_{Q,k}$ at a given level k is then given as:

$$c_{Q,k, \mathcal{P}^*} = 1 - q_k = 1 - Q(\lambda_k(\mathcal{P}), \lambda_k(\mathcal{P}^*))$$

For a proximity-based query, the optimised version of the plan \mathcal{P}^* may involve the evaluation of an approximate index while the unoptimised version \mathcal{P} does not. However, any type of query plan that involves an approximate query method can be compared that way (e.g., stochastic scanning proposed by [Gia18]). In practice, the obtained numbers for $c_{Q,k}$ can be stored in a lookup-table and be updated on a regular basis (e.g., when querying the index anyway). In case no observation is available, one can provide a cost prior with the index implementation, which is an index-dependent estimate of the quality cost incurred. Furthermore, the DBMS could also pre-calculate certain cost factors based on the known classes of high-dimensional index replacements described in Section 6.1.3 and queries that are likely to be executed. Whenever a query is being planned, the DBMS can look up the estimate that best fits the query to perform the cost calculation.

As indicated by the subscripts k and \mathcal{P}^* , we estimate c_{Q,k,\mathcal{P}^*} separately for multiple recall levels k and for different families of plans. We propose a logarithmic progression of the ranking levels, e.g., $k = \{1, 2, 4, 8, \dots\}$ for two reasons: Firstly, the likelihood of an extreme (outlier) value is higher for lower than for higher levels for purely stochastical reasons. This is also visualised in Figure 6.2. Secondly, items with a lower rank are often considered more relevant than items that exhibit a high rank [JK02], i.e., we require a more accurate estimate for the lower levels. We also expect the impact of approximation to be different depending on the type of query that is being executed. For example, [JDS10] report that the distortion of the distance estimate for PQ is not very relevant for NNS but ought to be considered, if one is interested in the actual distance values.

Due to its simplicity, the model obviously only provides a rough estimate. Furthermore, there are some important implicit assumptions we operate upon: Firstly, we assume that the quality of an execution plan depends on the type of plan and not on the query parameters provided by the user (e.g., the query vectors used for proximity-based search). The underlying assumption is that the query parameters are being drawn from the same (albeit unknown) probability distribution that has produced the indexed data. We believe this to be a reasonable assumption, because as we have argued in Chapter 4, these vectors are usually produced by some well-defined feature transformation. The second assumption is, that the quality of the results at lower, constrained ranking levels provides a valid estimator for quality at higher levels, i.e., that the agreement is continued indefinitely. Both assumptions require some well-behavedness of the index structure, which may not necessarily be a given.

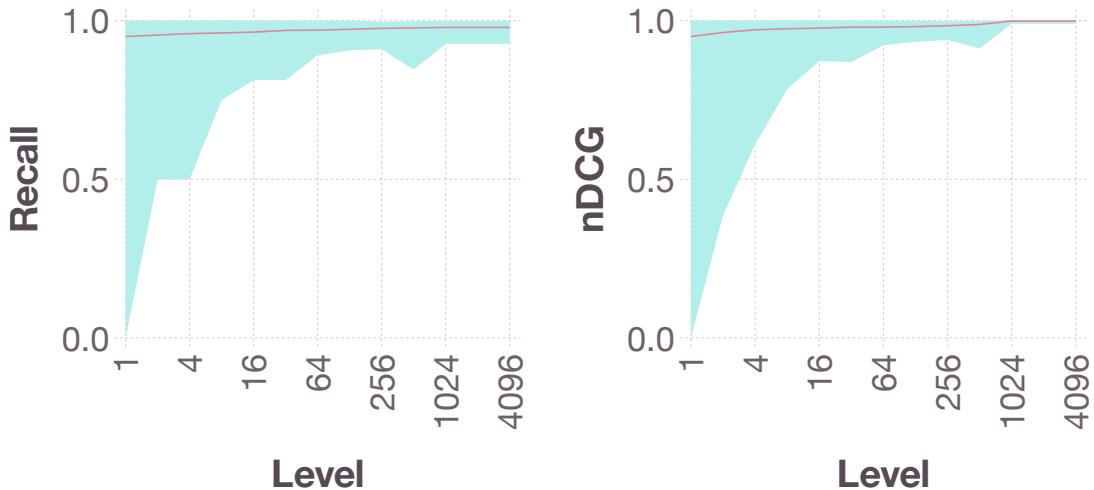


Figure 6.2 Quality of results produced by a PQ index in a NNS query in terms of recall and normalised DCG at different levels observed and averaged over a 1000 queries using different query vectors. The mean value remains more or less constant while minimum and maximum exhibit a large spread for smaller levels.

6.3 High-Dimensional Index Maintenance

As collections grow beyond the capacity for main memory architectures, the on-disk storage of indexes becomes an important aspect to consider [LÁJ⁺09; JSDS⁺19; CZW⁺21]. Yet, while the challenge of maintaining durable index structures and handling incremental as well as concurrent updates to them is a well-studied problem for, e.g., the B^+ -tree index [GUW09; Pet19], the on-disk storage and maintenance of high-dimensional indexes is not [Ams14], with the exception of a few examples [ÓJA⁺11; LÁJ⁺18; HJB19; LWW⁺20]. In addition to practical considerations, e.g., how data structures designed for use in main memory can be optimised for disk access to guarantee durability, there is the fundamental problem of an index's ability to cope with incremental changes to the data in a consistent manner, especially when concurrent read and write access is involved. What makes the issue challenging is that, as we have demonstrated in Section 4.2.2, many different types of algorithms exist for high-dimensional indexing, all of which bring their own characteristics. Since the problem of dynamic, high-dimensional indexing on secondary storage has not been widely considered in literature and is in fact often ignored by the inventors of new indexing methods,³ there are no established best practices on the handling of the issue at a systems level [ÓJA⁺11; Ams14].

³ It is either assumed that data is static or that re-indexing the entire collection is an option.

In this section, we propose a model to accommodate and maintain different types of durable high-dimensional index structures in a consistent manner. The model is supposed to be applicable to any type of index and considers it to be an auxiliary data structure, that provides certain (potentially imperfect) operations, that enable us to make incremental changes to it. Furthermore, indexes in this model exhibit an explicit state that allows the DBMS to make decisions as to whether and how the index should and can be utilised. The failure of an index due to accumulating changes is an explicit scenario that must be handled by the system. The proposed model consists of a series of components, based on which it can operate. Those are: (i) A *storage model*, (ii) a *write model* and, (iii) a *failure-and associated cost model* for every type of index, (iv) a system-wide notion of transaction isolation and (v) the ability to (re-)build an index from scratch. The model can be applied for any type of index, and we will see, that a traditional index structure can be considered a specific corner-case of the model.

6.3.1 Storage Model

The basic assumption and prerequisite is that all the high-dimensional index structures offered by the DBMS always act as secondary indexes. Consequently, the primary data (i.e., the original vectors) is always stored without loss of information and in a manner that allows for seamless propagation of changes, e.g., as index- or heap-organised tables in the case of a relational database.⁴ At this point, we reiterate the formal definition of a secondary index $\mathcal{I}_I^{\mathcal{R}}$ with respect to the relation \mathcal{R} it indexes from Definition 3.12 and the related Equation (3.10) that enables us to reconstruct attributes that are missing in the index.

The fallback strategy that the DBMS therefore can always resort to, in case an index becomes unavailable, is that of a sequential scan of the primary data, which results in a higher latency for reasons discussed in Section 4.2 and should therefore be avoided. Furthermore, the primary data can be used to reconstruct the index or parts thereof in larger batches – in a (re-)indexing process. Since, to the best of our knowledge, all relevant publications on high-dimensional indexes outline the steps required to construct the index from a data collection, we can safely assume that this is always possible. Moreover, the primary data can be used to correct missing or outdated information on the fly at query time, e.g., using random lookups.

⁴ This is an approach (implicitly) taken by methods such as DiskANN [JSDS⁺19] or SPANN [CZW⁺21] for the simple reason that in-memory storage is not feasible for very large collections.

In addition to this basic requirement, one must consider how the in-memory representation of a particular index can be efficiently mapped to disk. Referring to the list of indexes presented in Table 4.2, a naive mapping is often straightforward. Many indexes are organised as lists, trees or inverted (multi-)indexes [SZ03; BL14] and can therefore be organised using the persistent version of the respective data structure. At a high level, if an index is always scanned in its entirety, an array or list may be the easiest and most efficient form of storage. One could, for example, use a Binary Association Table (BAT)-like structure as used by MonetDB [IGN⁺12] to map TID to signature. If an index can benefit from filtering, e.g., to restrict a scan to a subset of the index in an inverted index [CPR⁺07; GJA10; JDS10], then a (B^+ -)tree structure might be more beneficial. In practice, there are many caveats to sort out, such as, balancing the size of clusters in inverted indexes [HJB19] or deciding which part of the data to keep in memory [LÁJ⁺09; JSDS⁺19; CZW⁺21]. We consider these to be important but index-specific aspects, which are, however, out of scope for this thesis.

6.3.2 Write Model of an Index

An index's write model specifies its ability to incorporate incremental changes into the data structures outlined by its storage model. To lay the groundwork, we formalise the notion of a change to a relation in the form of three operations listed in Definition 6.19.

Definition 6.19 Change Operations to Relation \mathcal{R}

Let \mathcal{R} be a relation and $t \in \mathcal{R}$, $t' \in \mathcal{R}'$ tuples with $SCH(\mathcal{R}) = SCH(\mathcal{R}')$. We call $OP_c(\mathcal{R}, \cdot) \in [\text{INSERT}(\mathcal{R}, \cdot), \text{DELETE}(\mathcal{R}, \cdot), \text{UPDATE}(\mathcal{R}, \cdot)]$ a change operation on \mathcal{R} and the quadruple $C = (OP_c, \mathcal{R}, t, t')$ a *change*.

$$\text{INSERT}(\mathcal{R}, t') = \mathcal{R} \cup \{t'\}$$

$$\text{DELETE}(\mathcal{R}, t) = \mathcal{R} \setminus \{t\}$$

$$\text{UPDATE}(\mathcal{R}, t, t') = ((\mathcal{R} \setminus \{t\}) \cup \{t'\})$$

That is, $OP_c(\mathcal{R}, \cdot)$ can be used to add, remove or replace tuples in \mathcal{R} . These operations correspond to the basic CRUD operations available to a DBMS. It is worth noting that formally, we can decompose any UPDATE operation into a consecutive INSERT and DELETE.

The write model WRITEM_I is now a function as given by Definition 6.20. It decides for every change C if and how it can be propagated to index I and applies it, if possible. In case a change C can be applied successfully, the write model returns `true`. If the change cannot be propagated, the write model returns `false`. In both cases, the change is forwarded to the index's failure model alongside with the outcome W of the write model.

Definition 6.20 Write Model for Index I

Let \mathcal{R} be a relation and I an index on attributes $I \subset \text{SCH}(\mathcal{R})$. The write model WRITEM_I is a function that determines, if a change $C = (\text{OP}_c, \mathcal{R}, t, t')$ can be applied to the index given additional parameters $P = \{p_1, p_2, \dots, p_n\}$:

$$\text{WRITEM}_I(C, P) = \begin{cases} \text{true}, & \text{if } \text{OP}_c(\mathcal{R}, t, t') \text{ can be applied to } I \\ \text{false}, & \text{otherwise} \end{cases}$$

The application of WRITEM_I always takes place within the context of the transaction \mathcal{T}_c that has issued the change and it therefore inherits all guarantees provided by it. Most importantly, we assume that all effects of WRITEM_I are made visible to other transactions upon `COMMIT` and are negated by a `ROLLBACK`.

6.3.3 Failure Model and State of an Index

The failure model FAILM_I is another function as given by Definition 6.21 that determines the new state of the index given change C , the result of the write model W and optional parameters $P = \{p_1, p_2, \dots, p_n\}$ (e.g., up-to-date index statistics).

Definition 6.21 Failure Model for Index I

Let \mathcal{R} be a relation and I an index on attributes $I \subset \text{SCH}(\mathcal{R})$. The failure model FAILM_I is a function that determines the index's new state $S \in \{\text{CLEAN}, \text{DIRTY}, \text{STALE}\}$ given a change $C = (\text{OP}_c, \mathcal{R}, t, t')$, the outcome of the write model W , and parameters $P = \{p_1, p_2, \dots, p_n\}$:

$$\text{FAILM}_I(C, W, P) = \begin{cases} \text{CLEAN}, & \text{if change can be applied without impairing the index} \\ \text{DIRTY}, & \text{if change may impair the index} \\ \text{STALE}, & \text{if change leads to index failure} \end{cases}$$

Based on the outcome of the failure model, the index's state and cost characteristics can be updated. The state determines, if and how the index can be used for query execution and thereby enforces consistent results within the quality parameters provided by the DBMS. The cost characteristic allows for more fine-grained control, in case certain cost factors are directly influenced by a change (or an accumulation thereof), for example, if changes lead to a deterioration in quality. We distinguish between three states as illustrated in Figure 6.3:

Clean indexes are ready for use by the DBMS and able to produce results within the speed and quality parameters they naturally provide.

Dirty indexes are ready for use by the DBMS. However, the flag signals that an index's ability to produce results within speed and quality parameters is impaired. This flag is usually accompanied by adjustments to the cost characteristic of the index that account for this.

Stale indexes have stopped functioning within their defined parameters and are therefore removed from consideration for query planning. The reason for this state is often an inconsistency, e.g., due to a missing record.

Freshly (re-)built indexes start in the CLEAN state. If the write model fails for a change C , i.e., $\text{WRITEM}(C, P) = \text{false}$, the state is updated to STALE and the index becomes unavailable. However, even if the write model succeeds, the change may still trigger a transition to either the DIRTY or the STALE state, if the changes that have accumulated in the meanwhile are expected to lead to index deterioration that goes beyond a certain threshold. Internal statistics can be consulted and updated to, e.g., adjust the index's cost model.

Again, we require that the application of WRITEM_I – including all updates to state and cost model – take place within the context of the transaction \mathcal{T}_c that has issued the change. Consequently, it also inherits all guarantees and is subject to its COMMIT and ROLLBACK semantics.

6.3.4 Concurrent Index Rebuilding

Once an index has been marked as STALE it can no longer be used by the DBMS until it is rebuilt by the $\text{REBUILD}(\mathcal{R}, \mathcal{I}, P)$ function, where $P = \{p_1, p_2, \dots, p_n\}$ again denotes a list of additional parameters. We can thereby ensure consistency as unprocessable changes accumulate, by preventing potentially faulty indexes from generating erroneous results due to wrong or missing entries. However, the failure of an index means that the DBMS must resort to a sequential scan,

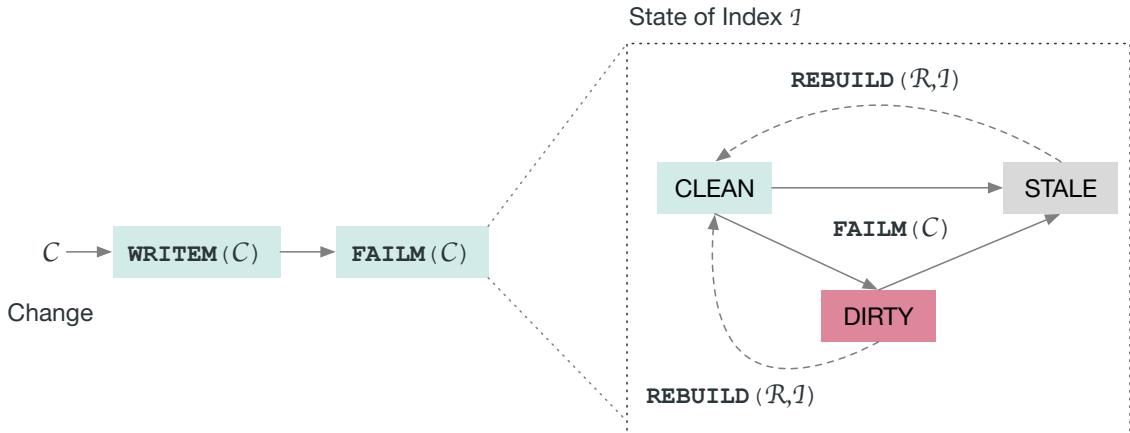


Figure 6.3 Illustration of a change C to the data that is processed by an index \mathcal{I} 's write and failure model and leads to potential state changes. An index knows the three states CLEAN, DIRTY, and STALE.

which impacts query performance negatively. This issue can only be resolved by a full index rebuild, which is often a long-running process that may take several minutes to hours to complete for large datasets. A time during which concurrent queries (read-only) as well as new changes (write) may arrive, which will inevitably lead to conflicts with a long-running rebuild operation that may cause operations to stall.

To prevent an index rebuild operation from blocking the entire DBMS instance for an extended period of time, we propose a strategy that minimises such conflicts. As a pre-requisite for this strategy to work, we assume that every transaction \mathcal{T}_t operates on a consistent snapshot \mathcal{S}_t of the database state at the time t when the transaction started. This is a concurrency control mechanism known as Snapshot Isolation (SI) [BBG⁺95], which is quite common in modern SQL and NoSQL systems such as PostgreSQL and MongoDB. For the remainder of this section, we will refer to snapshots of relations and indexes with subscripts, as illustrated in Example 6.10.

Example 6.10 Nomenclature for Snapshot Isolation

Let $\mathcal{T}_{t=1}, \mathcal{T}_{t=2}$ be two transactions that started at timepoint $t = 1$ and $t = 2$ respectively. We call $\mathcal{S}_{t=1}$ the database snapshot that is associated with the start of $\mathcal{T}_{t=1}$ and $\mathcal{S}_{t=2}$ the database snapshot that is associated with the start of $\mathcal{T}_{t=2}$. We further refer to the states of relation \mathcal{R} and \mathcal{I} at the a given timepoint with the same subscript, e.g., $\mathcal{R}_{t=1}$ and $\mathcal{I}_{t=1}$ respectively. Both are member of the respective database snapshot \mathcal{S} .

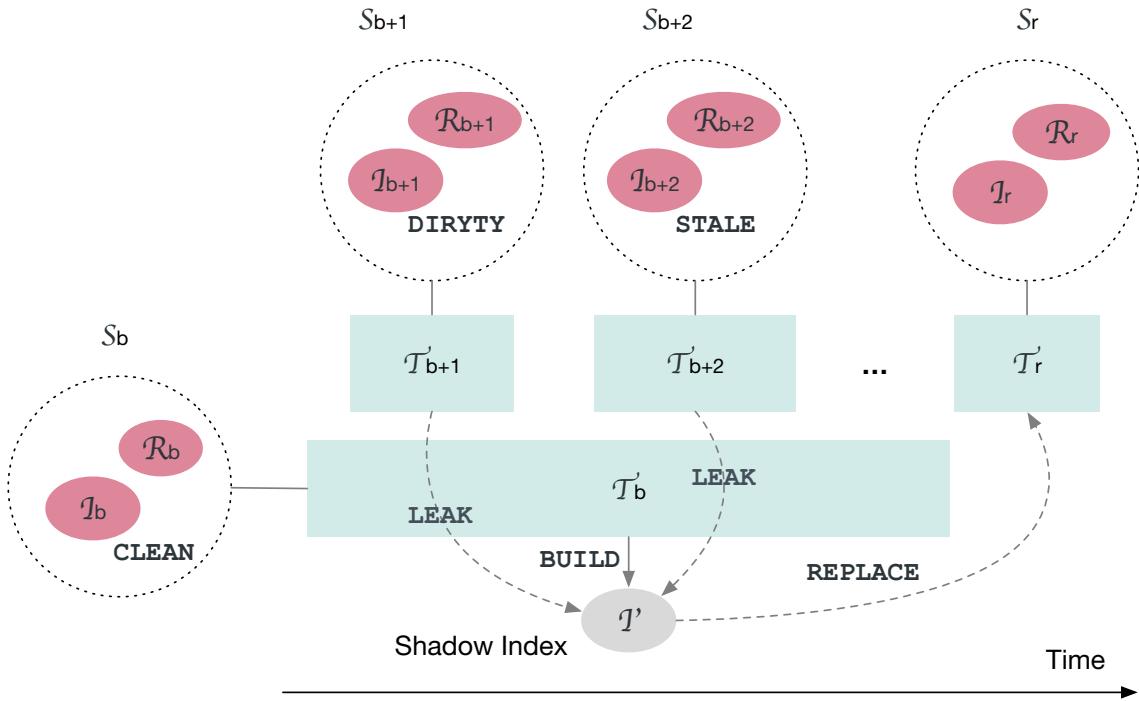


Figure 6.4 Illustration of asynchronous index rebuilding, which takes place in two separate phases BUILD and REPLACE. Information from transactions that take place concurrently to the BUILD phase are LEAK-ed through a side-channel upon commit of the respective transaction.

The method we propose tries to rebuild the (high-dimensional) index in two phases: The BUILD and the REPLACE phase, which start at timepoints b and r respectively (with $b < r$) and which are orchestrated by a dedicated REBUILD operator that exists and operates outside of a specific transactional scope. The basic assumption behind separating these two steps is, that building up the index structure is the time-consuming part, while writing it to disk is comparatively fast, since it does neither involve access to the primary data or any complex computations. The aim of our approach is to minimise exclusive access to the data structures maintained by the DBMS by exploiting the properties of SI. The approach is also illustrated in Figure 6.4.

The algorithm starts with the BUILD phase, during which the primary data in $\mathcal{R}_{t=b}$ is scanned in a read-only transaction $\mathcal{T}_{t=b}$ and the obtained data is used to build up the data structures required by the index. The so constructed *shadow index* I' exists outside of a proper database snapshot and is therefore not part of any visible database state and only accessible by the REBUILD operator. Since transaction $\mathcal{T}_{t=b}$ is read-only, SI allows for concurrent reads and writes to the database to take place and all transactions $\mathcal{T}_{t \in (b,r)}$ will see the state of the original $I_{t \in (b,r)}$ at that given point in time.

The shadow index construction is followed by the REPLACE phase, which comes directly after $\mathcal{T}_{t=b}$ has concluded in a second transaction $\mathcal{T}_{t=r}$. We assume that all write-transactions preceding $\mathcal{T}_{t=r}$ have finished once $\mathcal{T}_{t=r}$ starts, i.e., $\mathcal{T}_{t=r}$ must wait for these to complete. It then writes the shadow index \mathcal{I}' created during the BUILD phase to the snapshot $S_{t=r}$ and commits that snapshot, thereby replacing the $\mathcal{I}_{t=r-1}$ with an updated version $\mathcal{I}_{t=r}$ that can take over after $\mathcal{T}_{t=r}$ has committed. All transactions $\mathcal{T}_{t \geq r}$ will see the state of the newly built index $\mathcal{I}_{t \geq r}$. Since $\mathcal{T}_{t=r}$ is a write transaction, it may need exclusive access to all the data structures involved and thus acquire locks, i.e., concurrent write transactions that start after $\mathcal{T}_{t=r}$ may be required to wait until $\mathcal{T}_{t=r}$ concludes.

Separating these two steps obviously ignores one important aspect, which is that of change operations that take place between timepoints b and r . Since the SCAN phase operates on snapshot $\mathcal{R}_{t=b}$, those are not taken into account when re-building the index. We address this, by introducing a side-channel between $\mathcal{T}_{t=b}$ and concurrent transactions $\mathcal{T}_{t \in (b,r)}$ which LEAK relevant information through that channel. Hence, every concurrent transaction that commits successfully, reports all changes to the REBUILD operator, which will be responsible for merging these changes with the shadow index \mathcal{I}' using its write and failure model.

While there may be multiple ways to facilitate the LEAK process, depending on the type of DBMS, we propose to realise it using an internal event bus that implements a publish-subscribe-like communication model [Cha98a]. Upon start of $\mathcal{T}_{t=b}$, the REBUILD operator can subscribe to events that affect the relation \mathcal{R} it is currently re-indexing. It must be made sure, that all write-transactions that precede $\mathcal{T}_{t=b}$ conclude before $\mathcal{T}_{t=b}$ starts, which can lead to a very brief, system-wide halt. Relevant events that take place in a concurrent transaction $\mathcal{T}_{t \in (b,r)}$ are then recorded in a log and multicasted to all subscribers once $\mathcal{T}_{t \in (b,r)}$ commits. In addition to the changes discussed thus far, those events may also include DDL operations such as dropping the index or the relation. The operator then processes the incoming events in the order they arrive and writes the result to a second log. Similarly to the BUILD phase itself, this processing can and should include expensive computations. Once the BUILD phase has concluded, the content of the second log is replayed into the shadow index as part of the exclusive transaction $\mathcal{T}_{t=r}$. We must again synchronise the start of $\mathcal{T}_{t=r}$ with the end of all write-transactions that precede it.

We do neither assume nor require guaranteed delivery of messages nor do we need to block $\mathcal{T}_{t \in (b,r)}$ until all leaked events have been processed. However, we expect both to contribute to an improved stability.

6.3.4.1 Failure Cases and Limitations

The proposed algorithm has several failure cases that must be handled, and limitations that must be addressed. The first point of failure is message delivery during the LEAK process. Depending on how changes are recorded by the individual transactions, they may not be delivered to the REBUILD operator. If one, for example, imagines an in-memory event log, then events may be dropped for very long-running transactions simply because of memory limitations. The same can happen for the logs maintained by the REBUILD operator itself. In both instances, the shadow index must be discarded, and rebuilding of the index must be resumed from scratch in a new, up-to-date transaction. Both issues can be prevented by using a persistent messaging mechanism (e.g., a log on disk).

The REBUILD process will also fail, if at any time during the BUILD phase, the shadow index becomes unusable, i.e., if its failure model signals a problem with some of the indexed information and the shadow index becomes STALE. Again, in this case, the shadow index must be discarded and rebuilding of the index must be restarted. Therefore, for this algorithm to work, the index's write model should bring a certain degree of robustness to accumulating changes, otherwise the system will try endlessly to rebuild the index (which can be addressed by limiting the number of retries).

Since the likelihood of accumulating enough changes for the rebuild to fail increases with runtime, the algorithm can also be combined with other strategies to reduce time spent in the critical BUILD phase, e.g., by having multiple, smaller partitions of the primary data in \mathcal{R} being (re-)indexed independently. This is a strategy, e.g., employed by Milvus [WYG⁺21] and formally this leads to an index that is composed of multiple sub-indexes, each with a shared write and failure model but their own state. The state of the index is then determined by the state of its sub-indexes.

Last but not least, it is important to note that the proposed solution becomes a powerful tool if index deterioration and failure can be (approximately) anticipated. Presuming this is possible, the index rebuild can be started ahead of time while the original index \mathcal{I} is still functional. Due to SI, concurrent read transactions can continue to use the original index until the new version becomes available. In addition to heuristics (e.g. after x change operations the index is being rebuilt) one can also leverage the cost model proposed in Section 6.2 to determine if index quality is deteriorating and re-indexing is required.

6.3.5 Case Studies

We illustrate the idea of the storage, write and failure model using two index structures VAF [WSB98] and PQ [JDS10]. These case studies serve as mere examples and the concepts themselves can be applied to other indexes as well.

6.3.5.1 Vector Approximation File (VAF)

The VAF index (see Section 4.2.2.1) generates a signature SIG for every entry in a relation and an index search always results in a scan of all signatures $SIG \in \mathcal{I}_{VAF}$. Therefore, those signatures can be easily organised in a list or a tree-structure, e.g., by mapping TID to SIG in a B^+ -tree, which makes insertions and deletions straightforward. Alternatively, one can also map another primary key to the signature. A completely different type of VAF index could also try to map any other attribute to the signatures, to allow for the fast evaluation of proximity-based queries that restrict themselves to a subset of the data based on filtering.

To generate the signature from a vector, the VAF algorithm assigns that vector to a cell in a high-dimensional lattice that is determined by a partitioning of the vector space. That lattice – which acts as a scalar quantiser that must be stored with the index metadata – is derived from the component-wise minimum and maximum f_l and f_u of all feature vectors $f \in \mathcal{D}_f$. From this we can derive the write model for a relation \mathcal{R} with the attribute \mathcal{A}_f of data domain $\mathcal{D}_f \subset \mathbb{R}^d$ with $d \in \mathbb{N}_{>0}$. The `INSERT` is given by Equation (6.13) and illustrated in Figure 6.5.

$$\text{WRITEM}_{VAF}((\text{INSERT}, \mathcal{R}, t'), f_l, f_u) = \begin{cases} \text{true} & \text{if } f_{l,i} \geq t'[\mathcal{A}_f]_i \geq f_{u,i} \forall i \in [1, d] \\ \text{false} & \text{otherwise} \end{cases} \quad (6.13)$$

When a new tuple t' is inserted, the write model for the VAF index must check if the vector $t'[\mathcal{A}_f]$ falls within bounds of f_l and f_u . If it does, then the signature can be calculated directly and appended to the list – the write model returns true. Otherwise, the vector's position in the grid is undefined and the write model should return false. However, to harden the index for those types of failures, one can simply assign the closest cell (i.e., the first or the last one) instead, as illustrated in Figure 6.5. When the index is scanned and the bounds are not sufficient for filtering out the vector, the original entry must be fetched from the primary data and the exact distance must be obtained. Over time, we expect this to result in a deterioration of query latency, since more vectors must be read from disk.

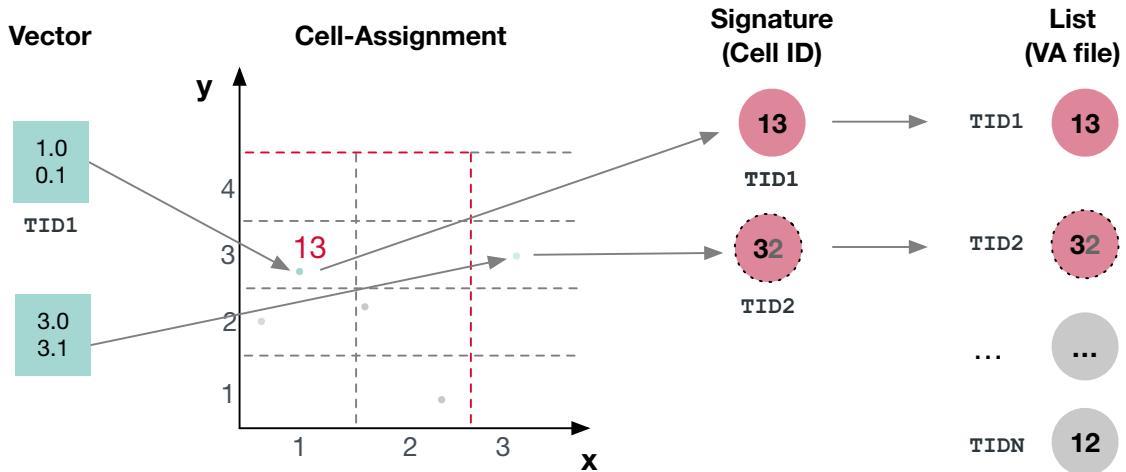


Figure 6.5 Illustration of the VAF index's write model (insert). If a vector lies within the bounds spanned by the minimum and maximum, the signature can be appended. If not, the closest cell is assigned, leading to a poor estimate of bounds during scanning.

If an existing tuple $t \in \mathcal{R}$ is removed, then this change can be applied in any case. The signature is simply looked up and removed from the list. However, as deletes accumulate, the component-wise minimum and maximum of the \mathcal{D}_f may shift, which potentially leads to an overly large and therefore sparsely populated lattice.

$$\text{WRITEM}_{\text{VAF}}(\text{DELETE}, \mathcal{R}, t) = \text{true} \quad (6.14)$$

Formally, the ability to apply an update can be directly expressed as a combination of $\text{WRITEM}((\text{INSERT}, \mathcal{R}, t'), \cdot)$ and $\text{WRITEM}((\text{DELETE}, \mathcal{R}, t), \cdot)$, as outlined in Definition 6.19. In practice, the two steps would be merged into a dedicated operation to optimise performance.

The failure model must now determine, how to handle the aforementioned changes. Both the sparseness of the overlay grid as well as the accumulation of out-of-bound entries potentially lead to a deterioration of the index's ability to produce results in a reasonable amount of time. Hence, as these changes accumulate, the need for an index rebuild may become more pressing. According to [WSB98], a typical VA file should be able filter out between 90 and 99% of the candidates without accessing the original vectors. That number can be used as a baseline to determine if an index rebuild is needed. Every time an index is scanned when executing a query, the actual value can be used to update the index's cost model. Eventually, the index's filtering capacity will drop below a threshold, in which case it can be marked as DIRTY.

6.3.5.2 Product Quantisation (PQ)

The PQ index (see Section 4.2.2.3) also generates a signature SIG for every entry in a relation. Depending on the type of PQ index, different storage strategies may be chosen. Exhaustive search always results in a scan of all signatures $SIG \in \mathcal{I}_{PQ}$. Therefore, the index can be organised similarly to the VAF index, e.g., by mapping TID (or any other primary key) to SIG in a B^+ -tree.

In case an IVFPQ index for non-exhaustive search is constructed, every signature SIG is additionally assigned to one (or multiple) clusters $c \in [1, C]$. During search, the scan is restricted to a subset of the clusters. A naive approach would be to simply prefix the signature SIG with the assigned cluster identifier c and then map the resulting, extended signature $CSIG$ to the TID in a B^+ -tree. This strategy can also handle the assignment of the same TID to multiple clusters, which will simply lead to multiple mappings for a given TID. At query time, c can then be used to quickly move to the clusters that should be considered and then scan all entries that share the same prefix. Alternatively, one could use a more specialised tree structure or could try to maintain separate lists per cluster.

To generate the signature SIG (and the cluster number c), PQ involves quantisation of a vector based on the assignment to a list of centroids in a codebook. For both PQ and IVFPQ, the quantiser used to construct the original index can be stored in the index metadata and re-used to process changes. Consequently, an incoming vector is simply quantised and appended to the respective data structure using the stored quantiser. The write model for a (IVF-)PQ should therefore always return true, as indicated by Equation (6.15).

$$\text{WRITEM}_{PQ}(\text{INSERT}, \mathcal{R}, t') = \text{true} \quad (6.15)$$

Similarly, for deletes, the entry associated with the deleted TID can be removed from the list. In the case of an IVFPQ index, the procedure incurs some additional overhead, since for a successful delete, the original vector must be read and its position in the inverted index must be determined by applying the two quantisers. In any case, the write model should return true as indicated by Equation (6.15).

$$\text{WRITEM}_{PQ}(\text{DELETE}, \mathcal{R}, t) = \text{true} \quad (6.16)$$

Again, formally, the ability to apply an update can be directly expressed as a combination of $\text{WRITEM}(\text{INSERT}, \mathcal{R}, t')$ and $\text{WRITEM}(\text{DELETE}, \mathcal{R}, t)$, as outlined in Definition 6.19.

The failure model of a PQ must potentially take multiple aspects into account. Firstly, a quantiser learned for a given size of the collection may become non-representative of the collection as it grows [LLN⁺20]. Secondly, a growing number of potentially very similar vectors may lead to a degeneration of the indexes ability to distinguish between them (part of the reason, why [JDS10] proposed the quantisation of residuals). And lastly, for IVFPQ, the different clusters constructed by coarse quantisation may become imbalanced over time, a phenomenon also observed in eCP [HJB19]. In all these cases, re-indexing with potentially different parameters may be required. It is to be noted, that finding the optimal choice of parameters for PQ for a given collection is still a poorly understood problem [GHK⁺14].

7

[...] But first they must catch you,
digger, listener, runner, prince
with the swift warning. Be
cunning and full of tricks and
your people shall never be
destroyed.

— Richard Adams, 1972

Cottontail DB

Cottontail DB [GRH⁺20] is our reference implementation for the query, cost, and index maintenance models described in Chapter 6. Starting out as a drop-in replacement for ADAM_{pro} [GS16], it has grown to be a full-fledged DBMS with support for classical Boolean, fulltext (powered by Apache Lucene¹) as well as proximity-based search used in multimedia retrieval and analytics applications in the *vitrivr* [RGT⁺16; GRS19a] context. All three aspects are combined in a unified query and execution model and many explicit and implicit decisions, which are described in this chapter, went into its conception and design.

Cottontail DB is written in Kotlin² and runs on the Java Virtual Machine (JVM). Cottontail DB’s source code can be downloaded from GitHub.³ It is freely available under a permissive open-source license and can be used as a standalone application, Docker container, or in an embedded mode.

7.1 Data Model and Nomenclature

Cottontail DB uses a data model very similar to that of other relational DBMS: All data is organised into *entities* – which correspond to tables or relations. Every entity consists of individual, strongly typed *columns*. The different data types currently supported are listed in Table 7.1. Columns can hold values of the declared type, or NULL, to indicate the absence of information, if the column has been declared as being nullable. To support the features used in multimedia retrieval applications, vector types are first-class members of the type system.

¹ See <https://lucene.apache.org/>; accessed September 2022.

² See <https://kotlinlang.org/>; accessed July 2022.

³ See <https://github.com/vitrivr/cottontaildb/>; accessed July 2022.

Table 7.1 Data types supported by Cottontail DB. Types in the numeric, vector, and complex class allow for domain specific arithmetics.

Name	Scalar	Numeric	Vector	Complex	JDBC / SQL
String	✓	✗	✗	✗	VARCHAR
ByteString	✓	✗	✗	✗	BLOB
Date	✓	✗	✗	✗	TIMESTAMP
Boolean	✓	✓	✗	✗	BOOLEAN
Byte	✓	✓	✗	✗	TINYINT
Short	✓	✓	✗	✗	SMALLINT
Int	✓	✓	✗	✗	INT
Long	✓	✓	✗	✗	BIGINT
Float	✓	✓	✗	✗	REAL
Double	✓	✓	✗	✗	DOUBLE
Complex 32	✓	✗	✗	✓	-
Complex 64	✓	✗	✗	✓	-
Boolean Vector	✗	✗	✓	✗	ARRAY[BOOLEAN](1, d)
Integer Vector	✗	✗	✓	✗	ARRAY[INT](1, d)
Long Vector	✗	✗	✓	✗	ARRAY[BIGINT](1, d)
Float Vector	✗	✗	✓	✗	ARRAY[REAL](1, d)
Double Vector	✗	✗	✓	✗	ARRAY[DOUBLE](1, d)
Complex32 Vector	✗	✗	✓	✓	-
Complex64 Vector	✗	✗	✓	✓	-

An entity can host multiple *records*, which correspond to tuples in the relation. Since Cottontail DB is a column store, records are only logical constructs in memory and do not reflect the physical data representation on disk. Consequently, records are assembled on-the-fly as queries are being executed. Similarly to columns, every record is also strongly typed, wherein a record is a tuple type of its elements. Internally, every record is uniquely identified by a *tuple identifier* (TID), which is a long value that can be used to address the record within an entity or an (in-memory) *recordset*. This TID is not exposed to the outside because it remains at the discretion of the storage and execution engine to generate, assign, change, and (re-)use it as data gets (re-)organised.

Furthermore, Cottontail DB allows for multiple entities to be grouped into *schemata*, which currently serves an organisational purpose only. Every entity can also host one or multiple secondary indexes that allow for more efficient data access to a single or multiple columns. In addition to indexes for Boolean and fulltext search, Cottontail DB also hosts different types of high-dimensional index structures used for proximity-based queries, for example, PQ [JDS10], VAF [WSB98], and LSH-based [IM98] indexes.

To address database objects, Cottontail DB uses a hierarchical namespace, i.e., every schema, entity, column, and index must be uniquely identified by a *fully qualified name*. All the information about the database objects – including all column and index statistics – is tracked in an internal *catalogue*, which is backed by the main storage engine.

7.2 Architecture

The main components of Cottontail DB are depicted in Figure 7.1. We use the path a query takes within the system, as indicated by the directed arrows, to illustrate the components involved in its implementation.

At a high level, every query must undergo *parsing*, *binding*, *planning* and *execution* in that order, even though, planning may be skipped for certain types of queries (e.g., DDL statements) or upon request. Starting with the binding step, all the information required for query processing is accumulated in a *query context* object, which is passed between the steps. In the following sections, we will describe the aforementioned steps in reverse order, since important concepts can be introduced more naturally that way.

7.2.1 Query Implementation and Execution

Cottontail DB implements an *iterator model* for query execution. This means that an implemented query is like a pipeline of operators, wherein each operator processes single records it receives from its upstream (input) operator(s) and passes single records to the next operator in the pipeline. Most operators have a direct correspondence to the relational operators (Chapter 3) and extensions (Chapter 6) introduced thus far. However, not all operators behave exactly as required by the model due to implementation details. For example, the `FETCH` operation corresponds to a variant of the projection π but instead of projecting onto a subset of the incoming attributes, it extends them, by looking up values in the database. Logically, `FETCH` corresponds to the operation $\text{OP}(\mathcal{R}) \bowtie_{\text{TID}} \pi_{\mathcal{A}_+}(\mathcal{R})$ for the additional attribute $\mathcal{A}_+ \in \text{SCH}(\mathcal{R})$, wherein the join acts on the tuple identifier. Such an operation can be beneficial since Cottontail DB allows for access to individual columns and it can thus reduce the cost of IO by deferring access to columns until, e.g., after a (s,k)-selection.

Conceptually, Cottontail DB distinguishes between *Nullary*, *Unary*, *Binary* and *N-Ary* operators, which differ in the number of inputs they can accept (0,

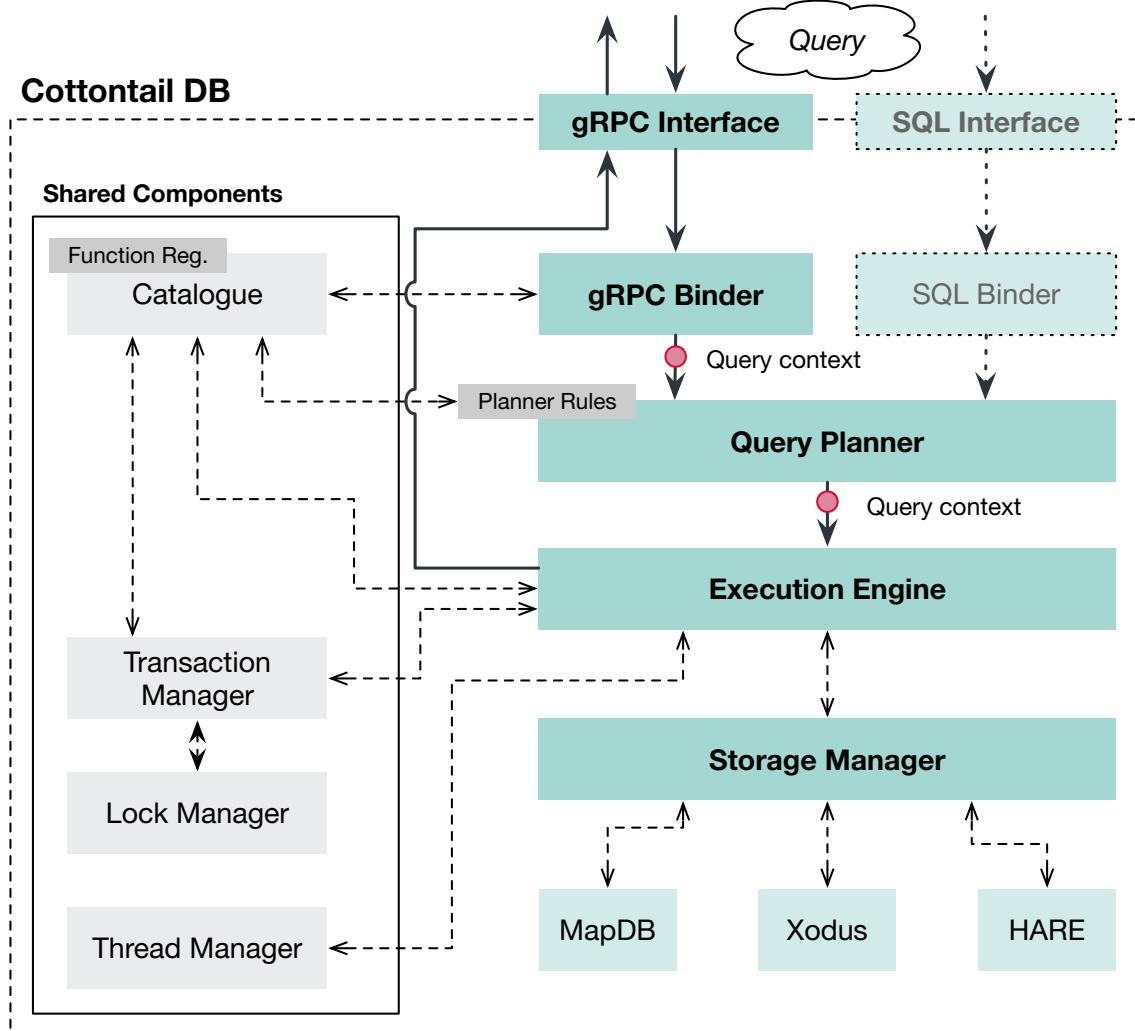


Figure 7.1 Architecture diagram of Cottontail DB’s main components. The directed arrows indicate the path a query takes within the system (dotted for a potential SQL interface). The dashed double arrows indicate interactions between components.

1, 2 or N). Furthermore, some operators require materialisation of intermediate resultsets and therefore act as pipeline breakers since they must collect all inbound records before they can start handing them down (e.g., sort operators). The most important types of operators implemented by Cottontail DB along with their correspondence to the relational operators are listed in Table 7.2. Every operator is implemented in Kotlin and pre-compiled to bytecode.

7.2.1.1 Execution Model

Operators in Cottontail DB are implemented as *coroutines*, i.e., every operator is a suspendable function that yields to the next operator in the pipeline upon emission of a value. Each operator collects a single record from its input opera-

Table 7.2 Main types of physical operators implemented by Cottontail DB alongside with their arity, their correspondence to relational operators and whether they require materialisation.

Type	Arity	Rel. Op.	Mat.	Description
SCAN	0	$\pi_P(\mathcal{R})$ $\rho_{\mathcal{A}_A \rightarrow \mathcal{A}_B}(\pi_P(\mathcal{R}))$		Act as a data source. Scans an <i>entity</i> and interacts directly with the storage manager.
INDEX	0	$\sigma_S(\mathcal{R})$ $\pi_{\mathbf{d}}(\mathcal{R})$ $\omega_{\mathcal{A}_{d\downarrow}}(\pi_{\mathbf{d}}(\mathcal{R}))$ $\lambda_k(\omega_{\mathcal{A}_{d\downarrow}}(\pi_{\mathbf{d}}(\mathcal{R})))$		Acts as a data source. Scans an <i>index</i> and interacts directly with the storage manager. Can be any type of operator sequence outlined in Section 6.1.4.
FETCH	1	$\bowtie_{\text{TID}}(\cdot, \pi_{\mathcal{A}_+}(\mathcal{R}))$		Extends every incoming record by fetching values for one or multiple columns and appending them to the record. This is a specialised projection that extends the columns present in the outgoing relation.
FUNCTION	1	$\bowtie_{\text{TID}}(\cdot, \pi_{f(\cdot)}(\mathcal{R}))$		Extends every incoming record by evaluating a specific function and appending the result as a new column to the record. This is a specialised projection that extends the columns present in the outgoing relation.
FILTER	1	$\sigma_S(\cdot)$		Filters incoming records by evaluating a given predicate. Records that do not match the predicate are not handed down the pipeline.
SORT	1	$\omega_O(\cdot)$	✓	Sorts the incoming records based on the specified columns in the specified direction.
LIMIT	1	$\lambda_k(\cdot)$		Skips and drops incoming records according to specification and thus limits the number of out-bound records to a specified number.
PROJECT	1	$\pi_P(\cdot)$ $\rho_{\mathcal{A}_A \rightarrow \mathcal{A}_B}$		A terminal operator that actively removes and/or renames columns in the incoming records. This is sometimes necessary because columns that are fetched for processing may not be desired in the final result.
MERGE	N	-		An operator that merges tuples from incoming strands of query execution in an arbitrary order (union semantics). Used for intra-query parallelism.

tor(s), performs processing, and forwards the resulting record by calling `emit()`, thus yielding to the next operator in the pipeline. Source operators usually iterate some data collection through a cursor and thus only `emit()` values, whereas sink operators `collect()` values without emitting anything (practically, the collected values are usually sent to the client that issued the query). Pipeline breaking operators are required to `collect()` all inbound values before they can start to `emit()`, i.e., they materialise a recordset. The suspension and continuation of operator calls is orchestrated by *Flows*, which are part of Kotlin's coroutines framework. An example of a naive operator pipeline is provided in Figure 7.2.

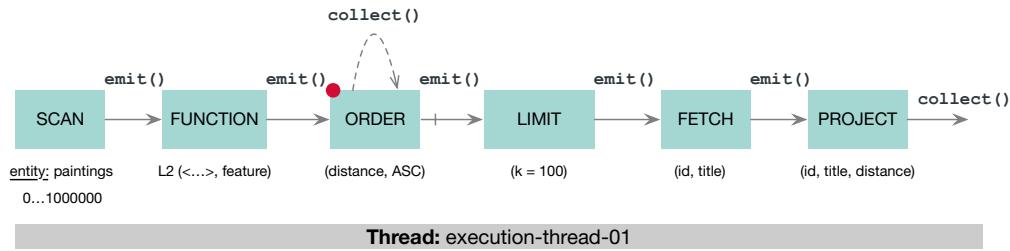


Figure 7.2 Operator pipeline for a proximity-based query that scans the entity *paintings*, calculates the L2 distance between a *feature* column and a query, orders by the obtained distance, and limits the number of outputs (kNN). Pipeline breaking operators are indicated by the red dot.

Inter-query parallelism is achieved by collecting different operator pipelines on different threads. The assignment of operator pipeline to thread is delegated to a *dispatcher*, which has a defined thread pool at its disposal and is also part of Kotlin's coroutines framework. Transaction isolation is guaranteed at all times by the *transaction manager* and is achieved through *snapshot isolation*. In addition, this execution model also allows for intra-query parallelism, which is used for queries that incur a high CPU cost. Intra-query parallelism is realised in a *data parallel* fashion by partitioning on the input. The decision whether to allow for parallelism is made upon implementation of a query – i.e., during the last step after the final plan has been selected – since it requires structural changes to the operator pipeline. The degree of parallelism depends on the expected CPU cost of the plan and the available CPUs on the host machine.

To prepare for intra-query parallelism, part of the operator pipeline is partitioned and every partition operates on a fraction of the input data. Whether partitioning is possible at a given point in the pipeline depends on whether the operator itself as well as all upstream operators allow for partitioning, which is queried upon implementation of the plan. An additional concurrency aware MERGE operator is introduced after such a partitioning point, which synchronises and merges all the incoming strands using a FIFO scheme. Since MERGE operators do not retain ordering, the sort operations must take place afterwards. An example for this is given in Figure 7.3. In practice, the three steps are combined into a single operator.

With an operator pipeline prepared in this manner, it again remains up to the dispatcher to allocate parts of the pipeline to concrete threads. Cottontail DB only declares how allocation should take place by choosing the appropriate type of dispatcher. Scheduling then depends on available hardware and system load. Consequently, the number of partitions is merely an upper bound on parallelism without guarantee that every strand will be processed in its own thread.

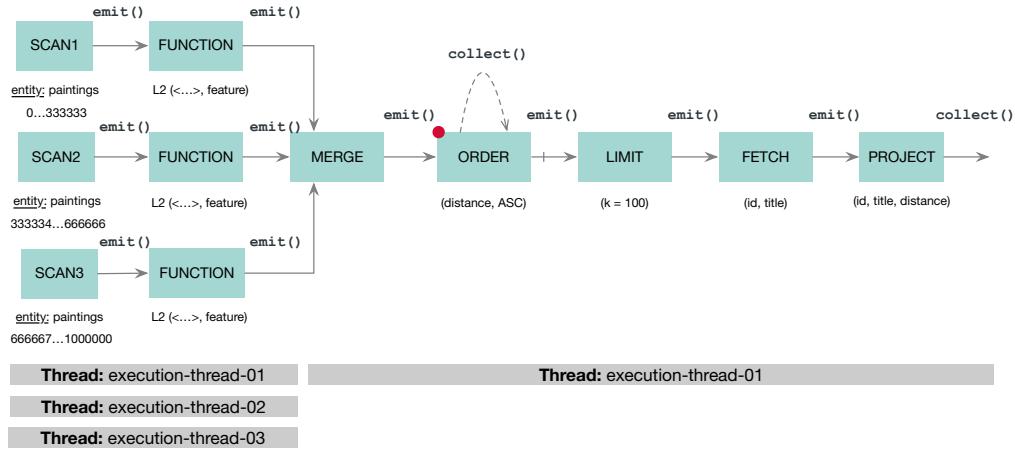


Figure 7.3 Operator pipeline that performs kNN on the entity *paintings*. The *scan* and distance *function* steps are partitioned and merged before order, limit, and projection is applied, which allows for parallel execution.

7.2.2 Query Planner

The goal of *query planning* is to find a cost-optimal execution plan for a query given any *policy* valid for the current query context and Cottontail DB’s cost model outlined in Section 6.2. The input into the query planner is a logical representation of the operators that should be executed as part of the operator pipeline, which we call the *canonical operator tree*. It results directly from the *query binding* stage. Query planning in Cottontail DB then takes place in three steps: Firstly, and optionally, the canonical operator node tree is compared to the *plan cache*. If the cache contains a hit for that tree, the already optimised plan in the cache can be reused and all the subsequent steps can be skipped. Secondly, if there is no hit in the cache or the cache is bypassed, the canonical operator tree undergoes *optimisation*. The optimisation step generates different versions of the input plan by applying rules that act on and manipulate the tree. Finally, the plan that minimises the expected cost is selected and implemented.

7.2.2.1 Optimisation Algorithm

The query planner used by Cottontail DB can be categorised as a rule-based query planner. It is loosely inspired by the Cascades Framework [Gra95] that was built as part of the Volcano project [GM93]. The planner generates new plans by enumerating and applying a set of rules to the input plan(s) in a recursive fashion. These rules are transformations that recognise patterns in the structure of the operator tree and try to rearrange and/or replace nodes to generate a more optimal but equivalent version.

The optimisation follows the same, basic algorithm: The available list of rules is enumerated, and every rule is applied to every node starting from the base of the tree. The application of rules takes place recursively, i.e., every node propagates a rule up the tree to its inputs. At every node, the planner first checks if the rule can be applied. This check usually incurs little cost and simply considers the type of node and other, available properties. If the check is successful, the actual application of the rule follows. This application may or may not yield a new version of the query plan, depending on more complex and more costly factors, such as the availability of an index, which requires a catalogue lookup. If a new plan results from the application of a rule, that new query plan is stored in a *memoization table* for future reference.

Every planning stage involves a defined number of passes, and the sum of all plans generated by one pass acts as the input for the next. All the intermediate plans are therefore run through the same algorithm, and all the rules are applied again. Consequently, the search space grows exponentially with the complexity of the input plan, the number of rules to apply, and the number of passes. To address this issue, Cottontail DB's query planner employs certain heuristics to keep the search space manageable:

- Very large, complex plans (e.g., stemming from sub-selects) are broken into groups and every group is optimised and treated as an isolated plan. The assumption being that the combination of smaller, near-optimal plans must be close to optimal as well.
- The planning of every group is done in two stages – a *logical* and *physical* stage – with distinct and disjoint sets of rules. The results of the logical stage act as an input for the physical stage.
- Logical and physical query plans are uniquely identified by a hash. Memoization in combination with the aforementioned hash is used to track trees that have already been optimised and skip them.
- Physical execution plans generated in the second stage are actively pruned by the planner between the different passes based on the cost estimate.

Between the logical and the physical planning, there is a step that maps every logical operator node to its naive, physical counterpart. The number of passes per stage and active pruning of plans during the physical phase are used to steer the high-level behaviour of the query planner. During the logical phase, the goal is to generate as many equivalent, logical query representations as possible to

have a corpus of plans to optimise on (expansion phase). In contrast, during the physical phase, the planner tries to limit the number of intermediate plans to prevent the search space from exploding.

Once a collection of potential plans has been generated for each group, the query planner enumerates the (sub-)plans per group, determines the expected cost, and selects the (sub-)plan that minimises it. It then reconstructs the full plan from the selected (sub-)plans, which is then implemented and executed. Furthermore, the mapping of the input plan to the resulting output plan is written to the plan cache for future reference.

7.2.2.2 Plan Caching

The plan cache is a mechanism that amortises the cost of the potentially expensive query planning if a certain type of query is encountered multiple times. At its core, the cache maps the unique digest of every incoming canonical operator tree to the resulting, optimised physical plan. If the same query is encountered again, that plan can be looked up and reused directly, thereby, avoiding another round of planning. Currently, there are certain limitations to the plan cache mechanism in Cottontail DB, which is why it is disabled by default. Firstly, the actual execution performance of a plan is not being evaluated against the cost model, i.e., the level of optimality of an execution plan is not assessed. Secondly, there is no mechanism that invalidates cached entries as changes to the data accumulate, e.g., due to indexes becoming unavailable as a result of its write and failure model (see Section 6.3). And finally, the optimal plan given a query does not only depend on the plan itself but also on the employed cost policy, which currently cannot be reflected by the cache (see Section 6.2).

7.2.2.3 Logical Optimisation

The logical optimisation step acts on structural properties of the operator node tree and aims to expand the list of query plans in a way that increases the likelihood of finding the optimal plan during the physical phase. It basically leverages the algebraic properties of the operators involved, e.g., the ability to decompose a FILTER operator evaluating a conjunction into two consecutive FILTER operators – one for each side. This is illustrated in the example given in Figure 7.4. The application of this rule can be seen as a preparation for the physical optimisation phase, where specific predicates may be satisfied by a secondary index.

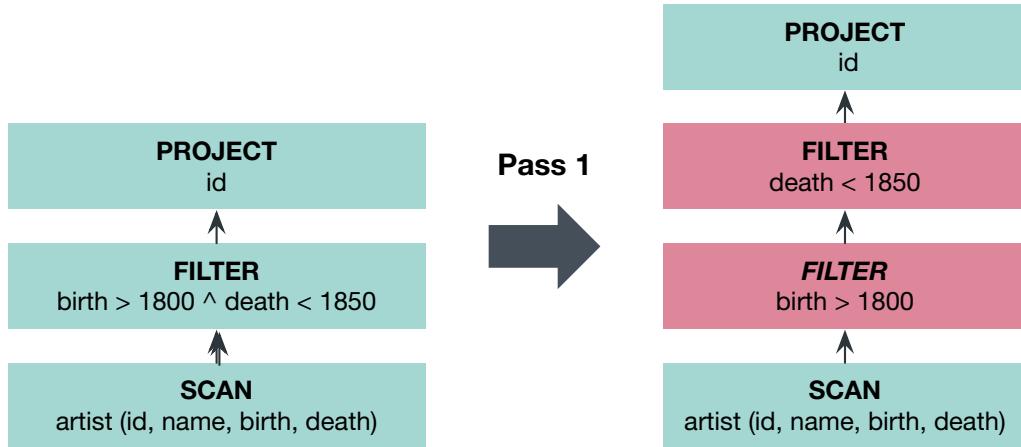


Figure 7.4 Decomposition of conjunctive predicates. The filter on the column `death` may be pushed down at a later stage during physical optimisation if an index were to be available.

7.2.2.4 Physical Optimisation

The physical optimisation step tries different implementations of operators to arrive at a more cost-effective plan. While structural changes to the operator tree may be possible, the focus lies on implementation aspects. At a high level, we distinguish between the following types of rules:

Index Replacement Rules try to replace FILTER predicates or FUNCTION invocations with an INDEX operation. For example, an entity SCAN followed by a FILTER can be replaced by a more efficient scan of a B^+ -tree index (see Section 3.3.5.2) followed by a FETCH operation or proximity-based operations can be delegated to an available, high-dimensional index using the classes of index replacements listed in Definitions 6.9, 6.10, and 6.11.

Deferral Rules try to defer the evaluation of functions and/or the fetching of columns to later stages of the query plan, which can be beneficial if a LIMIT or a FILTER leads to a reduction of the output cardinality. An example is given in Figure 7.5. This can significantly reduce IO and/or CPU costs. These rules also make sure, that only columns that are actually accessed are read from disk (i.e., they defer access until a column is removed from the plan).

Implementation Rules simply replace different implementations of an operation one-to-one, e.g., hash-join vs. nested-loop join for a theoretical JOIN operation or a scalar version of a FUNCTION invocation by a vectorised version that uses SIMD instructions.

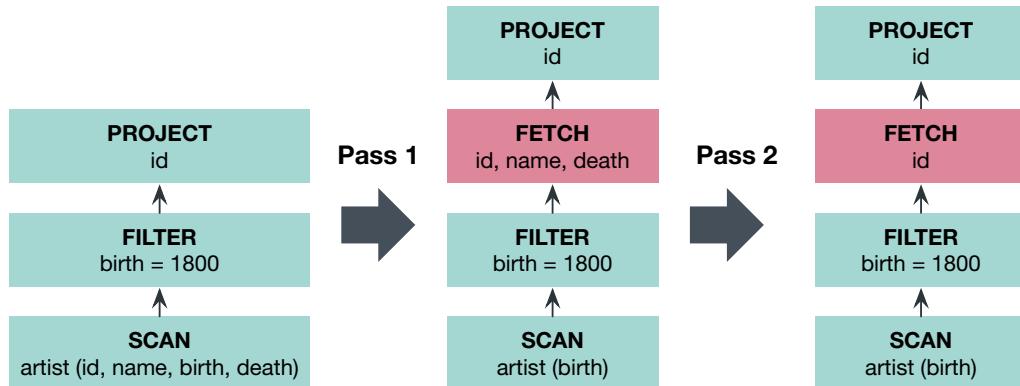


Figure 7.5 Illustration of deferring `FETCH` operations. Over multiple passes, the `FETCH` of the columns `id`, `name`, and `death` is pushed up until after the `FILTER` has been applied. Columns `name` and `death` are eventually eliminated because they are not required.

Summarisation Rules try to merge operations to arrive at a more efficient execution. For example, a `SORT` operation followed by a `LIMIT` operation can be replaced by a single, *limiting heap sort* operation, which applies sorting and limiting in a single step. Combining these two operations allows for more effective sorting, as compared to naively sorting all records first and thereafter limiting to the top k entries.

After each pass in the physical optimisation step, the list of resulting plans is pruned to the n best performing plans before handing them to the next pass. By this, we limit the size of the search space that is being explored. Once all passes have concluded, the best plan in terms of cost is selected and implemented. We directly apply the cost model introduced and described in Chapter 6.

7.2.3 Query Parsing and Binding

Cottontail DB uses a *gRPC*-based⁴ query language that allows for interaction with other systems. Hence, all queries are expressed as *Protocol Buffer*⁵ messages. These interactions include data definition, data management, transaction management, and the formulation of queries. Internally, the gRPC endpoint is a collection of four services for handling DDL, DML, DQL, and transaction management messages. The parsing of incoming messages, the mapping of messages to the correct service and endpoint and low-level handling of communication with clients is delegated to the gRPC Kotlin library.

⁴ See <https://grpc.io/>; accessed May 2022.

⁵ See <https://developers.google.com/protocol-buffers/>; accessed May 2022.

The reason for the choice of gRPC is threefold: Firstly, Cottontail DB’s predecessor – ADAMpro [GS16] – already used and introduced gRPC into the *vitriv* ecosystem. Consequently, building on this work greatly simplified the task of integrating Cottontail DB. Secondly, gRPC – being a platform independent remote procedure call framework – introduces out-of-the-box support for a large range of programming environments.⁶ And finally, the gRPC Kotlin library greatly simplified the task of structuring query endpoints, parsing incoming query messages, and handling the query responses.

Even though Cottontail DB uses gRPC, it can be seen as compatible with the SQL standard insofar as it supports a specific feature. This was demonstrated by the successful integration of Cottontail DB as a storage engine for *Polypheny-DB* [VSS18; VHS⁺20]. Consequently, it would be possible to add additional endpoints that can accept, for example, SQL queries. However, for complete SQL compliance, one would have to add language features that are currently not supported by Cottontail DB (e.g., JOINS). Furthermore, handling the technical details of parsing the incoming query, and communication with the client must be handled as well, which would lead to the implementation of many different standards such as JDBC or ODBC.

Irrespective of its source, any incoming query message undergoes a *binding* step in Cottontail DB. Query binding includes three aspects:

1. Names that reference database objects such as entities, columns, and indexes are resolved using Cottontail DB’s catalogue and replaced by concrete pointers to the respective object.
2. Literal values, function calls, and column references, e.g., in predicates, are extracted and replaced by *bindings* and attached to a *binding context*.
3. The parsed message is converted into a *canonical operator node tree* that logically represents the query.

The result of the parsing and binding step is the canonical operator node tree, which acts as a logical and unoptimised query representation. Every node in the tree represents an operator that acts on a *record* or *tuple* and that must be executed in order to generate the result specified by the query. The conversion of a message to a tree is performed in a naive manner and only syntactic checks and optimisation are performed at this stage (e.g., type checks, removal of trivial predicates such as “ $1 = 1$ ”).

⁶ Furthermore, there are specialised client libraries for Java, Kotlin and Python.

7.2.3.1 Binding Context

The binding context is a data structure that is part of the query context and a very important piece of functionality for query planning and execution. At their core, bindings are proxies for values used in a query. Cottontail DB distinguishes between *literal*, *column*, and *function* bindings. At any point during query execution, the binding context holds the current values for all the bindings associated with it, so that these can be accessed as the query is being evaluated.

The layer of indirection introduced by bindings enables many of Cottontail DB's functionalities, e.g., in the evaluation of complex expressions for predicates or functions. Most importantly, however, it allows for reuse of query plans, as bindings can simply be rebound to another binding context. Plan reuse then simply comes down to copying a plan with all its bindings and assigning it a new binding context afterwards.

With regards to intra-query parallelism, it is important to note that binding contexts are not thread-safe, which is why they are copied upon partitioning of a query plan. Consequently, every branch of an operator tree acts on its own copy of the same, root binding context.

7.2.4 Functions and Function Generators

In compliance with the model proposed in Section 6.1, Cottontail DB implements a *function registry*, which is part of the catalogue and can be used to obtain function implementations during query binding and planning.

The `FunctionRegistry` is a singleton object that is kept in memory for an instance of Cottontail DB. It provides the interface to register `FunctionGenerators` and resolve `Functions` given a `ClosedSignature`, which specifies the name of the requested `Function` and the types of arguments it can take. Such a request is then forwarded to a `FunctionGenerator` with a matching `OpenSignature`, which are again objects that exist per type of `Function` – e.g., a Euclidean distance or a scalar multiplication – and act as factories for concrete `Function` implementations. The `OpenSignature` consists of a `Function`'s name and *open* arguments that specify the class (e.g., number, vector) but not the concrete type (e.g., float number, double vector). It is up to the `FunctionGenerator` to generate and return a concrete `Function` implementation based on a `ClosedSignature`, if such an implementation exists. This generation process can be a simple lookup of a pre-existing piece of logic or a more complex operation (e.g., for code generation at runtime).

Given this facility, function resolution takes place in two steps: Firstly, a `ClosedSignature` of the requested function is converted to an `OpenSignature` to look up the responsible `FunctionGenerator` in a map. That generator then produces the concrete implementation, based on the `ClosedSignature`. This two-step process allows for easy extension, since adding new functions comes down to adding a generator and the type-specific implementations. Furthermore, the mechanism could be adapted to generate code at compile- or runtime, e.g., to automatically generate optimised, type specific versions of the same logic, which Cottontail DB currently does not support.

The equivalence between the execution of DFCs and high-dimensional indexes are implemented as dedicated planner rules that support the proposed classes of index replacements listed in Definitions 6.9, 6.10, and 6.11. Furthermore, some `Functions` implement the `VectorisableFunction` interface. This is a signal to Cottontail DB, that a vectorised version of the `Function` exists and can be obtained through a call to the `vectorised()` method. However, this functionality relies on the still incubating Java Vector API and is highly experimental.

7.2.5 Storage

Over time, Cottontail DB has seen multiple different layers for low-level data organisation and storage. Each of these layers addressed certain requirements that were valid at the time of conception. To accommodate these different engines, the *storage manager* offers a unified interface to interact with underlying *storage engines*, regardless of the concrete implementation. The main abstraction is that of a Database Object (DBO), namely, `Schema`, `Entity`, `Index`, and `Column`, in the form of interfaces. An instance of a transaction `Tx` can be obtained through the respective DBO to initiate interaction (i.e., `Schema.Tx`, `Entity.Tx`, `Index.Tx`, and `Column.Tx`). For that, each `Tx` implementation provides low-level primitives, e.g., `Column.Tx` offers `read()`, `update()`, `insert()`, and `delete()` for individual `Values`. A list of primitives for the different DBOs can be found in Table 7.3.

Even though the layer of storage manager abstraction could be used to combine different storage engines within a single instance, this is currently not done in the sense proposed by [Gia18]. This is mainly because replacement of a storage engine was often triggered by a very specific requirement the other engines could not fullfil. On an implementation side, however, Cottontail DB also relies on low-level aspects of the engines to ensure atomicity and (sometimes) transaction isolation. It would therefore take additional engineering to maintain these guarantees across different storage engines.

Table 7.3 Primitive operations offered by the storage engines' Tx abstractions to interact with DBOs including the argument and return types (arg → ret).

Type	Schema	Entity	Index	Column
create()	✓	✓	✓	✗
drop()	✓	✓	✓	✗
read()	✗	TID → Record	✗	TID → Value
update()	✗	Record	✗	TID, Value → Value
insert()	✗	Record → TID	✗	Value → TID
delete()	✗	TID → Record	✗	TID → Value
scan()	✗	Cursor < Record >	Cursor < Record >	Cursor < Value >
filter()	✗	✗	Cursor < Record >	✗
count()	✗	✓	✓	✓
largestTupleId()	✗	TID	✗	TID
smallestTupleId()	✗	TID	✗	TID

It is important to note, that Tx objects simply facilitate the interaction with a DBO within a mandatory transactional context managed by the transaction manager (see Section 7.2.6). The ability to commit or rollback changes is only supported at the level of this overarching context.

7.2.5.1 Storage Engine: MapDB

The *MapDB* storage engine is based on version 3.0.8 of a library with the same name⁷ and was the first storage engine for Cottontail DB. The main requirement at that time was speed and ease of integration. At its core, *MapDB* offers persistent versions of the Map interface in Java, backed by segmented hash trees. It can thus be regarded as a simple key-value store. *MapDB* uses memory mapped files, which allow for fast IO, and a generic and extendible serialisation engine to read and write arbitrary types from and to disk. Atomicity and durability of transactions are guaranteed by write-ahead logging, and in-memory latches mediate concurrent access to the underlying tree structure.

Cottontail DB integrates with *MapDB* by storing every column in a dedicated page-file that maps TID to record value. To do so, it bypasses the official high-level APIs offered by *MapDB* and instead uses an internal, lower-level storage mechanism to write records to pages directly. This has proven to be significantly faster than using the higher-level abstractions, especially for linear scans. Only the catalogue and all metadata are stored in persistent maps as are some of the data structures required for secondary indexes.

⁷ See <https://mapdb.org/>; accessed May 2022.

While easy to use in practice and reasonably fast by default, *MapDB* has several disadvantages. Firstly, the use of memory-mapped files is highly discouraged for DBMS as argued by Crotty et al [CLP22]. Secondly, the implementation of memory mapped files provided by the JVM is far from optimal and plagued by issues with resource retention and file locks on certain platforms. And finally, *MapDB* does not offer any MVCC, which is why isolation is provided by S2PL.

7.2.5.2 Storage Engine: *Xodus*

The *Xodus* storage engine is based on version 2.0.1 of the open-source, transactional, and schema-less, embedded database for Java and Kotlin with the same name, developed and maintained by JetBrains.⁸ As opposed to *MapDB*, *Xodus* does not rely on the use of memory mapped files. Atomicity and durability in *Xodus* are guaranteed through write-ahead logging. A property that is very desirable in the context of high-dimensional index maintenance described in Section 6.3, is *Xodus*'s support for non-blocking reads through the use of MVCC and snapshot isolation, which was one of the reasons to experiment with this engine. *Xodus* offers three levels of abstraction over the same core IO facility:

Environments encapsulate multiple, named key-value stores that can hold arbitrary keys and values. They allow for transactional reading and modification of data across stores. The interface for stores provides simple primitives such as `read()`, `put()`, and `delete()` based on the key as well as `scan()` operations using `Cursors`.

Entity Stores use environments to store more complex entities with attributes and links to other entities. This higher-level abstraction allows for more advanced data modelling as well as the formulation of complex queries.

Virtual File Systems use the environments to provide a transactional file system within *Xodus*.

Cottontail DB integrates directly with *Xodus*'s environment API, by having a dedicated store for each column within a single environment. In a store, the TID is mapped to the respective value. Furthermore, Cottontail DB maintains dedicated stores for all the necessary catalogue information as well as data structures required for indexes. Some indexes, e.g., the fulltext search index, rely on the higher-level virtual file system abstraction.

⁸ See <https://github.com/JetBrains/xodus/>; accessed May 2022.

While *Xodus* offers many advantages for Boolean retrieval and transactional workloads, it is slightly slower than *MapDB* for linear scans of columns containing large values in a read-only setting. This is due to the tree traversal that is involved, which incurs more overhead than the scanning of *MapDB*'s page files. However, we were able to compensate for this drawback by the application of *Snappy*⁹ compression on individual, high-dimensional vectors, leading to even higher throughput. We expect that even better compression ratio could be achieved, by compressing ensembles of values, which might provide an interesting direction for future optimisations.

7.2.5.3 Storage Engine: HARE

HARE is an experimental storage engine for Cottontail DB that so far has not made it to a final version and has thus never been used in an actual setup. The idea of *HARE* was to build the storage engine end-to-end to have more fine-grained control over data access patterns. *HARE* features all components required for persistent data management, including a dedicated *disk manager* and a *buffer pool manager*. Atomicity and durability are provided by maintaining undo and redo logs on a page level, which is reasonably straightforward to implement, but too slow for write-heavy workloads.

Structurally, *HARE* is built around page files, which hold the values for an entry. It interfaces with Cottontail DB's type system through a serialisation subsystem not unlike that of *MapDB*. Internal data organisation in *HARE* comes in two forms: Fixed-length columns essentially organise data in a persistent array, which enables addressing of values based on the TID by simple arithmetics. This allows for very fast random reads as well as extraordinary scan performance and is inspired by the BAT introduced by Monet DB [IGN⁺¹²]. However, this form of storage comes with challenges for transactional workloads, for which solutions have only emerged in highly experimental form.

Variable-length records use a skip-list structure with index and slotted data pages for lookup, which is slower for random reads but still reasonably fast for full table scans, since index pages can simply be skipped. In summary, *HARE* can be seen as a proof-of-concept for building a dedicated storage engine for Cottontail DB, which offers a very interesting foundation for future research but falls short of existing engines, in its current form, since it still lacks required features.

⁹ See <https://github.com/google/snappy/>; accessed May 2022.

7.2.6 Transactional Guarantees

The transaction manager provides transactional semantics and guarantees in Cottontail DB. A transaction formally starts with the creation of a transaction context and ends with either a call to `commit()`, `rollback()`, or `kill()`, which is a rollback for deadlocked transactions. The transaction context is always attached to the query context. However, unlike the query context, the transaction context may live beyond the scope of a single query. This depends on whether a transaction was explicitly started through the respective endpoints, or whether it is an implicit transaction, spawned by the issuing of a single query. Transactions of the first type must be explicitly terminated by calling `commit()` or `rollback()` whereas transactions of the second type end with the execution or abortion of the query that created them.

The transaction context provides access to the DBOs through Tx instances (see Section 7.2.5), which are used during query execution. Since all these accesses are routed through the transaction context, it can take necessary steps to ensure transaction isolation. Irrespective of the guarantees provided by the storage engine itself, isolation can always be guaranteed through Cottontail DB's *lock manager*, which provides DBO-level locking through the S2PL algorithm. While this is not the most efficient solution, it is applicable even for storage engines that do not provide any means of transaction isolation. In case they do the lock manager can be bypassed.

To guarantee atomicity and durability of transactions, Cottontail DB relies on the capabilities of the underlying storage engine. The transaction manager simply assures that calls to `commit()` or `rollback()` are propagated accordingly. As mentioned, both properties are assured by some form of write-ahead logging for all storage engines, which is state-of-the-art for current DBMS.

There are only few guarantees with respect to consistency, since Cottontail DB currently does not support any advanced constraints on columns. What Cottontail DB enforces, however, is that non-null columns are not nulled, and that changes to indexes are propagated to the underlying index structures as described in Section 6.3. Failure to fulfil this leads to a rollback of a transaction.

PART IV

Discussion

8

*I counted everything. I counted
the steps to the road, the steps up
to church, the number of dishes
and silverware I washed...
anything that could be counted, I
did.*

— Katherine Johnson, 2005

Evaluation

In this chapter we present the evaluation of Cottontail DB [GRH⁺20], which implements the concepts described in Chapter 6 and has become an integral part of the *vitrivr* stack [RGT⁺16; GRS19b; GRS19a]. This quantitative evaluation is structured around three experiments that highlight three different aspects:

Multimedia Analytics Workloads This experiment tries to simulate multimedia analytics workloads on million-scale datasets, for which Cottontail DB was primarily designed. We demonstrate Cottontail DB’s versatility and discuss and compare different execution strategies. We use the Vimeo Creative Commons Collection 1 & 2 (V3C) [BRS⁺19; RSB21], which currently serves as the reference dataset for both VBS [Sch19] and TRECVID AVS [CHW⁺21]. The dataset consists of 17235 videos, which amounts to 2300 h of content. We use a series of features derived by *vitrivr*’s feature extraction engine Cineast.

Large-Scale Similarity Search In this experiment, we test Cottontail DB’s ability to scale out to very large datasets as employed in challenges such as the (Big-) ANN-Benchmark [ABF17; SWA⁺21]. While not the primary focus of Cottontail DB’s design, this is still an interesting task to obtain a baseline for future efforts. We use this opportunity to compare Cottontail DB to Milvus (see Section 5.3.1) in a series of large-scale NNS queries. All queries are run against shards of the Yandex Deep1B dataset [BL16], which is a collection of descriptors derived from a GoogLeNet [SLJ⁺15] neural network.

High-Dimensional Index Maintenance In this experiment we test how Cottontail DB copes with incremental changes to the dataset that take place concurrently to query activities. We demonstrate its ability to implement such changes in the presence of indexes. Finally, we test the proposed mechanisms for concurrent index rebuilding.

8.1 Evaluation Setup

The entire evaluation is centred around database queries that are sent to the system under testing. We generate measurements from these queries and the results they return. Typically, we average the obtained values over 10 repetitions to compensate for anomalies and we execute a single warm-up query to give the systems the ability to initialise necessary caches (exceptions from this rule are mentioned explicitly).

Rather than working with randomly generated data, we decided to use existing data corpora that reflect the structure of features as generated by different models. For that, we use the Vimeo Creative Commons Collection (V3C) [BRS⁺19; RSB21] and the Yandex Deep1B dataset [BL16]. For V3C, we rely on features generated by our own extraction engine Cineast [RGT⁺16] and from the Deep1B dataset we have prepared different shards that contain the first 5 million, 10 million, 100 million, and all 1 billion pre-generated feature vectors. All data is imported as a dedicated entity (Cottontail DB) or collection (Milvus) respectively, prior to executing the actual workloads. A list of all datasets and the resulting entities is provided in Table 8.1 and we refer to the collections and entities by their name throughout this chapter. Necessary index structures are prepared beforehand if not stated otherwise. Every entity exhibits the columns `id` and `feature`. The `id` column serves as a primary key and is either a `string` (all V3C-based datasets) or a `long` value (all Deep1B-based datasets). The Deep1B shards also exhibit an additional `category` column, which holds a randomly assigned `int` value that is used for Boolean filtering. The state of all collections is listed in Appendix A.

Both Cottontail DB and Milvus are deployed on the same physical server (but they do not run at the same time). This server exhibits two NUMA nodes with an Intel Xeon CPU E5-2630 v4 (20 cores@2.2 GHz) and 192 GB RAM each. Therefore, the machine provides 40 compute cores and a total of 384 GB of RAM. The CPU supports the AVX2 instruction set extension, which allows for vectorised execution. The data directories from which Milvus and Cottontail DB access their data resides on three 500 GB SSDs that have been combined into a single, logical disk using RAID0 (striping). The server runs Ubuntu 20.04 and the OpenJDK Java version 17.0.3. All benchmarks are orchestrated by and sent from a separate node on the same network to minimise the impact on query execution. The evaluation scripts we used are available online.¹

¹ See <https://github.com/ppanopticon/cottontaildb-evaluation/>; accessed July 2022.

Entity	Source	d	N	Description
cineast_segment	V3C1&2	-	2512715	Segment metadata for V3C2. Does not contain any vectors.
averagecolor	V3C1&2	3	2512715	Average colours derived from video segments.
visualtextcoembedding	V3C1&2	25	2506273	Video to text co-embedding [SGH ⁺ 21a].
hogmf25k512	V3C1&2	512	2500943	HOG [BTG06] descriptors derived from video segments.
inceptionresnetv2	V3C1&2	1536	2508358	Vector derived from the last fully connected layer of a pre-trained InceptionResNetV2 applied on video segments.
conceptmaskssade20k	V3C1&2	2048	2469844	Embedding used for query by semantic sketch [RGS19].
yandex_deep5M	Deep1B	96	5000000	See [BL16] for more details.
yandex_deep10M	Deep1B	96	10000000	See [BL16] for more details.
yandex_deep100M	Deep1B	96	100000000	See [BL16] for more details.
yandex_deep1B	Deep1B	96	1000000000	See [BL16] for more details.

Table 8.1 List of the data collections used for this evaluation.

The binary version of Cottontail DB is started with a minimum and maximum heap size of 64 GB and 256 GB respectively. We have set up Cottontail DB’s cost model so that it parallelises workloads aggressively, whereas, by default, it tries to balance resource use per query and expected speed-up gained. The set of configuration parameters including explanation can also be found in Appendix A. For Milvus we use the standalone version and we followed the official tutorial for setting it up.² All benchmarks were conducted with the latest version 2.0.2 without further adjustments to the default configuration. In this default state, Milvus is allowed to use all available resources.

8.1.1 Metrics

We assume that for every query, we receive a list of results R from the respective database. That list contains K items $r_i \in R, i \in [1, K]$ that match the query. The returned items r_i are ranked either explicitly based on score or distance (for NNS, FNS or range search), or just arbitrarily, i.e., r_1 comes at position one, r_2 at position two and so forth. The ranking is solely dependent on the database and query. Every item r_i is simply a tuple containing the primary key of the retrieved database entry, which we always use to test for equality.

² See <https://milvus.io/docs/v2.0.x/>; accessed July 2022.

For every query, we are interested in assessing the efficiency and effectiveness of its execution. Efficiency can be easily gauged in terms of query execution time or latency, which is the elapsed real time in seconds between issuing the query and having received all results. This metric is simple enough to obtain and reason about and does not require further elaboration. The effectiveness of the query, i.e., the quality of the results it produces, is a bit more complicated. At a high-level, the results obtained by a query are typically compared to a resultset known to be correct – the *ground truth*. In our case, we do not have an absolute ground truth and we therefore rely on the results produced by a sequential, non-parallelised scan of the data.³ Given a list of results R produced by the database and the ground truth G , there are different ways to obtain a quality metric as, e.g., described in [WMZ10]. Those include but are not limited to precision and recall (sometimes at fixed levels), average precision, Mean Average Precision (MAP), Spearman’s rank correlation, and Discounted Cumulative Gain (DCG) [JK02], to name a few. Some of those metrics are merely set-based (e.g., recall and precision) while others take the position of the individual items, i.e., the ranking, into account (e.g., DCG or Spearman’s rank correlation).

We limit ourselves to obtaining the recall as well as the normalised DCG [JK02] of the result R compared to the ground truth G at a given level $k \leq K$ with $K = |G|$.⁴ Using these two metrics, we can assess both the completeness as well as the ranking quality of the results produced by different execution strategies.

Recall at a fixed level k , for which a definition is provided in Equation (8.1), is purely set-based and simply checks for the existence of an element from the ground truth G_k in the result R_k , without taking the exact positioning of the items into account. It provides a good estimate as to whether an access method is able to produce all the relevant items.

$$\text{REC}_k(R_k, G_k) = \frac{|R_k \cap G_k|}{k} \quad (8.1)$$

If R_k contains items that are not contained in the ground truth (false positives) or G_k contains documents that are missing in R_k (false negative), this directly results in a drop in REC_k . Therefore, if all of the items $r_i \in R_k$ were to be contained in G_k , i.e., $R_k \cap G_k = G_k$, the recall would become 1.0 and a result could be considered a perfect match. In contrast, if none of the items in $r_i \in R_k$ were contained in G_k , i.e., $R_k \cap G_k = \emptyset$ then recall drops to 0.0.

³ Having verified the correctness of Cottontail DB’s search algorithms on multiple occasions over the years, we believe this to be a reasonable choice.

⁴ We use k as a subscript to indicate, that a set has been limited to the first k items.

However, recall does not provide us with any information about the ranking of the individual items. In an extreme case, two lists containing the same items but in reverse order would produce the same recall value of 1.0. Often, though, we find that items with a low rank are more important than items with a very high rank. This is true both when considering a human user browsing a list from top to bottom, typically paying attention only to the top entries, or a use case in which we are only interested in the top item(s) (see, for example, Section 2.3). Nevertheless, and despite these limitations, the recall metric and variants thereof are very popular in ANNS evaluations [ABF17; SWA⁺21].

To compensate for the absence of information about the ranking quality, we turn to a variant of the DCG at a given level k , which in its original form was proposed by [JK02]. Our adapted version is specified in Equation (8.2). It builds the sum over all items $r_i \in R_k$. The numerator of the expression specifies the assigned relevance of an item based on its position in the ground truth G_k , which is expressed by the function $\text{rel}(r_i)$ in Equation (8.3). If r_i has rank 1 in G_k , it is considered most relevant and thus receives the relevance k . If r_i has rank k in G_k , it is not considered very relevant and therefore receives the relevance 1. If r_i is not contained in G_k at all (false positive), then $\text{rank}(g_i)$ returns the lowest possible relevance 0. The denominator is the logarithm of the actual rank of $r_i \in R_k$ and therefore the farther down it appears, the smaller its impact on the overall score. Consequently, items that are highly relevant but appear far down in the list are penalised.

$$\text{DCG}_k(R_k, G_k) = \sum_{i=1}^k \frac{\text{rel}(r_i) + 1}{\log_2(i + 1)} \quad (8.2)$$

$$\text{rel}(r_i) = \begin{cases} k - \text{rank}_{G_k}(r_i), & \text{if } r_i \in G_k \\ 0, & \text{if } r_i \notin G_k \end{cases} \quad (8.3)$$

The reasoning behind using this DCG variant is that items with a low rank in G_k are assumed to be more important than items with a very high rank. The numerator accounts for this by assigning high relevance to items that appear high up on the list, which quantifies the gain of inspecting such an item. That gain is discounted by an item's actual rank in R_k , i.e., the further down in the list an item appears the larger the discount. If an item in the result R_k does not appear in the ground truth (false positive), then there is no gain at all. What is not directly captured by the DCG, however, are entries that appear in G_k but that are missing in R_k (false negatives).

To make the DCG values comparable across queries (potentially with different values of k), we normalise it by the ideal iIDCG_k (Equation (8.4)) to obtain the normalised nDCG_k (Equation (8.5)). The iIDCG_k simply quantifies the maximum DCG that could be obtained if the ranking of R_k was perfect with respect to the ground truth G_k . Therefore, the normalised nDCG_k always assumes values between 0.0 (poor) and 1.0 (perfect). All the metrics described here are also illustrated in Example 8.1.

$$\text{iIDCG}_k(G_k) = \sum_{i=0}^K \frac{k+1-i}{\log_2(i+1)} \quad (8.4)$$

$$\text{nDCG}_k(R, G_k) = \frac{\text{DCG}_k(R_k, G_k)}{\text{iIDCG}_k(G_k)} \quad (8.5)$$

Example 8.1 Result R , ground truth G and obtained metrics.

We consider the following results R and the associated ground truth GT ($k = 7$), which contain the same items but in exact reverse order.

Rank	GT	$k + 1 - i$	R	$R_k \cap G_k$	$\text{rel}(r_i)$	$\log_2(i+1)$
1	87	7	597	✓	2	1
2	123	6	331	✓	3	1.58
3	542	5	3213	✓	4	2
4	3213	4	542	✓	5	2.32
5	313	3	123	✓	6	2.58
6	597	2	87	✓	7	2.81
7	757	1	888	✗	0	3

The values for REC_7 , DCG_7 , IDCG_7 , and NDCG_7 according to Equations 8.1, 8.2, and 8.5 are given as follows:

$$\text{REC}_7 = \frac{6}{7} = 0.86$$

$$\text{DCG}_7 = \frac{2}{1} + \frac{3}{1.58} + \frac{4}{2} + \frac{5}{2.32} + \frac{6}{2.58} + \frac{7}{2.81} + \frac{0}{3} = 12.87$$

$$\text{IDCG}_7 = \frac{7}{1} + \frac{6}{1.58} + \frac{5}{2} + \frac{4}{2.32} + \frac{3}{2.58} + \frac{2}{2.81} + \frac{1}{3} = 17.23$$

$$\text{NDCG}_7 = \frac{12.87}{17.23} = 0.75$$

8.2 Simulated Multimedia Analytics Workloads

This series of experiments is about demonstrating Cottontail DB's ability to execute different types of query workloads. For that purpose, we have prepared a test protocol consisting of queries that we execute and that leverage the proposed relational algebra extensions and the generalised proximity-based operations described in Section 6.1. In addition to the basic performance measurements, we also try to highlight different aspects of query execution and how they influence the results.

The test protocol starts by obtaining a random vector from the test collection, which will then act as a query vector q . It then goes on to obtain the mean distance m of all vectors in the collection from that vector q to then perform a range search to find all entries that are between m and $\frac{m}{2}$ away from q , which results in result r_1 . Subsequently, it performs a simple NNS query using q as a query vector, which results in r_2 . Finally, it obtains the segment entries that are associated with the returned entries in r_1 and r_2 . In summary, we execute two ordinary Boolean queries and three different types of proximity-based queries. The queries are expressed as Pseudo-SQL in Listing 8.1.

```

/* Q1a: Select a random vector from collection --> q. */
select feature from <collection> skip <random> limit 1

/* Q1b: Select mean distance from the query vector --> m*/
select mean(euclidean(feature, <q>)) from <collection>

/* Q1c: Perform a range search query --> r1. */
select id, euclidean(feature, <q>) as dst from <collection>
where dst between m/2.0, m order by dst asc limit 1000

/* Q1d: Perform a nearest neighbour search query --> r2. */
select id, euclidean(feature, <q>) as dst from <collection>
order by dst asc limit 1000

/* Q1e: Obtain segment entries associated with the IDs. */
select * from cineast_segment where segmentid in <r1 + r2>
```

Listing 8.1: Pseudo-SQL of the queries executed for the analytics workload.

We run all the queries against the V3C-based collections listed in Table 8.1, which are instances of the dataset that have been prepared for *vitrivr*'s [RGT⁺¹⁶; GRS19a] participation to VBS 2022 [HAG⁺²²]. *vitrivr*'s data model has been explained in Section 4.1.1.

For the benchmark, we prepared the following indexes on the respective entities in Cottontail DB: The PQ index variant for exhaustive search (8 subspaces, 2048 centroids) and a VAF index (35 marks per vector component). In addition to the obtained metrics, we also let Cottontail DB report the execution plan it selected for a given query using its EXPLAIN query endpoint. These execution plans are sometimes illustrated for the purpose of discussion.

8.2.1 Experiment 1: Influence of Parallelism & Use of Indexes

For the first experiment, we nudge Cottontail DB into using specific, high-dimensional index structures in order to be able to compare them directly. Furthermore, we fix the level parallelism to a given value between 2 and 16. Both can be done by providing *hints* with the query, which are respected by the query planner as far as possible given the solution space. The results of these measurements are presented in Figure 8.1. Unsurprisingly, the total execution time for the entire workload goes up as the dimensionality increases. There is one anomaly for *inceptionresnetv2* ($d = 1536$) and *conceptmasksade20k* ($d = 2048$), where a sequential scan is faster for the latter for all three types of proximity-based queries (i.e., Q1b, Q1c and Q1d). While somewhat counterintuitive, this can be explained by the structure of the data itself: The vectors contained in *conceptmasksade20k* are very spare and can therefore be compressed very efficiently by Cottontail DB’s rather naive compression scheme. In a quick follow-up experiment, we observed a compression to 550 B on average, which is a reduction of roughly 95%. In contrast, *inceptionresnetv2* contains very dense vectors and compression is not possible at all. Therefore, the full 6144 B plus an additional 5 B overhead introduced by the compression must be read from disk. While this is a glitch in the implementation, its impact on query performance demonstrates the importance of IO and efficient on-disk representations.

We can also see from Figure 8.1, that all the workloads, with the exception of the Fetch (Q1a) and the Select (Q1e) query, clearly benefit from an increased level of parallelism both for full entity scans as well as index scans, which confirms the need for an efficient parallelisation and distribution model [Gia18]. The two only exceptions are rather simple queries that cannot be parallelised effectively because they are either inherently sequential (Q1a) or because the partitioning model of Cottontail DB does not allow for parallelisation (Q1e).⁵ The overall execution time of the workload is clearly determined by the proximity-based queries.

⁵ Q1e leverages a B^+ -tree index, for which a scan cannot be parallelised in its current version.

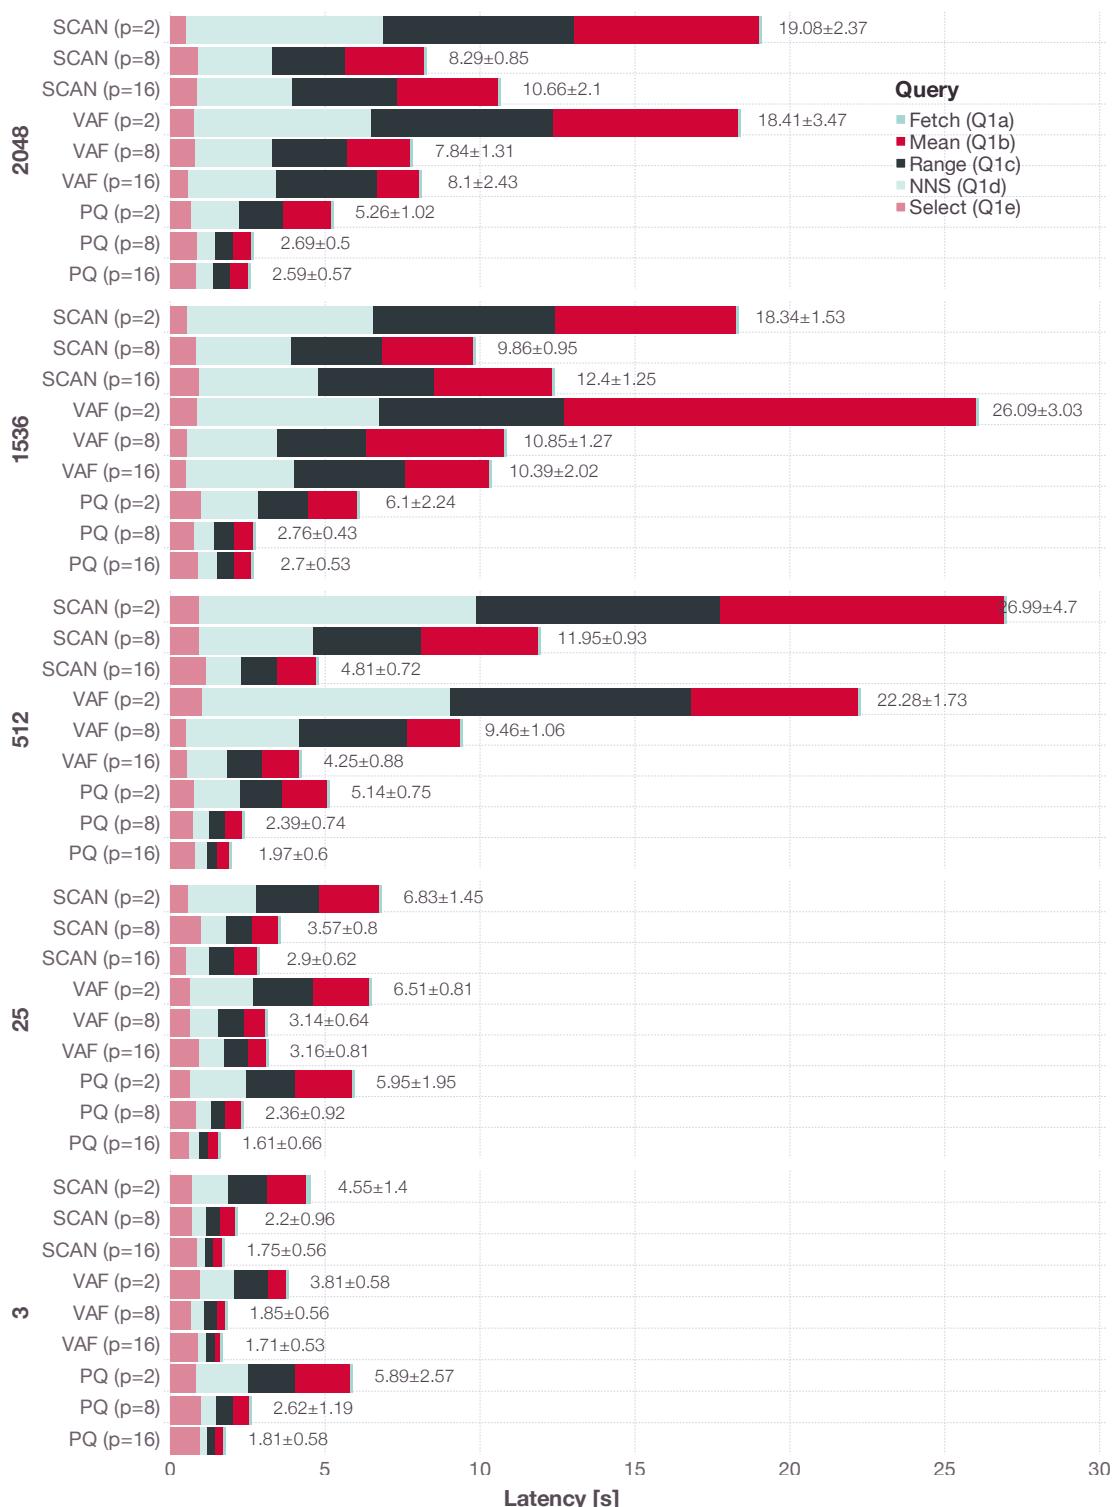


Figure 8.1 Latency in seconds (x-axis) on different entities using different access methods and levels of parallelisation (y-axis) for the analytics workloads during the first experiment. The individual queries Q1a-Q1e are highlighted in different colours. The use of the VAF and PQ index was enforced using query hints.

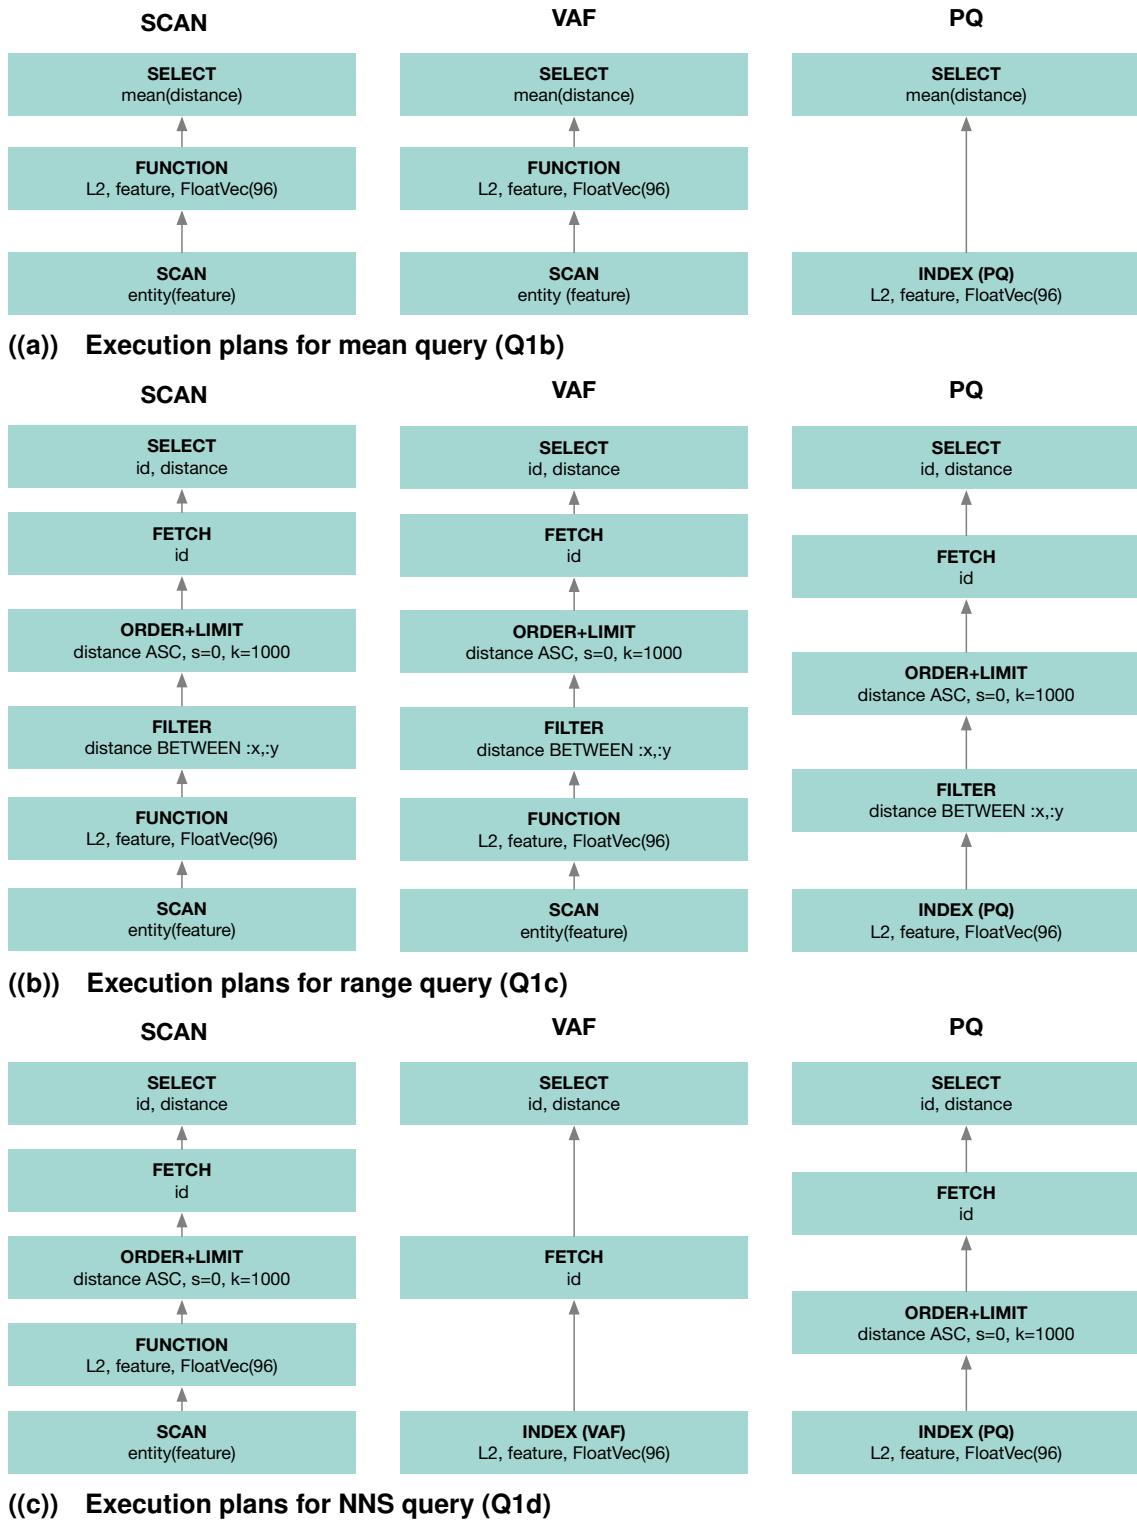


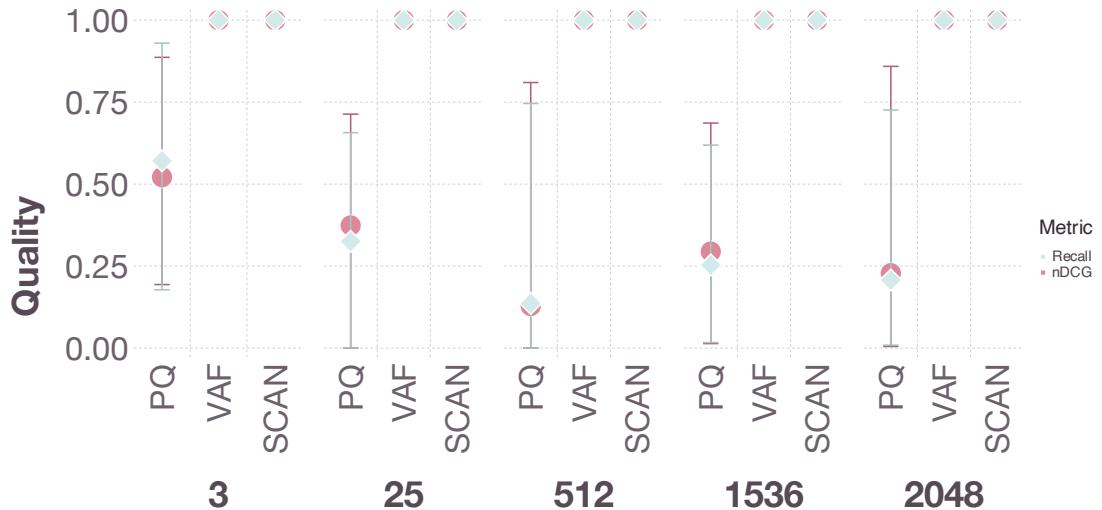
Figure 8.2 Execution plans produced by Cottontail DB for the mean (Q1b), range (Q1c), and NNS (Q1d) query workloads during the first experiment. The use of the VAF and PQ index was enforced using query hints.

An aspect that we would like to draw attention to is the influence of indexes: The Mean (Q1b), Range (Q1c), and NNS (Q1d) queries all seem to benefit from the PQ index, which, according to Figure 8.1, provides the lowest latency across all entities and for every degree of parallelisation. This is confirmed when cross-examining the query plans produced by Cottontail DB, which are listed in Figure 8.2. The plans for the Q1b (8.2(a)), Q1c (8.2(b)) and Q1d (8.2(c)) indeed use the PQ, if the respective hint is present. This is possible because PQ can be used to perform a class 1 index replacement (see Definition 6.9).

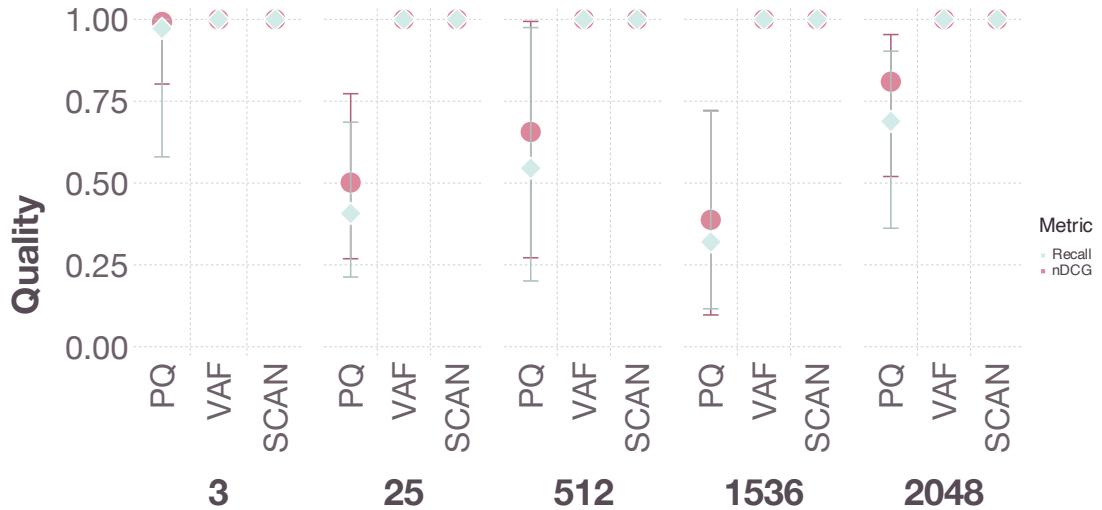
However, the observed speed-up is traded for a non-negligible, negative impact on quality in terms of recall and NDCG, as can be seen in Figure 8.3. Slightly higher values for NDCG provide us with an indication that mismatches occur further down in the ranking and may therefore be less relevant in practice. What is striking, though, is the very large spread for both recall and NDCG across all collections, which indicates that the performance of an index is highly dependent on the concrete query vector and thus hard or even impossible to reliably predict. Furthermore, we can observe that the negative impact on quality seems to be larger for range (Q1c, 8.3(a)) than for NNS (Q1d, 8.3(b)) queries for both recall and NDCG. This can be explained by the distortion of the approximate distance as a side-effect of PQ, which has been reported by [JDS10] and which is obviously very relevant for range queries. We see this as a confirmation of our proposition, that quality metrics and statistics for the purpose of query planning should be obtained for a type of execution plan rather than globally per index (see Section 6.2.2). Additionally, we can observe that the same index configuration for PQ yields very different results for the different collections. Unfortunately, there does not seem to be an apparent correlation between the dimensionality and the quality of PQ index results, which implies that the quality depends on more than just the size of the vectors.

In contrast, the VAF index can only provide speed-up for NNS queries (Q1d), since its use is limited to class 3 replacements (see Definition 6.11) and the current implementation does not support range queries.⁶ Again, this is confirmed by the query plans shown in Figure 8.2, where Cottontail DB resorts to a sequential scan in cases where the VAF cannot be employed. The speed-up gained by using VAF seems rather small, and it is directly related to the efficiency of the filtering. Over the different runs, we observed VAF filter efficiency values of between 80% and 99%. Again, there seems to be no apparent connection between dimensionality and the effectiveness of the filtering step.

⁶ The latter is a limitation of our implementation and could be added, according to [WSB98].



((a)) Quality for range query (Q1c)



((b)) Quality for NNS query (Q1d)

Figure 8.3 Quality of results (x-axis) in terms of recall (mint) and DCG (red) on different entities using different access methods (y-axis) for the range (Q1c) and NNS (Q1d) query workloads during the first experiment. The use of the VAF and PQ index was enforced using query hints.

We must assume, that the quality of all the tested index structures can be fine-tuned by choosing appropriate values for the hyperparameters. Since these values seem to depend on a multitude of factors (e.g., collection size, distribution of the data, dimensionality), this may not be a straightforward undertaking. However, learning an optimal set of hyperparameters for a given collection and index could be an interesting research direction for the future, very similar to learned index structures proposed by [KBC⁺18].



Figure 8.4 Score and rank of different query plan options for the mean query (Q1b). The preferred option for $w_Q = 0.0$ involves the scan of a PQ index. As w_Q goes up, the cost of impaired quality increases and a full entity scan (P_2) takes over, since none of the available indexes can satisfy the query. The scores for P_2 and P_3 do not depend on the quality parameter and incur almost the same cost, with P_2 being the less expensive variant.

8.2.2 Experiment 2: Influence of Cost Policy

So far, we have forced Cottontail DB into using one index over another by providing query hints that make an explicit choice. For this second experiment, we delegate the decision to Cottontail DB and observe how changing parameters in the cost policy influence plan selection. The idea behind the cost policies was explained in Section 6.2. Basically, they assign weight to the parameters of the cost model – the cost of IO, CPU, memory use, and reduction in quality – and can be used to steer the planner’s behaviour depending on the use case. For the experiment, we fixed the values of $w_{CPU} = 0.3$, $w_{IO} = 0.6$, and $w_{MEM} = 0.1$ (Cottontail DB’s default values) via the respective query hint. We then used the EXPLAIN endpoint to generate and print the execution plans as we increased the weight of the quality parameter w_Q from 0.1 to 0.9 in steps of 0.1. The results are depicted separately for the mean (Q1b, 8.4), range (Q1c, 8.5), and NNS (Q1d, 8.6) queries, for which we plot the scores and the ranking assigned to the individual plans by Cottontail DB’s query planner. Furthermore, we visualise all query execution plans that appear in the top three ranks.

For Q1b (Figure 8.4) we see that at least three different options exists and that by default, i.e., for $w_Q = 0.0$, the option involving the PQ index (P_1) is favoured. This is consistent with the findings from our first benchmark, where we demonstrated that PQ is the fastest of the available options. As long as the quality weight parameter is low ($w_Q \in [0.1, 0.2]$), P_1 remains the plan with the lowest score and thus the selected candidate. As w_Q increases, the score of P_1 increases as well and it moves up in the ranking, and plan P_2 , which involves a full table scan, takes over the lead. The two alternatives P_2 and P_3 are almost equivalent with the only difference that P_3 includes the `id` column in the initial scan, which is not required to calculate the mean and can therefore be optimised out (the application of the deferral rules described in Section 7.2.2). Consequently, P_3 is slightly more expensive and P_2 is selected instead.

A very similar behaviour can be observed for Q1c (Figure 8.5). Over the course of the measurements, six options appear in the top three ranks. Again, we start with $w_Q = 0.0$, where P_1 is the fastest option and again, it involves a scan of the PQ index. P_2 and P_3 are different variants involving the same index scan but use different intermediate operations and therefore differ in the fetching of required columns and sort algorithms (application of the deferral and implementations rules described in Section 7.2.2). As w_Q increases, P_1 , P_2 , and P_3 disappear from the top three ranks, since all three plans use the same index and therefore result in impaired quality. P_4 – which is the query plan involving the evaluation of a full entity scan – takes over and keeps the lowest rank. Again, P_5 and P_6 are less optimal variants of P_4 and are therefore not considered.

For both Q1b and Q1c using PQ or a full entity scan were the only available options due to the constraints discussed in Section 6.1.4. For Q1d (Figure 8.6) we see a third option appear, which is the use of a VAF index (P_3). As w_Q exceeds 0.1, the scan of the VAF index becomes the preferred choice and P_1 and P_2 become the second and third options. As w_Q moves beyond 0.2, P_1 and P_2 disappear from the top 3 and the entity scan becomes the preferred alternative to the VAF index scan. However, since the VAF index does not incur quality cost, it remains the favoured option.

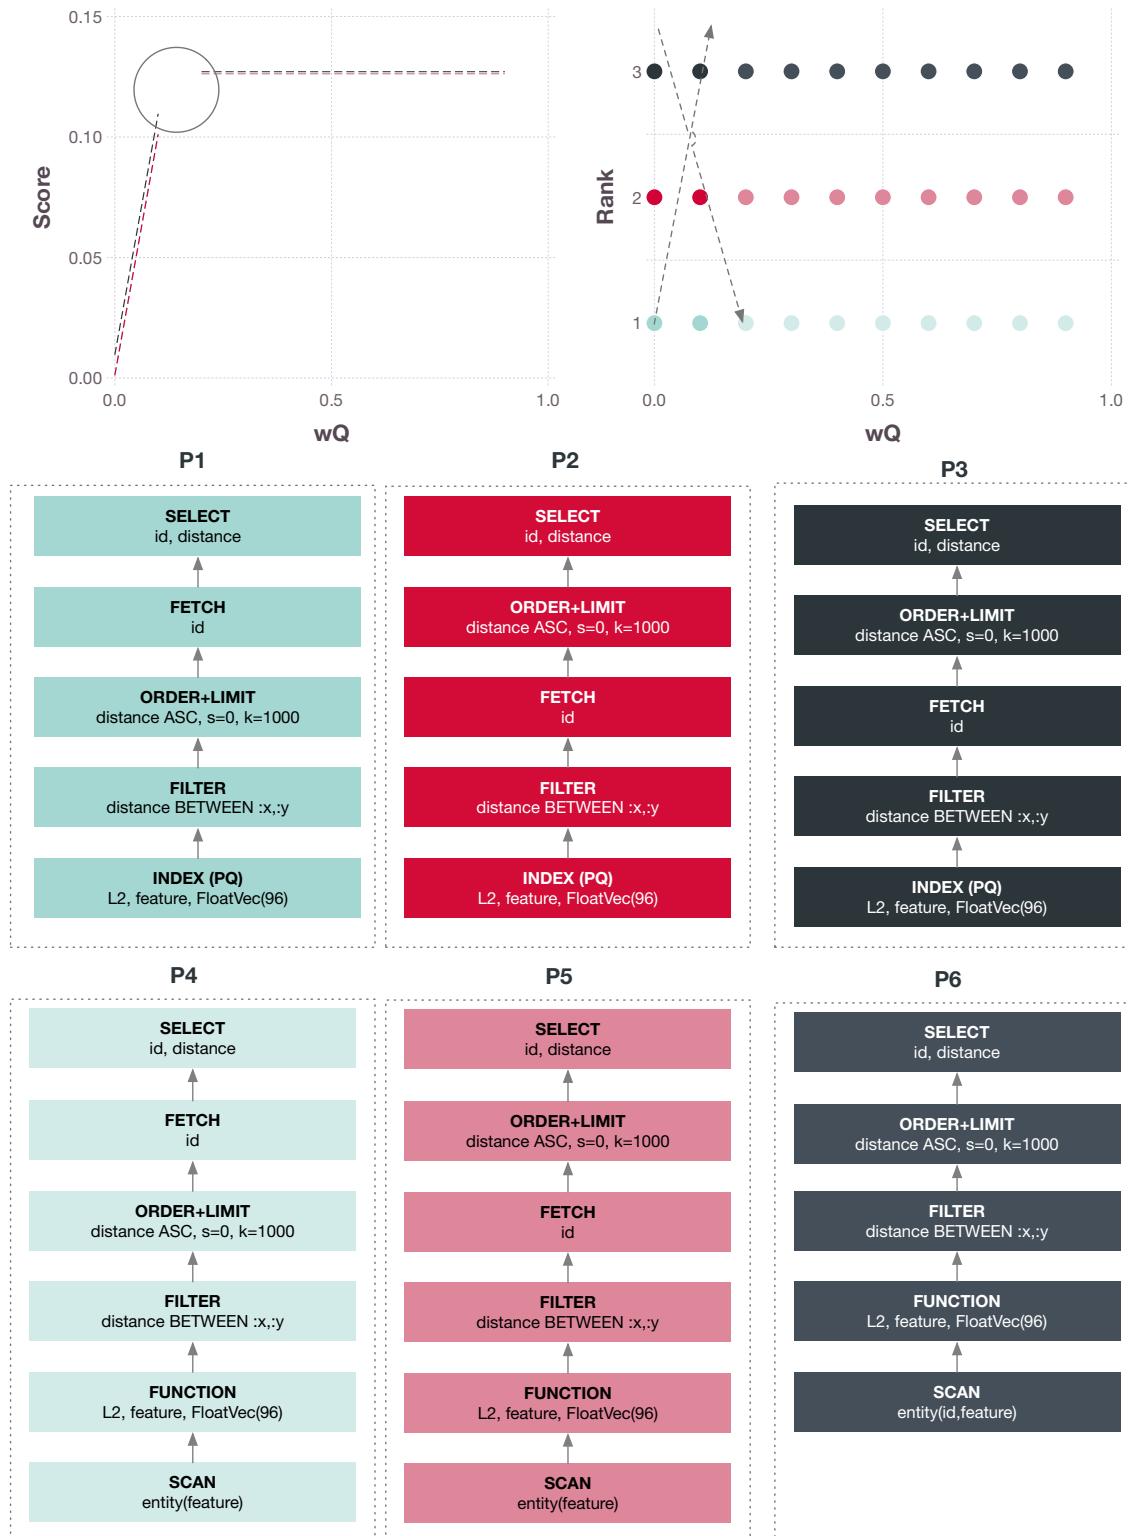


Figure 8.5 Score and rank of different query plan options for the range query (Q1c). The preferred option for $w_Q = 0.0$ involves the scan of a PQ index. As w_Q goes up, the cost of impaired quality increases and a full entity scan takes over, since none of the available indexes can satisfy the query. The scores for P_4 through P_6 do not depend on the quality parameter and incur almost the same cost, with P_4 being the least expensive.

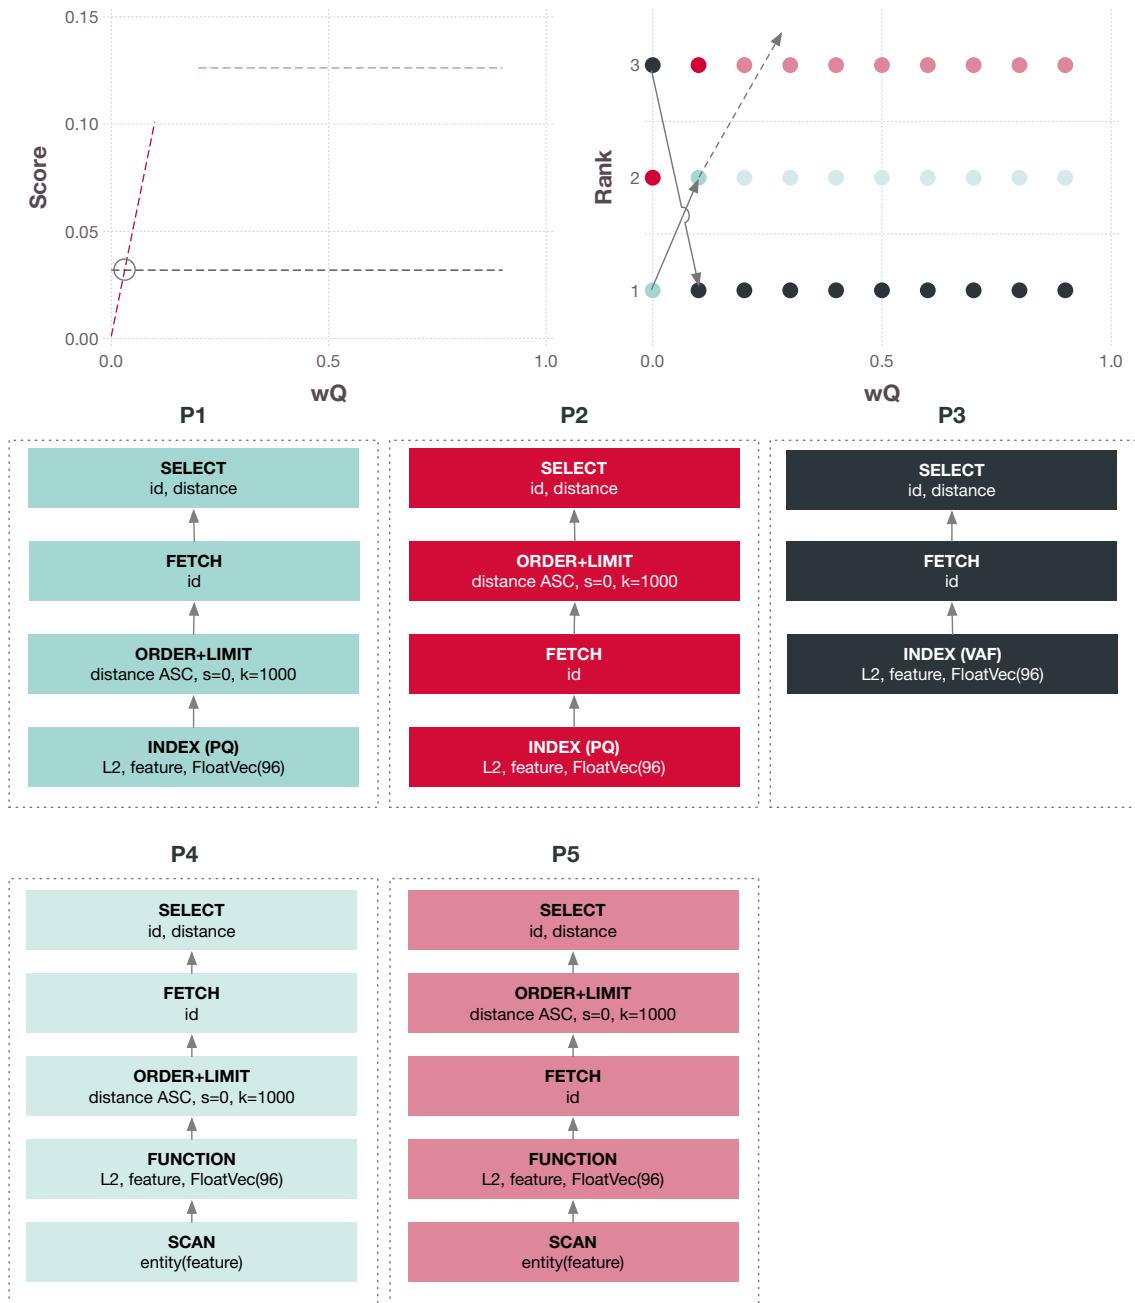


Figure 8.6 Score and rank of different query plan options for the NNS query (Q1d). The preferred option for $w_Q = 0.0$ involves the scan of a PQ index. As w_Q goes up, the cost of impaired quality increases and a VAF index scan (P_3) takes over, which can also satisfy NNS queries. Since VAF does not incur quality costs, it remains the favoured option.

8.2.3 Benchmark 3: Influence of Optimisation

We have demonstrated in the first experiment, that high-dimensional index selection has a considerable impact on query execution performance. However, as described in Chapter 7, in addition to selecting indexes, the query planner also performs other optimisations of the plan before executing it. For this experiment, we completely bypassed plan optimisation and observed the execution time for an unoptimised plan – which is a naive implementation of the canonical operator tree. All experiments were performed with parallelisation switched off, i.e., single-threadedly. The results are depicted in Figure 8.7.

We can see from this graph, that optimisation seems to have a larger effect on latency for proximity-based queries (Q1b to Q1d) involving high-dimensional vectors than it does for low dimensional ones. However, optimisation also seems to introduce a larger variance in latency, making it less predictable. Both effects can be attributed to the column-access deferral optimisation rule, which pushes access to columns that are not required for query execution up in the tree until after a (s, k) -selection has been executed. The idea is to reduce or even avoid access to data that is not needed and therefore the number of IO operations. This optimisation seems to be highly relevant for large vectors, but it diminishes as dimensionality decreases. We suspect that the page cache, which – given the huge amount of memory available – can keep many of the requested pages in memory, is responsible for this behaviour. When larger entities are being scanned, the cache becomes more saturated, and the likelihood of a cache miss becomes higher resulting in a higher overall latency. This effect is of course amplified if more than one column is being scanned, as in the unoptimised case.

As a second observation, we see that optimisation consistently benefits the Select (Q1e) query, where we load segment entries that match the preceding query results. The explanation of this is as simple as it is unexciting: The select query (Q1e) benefits from a B^+ -tree index scan in the optimised case, while the unoptimised case uses a full entity scan. Last but not least, we see that the impact of optimisation is slightly larger for range (Q1c) and NNS (Q1d) queries than it is for the mean query. This can be attributed to a second optimisation, which combines the sorting and limiting used for Q1c and Q1d into a single step that uses a specialised heap data structure. This sort algorithm can be executed more efficiently and uses less memory than naive sorting. We expect the impact to be even larger for setups that do not have the same amount of memory and must therefore resort to on-disk sorting.

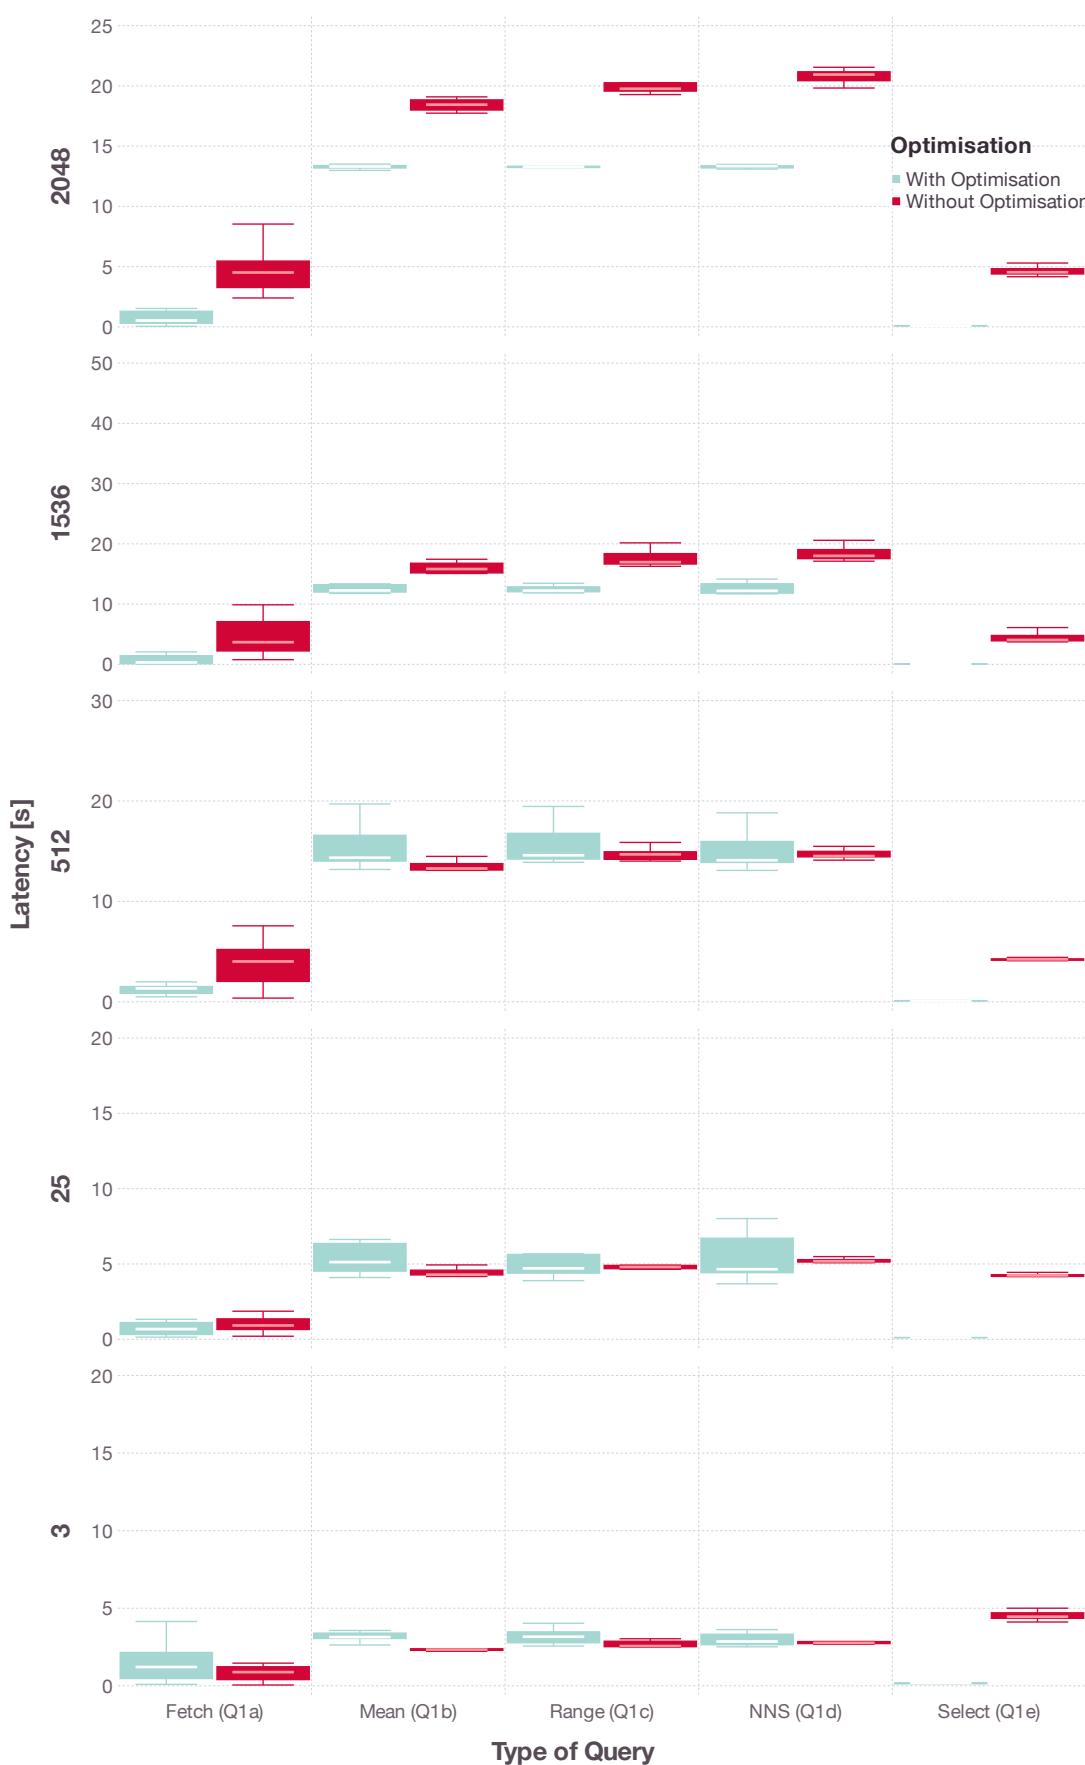


Figure 8.7 Latency in seconds on different collections (y-axis) for the query workloads Q1a through Q1e (x-axis) with (mint) and without (red) query optimisation.

8.3 Large-Scale Similarity Search

This experiment is a direct comparison of Milvus and Cottontail DB and is about mere performance. We execute all queries on prepared shards of the Deep1B [BL16] dataset ($d = 96$), with $k = 1000$ using different execution strategies and compare the obtained metrics. The shards contain 5 million, 10 million, 100 million and 1 billion 96-dimensional float vectors and are stored in dedicated collections (Milvus) or entities (Cottontail DB). The three types of queries are: (i) A simple NNS query that only returns the primary key and the distance (Q2a), (ii) a NNS query that additionally returns the result vector (Q2b), and, (iii) a NNS query with a Boolean filter (Q2c, hybrid query in Milvus terminology). The Pseudo-SQL is provided in Listing 8.2. We use the query vectors provided with the Deep1B dataset and establish a ground truth by executing a brute-force search.

```

/* Q2a: Simple NNS without feature vectors. */
select id, euclidean(feature, <query>) as dst from <collection>
    order by dst limit 1000

/* Q2b: NSS that returns feature vectors. */
select id, feature, euclidean(feature, <query>) as dst
    from <collection> order by dst limit 1000

/* Q2c: Hybrid query without feature vector. */
select id, euclidean(feature, <query>) as dst from <collection>
    where category = <category> order by dst limit 1000

```

Listing 8.2: Pseudo-SQL of the queries executed for this experiment.

8.3.1 Cottontail DB

For this benchmark, we prepared a selection of indexes on the respective entities in Cottontail DB: Two variants of a PQ index, one organised as a list for exhaustive search (8 subspaces, 256 centroids) and one organised as an inverted list for approximate search (8 subspaces, 256 centroids and 256 coarse centroids) and a VAF index. Furthermore, we created a B^+ -Tree index on the column `categories`, which we expect to be beneficial for Boolean search. We again use query hints to nudge Cottontail DB into using certain high-dimensional index structures to be able to compare the different query execution strategies. Across all workloads, we compare the performance of four different types: Sequential scan and the scan of the VAF and the (IVF-)PQ indexes.

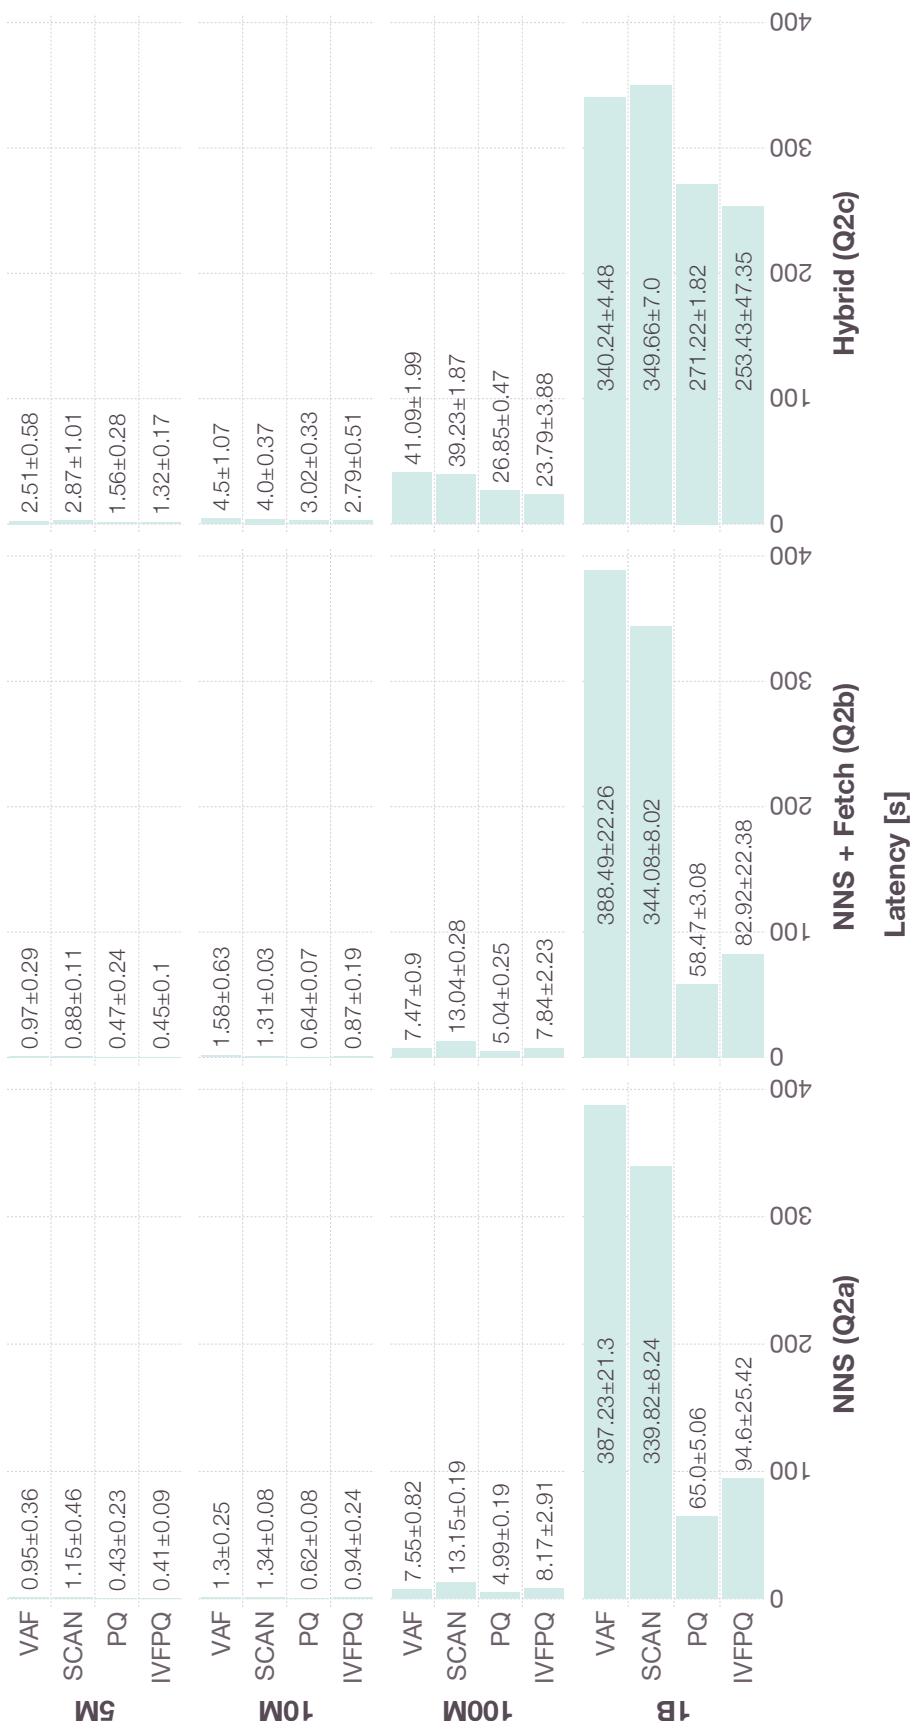


Figure 8.8 Cottontail DB's latency in seconds for different workloads (x-axis) on different shards of the Deep 1B dataset using different execution strategies (y-axis).

In our first experiment, we executed the simple NNS workload (Q2a). The results are depicted in Figure 8.8. Brute-force search took between 0.95 ± 0.36 s (5 million) and 339.82 ± 8.24 s (1 billion). We can observe, that the VAF index brings only minor advantages for the smaller datasets but reduces execution time by almost 4 s for the 100 million entries dataset. This is an artifact of the query execution engine, which uses less than the assigned 32 workers for the smaller collections because of the derived CPU costs. Unfortunately, VAF seems to perform poorly for the 1 billion entries dataset, where it is outperformed by an entity scan by more than 40 s. We believe this to be due to the many random accesses necessary for entries that could not be filtered out. If one considers a filter efficiency of approximately 90% as advertised by [WSB98], one must still fetch 100 million entries through random access, which seems more expensive than simply scanning the entire collection.

We can also clearly see, that using the PQ and IVFPQ index reduces execution time significantly, but at the cost of impaired quality, which is below 0.5 on average for both recall and NDCG, as can be seen in Figure B.1 (see Appendix B). Similarly to the analytics workload, we observe quite a large variance between individual queries. We expect, however, that the overall performance can be optimised by tuning the hyperparameters used upon index construction, specifically, the number of coarse and fine centroids. It is also worth noting, that the IVFPQ index does currently not support intra-query parallelism due to Cotontail DB’s partitioning model, i.e., the execution times we see for the IVFPQ index (94.6 ± 25.42 s for 1 billion entries) are always single-threaded. The execution plans are depicted in Figure 8.9(a) and are fairly unsurprising: The use of VAF and PQ constitutes a class 3 resp. class 1 index replacement according to Definitions 6.11 and 6.9. The fetching of the id columns is pushed down to after the sort and limit operations, since this significantly reduces the amount of IO (only 1000 fetches instead of scanning millions of entries).

Our second experiment involved the combined NNS and fetching of the resulting vectors. One can see in Figure 8.9(b) that in terms of execution plan, accessing this additional column does not make a difference. Since both the plan involving the entity scan as well as the VAF index scan produce the feature column early on, there is no need to fetch it later. Only for the PQ index – which uses a distance approximation that is obtained without accessing the feature column – an additional fetch is required. Consequently, the execution times are very similar to the simple NNS as one can see in Figure 8.8 (NNS + Fetch). The numbers for a sequential scan range between 0.68 ± 0.06 s (5 million) and

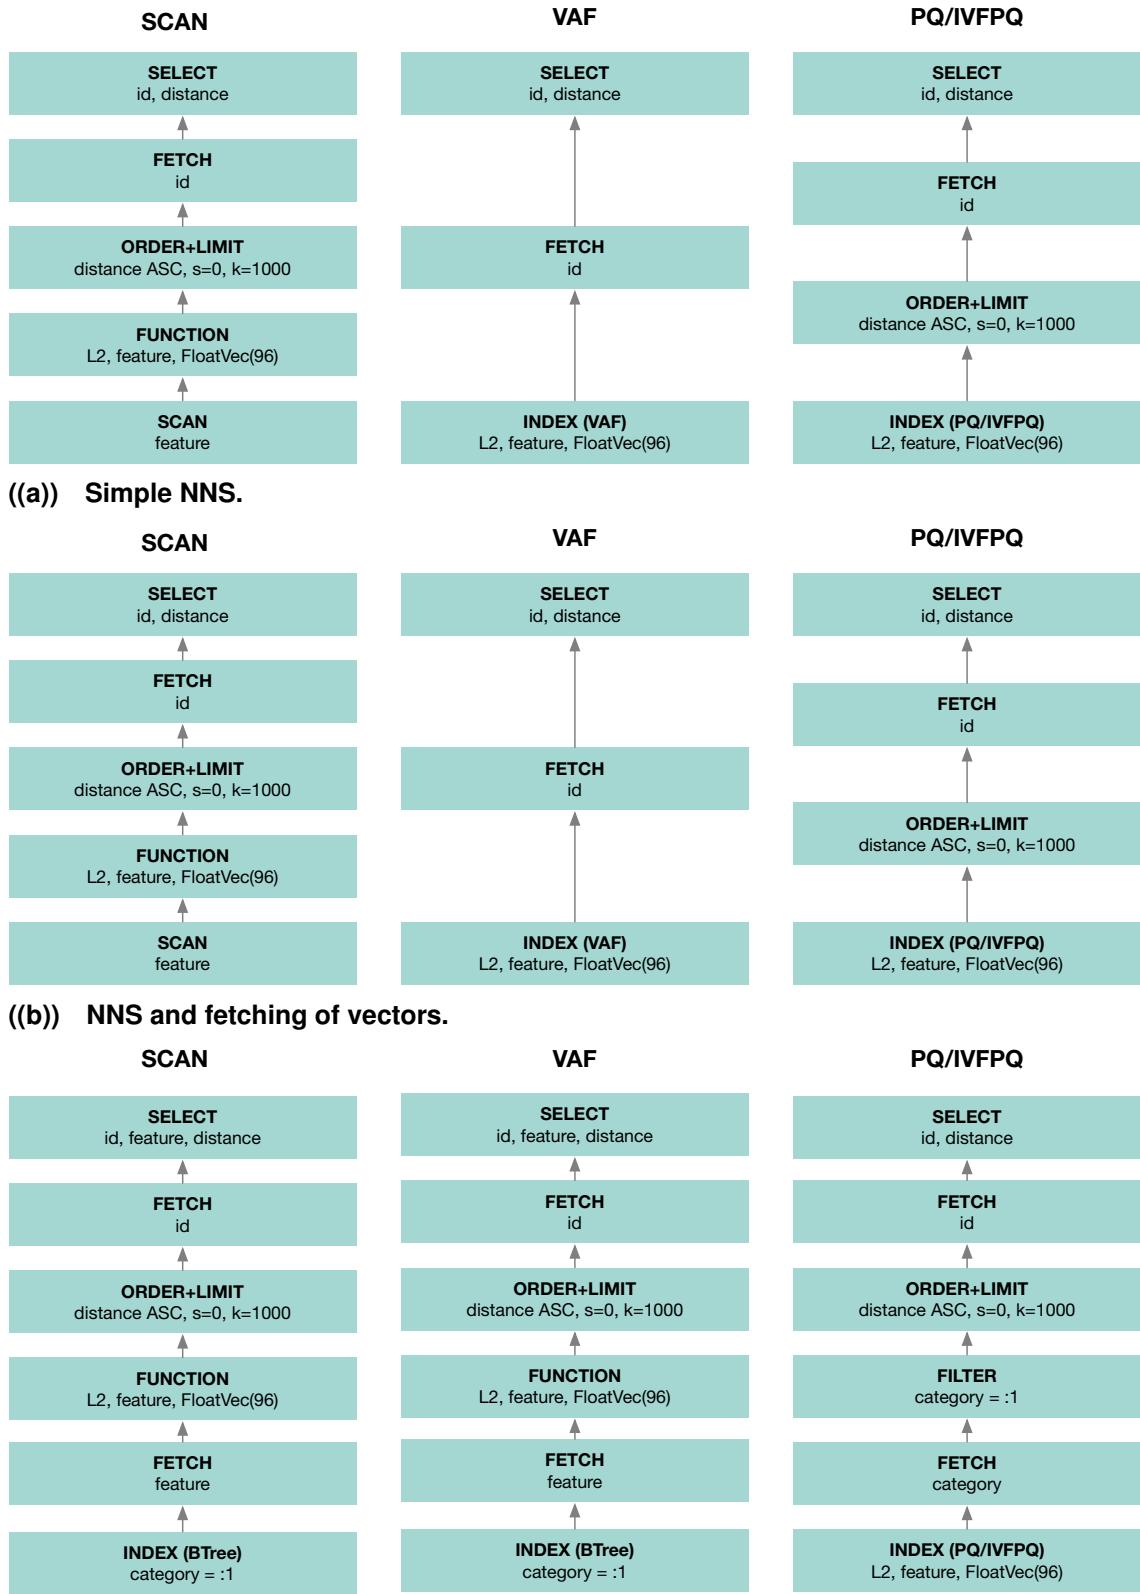


Figure 8.9 Execution plans produced by Cottontail DB for different query workloads prior to intra-query parallelisation. The use of the VAF and PQ index was enforced using query hints.

344.08 ± 8.02 s (1 billion). The impact of the additional fetching of 1000 feature columns is negligible for the PQ index, which can be attributed to the fact that the fetch is pushed until after the s, k -selection by the optimiser. We expect the impact of this fetching step to become more important for larger values of k but it should remain negligible since typically, $k \ll N$. The quality metrics shown in the appendix (Figure B.2), are comparable to those of the first experiment, exhibiting the same low numbers and high variance.

For the final experiment, we executed hybrid queries, i.e., NNS was restricted to a subset of the data that matches a Boolean predicate. The predicate involves a simple equality check that should effectively limit the exhaustive NNS to roughly $\frac{1}{10}$ of the collection. The values for a sequential scan range between 1.47 ± 0.15 s (5 million) and 349.66 ± 7.0 s (1 billion) with very similar numbers for the VAF and PQ indexes. To make sense of these values, we again turn to the query plans illustrated in Figure 8.8 (Hybrid). The sequential scan was executed with the help of a B^+ -tree index, which speeds up the Boolean filtering. Unfortunately, the current implementation of the VAF index cannot natively accommodate Boolean predicates, since it constitutes a class 3 index replacement according to Definition 6.11 and therefore, a special implementation of the index would be required.⁷ Consequently, a B^+ -tree scan was executed instead. In contrast, the PQ index can be combined with Boolean predicates, since it constitutes a class 1 index replacement according to Definition 6.9. However, one can see in Figure 8.8 (NNS + Fetch) that the advantage of using the PQ index is negligible. This can be attributed to the fact that a fetch operation for the category column must be executed for every entry before executing the filtering. The IO overhead is much larger than the speed-up gained through the PQ index. This is confirmed by the numbers we see for the IVFPQ index, which restricts the scan to a small subset of the collection. Similarly to the IVFPQ index, the B^+ -tree index does currently not allow for parallel evaluation due to Cottontail DB's partitioning model. This is severely limiting for the hybrid queries and something that should be addressed in the future. The NDCG and recall values are again shown in Appendix B (Figure B.3).

In summary, all modes of operation fail to deliver interactive latency for very large collections, which currently is the price that must be paid for being bound to disk. It is to be noted, however, that Cottontail DB does not employ any caching other than a page cache. We presume that such advanced caching of, e.g., signatures could result in a considerable speed-up.

⁷ It is, however, unclear if such an implementation would be beneficial in practice.

8.3.2 Milvus

For executing queries in Milvus, the mode of operation is slightly different than it is for Cottontail DB, since Milvus requires a data collection to be available in main memory, to be able to query it. Therefore, every query consists of two steps: Loading the collection and then executing the query once the collection is ready. We have obtained the elapsed time for loading and query execution separately, so that we can consider the individual components. The way queries can be specified is outlined in Listing 8.3. We show the Python instead of the Java syntax, because it is less verbose and more readable, assuming, that the functionality of the two client libraries is identical.

```
from pymilvus import Collection

# Load an existing collection.
collection = Collection("images")
collection.load()

# Perform search.
query = [[0.0, 0.0]]
params = {"metric_type": "L2", "params": {"nprobe": 10}}
results = collection.search(
    data=query, anns_field="feature", param=params, limit=10
)
```

Listing 8.3: Example of a similarity search query in a collection “images” using a two-dimensional query vector and the Euclidean distance. Before executing the query, the data collection must be loaded.

Across all workloads, we compared the performance of two different types of indexes: The FLAT index is the equivalent of a sequential scan and it is the only index in Milvus that guarantees a recall of 1.0. The IVF_SQ8 index uses an inverted file of clusters and limits the search to a subset of these clusters based on the parameters provided by the user and an initial distance calculation between the query and the cluster centres. This approach is comparable to cluster pruning [CPR⁰⁷] or the two-stage quantisation process described in [JDS10], where each vector is mapped to an inverted list using a coarse quantiser. In addition to limiting the search space, the IVF_SQ8 index also compresses the 4 B (float) into a 1 B (int8) presentation, which significantly reduces memory and CPU usage by up to 70%. We built the IVF_SQ8 index beforehand using 4096 clusters and instructed Milvus to consider 1024 and 2048 (query parameter nprobe) clusters respectively, when searching the collection.

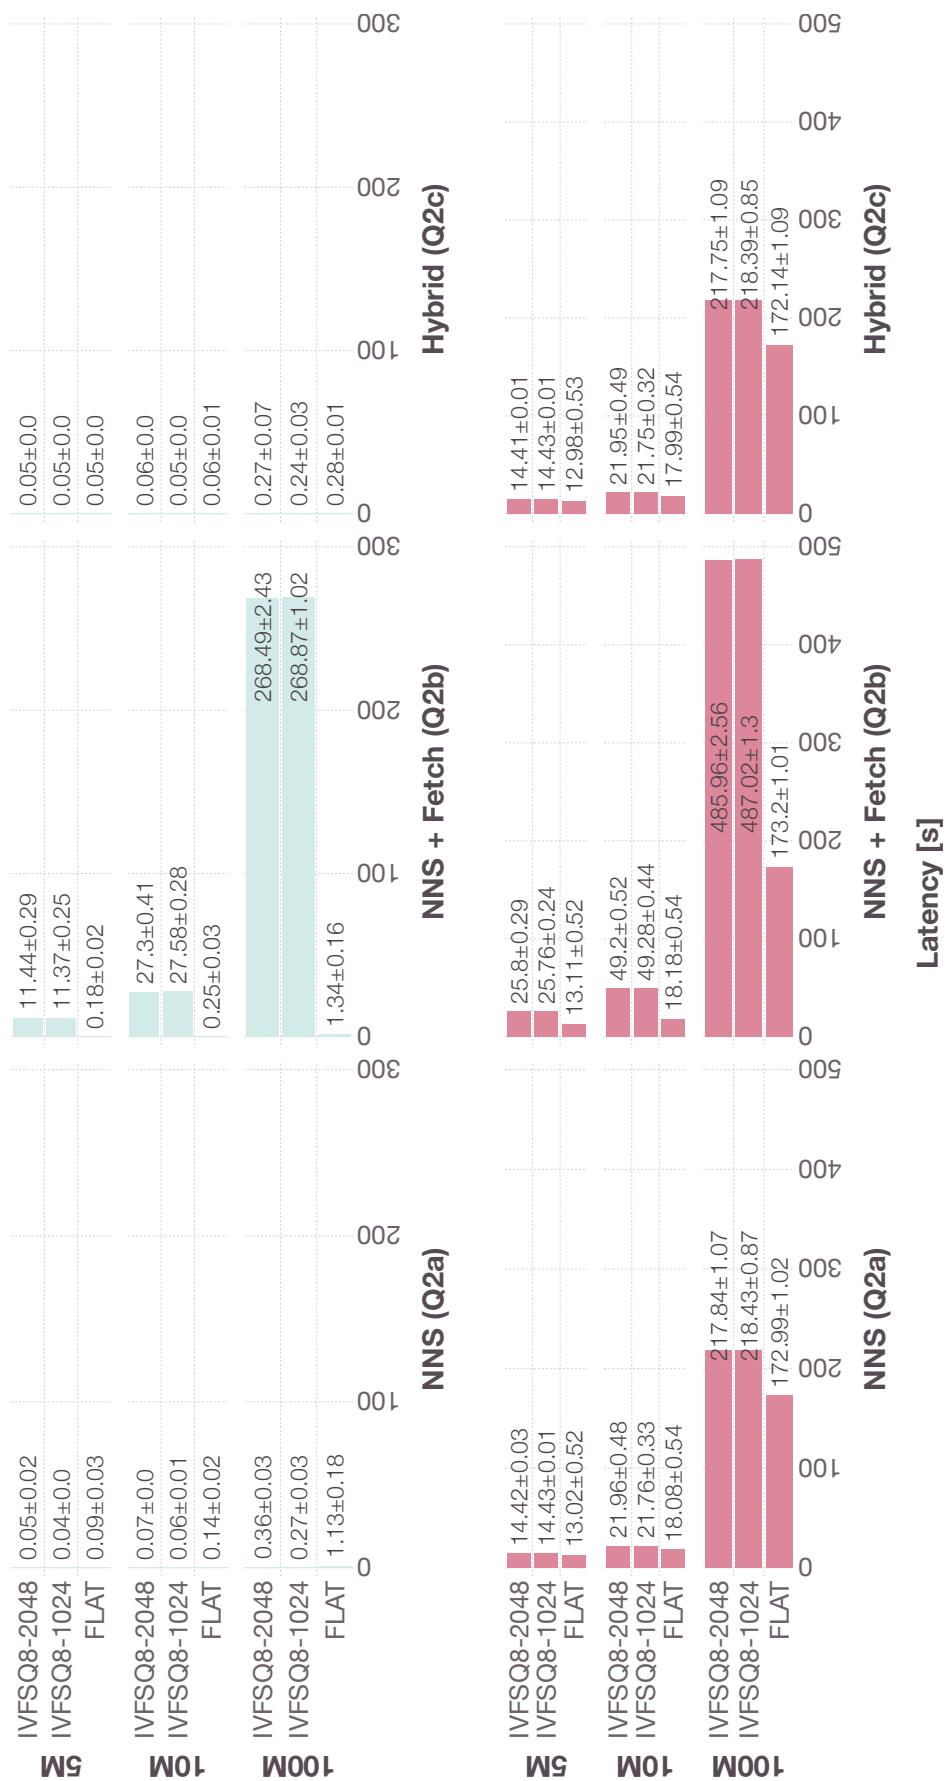


Figure 8.10 Milvus's latency in seconds for different workloads (x-axis) on different shards of the Deep 1B dataset using different access methods (y-axis), with (red) and without (mint) time used for loading the data from disk.

In our first experiment we executed the simple NNS workload (Q2a). The results are depicted in Figure 8.10. We can summarise that even for brute-force search (FLAT), query execution speed is very impressive once a collection has been loaded into main memory. For our test collection, it took between 0.09 ± 0.03 s (5 million) and 1.3 ± 0.18 s (100 million) to execute a query, all the while retaining a perfect recall and DCG of 1.0. However, if one considers time required to load a collection from disk to be part of the query time, these numbers rise to between 13.02 ± 0.52 s and 172.99 ± 1.02 s. If one loads a collection once to then execute many queries thereafter, the performance provided by Milvus is very desirable as the cost of loading the collection can be amortised over time. However, if a large number of collections must be queried without the ability to predict which one, and not all collections can be pre-loaded, query execution time is expected to deteriorate. Furthermore, the requirement to load the collection prior to querying it turned out to be an insurmountable roadblock for the shard that contained 1 billion entries. Milvus was unable to load the collection as it ran out of available memory. We can confirm that Milvus indeed used up the entire 376 GB of free RAM on the node.

In an attempt to alleviate the memory pressure, we also considered the IVF_-SQ8 index, which seems a logical choice due to its data compression characteristics. Unfortunately, using this index instead did not resolve the problem of collection loading, since, Milvus always loads the indexes and the original data. Nevertheless, there are two interesting insights from the measurements: (i) IVF_SQ8 brings considerable speed-up for raw query execution time, especially for larger collections (roughly 1 s faster on 100 million shard), (ii) however, it exacerbates the performance impact of collection loading, since the indexes must be loaded in addition to the raw data.

The second experiment involved the combined NNS and fetching of the resulting vectors. Unfortunately, Milvus currently does not support the returning of vectors in an NNS, which is why this is a two-step process. First, the NNS is executed to subsequently look up the vectors for the resulting primary keys in a second query. We measure the time for both steps. The results are depicted in Figure 8.10 (Q2b). The additional fetching step, while cumbersome, does not seem to have a negative impact in cases where the FLAT index is used. However, use of the IVF_SQ8 index seems to lead to a significant deterioration of the query performance for the fetching step, adding between 10 s (5 million) and more than 250 s (100 million) to the total query execution time. We do not have an explanation for this and consider it to be abnormal behaviour.

Last but not least, we did execute hybrid queries, wherein NNS was restricted to a subset of the data that was filtered by a simple predicate. In Milvus, this is a query primitive that is part of the similarity search API. The results are visualised in Figure 8.10 (Q2c) and do not hold any surprises. In-memory query execution speed is again outstanding whereas the time required to load a collection is significant. Interestingly, the IVF_SQ8 did not seem to have a negative impact on query execution speed, which confirms our suspicion that the deterioration observed in the second experiment must be considered a bug.

We refrain from visualising NDCG and recall values but can report that both always remained at 1.0 for FLAT and around 0.99 for all queries that used the IVF_SQ8 index.

8.3.2.1 Qualitative Assessment

While very convincing in terms of sheer execution speed, the current version of Milvus also has limitations that we list here for future reference:

- Milvus currently only supports the Euclidean and Inner Product distance for single-precision floating point embeddings. There is no way to use other metrics or data types.
- Milvus does not support proximity-based search strategies other than NNS. However, according to the official GitHub issue tracker,⁸ range search is due for the 2.2.0 release.
- Milvus cannot retrieve the feature vectors as part of an NNS query, i.e., they must be fetched in an additional lookup step based on the primary key. This issue is also due to be addressed in the 2.2.0 release.⁹
- Milvus can only maintain a single index per feature column. This effectively limits the set of available distance functions to one, since most indexes must be trained for a specific distance.

In addition, we have observed that as Milvus uses up all the memory during collection loading, it behaves erratically making it difficult to unload the collection again. In some of our runs, Milvus started to reload the same collection after being restarted after a crash, effectively rendering the entire instance unusable for several hours.

⁸ See <https://github.com/milvus-io/milvus/issues/17599/>; accessed August 2022.

⁹ See <https://github.com/milvus-io/milvus/issues/16538/>; accessed August 2022.

8.4 High-Dimensional Index Maintenance

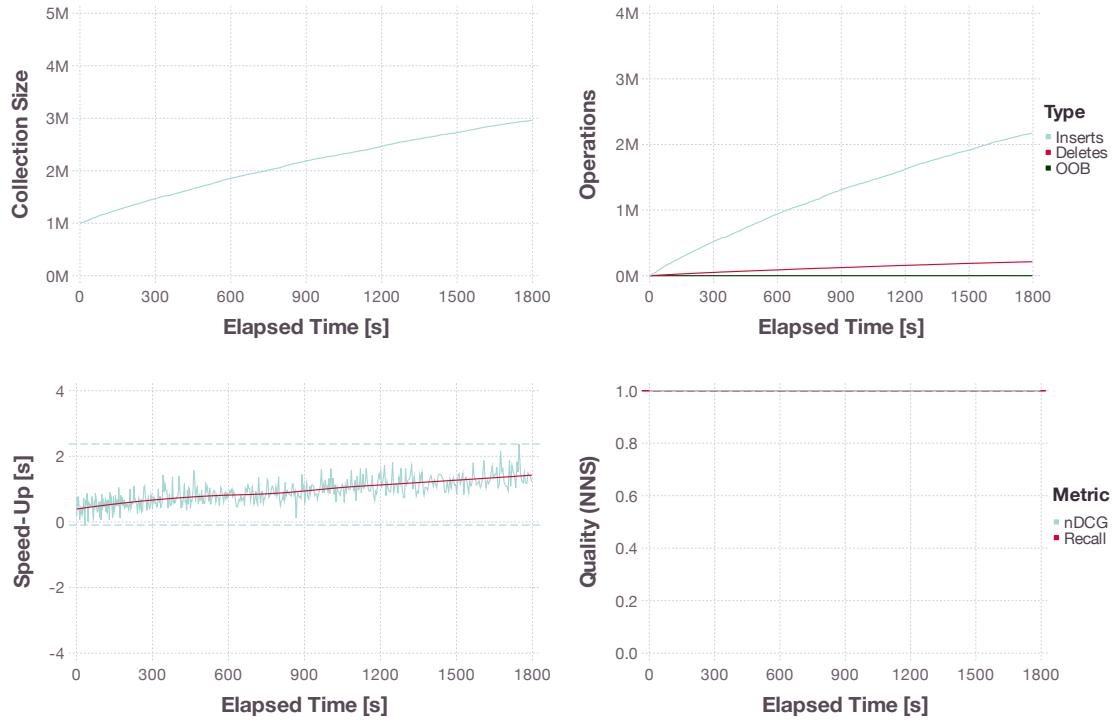
With this evaluation, we demonstrate Cottontail DB’s ability to process incremental changes to data collections that hold a high-dimensional index structure. It is a test of the index maintenance mechanism described in Section 6.3 and of the transaction support Cottontail DB provides. The setup of this experiment is slightly different than that of all previous ones, in that it involves multiple workloads that run concurrently. The evaluation starts by creating an entity and loading it with 1 million entries. We again use the Yandex Deep 1B ($d = 96$) dataset as a basis here and we start with the first one million vectors, i.e., the state of the collection at the beginning is identical for all runs. After data loading has completed, we generate and build an index – either a VAF (35 marks per vector component) or a PQ index (8 subspaces, 1024 centroids) – and start the actual measurement once it is ready. We run three threads simultaneously:

The first thread runs a loop that performs a batch insert operation per cycle. For a single insert, we randomly draw between 100 and 5000 entries from the Yandex Deep 1B dataset. The range can be calibrated by a user-defined factor $c_{\text{ins}} \in [0.0, 1.0]$. The second thread runs a loop that randomly deletes between 100 and 5000 entries per cycle based on their ID. We use the sequential nature of the ID and a blacklist to make sure that every ID is only selected once to prevent deletes that have no effect. Again, the range can be calibrated by a user-defined factor $c_{\text{del}} \in [0.0, 1.0]$. Both the insert and the delete takes place in a dedicated transaction. Between cycles there is a pause of 50 ms to 500 ms to simulate operations that arrive randomly.

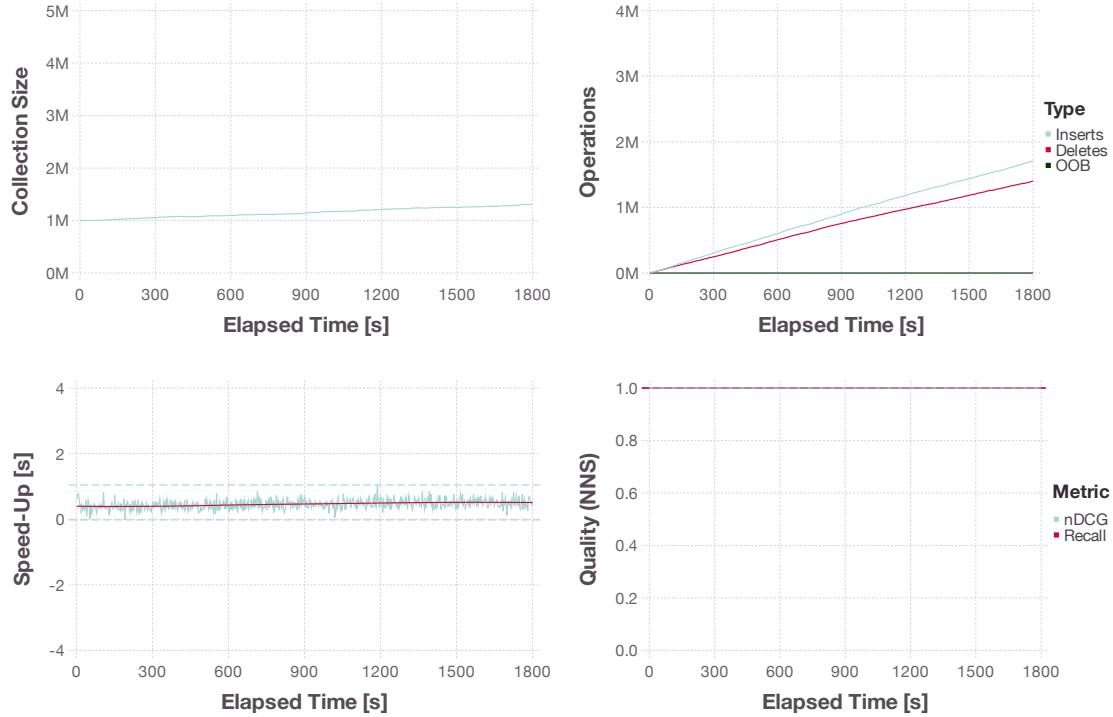
The third thread continuously executes NNS queries using the query vectors from the Yandex Deep 1B ground truth. One NNS query explicitly uses the previously created index whereas the other query employs a brute-force scan to obtain the ground truth. Furthermore, we prohibit intra-query parallelism to reduce noise. In contrast to previous experiments, we do not average over multiple queries and do not execute warm-ups. To make sure that the two queries operate on the same state of the database, they are executed in a single transaction.¹⁰ We also obtain the query plan to verify that the index is used.

In our first experiment, we run the workload for 30 min using different factors, c_{ins} and c_{del} , to steer the ratio of inserts and deletes. While it runs, we gauge how the collection, the speed-up, i.e., the difference in latency between brute-force and indexed search for NNS, and the quality of the results develop.

¹⁰ If we were to execute the queries without an explicit transaction, we would expect differences between the two results due to interleaving changes.

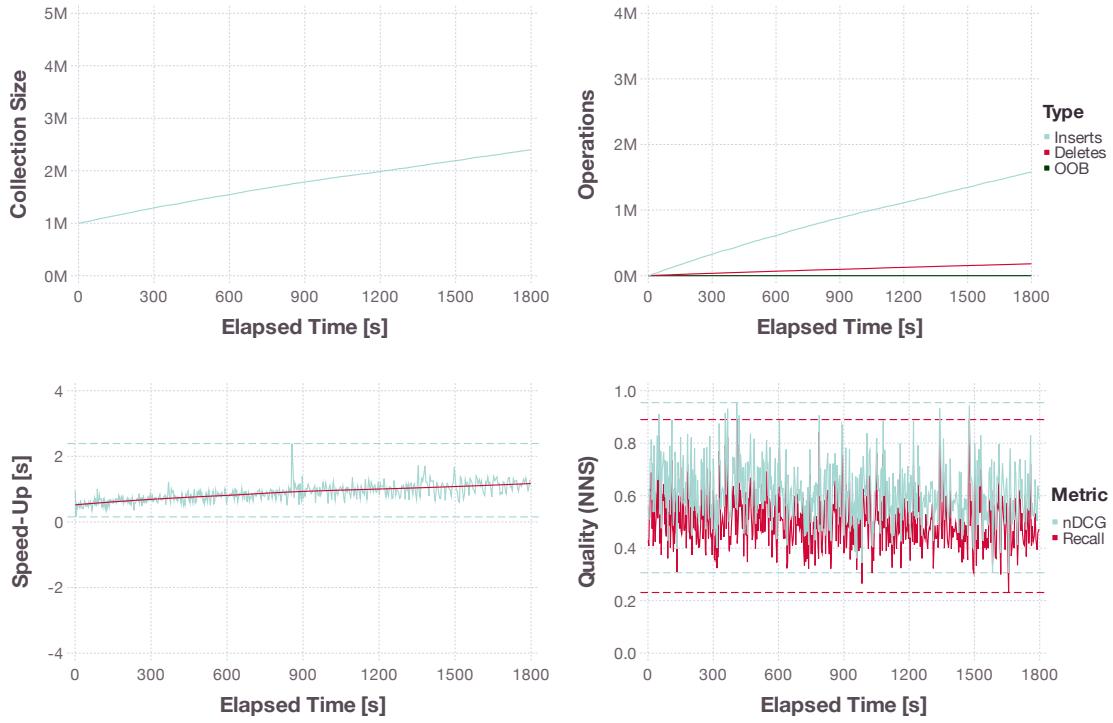


((a)) Index adaptiveness measurements for VAF index; for $c_{ins} = 0.9$, $c_{del} = 0.1$.

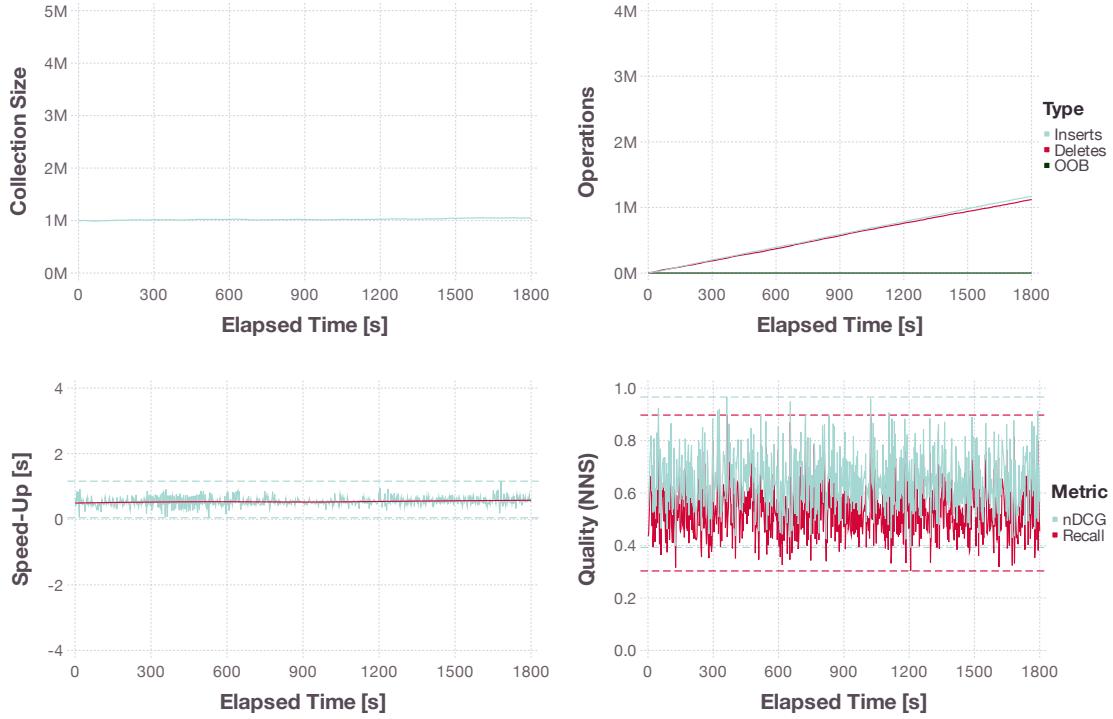


((b)) Index adaptiveness measurements for VAF index; for $c_{ins} = 0.5$, $c_{del} = 0.5$.

Figure 8.11 Index adaptiveness measurements for VAF index for two different values of c_{ins} and c_{del} , showing the size of the collection (top left), the number of accumulated insert and delete operations (top right), speed-up for NNS using the index (bottom left), and quality of the NNS query with respect to ground truth (bottom right) over time (x-axis).



((a)) Index adaptiveness measurements for PQ index; for $c_{ins} = 0.9$, $c_{del} = 0.1$.



((b)) Index adaptiveness measurements for PQ index; for $c_{ins} = 0.5$, $c_{del} = 0.5$.

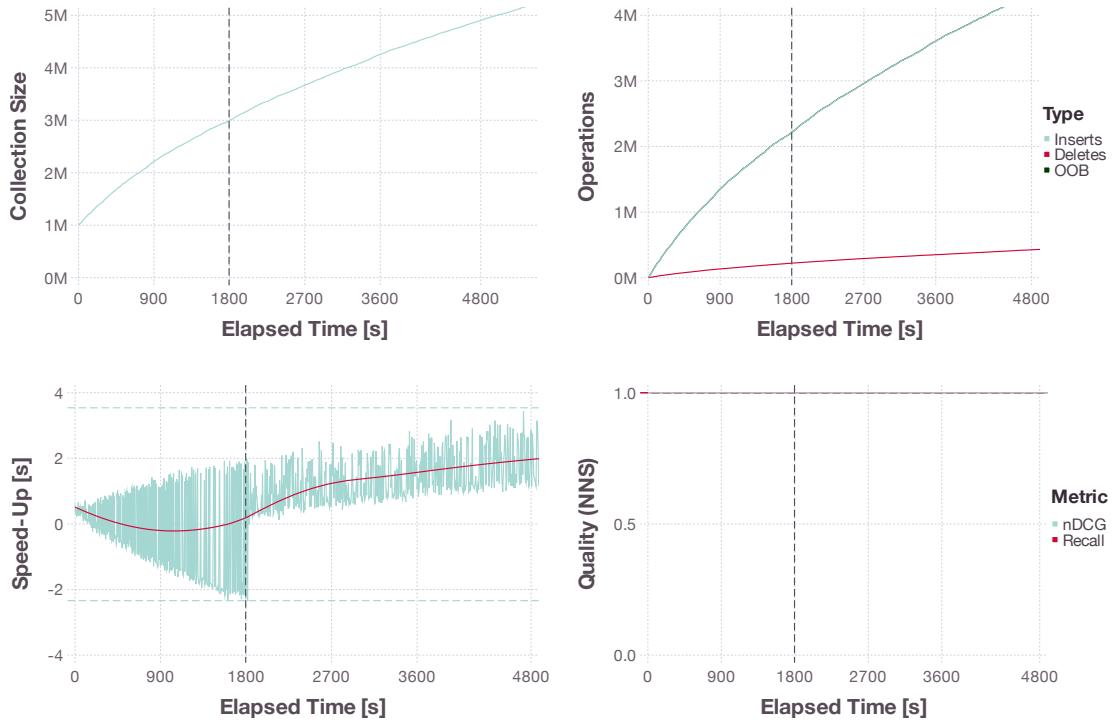
Figure 8.12 Index adaptiveness measurements for PQ index for two different values of c_{ins} and c_{del} , showing the size of the collection (top left), number of accumulated insert and delete operations (top right), speed-up for NNS using the index (bottom left), and quality of NNS with respect to the ground truth (bottom right) over time (x-axis).

The plots for the first experiment are depicted in Figure 8.11 (VAF) and Figure 8.12 (PQ) and we can make several observations: For the workload with $c_{\text{ins}} = 0.5$ and $c_{\text{del}} = 0.5$ we see that the collection size remains more or less constant and with it both speed-up and quality delivered by the VAF (Figure 8.11(b)) as well as the PQ (Figure 8.12(b)) index.

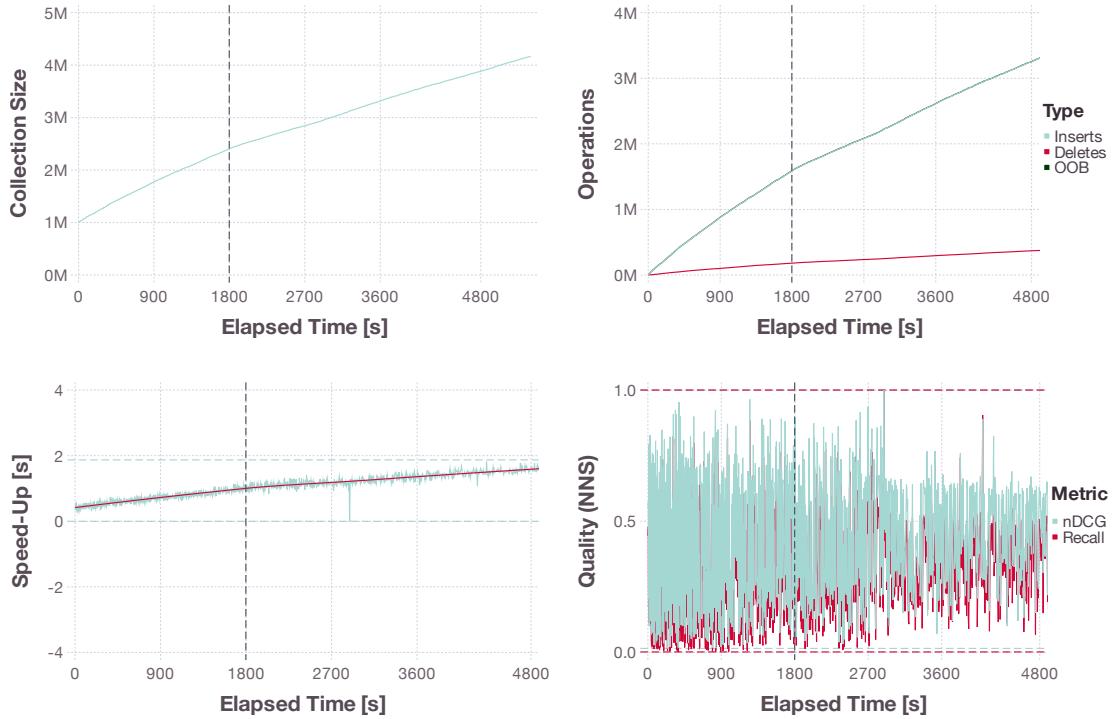
In contrast, speed-up for executing NNS shows a slight upward trend for both the VAF (Figure 8.11(a)) and PQ (Figure 8.12(a)) for $c_{\text{ins}} = 0.9$ and $c_{\text{del}} = 0.1$. This is to be expected, since the indexes are designed to scale better with the linear increase of collection size. We cannot observe any visible deterioration in terms of speed-up building up over time. Similarly, both recall and NDCG remain steady for both indexes over the course of the measurement despite the concurrent write operations, which is the desired and thus expected behaviour. On an unrelated note, we can also see how unpredictable and noisy the quality of the results produced by the PQ index is, which assumes values of between 0.9 and 0.2. It is also to be noted that growth of the collection is slower for PQ despite being governed by the same benchmark logic. We suspect this to be a consequence of the higher overhead incurred by generating the PQ signature upon insert. This also explains why there is a slight net-growth for VAF $c_{\text{ins}} = 0.5$ and $c_{\text{del}} = 0.5$ workload, as we generally expect deletes to be slower than inserts, which indeed seems to be the case for VAF (but not for PQ).

We have argued in Section 6.3, that we expect the performance of an index to depend on how well the data fits the (unknown) distribution of the corpus the index was built upon. To get an indication of this, we have also recorded the component-wise minimum and maximum at the start of the run (i.e., of the first 1 million entries). Every time a new vector is inserted, we check if a component exceeds that minimum and maximum and record such entries as out-of-bounds (OOB). For this first experiment, wherein we use the collection in its unaltered state, we can spot no discernible growth in these entries and indeed the data suggest that only a few 100 appear among the first 2 million inserts. This is an indication albeit not a proof of the validity of our assumption.

To test our method with data that does not exhibit this well-behavedness, we run a second experiment in which we introduce noise. This *jitter* is simply a random number of configurable magnitude M that is added to each vector component prior to insert. We run this experiment for 90 min and randomly apply the same transformation to 50% of the query vectors, i.e., some queries exhibit the same noise while others do not. After 30 min we trigger an asynchronous rebuild of the index structure to observe its effect.



((a)) Index adaptiveness measurements for VAF index; for for $c_{ins} = 0.9$, $c_{del} = 0.1$, $M = 100$.



((b)) Index adaptiveness measurements for PQ index; for for $c_{ins} = 0.9$, $c_{del} = 0.1$, $M = 100$.

Figure 8.13 Index adaptiveness measurements for the VAF and PQ index, showing the size of the collection (top left), the number of accumulated insert and delete operations (top right), speed-up for NNS using the index (bottom left), and quality of the NNS query with respect to the ground truth (bottom right) over time (x-axis). Data and queries exhibit a random noise and index rebuild is started after 30 min (grey, dashed line).

The plots for the second experiment are depicted in Figure 8.13 and we break them down for the two indexes: For VAF (Figure 8.13(a)), we see constant quality of 1.0 over the entire measurement – both for NDCG and recall. This is expected since VAF should retain quality guarantees even for very noisy data. However, speed-up provided by the index deteriorates significantly during the first 30 min to the point where it is actually slower than the linear scan. Furthermore, we observe a massive increase in variance between individual queries. This can be attributed to the random noise that is added to some query vectors. Queries that do not exhibit that noise are still filtered out by the index quite effectively (between 85% and 95% efficiency). In contrast, the filtering capacity continuously drops for noisy queries as noisy data accumulates. This can be explained by how the VAF works. Since the index still reflects the original distribution without noise, vectors that match this distribution can be distinguished with a constant resolution and will be assigned to one of the m^d cells (m denoting the number of marks per dimension) leads to 4.81×10^{16} possibilities. All the noisy vectors accumulate on the edge of the high-dimensional lattice, of which there are only $2 * m * d$ (which leads to 7720 possibilities). As soon as a noisy vector is being queried, it is very likely to end up in one of the $2 * m * d$ cells which, in turn, is likely to be occupied by a large number of vectors, all of which must be fetched from disk to perform a distance calculation. This is why this effect increases as the benchmark proceeds. After 30 min, we trigger the asynchronous re-indexing and it finishes shortly thereafter. The effect becomes visible immediately, as speed-up and especially the variance decrease significantly. The index rebuild itself has a slight effect on the concurrent DML operations, which take more time but can continue normally.

For the PQ index, the deterioration manifests itself in the quality, which now oscillates between 0.9 and 0.0. The explanation is similar to the one for the VAF: The old index is still capable of assigning query vectors that do not exhibit any noise to their proper centroid and can therefore reliably approximate the distance for them. Noisy vectors get assigned to some centroid but the distortion in distance becomes too big for reliable NNS. After 30 min, we also trigger the asynchronous re-indexing. In the case of PQ, the overhead incurred is clearly visible in the throughput for both insert and delete, which otherwise continue normally. After about 52 min, we see a normalisation in the quality metric in that the variance is less severe. However, quality is still far from optimal even after the rebuild has concluded, which can be attributed to the parametrisation of the index.

8.5 Summary, Related Work and Discussion

With the evaluation we have demonstrated the applicability of the concepts introduced in Chapter 6 to practical scenarios using our reference implementation, Cottontail DB. Even though Cottontail DB was unable to compete with Milvus in terms of raw processing speed, we consider this evaluation to be a success. Maintaining and accessing all the relevant data on disk incurs an overhead that makes it difficult or even impossible to reach the processing speed of pure in-memory systems like Milvus. This was pointed out by L. Amsaleg, who stated: *“When purposely dealing with secondary storage, it is hard if not impossible to beat high-dimensional indexing solutions that have been designed to run in main memory. This is unfortunately true even for collections that might fit in main memory, as enforcing durability etc. generates a substantial overhead”* [Ams14] (p. 12). Regardless of these limitations, Cottontail DB still seems to perform pretty well against Milvus, especially when the time required for loading data from disk is considered as well. In our experiments, during the time that Milvus required to load a collection, Cottontail DB could query that same collection between 15 (VAF) to 34 (PQ) times, all the while retaining the flexibility to switch between collections and/or execute different types of workloads that Milvus could not (e.g., range search). Our experiments have also demonstrated the limitations of pure in-memory processing in cases where collections simply do not fit into main memory, regardless of how optimised the execution engine may be.

8.5.1 Generalised Proximity-Based Queries

The notion of generalised, proximity-based operations as an extension to the relation algebra described in Section 6.1 provides us with the ability to construct and execute a wide range of query workloads in a much more flexible way than purpose-built systems such as Milvus [WYG⁺21] can provide. Furthermore, the proposed algebra enables the DBMS to decompose and optimise queries based on the needs expressed by a system user and the parameters of the system itself.

In the past decades, many information retrieval models have been proposed [Uma83; SFW83; WZW85]. Some of these models are set-based [SFW83; Uma83], and thus similar to the relational model [Cod70], while others are based entirely on the vector space model that considers documents to be high-dimensional vectors [WZW85]. Our approach uses the relational model as a foundation and integrates the notion of complex data types (e.g., vectors), functions that operate on these types, as well as scoring and ranking as explicit extensions in way that en-

ables its use beyond mere multimedia retrieval. In that it follows an established tradition of extending the expressive power of SQL and the relational model [Lib03], which has been practiced for many years [LCI⁺05; ZHC⁺06; BOV07], as we have argued in Section 3.2.3. Conceptually, we see the proposed extensions as a first, minimally invasive step towards the realisation of the *Extended Boolean Retrieval Model* as proposed in [SFW83]. In such a model, distances and scores can be employed to perform advanced operations such as fuzzy-set operations or similarity joins [Uma83; Zad96; BMS⁺01]. From our perspective and with an eye on analytics workloads, we would also like to see query primitives such a clustering or RNNS [KM00] in that list of supported operations.

The convergence of information retrieval and databases as well as the addition of multimedia support have been a recurring subject at many of the database self-assessment meetings since the 1980s [AAB⁺08]. Unsurprisingly, many attempts at convergence have been made over the years at a conceptual [MS96; AN97; WOK⁺98], logical [ZHC⁺06; BOV07], and physical [SAA⁺10; GAKS14; WLL⁺15; GS16; YLF⁺20] level. We follow that trend and, the model we propose provides a sound, theoretical foundation rooted in relational algebra and has been demonstrated to be useful in practical scenarios.

An important, mental step that enables such simple, yet expressive models is that of “separation of concerns”, or what [Gia18; Fer14] refer to as “pragmatics”, which is in fact realised throughout the *vitrivr* stack [RGT⁺16; GRS19a]. While similarity search and retrieval may be the ultimate aim, from a database perspective, we merely execute simple primitives such as the evaluation of distance functions, sorting, and selection. By carefully ignoring the noise introduced by higher-level concepts such as documents, features, scores, and distances – which may not be useful for data processing – the database system can focus on its core purpose which is that of providing reliable and fast data management and access and leave the interpretation and advanced data modelling to upper-tier systems. This is a common approach in database research, and with our model we further move in such a direction, similar to computational databases such as SciDB [SBZ⁺13], MonetDB [IGN⁺12], or HypeR [HPS⁺17], which focus on providing support for efficient computation close to where the data resides, without too strictly limiting the concrete application. Such systems form an active research field [LGG⁺18; BGL⁺22] and might hold interesting benefits for the multimedia community. They constitute a stark contrast to systems with very focused applications such as Milvus [WYG⁺21], Vald¹¹, or Qdrant.¹²

¹¹ See <https://github.com/vdaas/vald/>; accessed August 2022.

¹² See <https://github.com/qdrant/qdrant/>; accessed July 2022.

Since ADAM_{pro}[GS16] is Cottontail DB’s predecessor and one of the inspirations for the work presented in this Thesis, we would also like to highlight the differences between its query model and our own. As outlined in Section 5.5, [Gia18] proposes a new similarity operator $\tau_{\delta(\cdot,\cdot),a,q}(\mathcal{R})$ that “performs a similarity query under a distance $\delta(\cdot,\cdot)$ applied on an attribute a of relation \mathcal{R} and compared to a query vector q ” [Gia18] (p. 138). It then uses this operator to derive more specific versions that basically consider NNS and range search (ϵ NN) to be dedicated query primitives. As we have argued in Section 6.1, this model is very restrictive in its assumptions and to some extent violates the aspect of pragmatics as proposed by both [Fer14] and [Gia18].

In contrast, our model considers the distance computation to be the result of an arbitrary function invocation. We restrict ourselves to certain distance functions in practice but in theory, any function could be employed. Furthermore, we use the extended projection to make the derived distance explicit and we argue, that making the attribute explicit is what allows downstream processing such as normalisation (from distance to score) or combination of multiple scores to express score fusion or variants thereof [BMS⁺01]. Ultimately, we consider it to be at the discretion of the user to decide whether and how calculated attributes should be used and included and it should be the task of the query optimiser to take these decisions into account when preparing an efficient execution plan.

Consequently, the relational algebra extensions we propose are rather limited and lie well within the capabilities most modern DBMS possess today. In line with [AAA¹⁴; AAA²⁰], we identify the efficient execution of the functions and computations involved using the available means in terms of hard- and software (e.g., SIMD, GPU), as important, future research directions. In order to enable fast execution, functions must be enriched with metainformation that allows for optimisation at different stages of the query planning and execution process. In addition, we also believe the relationship between functions and indexing techniques to be a promising avenue to explore.

Nevertheless, other systems, such as Milvus [WYG²¹], take a similar approach to [GS16] as they add query primitives for the different types of workloads they accommodate – leading to a growing number of disjoint APIs for different purposes. We do not consider this a sustainable approach, since it ignores future workloads that may combine existing primitives in a different way (e.g., aggregating on distance or FNS). Furthermore, generating specialised blackbox operators for different types of high-level primitives unnecessarily limits the potential for optimisation by the DBMS and increases system complexity.

8.5.2 Query Planning and Cost Model

In Section 6.2 we have proposed a cost model that takes the expected quality of results into account, and we have discussed methods that can be used to estimate that quality for different, known query plans. Our reference implementation Cottontail DB uses a query planner that leverages the aforementioned cost model to determine the optimal execution path given a query and a cost policy specified by the user, as demonstrated in Section 8.2.2.

To the best of our knowledge, both aspects – the estimation of quality as well as its use in a cost model for query planning – are rather uncommon in multimedia search systems today. Even though the quality versus execution time trade-off is an implicit consequence of approximate solutions to proximity-based search in general and NNS in particular, there is very little research that considers this issue systematically, with a few, notable exceptions: [WB00; BHC⁺01; SJH⁺05]. This is reflected by the fact, that most state-of-the-art approximate index structures fail to provide any strong quality guarantees [EJP21]. Here we see a huge corpus for future research: On the one hand, we see a need for index structures that provide explicit quality guarantees such as [LWW⁺20]. On other the hand, we should continue to explore system-level mechanisms that continuously evaluate the efficiency and effectiveness by which results have been produced and adjust the cost model accordingly. Such an empirical approach to query planning has already been proposed by [Gia18], however, taking only execution time into account. This goes into the direction of adaptive query processing [DIR⁺07] and the cost model we propose can be seen as a foundation for a continuous control loop between the execution engine and the query planner.

8.5.3 High-Dimensional Index Maintenance

Maintaining a durable and consistent state and providing transactional guarantees for data structures required for executing queries – including secondary index structures – is a basic trait shared by most if not all traditional SQL and NoSQL systems. For some reason, however, durability and the notion of isolation and consistency in the face of dynamic data are aspects often overlooked for high-dimensional indexing [Ams14; HJB19], especially in combination [Ams14]. In Section 6.3, we have proposed a system-level mechanism for maintaining high-dimensional index structures on disk and for processing incremental changes to the data in a transactional manner. The viability and limitations of the proposed mechanism were demonstrated in Section 8.4.

Despite an already huge corpus of research that dates back to the 1970s [Ben75; Gut84; BKS⁺90; IM98; WSB98; JDS10], high-dimensional indexing for NNS is still a very active research topic today [SOV⁺21; KBC⁺18]. Since different index structures provide different properties in terms of scalability, performance, and quality (guarantees) and can potentially be employed for different tasks, we regard this work as pivotal for further extending the indexing capabilities of a multimedia DBMS such as Cottontail DB. One could argue that having a single index may be sufficient, however, we believe that in light of a DBMS’s ability to plan queries and of the different properties these indexes bring, having a choice is better. Furthermore, multiple indexes could also be combined to improve the quality of results and the speed of execution [Gia18].

In light of ever-growing data collections, the topic of disk-based ANNS has been getting more attention lately, with a few, early examples [LJA11; LÁJ⁺09; GJA10] and new, emerging techniques such as DiskANN [JSDS⁺19] or SPANN [CZW⁺21], which both advertise a hybrid in-memory/on-disk processing model and report impressive numbers for both latency and recall. However, in contrast to [LJA11; LÁJ⁺09; HJB19], both these state-of-the-art algorithms do not consider dynamic data and operate on the premise that the complete data collection can be (re-)indexed. This severely limits their usefulness. Nevertheless, such efforts lay important groundwork for disk-based multimedia index maintenance and must simply be extended in terms of their ability to process incremental changes.

In [LJA11; LÁJ⁺09], the authors propose the NV-tree, which is a disk-resident, high-dimensional index structure that provides transactional support [LÁJ⁺18] and thus allows for inserts, updates, and deletes. The NV-tree would make a great extension to Cottontail DB and we expect it to fully fit into the proposed model for index maintenance in that it is an index structure that does not fail per-se and should not exhibit deterioration due to dynamic data (i.e., it can be considered a corner case, like the B^+ -tree). Unfortunately, however, the index structure itself seems proprietary and the available literature does not provide us with enough information to come up with our own implementation.

The subject of online indexing (i.e., processing incremental changes) has also been addressed by several authors, for indexes such as PQ [XTZ18], LSH [CS15], (e)CP [HJB19] and some graph-based methods [ZWN22]. While the proposed methods lack a notion of index state and transactionality, and (with the exception of [HJB19]) do not consider indexes to be durable data structures, these are still very important contributions towards robust write-models for different types of indexing techniques.

We have also discussed different vector database systems in Section 5.3, which all face the issue of dynamic data. Milvus [WYG⁺21], for example, handles inserts by marking them with timestamps and it allows for configurable consistency levels that differ in the timestamps of entries that are being considered. Hence, it uses a kind of MVCC-based consistency and concurrency control model, similarly to Cottontail DB. In contrast, deletes are not propagated to the indexes and instead enabled by setting bit mask for deleted entries and filtering them on the fly at query time. This method is reminiscent of the “tombstone markers” used in Apache Cassandra, which are thought to be culprits for major performance hits. However, it is not fully clear how the method employed by Milvus scales as changes accumulate. Nevertheless, some concepts employed by Milvus may be very interesting, e.g., the (re)-indexing of chunks rather than the entire data collection at once. Vald advertises an “asynchronous auto indexing” that uses “distributed index graphs” to avoid locking the entire graph during indexing. However, it is not clear what type of guarantees are provided by this mechanism other than preventing a system-wide halt, and generally it seems, that the method is very specific to the NGT index used [Iwa16]. Unfortunately, for Qdrant, we could not find any information as to how dynamic data and concurrency are handled.

8.5.4 A Remark on Benchmarking Multimedia Analytics

The experiment we have presented in Section 8.2 is supposed to be a simulation of what a simple analytics query workload could look like, in that it relies on different types of queries within a single session. Chinchor et al. [CTW⁺10] have already mentioned the “resounding call for datasets” that allow for standardised benchmarking of multimedia analytics systems and techniques. We tend to reinforce this call and add, that a standardisation of benchmarks for testing data management aspects might also be a requirement for future work in this area. Such benchmarks are very common for traditional OLAP and OLTP (e.g., YCSB¹³ or TPC¹⁴), and usually involve a series of queries that one expects in a typical workload. Such a baseline, which should go beyond existing benchmarks such as (Big-) ANN-Benchmark [ABF17; SWA⁺21], would simplify the testing of techniques and systems that tackle data management in the interactive multi-media retrieval and analytics domain.

¹³ See <https://github.com/brianfrankcooper/YCSB>; accessed September 2022.

¹⁴ See <https://www.tpc.org/>; accessed September 2022.

9

One never notices what has been done; one can only see what remains to be done.

— Marie Curie, 1894

Conclusion and Outlook

Inspired by the questions raised in [JWZ⁺16] and the solutions presented by [Gia18], we set out to challenge some assumptions about the storage, management, and processing of data used in multimedia retrieval and analytics workloads. Most importantly, we called into question that (i) multimedia retrieval and analytics only involve NNS in high-dimensional, real-valued vector spaces, (ii) that data collections remain static while they are being queried, and (iii) that the execution path of a query should be an explicit choice of the user. Throughout this Thesis, we tried to take the viewpoint of a database engineer who considers the problems at hand to be data management and processing issues that should be delegated to the underlying DBMS whenever possible [Fer14; Ams14]. At this (preliminary) end of our journey, it has led to several contributions:

1. We specified a query model based on the notion of generalised, proximity-based operations to support multimedia retrieval and analytics workloads on top of relational algebra (see Section 6.1). In our model, we put a special emphasis on the role of functions and Distance Function Classes in query formulation, planning, and execution, and we explore the relationship between these functions and high-dimensional indexes.
2. We proposed a cost model that takes the expected quality of results produced by approximate index structures into account (see Section 6.2). This cost model exposes the often implicit and opaque trade-off between quality and speed to the system and the user, who can act upon it.
3. We described a mechanism that can function as a template for integrating and maintaining high-dimensional indexes in a transactional and durable database system (see Section 6.3). Such an integration is crucial for supporting workloads that rely on changing data.

4. We introduced our open-source reference implementation Cottontail DB, which employs the aforenamed, theoretical contributions and more (see Chapter 7).
5. We presented an evaluation that uses Cottontail DB to demonstrate the viability of our approach (see Chapter 8). One focus of this evaluation was a direct comparison of Cottontail DB to the vector database Milvus [WYG⁺21].
6. We identified important issues and research questions to further the fields of multimedia retrieval and analytics on the one hand and database research on the other, as well as the convergence of the two. These insights also provide guidance for the continued development of Cottontail DB.

Cottontail DB can be considered a classical database in many respects in that it offers traditional database features such as durability and transactionality. Yet its main purpose, until today, has been the storage and management of multimedia metadata and descriptors, the processing of interactive multimedia retrieval query workloads and the serving of data to the user. Cottontail DB has been part of the *vitrivr* ecosystem [RGT⁺16; GRS19a] for several years and its versatile query model has enabled successful participation of *vitrivr* in many iterations of events like the VBS [RAPG⁺19; SAPG⁺20; SGH⁺21a; HAG⁺22] or LSC [SGH⁺21b; HAPG⁺20; HGP⁺21; SGH⁺22].

The work on Cottontail DB and thus the work presented in this Thesis contributes towards integrating efforts in data management aspects of multimedia retrieval and analytics ([VHC⁺19] refers to this as multimedia delivery) on the one hand and database research on the other. In its current form, it constitutes an important foundation for current and future research endeavours in various multimedia applications. In addition, we see a huge potential for using both *vitrivr* and Cottontail DB for practical applications and for educational purposes.

9.1 Motivating Scenarios

In Chapter 2 we have described three scenarios that motivate the work presented in this thesis, namely, interactive multimedia retrieval (Section 2.1), social media analytics (Section 2.2) and signal analysis in medical imaging through MRF (Section 2.3). In this section, we would like to elaborate on how we expect our contributions to benefit those use cases.

The proposed query model can be directly applied to all three applications. In [Hür20; Zih21] we have used it to solve the MRF problem by adding support for complex vectors and an inner product distance function defined for this domain. The same model has also enabled the pose-queries described in [HAG⁺22], which rely on the evaluation of nested (distance) functions and which facilitate a new type of multimedia retrieval query in *vitrivr*. Furthermore, Cottontail DB powered rather experimental systems such as LifeGraph [RBG⁺21], which speaks to its and the query model's versatility. We expect that the query model in its current state (and with future extensions) can be used to facilitate many additional types of workloads required by multimedia retrieval and analytics applications as, e.g., outlined in [SHS⁺17].

The cost model we propose can be used to configure an instance of Cottontail DB to match a specific application using cost policies. For solving the MRF problem, one would probably put more weight onto quality whereas execution speed remains the focus for interactive retrieval operations. We expect the aspect of quality cost to become more important in the future, especially as new methods allow for a more fine-grained control over the quality versus speed trade-off, e.g., with progressive querying proposed by [Gia18] combined with a user's ability to abort queries early and re-formulate them.

The maintenance of high-dimensional indexes in the face of incremental changes is probably most relevant for the social media analytics use case. In this scenario, changes to the data are inevitable since new content is being produced continuously and must therefore be considered during indexing. Furthermore, changes to the analytics models and insights gained during the process can also lead to changes to the data. However, the proposed mechanisms may also benefit more traditional retrieval scenarios especially when considering CRUD applications in which multiple users work on the same data corpus, which they query and alter.

9.2 Future Work

The process of writing this Thesis, the formulation of the theoretical models and the implementation of a system that employs them, has answered several important questions but it has raised even more. Over the course of this project, we came up with many (brave) new ideas concerning both the theoretical foundation and the implementation. We will conclude this chapter by elaborating on these ideas.

9.2.1 Query Model

For decades, much effort has been made in the development of different query models for information retrieval and their combination with the Boolean model. However, only few have actually been implemented and put to the test. We consider the query model we have presented to be a first step in combining two worlds in a pragmatic way that opens new avenues to explore. On the one hand, there are many potentially relevant extensions that lead towards a fully fledged Extended Boolean Retrieval Model [SFW83], such as fuzzy queries [Uma83; BMS⁺01] or similarity joins [YLK10]. One step in realising this could be to consider distances and scores explicit data types that, in addition to basic arithmetics, allow for advanced operations required in this context [SAA⁺10]. On the other hand, a general-purpose multimedia DBMS must also consider query primitives, such as RNNS [KM00], bichromatic RNNS [SRA⁺01] or clustering. Both could prove to be interesting extensions to the proposed framework.

In light of such potential extensions we would like to briefly touch upon the question “*Is a novel multimedia query language needed for the database to fully support multimedia analytics, or is an extension of classic query languages sufficient?*” [JWZ⁺16] (p. 299, RQ2). Considering the work in this thesis, we can state that all of the functionality required can be expressed by extended relational algebra and therefore be easily translated to SQL. The idea of expressing advanced functionality as mere function invocations can be further extended to other concepts such as fulltext search, as, e.g., done by MS-SQL. However, advanced query paradigms such as clustering and RNNS are difficult or impossible to express in standard SQL and may therefore require language extensions similar to, e.g., common table expressions for recursive queries. While some authors may advertise a dedicated “information retrieval query language” [Fer14], we tend to argue from personal experience, that an extension of SQL would be preferable to coming up with a new language, simply because SQL is well established, well known, and easy to use.¹

9.2.2 Multimedia Indexing

We have shown in our evaluation that the ability to cope with changes to the data at an index level may not be sufficient to maintain a constant quality of service, at least for the index structures considered in our experiments. However, we strongly believe that the argument can be extended to other indexes. We

¹ In fact, a SQL interface is among the top feature requests for Cottontail DB.

see a major challenge in tuning hyperparameters of an index to the concrete collections at hand. As it seems, the factors that determine the right choice are not only tied to the dimensionality of the vectors but also to the size of the collection and the distribution of data within it. Until now, this problem has not been systematically considered and finding the set of parameters that yields the best results in terms of performance and quality has been a manual task that relies on the experience of the user and empirical analysis. This may work well if a collection remains static, but fails once a collection is subject to accumulating changes that cannot be dealt with, given the set of hyperparameters. Hence, rather than regarding these parameters as fixed properties, a multimedia DBMS should move towards dynamically adapting them to its needs. In order for this to work, however, we require a better understanding of the relationship between those parameters, the properties of a collection, and the quality of the results. And we would like to take the opportunity to encourage researchers who work in high-dimensional indexing to put more emphasis on these aspects.

In addition, the traditional model applied to indexing is that of a system user or database administrator creating indexes based on their requirements. Due to the complex interdependencies we have observed, this might not work well for high-dimensional indexes which is why we also see multimedia indexing as a very interesting area for the application of Database Cracking techniques [IKM07; SJD13] wherein the DBMS tries to learn indexes and optimised storage views based on query workloads. To some extent, our adaptation of the cost model and state of existing indexes based on the outcome of queries can be regarded as the first small steps in such a direction.

9.2.3 Storage Model

Over the course of Cottontail DB's lifetime, we have gone through three different versions of storage engines (four, if one counts the experimental one). A key takeaway from these projects is that the organisation and layout of data on disk has a great impact on performance and that storage of and access to large, complex data types comes with unique requirements. At a high level, there are two ways to approach this issue, both equally interesting and challenging:

On the one hand, [Gia18] did propose and implement a polystore approach, which delegates storage to different systems that handle a certain type of data- or query-workload particularly well. Such an approach comes with the challenge that transactional guarantees must be provided across different systems. This is a problem tackled by PolyDBMS [VLC⁺21] systems such as Polypheny-

DB [VSS18; VHS⁺20]. In fact, first steps in such a direction have already been taken with the integration of Cottontail DB into Polypheny-DB and they could give rise to interesting new insights for mixed query workloads in the *vitrivr* context. However, the disadvantage of such an approach is the inherent complexity incurred by relying on different systems, which may complicate the configuration and parametrisation of the individual components for a use case.

On the other hand, it would also be worthwhile to devise efficient low-level storage solutions for the different types of data and workloads faced by a multimedia DBMS as this would allow for tighter integration and optimisation of different parts of a system and reduce dependency on external components. The relationship between storage and low-level data organisation on the one hand and performance on the other is well established [LÁJ⁰⁹; LJA11; HJB19] and the recent years have seen many interesting strategies to database storage that could be relevant for our purpose (e.g., [IGN¹²; SR12]). One concrete step we would like to take is serialisation and storage of multiple values in a single chunk or segment. We know from experiments that storing and compressing multiple vectors together can lead to significantly better compression and thus reduced IO for read workloads. Such a storage scheme could also be beneficial to a batched execution model. Furthermore, several specialised storage engines could and should be integrated by an abstraction layer, e.g., as proposed by [DJ11; AIA14] and (to some extent) implemented by Cottontail DB.

9.2.4 Cost Model and Query Planning

By making the cost of impaired quality explicit, we have opened an interesting avenue for future research. Thus far, we have only considered rather simple cases in which an approximate index structure is used to satisfy a NNS or range search query. In addition to making these estimates more accurate and exploring the impact of approximation on different types of primitive constructs, it may also be interesting to see how errors introduced into a complex execution plan propagate and influence a final result. If one considers, for example, the case of a (similarity-)join for which both incoming strands are generated by an approximate method – can we predict how these error (and thus the cost) are combined downstream given an accurate model of the initial error?

Furthermore, we see a potential in using query planning techniques for the optimisation of proximity-based queries and we would like to see how our implementation can cope with a larger solution space involving additional indexes, more complex operations, and functions.

9.2.5 Execution Model

In its early days, Cottontail DB implemented a materialisation model wherein each operator generated intermediate resultsets in memory. This strategy was quickly dropped in favour of the much simpler and less memory-intensive iterator model. However, due to the great potential we see in vectorisation and SIMD, especially in light of the recent launch of the Java Vector API², this model is no longer an optimal fit and a transition to a batch-model should be realised rather sooner than later. Such a transition must, of course, be accompanied and aligned with changes to the storage model, which should also be batched.

Another interesting aspect is that of distributing workloads. ADAM_{pro} used to be inherently distributed [GS16; Gia18] whereas Cottontail DB is inherently local. Adding the ability to distribute workloads among different nodes might go a long way in achieving better execution performance, especially for analytical workloads. However, such a step will inevitably raise new research and engineering questions with respect to partitioning and distribution of the data, indexing, and transactionality through, e.g., distributed snapshot isolation [BHF⁺14]. Generally speaking, we see this as a worthwhile opportunity to rethink Cottontail DB's data partitioning model from storage to processing with regards to both local parallelisation as well as distribution.

9.2.6 Implementation and Vision

On the implementation side, we see some room for work on the aspect of automated code generation. In order to attain the execution speed, much of the logic is re-implemented over and over again to make use of optimisations for specific data types and to avoid type casting in tight loops. This process could and should be automated either at compile or runtime (code generation). In addition, there are many quality of life features that have been requested for Cottontail DB, e.g., a SQL interface or a management user interface.

As a long-term vision and given the discussions thus far and potential applications outside the multimedia domain, one should carefully observe the progress in the domain of scientific databases and the convergence between linear and relational algebra. Moving towards a general-purpose, scientific database such as [SBZ⁺13; LGG⁺18; BGL⁺22] could bear many benefits for multimedia retrieval and analytics.

² See <https://openjdk.org/jeps/338> and JEPs 414 & 417; accessed August 2022.

Appendix

A

Evaluation Supplements

```
{  
    /* Path to data folder. */  
    "root": "/mnt/raid0/evaluation/cottontail-data",  
    "server": {  
        /* Port used by gRPC query interface. */  
        "port": 1865  
    },  
    "cost": {  
        /* Allows for aggressive parallelisation. */  
        "speedupPerWorker": 0.00,  
  
        /* Portion of \acrshort{io} cost considered  
        parallelisable. */  
        "parallelisableIO": 0.9  
    }  
}
```

Listing A.1: Cottontail DB configuration used during the entire evaluation (config.json).

Table A.1 State of the Yandex Deep 1B [BL16] shards used during the evaluation with Cottontail DB.

DBO	Class	Type	Rows	Size	Info
yandex_deep5m	ENTITY	-	5e6	-	-
yandex_deep5m.id	COLUMN	INTEGER	5e6	1	NOT NULL
yandex_deep5m.feature	COLUMN	FLOATVEC	5e6	96	NOT NULL
yandex_deep5m.category	COLUMN	INTEGER	5e6	1	NOT NULL
yandex_deep5m.idx_feature_vaf	INDEX	VAF	5e6	-	CLEAN
yandex_deep5m.idx_category_btree	INDEX	BTREE	5e6	-	CLEAN
yandex_deep5m.idx_feature_pq	INDEX	PQ	5e6	-	CLEAN
yandex_deep5m.idx_feature_ivfpq	INDEX	IVFPQ	5e6	-	CLEAN
yandex_deep10m	ENTITY	-	1e7	-	-
yandex_deep10m.id	COLUMN	INTEGER	1e7	1	NOT NULL
yandex_deep10m.feature	COLUMN	FLOATVEC	1e7	96	NOT NULL
yandex_deep10m.category	COLUMN	INTEGER	1e7	1	NOT NULL
yandex_deep10m.idx_feature_vaf	INDEX	VAF	1e7	-	CLEAN
yandex_deep10m.idx_category_btree	INDEX	BTREE	1e7	-	CLEAN
yandex_deep10m.idx_feature_pq	INDEX	PQ	1e7	-	CLEAN
yandex_deep10m.idx_feature_ivfpq	INDEX	IVFPQ	1e7	-	CLEAN
yandex_deep100m	ENTITY	-	1e8	-	-
yandex_deep100m.id	COLUMN	INTEGER	1e8	1	NOT NULL
yandex_deep100m.feature	COLUMN	FLOATVEC	1e8	96	NOT NULL
yandex_deep100m.category	COLUMN	INTEGER	1e8	1	NOT NULL
yandex_deep100m.idx_feature_vaf	INDEX	VAF	1e8	-	CLEAN
yandex_deep100m.idx_category_btree	INDEX	BTREE	1e8	-	CLEAN
yandex_deep100m.idx_feature_pq	INDEX	PQ	1e8	-	CLEAN
yandex_deep100m.idx_feature_ivfpq	INDEX	IVFPQ	1e8	-	CLEAN
yandex_deep1b	ENTITY	-	1e9	-	-
yandex_deep1b.id	COLUMN	INTEGER	1e9	1	NOT NULL
yandex_deep1b.feature	COLUMN	FLOATVEC	1e9	96	NOT NULL
yandex_deep1b.category	COLUMN	INTEGER	1e9	1	NOT NULL
yandex_deep1b.idx_feature_vaf	INDEX	VAF	1e9	-	CLEAN
yandex_deep1b.idx_feature_pq	INDEX	PQ	1e9	-	CLEAN
yandex_deep1b.idx_feature_ivfpq	INDEX	IVFPQ	1e9	-	CLEAN

Table A.2 State of the V3C-based [BRS⁺19; RSB21] entities used during the evaluation with Cottontail DB.

DBO	Class	Type	Rows	Size	Info
averagecolor	ENTITY	-	2512715	-	-
averagecolor.id	COLUMN	STRING	2512715	1	NOT NULL
averagecolor.feature	COLUMN	FLOATVEC	2512715	3	NOT NULL
averagecolor.idx_feature_vaf	INDEX	VAF	2512715	-	CLEAN
averagecolor.idx_feature_pq	INDEX	PQ	2512715	-	CLEAN
visualtextcoembedding	ENTITY	-	2506273	-	-
visualtextcoembedding.id	COLUMN	STRING	2506273	1	NOT NULL
visualtextcoembedding.feature	COLUMN	FLOATVEC	2506273	256	NOT NULL
visualtextcoembedding.idx_feature_vaf	INDEX	VAF	2506273	-	CLEAN
visualtextcoembedding.idx_feature_pq	INDEX	PQ	2506273	-	CLEAN
hogmf25k512	ENTITY	-	2512715	-	-
hogmf25k512.id	COLUMN	STRING	2512715	1	NOT NULL
hogmf25k512.feature	COLUMN	FLOATVEC	2512715	512	NOT NULL
hogmf25k512.idx_feature_vaf	INDEX	VAF	2512715	-	CLEAN
hogmf25k512.idx_feature_pq	INDEX	PQ	2512715	-	CLEAN
inceptionresnetv2	ENTITY	-	2508358	-	-
inceptionresnetv2.id	COLUMN	STRING	2508358	1	NOT NULL
inceptionresnetv2.feature	COLUMN	FLOATVEC	2508358	1536	NOT NULL
inceptionresnetv2.idx_feature_vaf	INDEX	VAF	2508358	-	CLEAN
inceptionresnetv2.idx_feature_pq	INDEX	PQ	2508358	-	CLEAN
conceptmasksade20k	ENTITY	-	2469844	-	-
conceptmasksade20k.id	COLUMN	STRING	2469844	1	NOT NULL
conceptmasksade20k.feature	COLUMN	FLOATVEC	2469844	2048	NOT NULL
conceptmasksade20k.idx_feature_vaf	INDEX	VAF	2469844	-	CLEAN
conceptmasksade20k.idx_feature_pq	INDEX	PQ	2469844	-	CLEAN

B

Additional Results

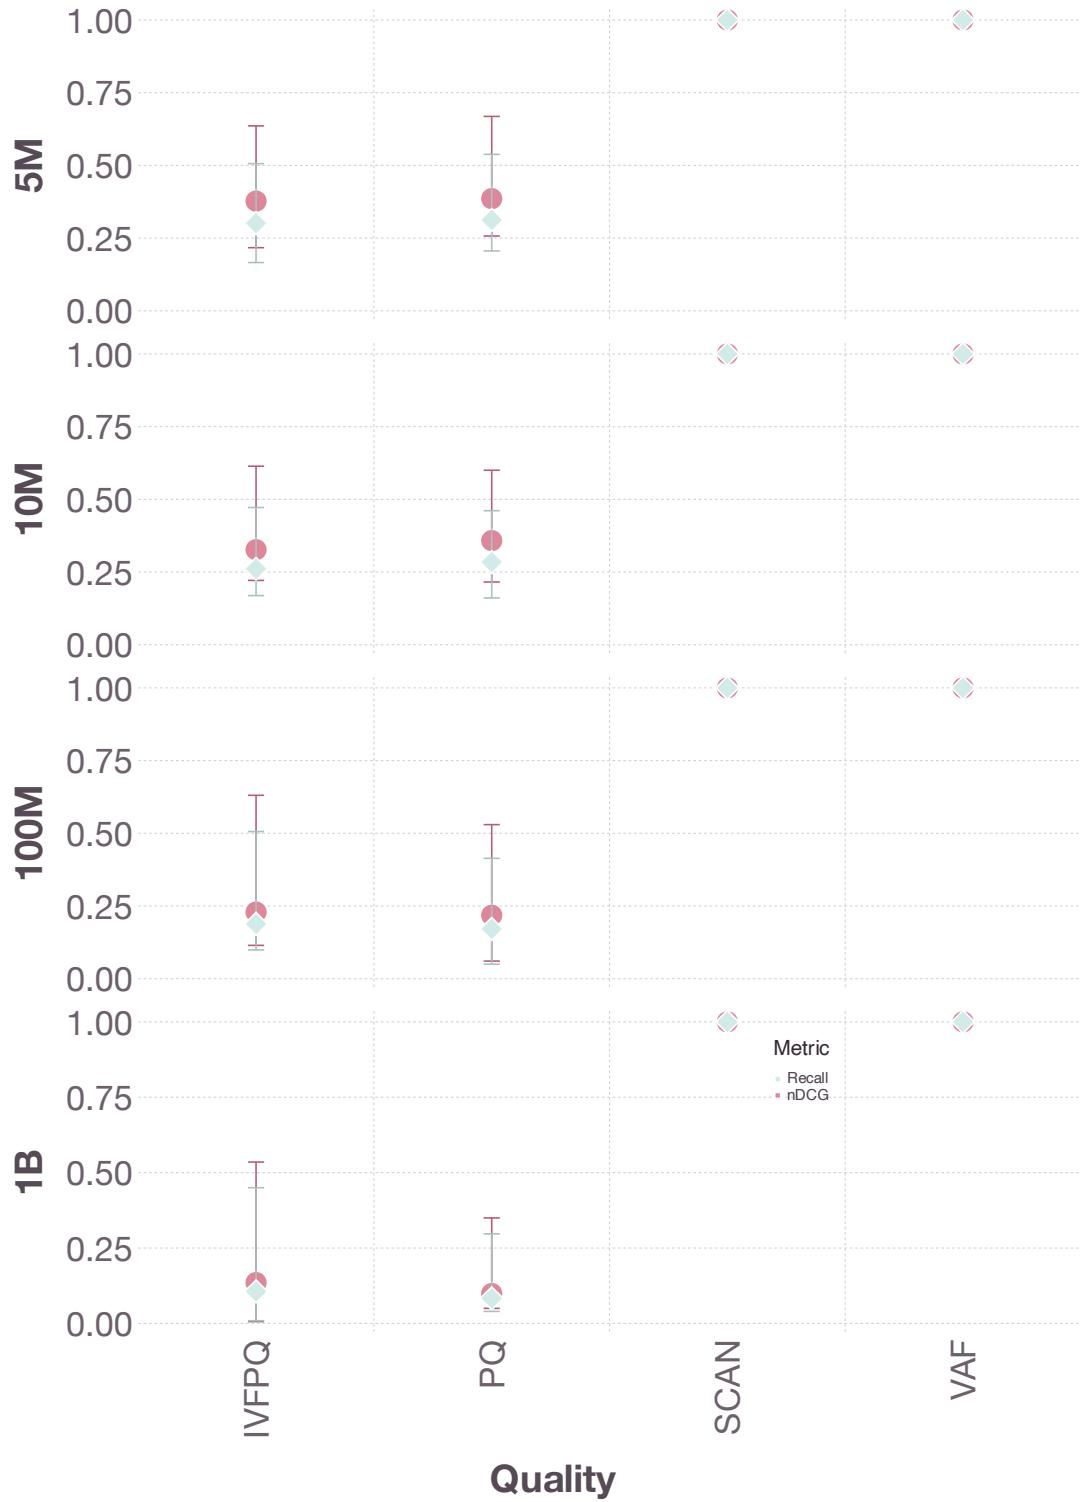


Figure B.1 Quality metrics for Cottontail DB during the simple NNS large-scale retrieval evaluation presented in Section 8.3.1. Recall and DCG are 1.0 for all execution strategies other than PQ, for which they exhibit rather low values with a large spread.

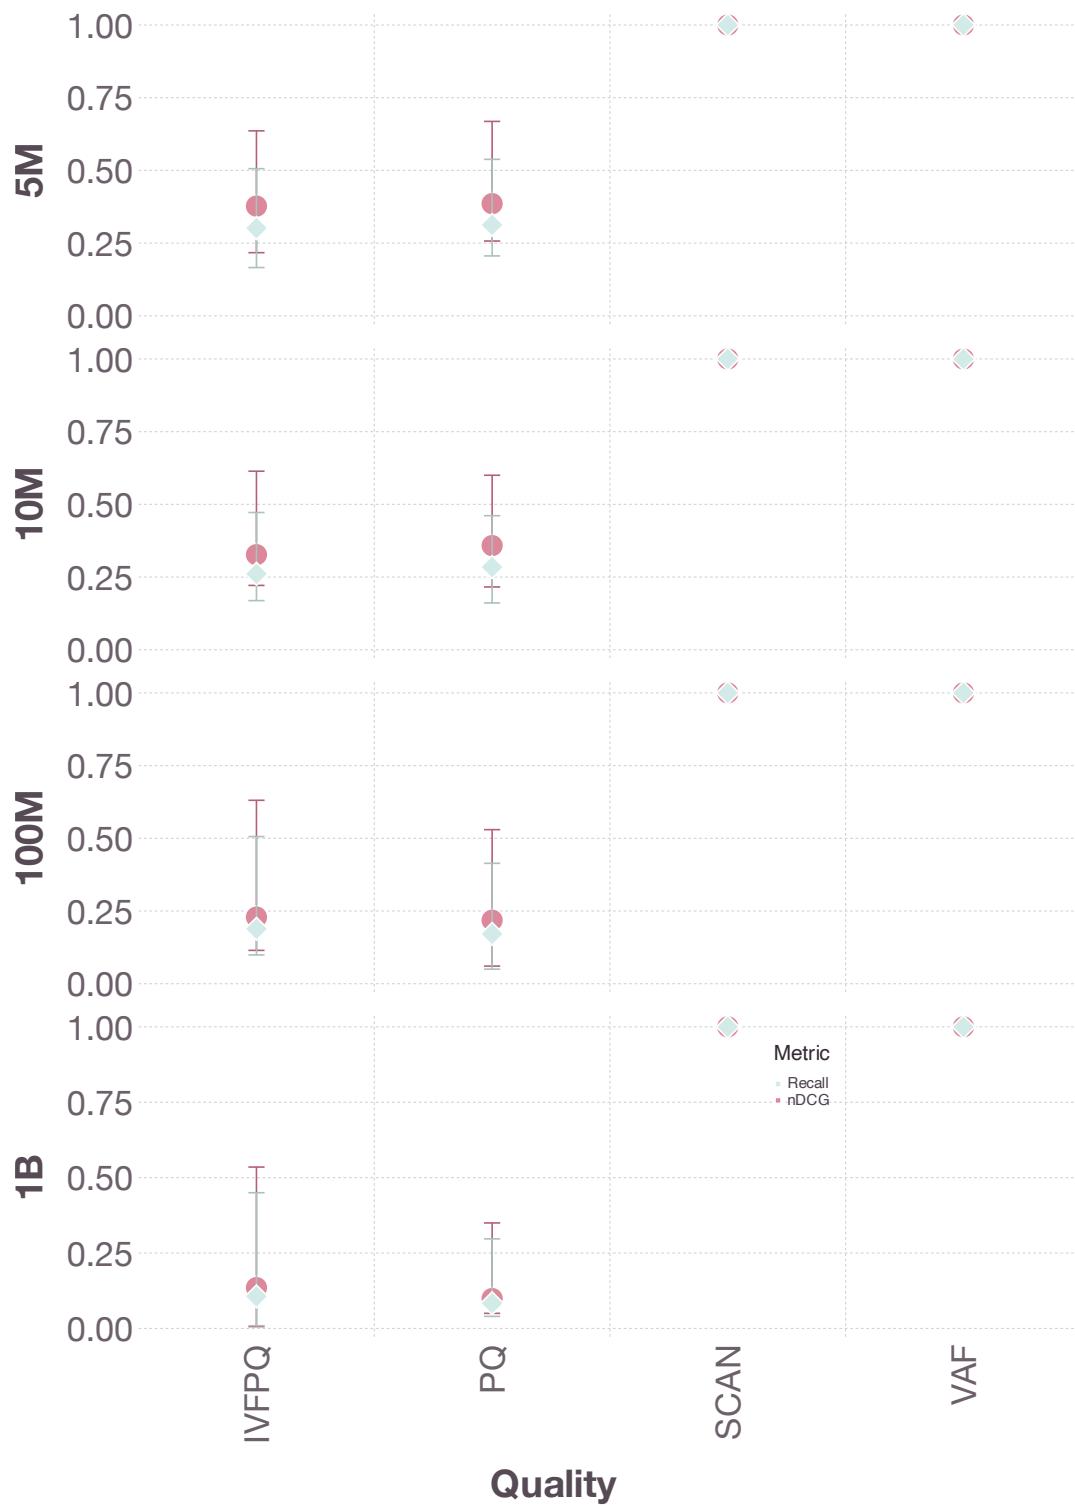


Figure B.2 Quality metrics for Cottontail DB during the NNS + fetch large-scale retrieval evaluation presented in Section 8.3.1. Recall and DCG are 1.0 for all execution strategies other than PQ, for which they exhibit rather low values with a large spread.

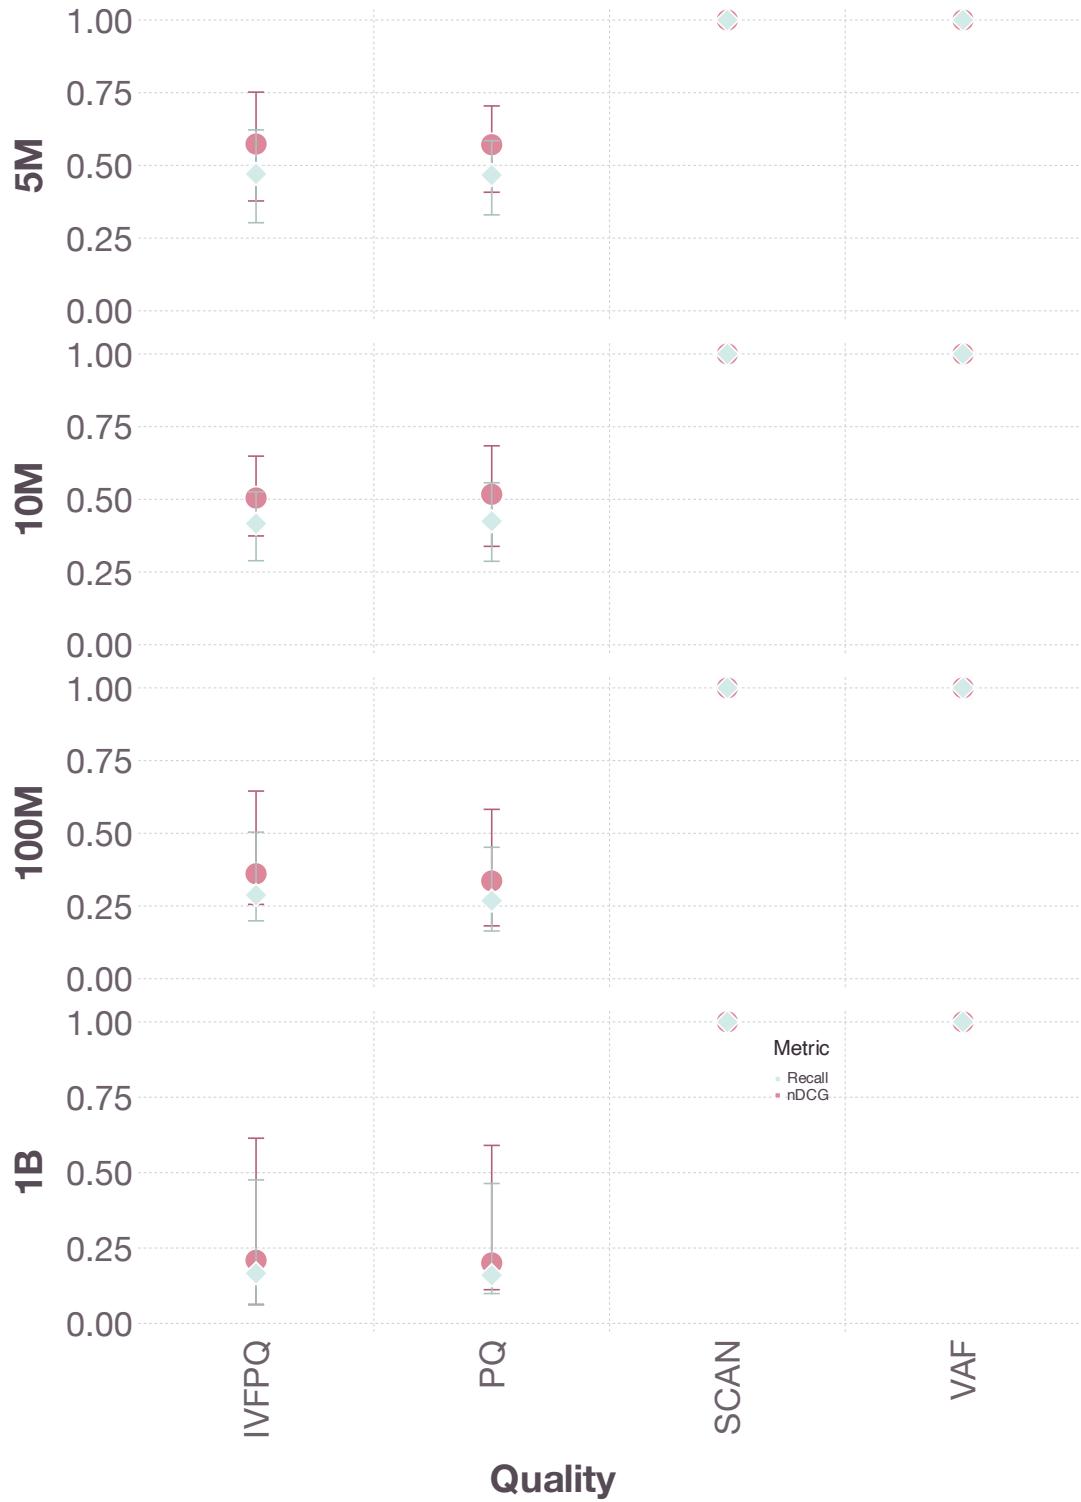


Figure B.3 Quality metrics for Cottontail DB during the hybrid query large-scale retrieval evaluation presented in Section 8.3.1. Recall and DCG are 1.0 for all execution strategies other than PQ, for which they exhibit rather low values with a large spread.

Bibliography

- [AAA^{+14]} Daniel Abadi, Rakesh Agrawal, Anastasia Ailamaki, Magdalena Balazinska, Philip A. Bernstein, Michael J. Carey, Surajit Chaudhuri, Jeffrey Dean, AnHai Doan, Michael J. Franklin, Johannes Gehrke, Laura M. Haas, Alon Y. Halevy, Joseph M. Hellerstein, Yannis E. Ioannidis, H. V. Jagadish, Donald Kossmann, Samuel Madden, Sharad Mehrotra, Tova Milo, Jeffrey F. Naughton, Raghu Ramakrishnan, Volker Markl, Christopher Olston, Beng Chin Ooi, Christopher Ré, Dan Suciu, Michael Stonebraker, Todd Walter, and Jennifer Widom. The Beckman Report on Database Research. *ACM Sigmod Record*, 43(3):61–70, December 2014. ISSN: 0163-5808. doi: 10.1145/2694428.2694441.
- [AAA^{+20]} Daniel Abadi, Anastasia Ailamaki, David Andersen, Peter Bailis, Magdalena Balazinska, Philip Bernstein, Peter Boncz, Surajit Chaudhuri, Alvin Cheung, AnHai Doan, Luna Dong, Michael J. Franklin, Juliana Freire, Alon Halevy, Joseph M. Hellerstein, Stratos Idreos, Donald Kossmann, Tim Kraska, Sailesh Krishnamurthy, Volker Markl, Sergey Melnik, Tova Milo, C. Mohan, Thomas Neumann, Beng Chin Ooi, Fatma Ozcan, Jignesh Patel, Andrew Pavlo, Raluca Popa, Raghu Ramakrishnan, Christopher Ré, Michael Stonebraker, and Dan Suciu. The Seattle Report on Database Research. *ACM Sigmod Record*, 48(4):44–53, February 2020. ISSN: 0163-5808. doi: 10.1145/3385658.3385668.
- [AAB^{+05]} Serge Abiteboul, Rakesh Agrawal, Phil Bernstein, Mike Carey, Stefano Ceri, Bruce Croft, David DeWitt, Mike Franklin, Hector Garcia Molina, Dieter Gawlick, Jim Gray, Laura Haas, Alon Halevy, Joe Hellerstein, Yannis Ioannidis, Martin Kersten, Michael Pazzani, Mike Lesk, David Maier, Jeff Naughton, Hans Schek, Timos Sellis, Avi Silberschatz, Mike Stonebraker, Rick Snodgrass, Jeff Ullman, Gerhard Weikum, Jennifer Widom, and Stan Zdonik. The Lowell Database Research Self-Assessment. *Communications of The ACM*, 48(5):111–118, May 2005. ISSN: 0001-0782. doi: 10.1145/1060710.1060718.

- [AN97] Donald A. Adjeroh and Kingsley C. Nwosu. Multimedia Database Management-Requirements and Issues. *IEEE MultiMedia*, 4(3):24–33, 1997. doi: 10.1109/93.621580.
- [AAB⁺08] Rakesh Agrawal, Anastasia Ailamaki, Philip A. Bernstein, Eric A. Brewer, Michael J. Carey, Surajit Chaudhuri, AnHai Doan, Daniela Florescu, Michael J. Franklin, and Hector Garcia-Molina. The Claremont Report on Database Research. *ACM Sigmod Record*, 37(3):9–19, 2008.
- [AIA14] Ioannis Alagiannis, Stratos Idreos, and Anastasia Ailamaki. H2O: A hands-free Adaptive Store. In *Proceedings of the 2014 ACM SIGMOD International Conference on Management of Data*, pages 1103–1114, New York, NY, USA. Association for Computing Machinery, 2014. ISBN: 978-1-4503-2376-5. doi: 10.1145/2588555.2610502.
- [Ams14] Laurent Amsaleg. *A Database Perspective on Large Scale High-Dimensional Indexing*. PhD thesis, Université Rennes, 2014.
- [AG08] Renzo Angles and Claudio Gutierrez. Survey of Graph Database Models. *ACM Computing Surveys (CSUR)*, 40(1):1–39, 2008.
- [Arm74] William Ward Armstrong. Dependency Structures of Data Base Relationships. In *Proceedings of the 1974 IFIP Congress*, volume 74, pages 580–583, 1974.
- [AP17] Joy Arulraj and Andrew Pavlo. How to Build a Non-volatile Memory Database ManagementSystem. In *Proceedings of the 2017 ACM International Conference on Management of Data*, pages 1753–1758, 2017.
- [ABC⁺76] Morton M. Astrahan, Mike W. Blasgen, Donald D. Chamberlin, Kapali P. Eswaran, Jim Gray, Patricia P. Griffiths, W. Frank King, Raymond A. Lorie, Paul R. McJones, James W. Mehl, Gi-anfranco R. Putzolu, Irving L. Traiger, Bradford W. Wade, and Vera Watson. System R: Relational Approach to Database Management. *ACM Transactions on Database Systems*, 1(2):97–137, June 1976. ISSN: 0362-5915. doi: 10.1145/320455.320457.
- [ABF17] Martin Aumüller, Erik Bernhardsson, and Alexander Faithfull. ANN-Benchmarks: A Benchmarking Tool for Approximate Nearest Neighbor Algorithms. In *Proceedings of the 2017 In-*

- ternational Conference on Similarity Search and Applications*, volume 10609 of *Lecture Notes in Computer Science*, pages 34–49, Cham. Springer International Publishing, 2017. doi: 10.1007/978-3-319-68474-1_3.
- [BL14] Artem Babenko and Victor Lempitsky. The Inverted Multi-index. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 37(6):1247–1260, 2014.
- [BL16] Artem Babenko and Victor Lempitsky. Efficient Indexing of Billion-scale Datasets of Deep Descriptors. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 2055–2063, 2016.
- [BK17] Tobias Baer and Vishnu Kamalnath. Controlling Machine-Learning Algorithms and their Biases. *McKinsey Insights*, 2017.
- [BPS⁺20] Abhishek Bagade, Ashwini Pale, Shreyans Sheth, Megha Agarwal, Soumen Chakrabarti, Kameswari Chebrolu, and S. Sudarshan. The Kauwa-Kaate Fake News Detection System: Demo. In *Proceedings of the 7th ACM IKDD CoDS and 25th COMAD*, CoDS COMAD 2020, pages 302–306, New York, NY, USA, 2020. ISBN: 978-1-4503-7738-6. doi: 10.1145/3371158.3371402.
- [BZ18] Pablo Barberá and Thomas Zeitzoff. The New Public Address System: Why do World Leaders Adopt Social Media? *International Studies Quarterly*, 62(1):121–130, 2018.
- [BHM16] Kai Uwe Barthel, Nico Hezel, and Radek Mackowiak. Navigating a Graph of Scenes for Exploring Large Video Collections. In *Proceedings of the 22nd International Conference on Multimedia Modeling*, volume 9517 of *Lecture Notes in Computer Science*, pages 418–423. Springer International Publishing, 2016. ISBN: 978-3-319-27674-8. doi: 10.1007/978-3-319-27674-8_43.
- [BTVG06] Herbert Bay, Tinne Tuytelaars, and Luc Van Gool. SURF: Speeded Up Robust Features. In *European Conference on Computer Vision*, pages 404–417, 2006.
- [BM70] Rudolf Bayer and Edward M. McCreight. Organization and Maintenance of Large Ordered Indexes. In *Proceedings of the 1970 ACM SIGFIDET Workshop on Data Description, Access and Control*, SIGFIDET '70, pages 107–141, New York, NY, USA. Asso-

- ciation for Computing Machinery, 1970. ISBN: 978-1-4503-7941-0. doi: 10.1145/1734663.1734671.
- [BKS⁺90] Norbert Beckmann, Hans-Peter Kriegel, Ralf Schneider, and Bernhard Seeger. The R*-Tree: An Efficient and Robust Access Method for Points and Rectangles. In *Proceedings of the 1990 ACM SIGMOD International Conference on Management of Data*, SIGMOD '90, pages 322–331, New York, NY, USA. Association for Computing Machinery, 1990. ISBN: 0-89791-365-5. doi: 10.1145/93597.98741.
- [BCH⁺18] Edmon Begoli, Jesús Camacho-Rodríguez, Julian Hyde, Michael J. Mior, and Daniel Lemire. Apache Calcite: A Foundational Framework for Optimized Query Processing over Heterogeneous Data Sources. In *Proceedings of the 2018 International Conference on Management of Data*, pages 221–230, New York, NY, USA. Association for Computing Machinery, 2018. ISBN: 978-1-4503-4703-7. doi: 10.1145/3183713.3190662.
- [BOV07] Radim Belohlavek, Stanislav Opichal, and Vilem Vychodil. Relational Algebra for Ranked Tables with Similarities: Properties and Implementation. In Michael R. Berthold, John Shawe-Taylor, and Nada Lavrač, editors, *Advances in Intelligent Data Analysis VII*, pages 140–151, Berlin, Heidelberg. Springer Berlin Heidelberg, 2007. ISBN: 978-3-540-74825-0.
- [Ben75] Jon Louis Bentley. Multidimensional Binary Search Trees Used for Associative Searching. *Communications of the ACM*, 18(9):509–517, September 1975. ISSN: 0001-0782. doi: 10.1145/361002.361007.
- [BBG⁺95] Hal Berenson, Phil Bernstein, Jim Gray, Jim Melton, Elizabeth O’Neil, and Patrick O’Neil. A Critique of ANSI SQL Isolation Levels. *ACM SIGMOD Record*, 24(2):1–10, 1995.
- [BS19] David Bernhauer and Tomáš Skopal. Non-metric Similarity Search using Genetic TriGen. In Giuseppe Amato, Claudio Gennaro, Vincent Oria, and Miloš Radovanović, editors, *Proceedings of the 12th International Conference on Similarity Search and Applications*, volume 11807 of *Lecture Notes in Computer Science*, pages 86–93, Cham. Springer International Publishing, 2019. ISBN: 978-3-030-32047-8. doi: 10.1007/978-3-030-32047-8_8.

- [BRS⁺19] Fabian Berns, Luca Rossetto, Klaus Schoeffmann, Christian Beecks, and George Awad. V3C1 Dataset: An Evaluation of Content Characteristics. In *Proceedings of the 2019 on International Conference on Multimedia Retrieval*, pages 334–338, 2019.
- [BDO95] Michael W. Berry, Susan T. Dumais, and Gavin W. O'Brien. Using Linear Algebra for Intelligent Information Retrieval. *SIAM Review*, 37(4):573–595, 1995.
- [BGR⁺99] Kevin Beyer, Jonathan Goldstein, Raghu Ramakrishnan, and Uri Shaft. When is “Nearest Neighbor” Meaningful? In *International Conference on Database Theory*, pages 217–235. Springer, 1999.
- [BHF⁺14] Carsten Binnig, Stefan Hildenbrand, Franz Färber, Donald Kossmann, Juchang Lee, and Norman May. Distributed Snapshot Isolation: Global Transactions Pay Globally, Local Transactions Pay Locally. *The VLDB Journal*, 23(6):987–1011, 2014.
- [BMCFM⁺19] Bhairav Bipin Mehta, Simone Coppo, Debra Frances McGivney, Jesse Ian Hamilton, Yong Chen, Yun Jiang, Dan Ma, Nicole Seiberlich, Vikas Gulani, and Mark Alan Griswold. Magnetic Resonance Fingerprinting: A Technical Review. *Magnetic Resonance in Medicine*, 81(1):25–46, 2019.
- [BGL⁺22] Mark Blacher, Joachim Giesen, Sören Laue, Julien Klaus, and Viktor Leis. Machine Learning, Linear Algebra, and More: Is SQL All You Need? In *Proceedings of the 11th Conference on Innovative Data Systems Research*, pages 9–12, 2022.
- [BdB⁺07] Henk M Blanken, Arjen P de Vries, Henk Ernst Blok, and Ling Feng. *Multimedia Retrieval*. Springer, 2007.
- [Blo46] Felix Bloch. Nuclear Induction. *Physical Review*, 70(7-8):460, 1946.
- [BHC⁺01] Henk Ernst Blok, Djoerd Hiemstra, Sunil Choenni, Franciska de Jong, Henk M. Blanken, and Peter MG. Apers. Predicting the Cost-quality Trade-off for Information Retrieval Queries: Facilitating Database Design and Query Optimization. In *Proceedings of the Tenth International Conference on Information and Knowledge Management*, pages 207–214, 2001.

- [BMS⁺01] Klemens Böhm, Michael Mlivoncic, Hans-Jörg Schek, Roger Weber, et al. Fast Evaluation Techniques for Complex Similarity Queries. In *Proceedings of the 27th International Conference on Very Large Databases*, volume 1, pages 211–220, 2001.
- [BKM08] Peter A. Boncz, Martin L. Kersten, and Stefan Manegold. Breaking the Memory Wall in MonetDB. *Communications of the ACM*, 51(12):77–85, December 2008. issn: 0001-0782. doi: 10 . 1145 / 1409360 . 1409380.
- [BGS⁺20] Samuel Börlin, Ralph Gasser, Florian Spiess, and Heiko Schuldt. 3D Model Retrieval Using Constructive Solid Geometry in Virtual Reality. In *Proceedings of the 2020 IEEE International Conference on Artificial Intelligence and Virtual Reality*, AIVR 2021, pages 373–374, 2020.
- [BGJ10] Nicolas Bruno, César Galindo-Legaria, and Milind Joshi. Polynomial Heuristics for Query Optimization. In *Proceedings of the 26th IEEE International Conference on Data Engineering*, pages 589–600, 2010. doi: 10 . 1109/ICDE . 2010 . 5447916.
- [BJZ13] Nicolas Bruno, Sapna Jain, and Jingren Zhou. Continuous Cloud-Scale Query Optimization and Processing. *Proceedings of the VLDB Endowment*, 6(11):961–972, 2013.
- [BBZ12] Petra Budikova, Michal Batko, and Pavel Zezula. Query Language for Complex Similarity Queries. In Tadeusz Morzy, Theo Härdter, and Robert Wrembel, editors, *Proceedings of the East European Conference on Advances in Databases and Information Systems*, pages 85–98, Berlin, Heidelberg, 2012. isbn: 978-3-642-33074-2.
- [BGS⁺17] Mateusz Budnik, Efrain-Leonardo Gutierrez-Gomez, Bahjat Safadi, Denis Pellerin, and Georges Quénot. Learned Features Versus Engineered Features for Multimedia Indexing. *Multimedia Tools and Applications*, 76(9):11941–11958, 2017.
- [CS15] Fatih Cakir and Stan Sclaroff. Adaptive Hashing for Fast Similarity Search. In *Proceedings of the 2015 IEEE International Conference on Computer Vision*, pages 1044–1052, 2015.
- [CWW⁺10] Yang Cao, Hai Wang, Changhu Wang, Zhiwei Li, Liqing Zhang, and Lei Zhang. MindFinder: interactive sketch-based image search on millions of images. In *Proceedings of the 18th ACM International Conference on Multimedia*, MM ’10, pages 1605–1608, New

- York, NY, USA. Association for Computing Machinery, 2010. ISBN: 978-1-60558-933-6. doi: 10.1145/1873951.1874299.
- [CHS⁺95] Michael J. Carey, Laura M. Haas, Peter M. Schwarz, Manish Arya, William F Cody, Ronald Fagin, Myron Flickner, Allen W. Luniewski, Wayne Niblack, and Dragutin Petkovic. Towards Heterogeneous Multimedia Information Systems: The Garlic Approach. In *Proceedings of the 5th International Workshop on Research Issues in Data Engineering-Distributed Object Management*, pages 124–131, 1995.
- [CSM⁺15] Stephen F. Cauley, Kawin Setsompop, Dan Ma, Yun Jiang, Hui-hui Ye, Elfar Adalsteinsson, Mark A. Griswold, and Lawrence L. Wald. Fast Group Matching for MR Fingerprinting Reconstruction. *Magnetic Resonance in Medicine*, 74(2):523–528, 2015.
- [Cha12] Donald D. Chamberlin. Early History of SQL. *IEEE Annals of the History of Computing*, 34(4):78–82, 2012. doi: 10.1109/MAHC.2012.61.
- [Cha98a] Arvola Chan. Transactional Publish/Subscribe: The Proactive Multicast of Database Changes. *ACM Sigmod Record*, 27(2):521, 1998.
- [Cha98b] Surajit Chaudhuri. An Overview of Query Optimization in Relational Systems. In *Proceedings of the 18th ACM SIGACT-SIGMOD-SIGART Symposium on Principles of Database Systems*, PODS ’98, pages 34–43, New York, NY, USA. Association for Computing Machinery, 1998. ISBN: 0-89791-996-3. doi: 10.1145/275487.275492.
- [CHW⁺21] Aozhu Chen, Fan Hu, Zihan Wang, Fangming Zhou, and Xirong Li. What Matters for Ad-hoc Video Search? A Large-scale Evaluation on TRECVID. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 2317–2322, 2021.
- [Che76] Peter Pin-Shan Chen. The Entity-Relationship Model — Toward a Unified View of Data. *ACM Transactions on Database Systems*, 1(1):9–36, March 1976. ISSN: 0362-5915. doi: 10.1145/320434.320440.

- [CZW⁺21] Qi Chen, Bing Zhao, Haidong Wang, Mingqin Li, Chuanjie Liu, Zengzhong Li, Mao Yang, and Jingdong Wang. SPANN: Highly-efficient Billion-Scale Approximate Nearest Neighborhood Search. *Advances in Neural Information Processing Systems*, 34:5199–5212, 2021.
- [CPR⁺07] Flavio Chierichetti, Alessandro Panconesi, Prabhakar Raghavan, Mauro Sozio, Alessandro Tiberi, and Eli Upfal. Finding Near Neighbors through Cluster Pruning. In *Proceedings of the 26th ACM SIGMOD-SIGACT-SIGART Symposium on Principles of Database Systems*, pages 103–112, 2007.
- [CTW⁺10] Nancy A. Chinchor, James J. Thomas, Pak Chung Wong, Michael G. Christel, and William Ribarsky. Multimedia Analysis + Visual Analytics = Multimedia Analytics. *IEEE Computer Graphics and Applications*, 30(5):52–60, 2010. doi: 10.1109/MCG.2010.92.
- [CPZ97] Paolo Ciaccia, Marco Patella, and Pavel Zezula. M-tree: An Efficient Access Method for Similarity Search in Metric Spaces. In *Proceedings of the 23rd International Conference on Very Large Data Bases*, volume 97, pages 426–435, 1997.
- [Ciu22] Dan Ciuriak. The Role of Social Media in Russia’s War on Ukraine. *Available at SSRN*, 2022.
- [Cod70] Edgar Frank Codd. Relational Model of Data for Large Shared Data Banks. *Communications of the ACM*, 13(6):377–387, June 1970.
- [Cod72] Edgar Frank Codd. Further Normalization of the Data Base Relational Model. *Data Base Systems*, 6:33–64, 1972.
- [Cod74] Edgar Frank Codd. Recent Investigations Into Relational Data Base. In *Proceedings of the 1974 IFIP Congress*, 1974.
- [CHN⁺95] William F. Cody, Laura M. Haas, Wayne Niblack, Manish Arya, Michael J. Carey, Ronald Fagin, Myron Flickner, Denis Lee, Dragutin Petkovic, and Peter M. Schwarz. Querying Multimedia Data from Multiple Repositories by Content: The Garlic Project. In *Proceedings of the 1995 Working Conference on Visual Database Systems*, pages 17–35, 1995.

- [CZR18] Ouri Cohen, Bo Zhu, and Matthew S. Rosen. MR Fingerprinting Deep Reconstruction Network (DRONE). *Magnetic Resonance in Medicine*, 80(3):885–894, 2018.
- [Cre20] Stefano Cresci. A Decade of Social Bot Detection. *Communications of the ACM*, 63(10):72–83, 2020.
- [CLP22] Andrew Crotty, Viktor Leis, and Andrew Pavlo. Are You Sure You Want to Use MMAP in Your Database Management System? In *Proceedings of the 2022 Conference on Innovative Data Systems Research*, Chaminade, CA, 2022.
- [CSW⁺19] Limeng Cui, Kai Shu, Suhang Wang, Dongwon Lee, and Huan Liu. DEFEND: A System for Explainable Fake News Detection. In *Proceedings of the 28th ACM International Conference on Information and Knowledge Management*, CIKM ’19, pages 2961–2964, New York, NY, USA. Association for Computing Machinery, 2019. ISBN: 978-1-4503-6976-3. doi: 10.1145/3357384.3357862.
- [DT05] Navneet Dalal and Bill Triggs. Histograms of Oriented Gradients for Human Detection. In *Proceedings of the 2005 IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, volume 1, 886–893 vol. 1, 2005. doi: 10.1109/CVPR.2005.177.
- [Xopa] *Data Management: SQL Call Level Interface (CLI)*. X/Open Company Ltd., U.K. ISBN: 1-85912-081-4.
- [Xopb] *Data Management: Structured Query Language (SQL) Version 2*. X/Open Company Ltd., U.K. ISBN: 1-85912-151-9.
- [DVF⁺16] Clayton Allen Davis, Onur Varol, Emilio Ferrara, Alessandro Flammini, and Filippo Menczer. Botornot: A system to evaluate social bots. In *Proceedings of the 25th International Conference Companion on World Wide Web*, pages 273–274, 2016.
- [dSP⁺17] Jonathan de Oliveira Assis, Vanessa CO Souza, Melise MV Paula, and João Bosco S Cunha. Performance Evaluation of NoSQL Data Store for Digital Media. In *Proceedings of the 12th Iberian Conference on Information Systems and Technologies*, pages 1–6, 2017.
- [dLG16] Gabriel de Oliveira Barra, Mathias Lux, and Xavier Giro-i-Nieto. Large Scale Content-based Video Retrieval with LIVRE. In *2016 14th International Workshop on Content-based Multimedia Indexing*, pages 1–4, 2016.

- [DG08] Jeffrey Dean and Sanjay Ghemawat. MapReduce: Simplified Data Processing on Large Clusters. *Communications of the ACM*, 51(1):107–113, 2008.
- [DBS19] Emir Demirel, Baris Bozkurt, and Xavier Serra. Automatic Chord-Scale Recognition using Harmonic Pitch Class Profiles. In *Proceedings of the 16th Sound & Music Computing Conference*, Málaga, Spain, 2019.
- [DBS⁺11] Oscar Déniz, Gloria Bueno, Jesús Salido, and Fernando De la Torre. Face Recognition using Histograms of Oriented Gradients. *Pattern Recognition Letters*, 32(12):1598–1603, 2011.
- [DM10] Adrien Depeursinge and Henning Müller. Fusion Techniques for Combining Textual and Visual Information Retrieval. In *ImageCLEF: Experimental Evaluation in Visual Information Retrieval*, pages 95–114. January 2010. ISBN: 978-3-642-15180-4.
- [DKN08] Thomas Deselaers, Daniel Keysers, and Hermann Ney. Features for Image Retrieval: An Experimental Comparison. *Information Retrieval*, 11(2):77–107, 2008. doi: 10.1007/s10791-007-9039-3.
- [DIR⁺07] Amol Deshpande, Zachary Ives, Vijayshankar Raman, et al. Adaptive Query Processing. *Foundations and Trends® in Databases*, 1(1):1–140, 2007.
- [Dha13] Vasant Dhar. Data Science and Prediction. *Communications of the ACM*, 56(12):64–73, 2013.
- [DA13] T. Dharani and I. Laurence Aroquiaraj. A Survey on Content Based Image Retrieval. In *Proceedings of the 2013 International Conference on Pattern Recognition, Informatics and Mobile Engineering*, pages 485–490, 2013.
- [DZJ⁺12] Shifei Ding, Hong Zhu, Weikuan Jia, and Chunyang Su. A Survey on Feature Extraction for Pattern Recognition. *Artificial Intelligence Review*, 37(3):169–180, March 2012. ISSN: 1573-7462. doi: 10.1007/s10462-011-9225-y.
- [DJ11] Jens Dittrich and Alekh Jindal. Towards a One Size Fits all Database Architecture. In pages 195–198, January 2011.

- [DJ21] Aaron Duane and Björn Þór Jónsson. ViRMA: Virtual Reality Multimedia Analytics at LSC 2021. In *Proceedings of the 4th Annual on Lifelog Search Challenge*, LSC '21, pages 29–34, New York, NY, USA. Association for Computing Machinery, 2021. ISBN: 978-1-4503-8533-6. doi: 10.1145/3463948.3469067.
- [DJ22] Aaron Duane and Björn Þór Jónsson. ViRMA: Virtual Reality Multimedia Analytics. In *Proceedings of the 2022 International Conference on Multimedia Retrieval*, pages 211–214, 2022.
- [EZP21] Karima Echihabi, Kostas Zoumpatianos, and Themis Palpanas. High-Dimensional Similarity Search for Scalable Data Science. In *Proceeding of the 37th IEEE International Conference on Data Engineering*, pages 2369–2372, 2021. doi: 10.1109/ICDE51399.2021.00268.
- [FKL⁺17] Franz Faerber, Alfons Kemper, Per-Åke Larson, Justin Levan-doski, Thomas Neumann, and Andrew Pavlo. Main Memory Database Systems. *Foundations and Trends in Databases*, 8(1-2):1–130, 2017.
- [Far09] Hany Farid. Image Forgery Detection. *IEEE Signal Processing Magazine*, 26(2):16–25, 2009.
- [Far16] Javeria Farooq. Object Detection and Identification using SURF and BoW Model. In *Proceedings of the 2016 International Conference on Computing, Electronic and Electrical Engineering*, pages 318–323, 2016.
- [FTA⁺00] Hakan Ferhatsmanoglu, Ertem Tuncel, Divyakant Agrawal, and Amr El Abbadi. Vector Approximation Based Indexing for Non-Uniform High Dimensional Data Sets. In *Proceedings of the 10th International Conference on Information and Knowledge Management*, pages 202–209, 2000.
- [FCC⁺20] Emilio Ferrara, Herbert Chang, Emily Chen, Goran Muric, and Jaimin Patel. Characterizing Social Media Manipulation in the 2020 US Presidential Election. *First Monday*, 2020.
- [Fer14] Nicola Ferro, editor. *Bridging Between Information Retrieval and Databases*, volume 8173 of *Lecture Notes in Computer Science*. Springer, 2014. ISBN: 978-3-642-54797-3.

- [FB74] Raphael A. Finkel and Jon Louis Bentley. Quad-Trees: A Data Structure for Retrieval on Composite Keys. *Acta Informatica*, 4(1):1–9, 1974.
- [FC13] Fausto C. Fleites and Shu-Ching Chen. Efficient Content-based Multimedia Retrieval using Novel Indexing Structure in PostgreSQL. In *Proceedings of the 2013 IEEE International Symposium on Multimedia*, pages 500–501, 2013.
- [FSA⁺95] Myron Flickner, Harpeet S. Sawhney, Jonathan Ashley, Qian Huang, Byron Dom, Monika Gorkani, Jim Hafner, Denis Lee, Dragutin Petkovic, David Steele, and Peter Yanker. Query by Image and Video Content: The QBIC System. *IEEE Computer*, 28(9):23–32, 1995. doi: 10.1109/2.410146.
- [Foo00] Jonathan Foote. Automatic Audio Segmentation Using a Measure of Audio Novelty. In *Proceedings of the 2000 IEEE International Conference on Multimedia and Expo*, volume 1, pages 452–455, 2000.
- [FGG⁺18] Nadime Francis, Alastair Green, Paolo Guagliardo, Leonid Libkin, Tobias Lindaaker, Victor Marsault, Stefan Plantikow, Mats Rydberg, Petra Selmer, and Andrés Taylor. Cypher: An Evolving Query Language for Property Graphs. In *Proceedings of the 2018 ACM SIGMOD International Conference on Management of Data*, pages 1433–1445, 2018.
- [FJK96] Michael J. Franklin, Björn Þór Jónsson, and Donald Kossmann. Performance Tradeoffs for Client-Server Query Processing. In *Proceedings of the 1996 ACM SIGMOD International Conference on Management of Data*, pages 149–160, New York, NY, USA. Association for Computing Machinery, 1996. ISBN: 0-89791-794-4. doi: 10.1145/233269.233328.
- [FXW⁺17] Cong Fu, Chao Xiang, Changxu Wang, and Deng Cai. Fast Approximate Nearest Neighbor Search with the Navigating Spreading-out Graph. *CoRR*, abs/1707.00143, 2017.
- [FT21] Henning Funke and Jens Teubner. Low-latency Compilation of SQL Queries to Machine Code. *Proceedings of the VLDB Endowment*, 14(12):2691–2694, July 2021. ISSN: 2150-8097. doi: 10.14778/3476311.3476321.

- [GLSJ⁺21] Liliana C. Gallegos, Giulian Luchini, Peter C. St. John, Seonah Kim, and Robert S. Paton. Importance of Engineered and Learned Molecular Representations in Predicting Organic Reactivity, Selectivity, and Chemical Properties. *Accounts of Chemical Research*, 54(4):827–836, 2021.
- [GS92] Hector Garcia-Molina and Kenneth Salem. Main Memory Database Systems: An Overview. *IEEE Transactions on Knowledge and Data Engineering*, 4(6):509–516, 1992.
- [GUW09] Hector Garcia-Molina, Jeffrey D. Ullmann, and Jennifer Widom. *DATABASE SYSTEMS The Complete Book, 2nd Edition*. second edition, 2009. ISBN: 1.
- [GRH⁺20] Ralph Gasser, Luca Rossetto, Silvan Heller, and Heiko Schuldt. Cottontail DB: An Open Source Database System for Multimedia Retrieval and Analysis. In *Proceedings of the 28th ACM International Conference on Multimedia*, ACM MM 2020, pages 4465–4468, 2020. doi: 10.1145/3394171.3414538.
- [GRS19a] Ralph Gasser, Luca Rossetto, and Heiko Schuldt. Multimodal Multimedia Retrieval with vitrivr. In *Proceedings of the 2019 International Conference on Multimedia Retrieval (ICMR 2019)*, ICMR 2019, pages 391–394, New York, NY, USA. Association for Computing Machinery, 2019. ISBN: 978-1-4503-6765-3. doi: 10.1145/3323873.3326921.
- [GRS19b] Ralph Gasser, Luca Rossetto, and Heiko Schuldt. Towards an All-purpose Content-based Multimedia Information Retrieval System. *CoRR*, abs/1902.03878, 2019.
- [GHK⁺14] Tiezheng Ge, Kaiming He, Qifa Ke, and Jian Sun. Optimized Product Quantization. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 36(4):744–755, 2014. doi: 10.1109/TPAMI.2013.240.
- [GTK01] Lise Getoor, Benjamin Taskar, and Daphne Koller. Selectivity Estimation using Probabilistic Models. In *Proceedings of the 2001 ACM SIGMOD International Conference on Management of Data*, pages 461–472, 2001.

- [GLC⁺95] Asif Ghias, Jonathan Logan, David Chamberlin, and Brian C. Smith. Query by Humming: Musical Information Retrieval in an Audio Database. In *Proceedings of the 3rd ACM International Conference on Multimedia*, pages 231–236, 1995.
- [Gia18] Ivan Giangreco. *Database Support for Large-scale Multimedia Retrieval*. PhD thesis, University of Basel, Switzerland, August 2018.
- [GAKS14] Ivan Giangreco, Ihab Al Kabary, and Heiko Schuldt. ADAM: A Database and Information Retrieval System for Big Multimedia Collections. In *Proceedings of the 2014 IEEE International Congress on Big Data*, pages 406–413, 2014.
- [GS16] Ivan Giangreco and Heiko Schuldt. ADAMpro: Database Support for Big Multimedia Retrieval. *Datenbank-Spektrum*, 16(1):17–26, 2016. doi: 10.1007/s13222-015-0209-y.
- [GBC16] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016.
- [GAR⁺16] Albert Gordo, Jon Almazán, Jerome Revaud, and Diane Larlus. Deep Image Retrieval: Learning Global Representations for Image Search. In *Proceedings of the 2016 European Conference on Computer Vision*, pages 241–257, 2016.
- [Gra93] Goetz Graefe. Query Evaluation Techniques for Large Databases. *ACM Computing Surveys (CSUR)*, 25(2):73–169, 1993.
- [Gra95] Goetz Graefe. The Cascades Framework for Query Optimization. *IEEE Data Engineering Bulletin*, 18(3):19–29, 1995.
- [GM93] Goetz Graefe and William J. MacKenna. The Volcano Optimizer Generator: Extensibility and Efficient Search. In *Proceedings of IEEE 9th International Conference on Data Engineering*, pages 209–218. IEEE, 1993.
- [GJA10] Gylfi Þór Gudmundsson, Björn Þór Jónsson, and Laurent Amaleg. A Large-scale Performance Study of Custer-based High-Dimensional Indexing. In *Proceedings of the 18th ACM International Conference on Multimedia*, VLS-MCMR '10, pages 31–36, New York, NY, USA. Association for Computing Machinery, 2010. ISBN: 978-1-4503-0166-4. doi: 10.1145/1878137.1878145.

- [GdR⁺09] Denise Guliato, Ernani V. de Melo, Rangaraj M. Rangayyan, and Robson C. Soares. PostgreSQL-IE: An Image-Handling Extension for PostgreSQL. *Journal of Digital Imaging*, 22(2):149–165, 2009.
- [GSL⁺20] Ruiqi Guo, Philip Sun, Erik Lindgren, Quan Geng, David Simcha, Felix Chern, and Sanjiv Kumar. Accelerating Large-scale Inference with Anisotropic Vector Quantization. In *Proceedings of the 37th International Conference on Machine Learning*, pages 3887–3896, 2020.
- [GHQ95] Ashish Kumar Gupta, Venky Harinarayan, and Dallan Quass. Generalized Projections: A Powerful Approach to Aggregation. Technical report, Stanford InfoLab, 1995.
- [GJS⁺21] Cathal Gurrin, Björn Pór Jónsson, Klaus Schöffmann, Duc-Tien Dang-Nguyen, Jakub Lokoč, Minh-Triet Tran, Wolfgang Hürst, Luca Rossetto, and Graham Healy. Introduction to the Fourth Annual Lifelog Search Challenge, LSC’21. In *Proceedings of the 2021 International Conference on Multimedia Retrieval*, pages 690–691, 2021.
- [GSJ⁺19] Cathal Gurrin, Klaus Schoeffmann, Hideo Joho, Andreas Leibetseder, Liting Zhou, Aaron Duane, Duc-Tien Dang-Nguyen, Michael Riegler, Luca Piras, Minh-Triet Tran, et al. Comparing Approaches to Interactive Lifelog Search at the Lifelog Search Challenge (LSC2018). *ITE Transactions on Media Technology and Applications*, 7(2):46–59, 2019.
- [Gut84] Antonin Guttman. R-Trees: A Dynamic Index Structure for Spatial Searching. In *Proceedings of the 1984 ACM SIGMOD International Conference on Management of Data*, SIGMOD ’84, pages 47–57, New York, NY, USA. Association for Computing Machinery, 1984. ISBN: 0-89791-128-8. doi: 10.1145/602259.602266.
- [HR83] Theo Haerder and Andreas Reuter. Principles of Transaction-Oriented Database Recovery. *ACM Computing Surveys (CSUR)*, 15(4):287–317, 1983.
- [HE10] Philippe Hamel and Douglas Eck. Learning Features from Music Audio with Deep Belief Networks. In J. Stephen Downie and Remco C. Veltkamp, editors, *Proceedings of the 11th International Society for Music Information Retrieval*, pages 339–344, Utrecht,

- Netherlands. International Society for Music Information Retrieval, 2010.
- [HR16] Hadi Hashem and Daniel Ranc. Evaluating NoSQL Document Oriented Data Model. In *Proceedings of the 4th IEEE International Conference on Future Internet of Things and Cloud Workshops*, pages 51–56, 2016.
- [HAPG⁺20] Silvan Heller, Mahnaz Amiri Parian, Ralph Gasser, Loris Sauter, and Heiko Schuldt. Interactive Lifelog Retrieval with vitrivr. In *Proceedings of the 3rd Annual Workshop on Lifelog Search Challenge*, pages 1–6, 2020.
- [HAG⁺22] Silvan Heller, Rahel Arnold, Ralph Gasser, Viktor Gsteiger, Mahnaz Parian-Scherb, Luca Rossetto, Loris Sauter, Florian Spiess, and Heiko Schuldt. Multi-Modal Interactive Video Retrieval with Temporal Queries. In *Proceedings of the 28th International Conference on Multimedia Modeling*, volume 13142 of *Lecture Notes in Computer Science*, pages 493–498. Springer International Publishing, 2022. ISBN: 978-3-030-98355-0. DOI: 10.1007/978-3-030-98355-0_44.
- [HGP⁺21] Silvan Heller, Ralph Gasser, Mahnaz Parian-Scherb, Sanja Popovic, Luca Rossetto, Loris Sauter, Florian Spiess, and Heiko Schuldt. Interactive Multimodal Lifelog Retrieval with vitrivr at LSC 2021. In *Proceedings of the 4th Annual on Lifelog Search Challenge*, pages 35–39, 2021.
- [HSS⁺20] Silvan Heller, Loris Sauter, Heiko Schuldt, and Luca Rossetto. Multi-stage Queries and Temporal Scoring in vitrivr. In *2020 IEEE International Conference on Multimedia & Expo Workshops (ICMEW)*, pages 1–5, 2020.
- [HSH07] Joseph M Hellerstein, Michael Stonebraker, and James Hamilton. *Architecture of a Database System*. Now Publishers Inc, 2007.
- [Hen19] Gavin Henry. Howard Chu on Lightning Memory-Mapped Database. *IEEE Software*, 36(06):83–87, 2019.
- [HSJ⁺22] Nico Hezel, Konstantin Schall, Klaus Jung, and Kai Uwe Barthel. Efficient Search and Browsing of Large-scale Video Collections with Vibro. In *Proceedings of the 28th International Conference on Multimedia Modeling*, volume 13142 of *Lecture Notes in Computer*

- Science*, pages 487–492. Springer International Publishing, 2022. ISBN: 978-3-030-98355-0. doi: 10.1007/978-3-030-98355-0_43.
- [Hir77] Daniel S. Hirschberg. Algorithms for the Longest Common Subsequence Problem. *Journal of the ACM (JACM)*, 24(4):664–675, 1977.
- [HJB19] Anders Munck Højsgaard, Björn Þór Jónsson, and Philippe Bonnet. Index Maintenance Strategy and Cost Model for Extended Cluster Pruning. In Giuseppe Amato, Claudio Gennaro, Vincent Oria, and Miloš Radovanović, editors, *Proceedings of the 12th International Conference on Similarity Search and Applications*, volume 11807 of *Lecture Notes in Computer Science*, pages 32–39, Cham. Springer International Publishing, 2019. ISBN: 978-3-030-32047-8. doi: 10.1007/978-3-030-32047-8_3.
- [HSS⁺19] M. D. Zakir Hossain, Ferdous Sohel, Mohd Fairuz Shiratuddin, and Hamid Laga. A Comprehensive Survey of Deep Learning for Image Captioning. *ACM Computing Surveys (CSUR)*, 51(6):1–36, 2019.
- [HXL⁺11] Weiming Hu, Nianhua Xie, Li Li, Xianglin Zeng, and Stephen Maybank. A Survey on Visual Content-based Video Indexing and Retrieval. *IEEE Transactions on Systems, Man, and Cybernetics, Part C (Applications and Reviews)*, 41(6):797–819, 2011.
- [HPS⁺17] Nina Hubig, Linnea Passing, Maximilian E. Schüle, Dimitri Vorona, Alfons Kemper, and Thomas Neumann. HyPerInsight: Data Exploration Deep Inside HyPer. In *Proceedings of the 2017 ACM Conference on Information and Knowledge Management, CIKM '17*, pages 2467–2470, New York, NY, USA. Association for Computing Machinery, 2017. ISBN: 978-1-4503-4918-5. doi: 10.1145/3132847.3133167.
- [Hür20] Manuel Hürbin. *Retrieval Optimization in Magnetic Resonance Fingerprinting*. PhD thesis, University of Basel, Basel, Switzerland, 2020.
- [IGN⁺12] S. Idreos, F. Groffen, N. Nes, S. Manegold, S. Mullender, and M. Kersten. MonetDB: Two Decades of Research in Column-Oriented Database. *IEEE Data Engineering Bulletin*, 35(1):40–45, 2012.

- [IKM07] Stratos Idreos, Martin L. Kersten, and Stefan Manegold. Database Cracking. In *Proceedings of the 3rd Biennial Conference on Innovative Data Systems Research*, volume 7, pages 68–78, 2007.
- [IM98] Piotr Indyk and Rajeev Motwani. Approximate Nearest Neighbors: Towards Removing the Curse of Dimensionality. In *Proceedings of the 30th Annual ACM Symposium on Theory of Computing*, pages 604–613, 1998.
- [Ioa03] Yannis Ioannidis. The History of Histograms (abridged). In *Proceedings of the 29th International Conference on Very Large Databases*, pages 19–30, 2003.
- [Iwa16] Masajiro Iwasaki. Pruned Bi-directed k-Nearest Neighbor Graph for Proximity Search. In *Proceedings of the 9th International Conference on Similarity Search and Applications*, volume 9939 of *Lecture Notes in Computer Science*, pages 20–33, Cham. Springer International Publishing, 2016. doi: [10.1007/978-3-319-46759-7_2](https://doi.org/10.1007/978-3-319-46759-7_2).
- [JK84] Matthias Jarke and Jürgen Koch. Query Optimization in Database Systems. *ACM Computing Surveys (CSUR)*, 16(2):111–152, June 1984. ISSN: 0360-0300. doi: [10.1145/356924.356928](https://doi.org/10.1145/356924.356928).
- [JK02] Kalervo Järvelin and Jaana Kekäläinen. Cumulated Gain-based Evaluation of IR Techniques. *ACM Transactions on Information Systems*, 20(4):422–446, 2002.
- [JLX21] Ehsan Javanmardi, Sifeng Liu, and Naiming Xie. Exploring the Philosophical Foundations of Grey Systems Theory: Subjective Processes, Information Extraction and Knowledge Formation. *Foundations of Science*, 26, June 2021. doi: [10.1007/s10699-020-09690-0](https://doi.org/10.1007/s10699-020-09690-0).
- [JSDS⁺19] Suhas Jayaram Subramanya, Fnu Devvrit, Harsha Vardhan Simhadri, Ravishankar Krishnawamy, and Rohan Kadekodi. DiskANN: Fast Accurate Billion-point Nearest Neighbor Search on a Single Node. *Advances in Neural Information Processing Systems*, 32, 2019.
- [JDS10] Herve Jegou, Matthijs Douze, and Cordelia Schmid. Product Quantization for Nearest Neighbor Search. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 33(1):117–128, 2010.

- [JDJ19] Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-Scale Similarity Search with GPUs. *IEEE Transactions on Big Data*, 7(3):535–547, 2019.
- [JGA⁺19] Björn Þór Jónsson, Gylfi Þór Guðmundsson, Laurent Amsaleg, and Philippe Bonnet. Data Storage and Management for Big Multimedia. *Big Data Analytics for Large-Scale Multimedia Search*:211, 2019.
- [JWZ⁺16] Björn Þór Jónsson, Marcel Worring, Jan Zahálka, Stevan Rudinac, and Laurent Amsaleg. Ten Research Questions for Scalable Multimedia Analytics. In *Proceedings of the 22nd International Conference on Multimedia Modeling*, volume 9517 of *Lecture Notes in Computer Science*, pages 290–302. Springer International Publishing, 2016. ISBN: 978-3-319-27674-8. doi: 10.1007/978-3-319-27674-8_26.
- [KKE⁺10] Daniel Keim, Jörn Kohlhammer, Geoffrey Ellis, and Florian Mansmann. Mastering the Information Age: Solving Problems with Visual Analytics, 2010.
- [KCH95] Patrick M. Kelly, Michael T. Cannon, and Donald R. Hush. Query by Image Example: The Comparison Algorithm for Navigating Digital Image Databases (CANDID) Approach. *Storage and Retrieval for Image and Video Databases III*, 2420:238–248, 1995.
- [Kha21] Firas Layth Khaleel. An Overview on Multimedia Big Data Analytics. In *Proceedings of the 2021 International Conference on Electrical Engineering and Informatics*, pages 1–6, 2021. doi: 10.1109/ICEEI52609.2021.9611118.
- [KUG14] M. Ali-ud-din Khan, Muhammad Fahim Uddin, and Navarun Gupta. Seven V’s of Big Data: Understanding Big Data to Extract Value. In *Proceedings of the 2014 Zone 1 Conference of the American Society for Engineering Education*, pages 1–5, 2014. doi: 10.1109/ASEEZone1.2014.6820689.
- [KDJ⁺21] Omar Shahbaz Khan, Aaron Duane, Björn Þór Jónsson, Jan Zahálka, Stevan Rudinac, and Marcel Worring. Exquisitor at the Lifelog Search Challenge 2021: Relationships Between Semantic Classifiers. In *Proceedings of the 4th Annual on Lifelog Search Challenge*, LSC ’21, pages 3–6, New York, NY, USA. Association for

- Computing Machinery, 2021. ISBN: 978-1-4503-8533-6. doi: 10 . 1145/3463948.3469255.
- [KSJ⁺22] Omar Shahbaz Khan, Ujjwal Sharma, Björn Þór Jónsson, Dennis C. Koelma, Stevan Rudinac, Marcel Worring, and Jan Zahálka. Exquisitor at the Video Browser Showdown 2022. In *Proceedings of the 28th International Conference on Multimedia Modeling*, volume 13142 of *Lecture Notes in Computer Science*, pages 511–517. Springer International Publishing, 2022. ISBN: 978-3-030-98355-0. doi: 10 . 1007/978-3-030-98355-0_47.
- [KS04] Hyo Young-Gook Kim and Thomas Sikora. Comparison of MPEG-7 Audio Spectrum Projection Features and MFCC Applied to Speaker Recognition, Sound Classification and Audio Segmentation. In *Proceedings of the 2004 IEEE International Conference on Acoustics, Speech, and Signal Processing*, pages 925–928, Montreal, Canada. IEEE, 2004. doi: 10 . 1109/ICASSP.2004.1327263.
- [KCG⁺03] Serkan Kiranyaz, Kerem Caglar, Esin Guldogan, Olcay Guldogan, and Moncef Gabbouj. MUVIS: A content-based Multimedia Indexing and Retrieval Framework. In *Proceedings of the 7th International Symposium on Signal Processing and Its Applications*, volume 1, pages 1–8, 2003.
- [Koh90] Teuvo Kohonen. The Self-Organizing Map. *Proceedings of the IEEE*, 78(9):1464–1480, 1990.
- [KC01] Irena Koprinska and Sergio Carrato. Temporal Video Segmentation: A survey. *Signal processing: Image communication*, 16(5):477–500, 2001.
- [KM00] Flip Korn and Suresh Muthukrishnan. Influence Sets Based on Reverse Nearest Neighbor Queries. *Proceedings of the 2000 ACM SIGMOD International Conference on Management of Data*, 29(2):201–212, 2000.
- [KBC⁺18] Tim Kraska, Alex Beutel, Ed H. Chi, Jeffrey Dean, and Neoklis Polyzotis. The Case for Learned Index Structures. In *Proceedings of the 2018 International Conference on Management of Data*, SIGMOD ’18, pages 489–504, New York, NY, USA. Association for Computing Machinery, 2018. ISBN: 978-1-4503-4703-7. doi: 10 . 1145/3183713.3196909.

- [KVM⁺20] Miroslav Kratochvíl, Patrik Veselý, František Mejzlík, and Jakub Lokoč. SOM-hunter: Video Browsing with Relevance-to-SOM Feedback Loop. In *Proceedings of the 26th International Conference on Multimedia Modeling*, volume 11962 of *Lecture Notes in Computer Science*, pages 790–795. Springer International Publishing, 2020. ISBN: 978-3-030-37734-2. doi: 10.1007/978-3-030-37734-2_71.
- [KVH21] Alexander Krolik, Clark Verbrugge, and Laurie Hendren. R3d3: Optimized Query Compilation on GPUs. In *Proceedings of the 2021 IEEE/ACM International Symposium on Code Generation and Optimization*, pages 277–288, 2021.
- [LBB⁺18] David MJ Lazer, Matthew A. Baum, Yochai Benkler, Adam J. Berinsky, Kelly M. Greenhill, Filippo Menczer, Miriam J. Metzger, Brendan Nyhan, Gordon Pennycook, and David Rothschild. The Science of Fake News. *Science*, 359(6380):1094–1096, 2018.
- [Lee06] Kyogu Lee. Automatic Chord Recognition from Audio using Enhanced Pitch Class Profile. In *Proceedings of the 2006 International Computer Music Conference*, New Orleans, USA. Michigan Publishing, 2006.
- [LCM⁺22] Yejin Lee, Hyunji Choi, Sunhong Min, Hyunseung Lee, Sangwon Beak, Dawoon Jeong, Jae W Lee, and Tae Jun Ham. ANNA: Specialized Architecture for Approximate Nearest Neighbor Search. In *Proceedings of the 2022 IEEE International Symposium on High-Performance Computer Architecture*, pages 169–183, 2022.
- [LÁJ⁺09] Herwig Lejsek, Friðrik Heiðar Ásmundsson, Björn Þór Jónsson, and Laurent Amsaleg. NV-Tree: An Efficient Disk-based Index for Approximate Search in Very Large High-Dimensional Collections. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 31(5):869–883, 2009. doi: 10.1109/TPAMI.2008.130.
- [LÁJ⁺18] Herwig Lejsek, Friðrik Heiðar Ásmundsson, Björn Þór Jónsson, and Laurent Amsaleg. Transactional Support for Visual Instance Search. In Stéphane Marchand-Maillet, Yasin N. Silva, and Edgar Chávez, editors, *Proceedings of the 11th International Conference on Similarity Search and Applications*, volume 11223 of *Lecture Notes in Computer Science*, pages 73–86, Cham. Springer International

- Publishing, 2018. ISBN: 978-3-030-02224-2. DOI: 10.1007/978-3-030-02224-2_6.
- [LJA11] Herwig Lejsek, Björn Þór Jónsson, and Laurent Amsaleg. NV-Tree: Nearest Neighbors at the Billion Scale. In *Proceedings of the 1st ACM International Conference on Multimedia Retrieval*, New York, NY, USA. Association for Computing Machinery, 2011. ISBN: 978-1-4503-0336-1. DOI: 10.1145/1991996.1992050.
- [Lev65] Vladimir I. Levenshtein. Binary Codes Capable of Correcting Deletions, Insertions, and Reversals. *Soviet Physics Doklady*, 10:707–710, 1965.
- [LCI⁺05] Chengkai Li, Kevin Chen-Chuan Chang, Ihab F. Ilyas, and Sumin Song. RankSQL: Query Algebra and Optimization for Relational Top-k Queries. In *Proceedings of the 2005 ACM SIGMOD International Conference on Management of Data*, SIGMOD '05, pages 131–142, New York, NY, USA. Association for Computing Machinery, 2005. ISBN: 1-59593-060-4. DOI: 10.1145/1066157.1066173.
- [Lib03] Leonid Libkin. Expressive Power of SQL. *Theoretical Computer Science*, 296(3):379–404, 2003.
- [LKD⁺88] Volker Linnemann, Klaus Küspert, Peter Dadam, Peter Pistor, R. Erbe, Alfons Kemper, Norbert Südkamp, Georg Walch, and Mechtilde Wallrath. Design and Implementation of an Extensible Database Management System Supporting User Defined Data Types and Functions. In *Proceedings of the 1988 International Conference on Very Large Databases*, pages 294–305, 1988.
- [LLN⁺20] Chong Liu, Defu Lian, Min Nie, and Xia Hu. Online Optimized Product Quantization. In *Proceedings of the 2020 IEEE International Conference on Data Mining*, pages 362–371, 2020.
- [LBB⁺22] Jakub Lokoč, Werner Bailer, Kai Uwe Barthel, Cathal Gurrin, Silvan Heller, Ladislav Peška, Luca Rossetto, Klaus Schoeffmann, Lucia Vadicamo, Stefanos Vrochidis, et al. A Task Category Space for User-centric Comparative Multimedia Search Evaluations. In *Proceedings of the 28th International Conference on Multimedia Modeling*, volume 13141 of *Lecture Notes in Computer Science*, pages 193–204. Springer International Publishing, 2022. ISBN: 978-3-030-98358-1. DOI: 10.1007/978-3-030-98358-1_16.

- [LBS⁺18] Jakub Lokoč, Werner Bailer, Klaus Schoeffmann, Bernd Münzer, and George Awad. On Influential Trends in Interactive Video Retrieval: Video Browser Showdown 2015–2017. *IEEE Transactions on Multimedia*, 20(12):3361–3376, 2018.
- [LKM⁺19] Jakub Lokoč, Gregor Kovalčík, Bernd Münzer, Klaus Schöffmann, Werner Bailer, Ralph Gasser, Stefanos Vrochidis, Phuong Anh Nguyen, Sitapa Rujikietgumjorn, and Kai Uwe Barthel. Interactive Search or Sequential Browsing? A Detailed Analysis of the Video Browser Showdown 2018. *ACM Transactions on Multimedia Computing, Communications, and Applications*. TOMM, 15(1):1–18, 2019.
- [Low99] David G Lowe. Object Recognition from Local Scale-Invariant Features. In *Proceedings of the 7th IEEE International Conference on Computer Vision*, volume 2, pages 1150–1157, 1999.
- [Lu01] Goujun Lu. Indexing and Retrieval of Audio: A Survey. *Multimedia Tools and Applications*, 15(3):269–290, 2001.
- [LWW⁺20] Kejing Lu, Hongya Wang, Wei Wang, and Mineichi Kudo. VHP: Approximate Nearest Neighbor Search via Virtual Hypersphere Partitioning. *Proceedings of the VLDB Endowment*, 13(9):1443–1455, 2020.
- [LGG⁺18] Shangyu Luo, Zekai J Gao, Michael Gubanov, Luis L Perez, and Christopher Jermaine. Scalable Linear Algebra on a Relational Database System. *IEEE Transactions on Knowledge and Data Engineering*, 31(7):1224–1238, 2018.
- [LC08] Mathias Lux and Savvas A. Chatzichristofis. LIRE: Lucene Image Retrieval: An Extensible Java CBIR Library. In *Proceedings of the 16th ACM International Conference on Multimedia*, pages 1085–1088, New York, NY, USA, 2008. ISBN: 978-1-60558-303-7. doi: 10.1145/1459359.1459577.
- [MGS⁺13] Dan Ma, Vikas Gulani, Nicole Seiberlich, Kecheng Liu, Jeffrey L. Sunshine, Jeffrey L. Duerk, and Mark A. Griswold. Magnetic Resonance Fingerprinting. *Nature*, 495(7440):187–192, 2013.
- [ML86] Lothar F Mackert and Guy M Lohman. R* Optimizer Validation and Performance Evaluation for Local Queries. In *Proceedings of the 1986 ACM SIGMOD International Conference on Management of Data*, pages 84–95, 1986.

- [MY18] Yu A. Malkov and Dmitry A. Yashunin. Efficient and Robust Approximate Nearest Neighbor Search using Hierarchical Navigable Small World Graphs. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 42(4):824–836, 2018.
- [MBK02] Stefan Manegold, Peter A. Boncz, and Martin L. Kersten. Generic Database Cost Models for Hierarchical Memory Systems. In Philip A. Bernstein, Yannis E. Ioannidis, Raghu Ramakrishnan, and Dimitris Papadias, editors, *Proceedings of the 28th International Conference on Very Large Databases*, pages 191–202. San Francisco, 2002. ISBN: 978-1-55860-869-6.
- [MCS88] Michael V. Mannino, Paicheng Chu, and Thomas Sager. Statistical Profile Estimation in Database Systems. *ACM Computing Surveys*, 20(3):191–221, September 1988. ISSN: 0360-0300. DOI: 10.1145/62061.62063.
- [MS96] Sherry Marcus and V. S. Subrahmanian. Foundations of Multimedia Database Systems. *Journal of the ACM*, 43(3):474–523, May 1996. ISSN: 0004-5411. DOI: 10.1145/233551.233554.
- [MPM⁺14] Debra F McGivney, Eric Pierre, Dan Ma, Yun Jiang, Haris Saybasili, Vikas Gulani, and Mark A Griswold. SVD Compression for Magnetic Resonance Fingerprinting in the Time Domain. *IEEE Transactions on Medical Imaging*, 33(12):2311–2322, 2014.
- [MB03] Martin F. McKinney and Jeroen Breebaart. Features for Audio and Music Classification. In *Proceedings of the 4th International Conference on Music Information Retrieval*, Baltimore, USA, 2003.
- [MHL⁺92] Chandrasekaran Mohan, Don Haderle, Bruce Lindsay, Hamid Pirahesh, and Peter Schwarz. ARIES: A Transaction Recovery Method Supporting fine-Granularity Locking and Partial Rollbacks using Write-ahead Logging. *ACM Transactions on Database Systems*, 17(1):94–162, 1992.
- [MT99] Danilo Montesi and Alberto Trombetta. Similarity Search Through Fuzzy Relational Algebra. In *Proceedings of the 10th International Workshop on Database and Expert Systems Applications*, pages 235–239, 1999.

- [MSL⁺14] Hannes Mühlisen, Thaer Samar, Jimmy Lin, and Arjen De Vries. Old Dogs are Great at New Tricks: Column Stores for IR Prototyping. In *Proceedings of the 37th International ACM SIGIR Conference on Research & Development in Information Retrieval*, pages 863–866, 2014.
- [MMB⁺04] Henning Müller, Nicolas Michoux, David Bandon, and Antoine Geissbuhler. A Review of Content-Based Image Retrieval Systems in Medical Applications - Clinical Benefits and Future Directions. *International Journal of Medical Informatics*, 73(1):1–23, 2004.
- [MU17] Henning Müller and Devrim Unay. Retrieval from and Understanding of Large-Scale Multi-Modal Medical Datasets: A Review. *IEEE Transactions on Multimedia*, 19(9):2093–2104, 2017. doi: 10.1109/TMM.2017.2729400.
- [MK18] Y. V. Srinivasa Murthy and Shashidhar G. Koolagudi. Content-Based Music Information Retrieval (CB-MIR) and its Applications Toward the Music Industry: A Review. *ACM Computing Surveys (CSUR)*, 51(3), June 2018. issn: 0360-0300. doi: 10.1145/3177849.
- [Neu11] Thomas Neumann. Efficiently Compiling Efficient Query Plans for Modern Hardware. *Proceedings of the VLDB Endowment*, 4(9):539–550, 2011.
- [NL14] Thomas Neumann and Viktor Leis. Compiling Database Queries into Machine Code. *Bulletin of the IEEE Computer Society Technical Committee on Data Engineering*, 37(1):3–11, 2014.
- [ÓJA⁺11] Arnar Ólafsson, Björn Pór Jónsson, Laurent Amsaleg, and Herwig Lejsek. Dynamic Behavior of Balanced NV-trees. *Multimedia Systems*, 17(2):83–100, March 2011. issn: 1432-1882. doi: 10.1007/s00530-010-0199-4.
- [OT93] Eitetsu Oomoto and Katsumi Tanaka. OVID: Design and Implementation of a Video-Object Database System. *IEEE Transactions on Knowledge and Data Engineering*, 5(4):629–643, 1993.
- [PSS⁺15] Rasmus Pagh, Francesco Silvestri, Johan Sivertsen, and Matthew Skala. Approximate Furthest Neighbor in High Dimensions. In *Proceedings of the 8th International Conference on Similarity Search*

- and Applications*, volume 9371 of *Lecture Notes in Computer Science*, pages 3–14, Cham. Springer International Publishing, 2015. doi: 10.1007/978-3-319-25087-8_1.
- [PST⁺04] Dimitris Papadias, Qiongmao Shen, Yufei Tao, and Kyriakos Mouratidis. Group Nearest Neighbor Queries. In *Proceedings of the 20th International Conference on Data Engineering*, pages 301–312, 2004.
- [PTM⁺05] Dimitris Papadias, Yufei Tao, Kyriakos Mouratidis, and Chun Kit Hui. Aggregate Nearest Neighbor Queries in Spatial Databases. *ACM Transactions on Database Systems*, 30(2):529–576, 2005.
- [PLH21] Johns Paul, Shengliang Lu, and Bingsheng He. Database Systems on GPUs. *Foundations and Trends in Databases*, 11(1):1–108, 2021.
- [PW10] Ofir Pele and Michael Werman. The Quadratic-Chi Histogram Distance Family. In *European Conference on Computer Vision*, pages 749–762, 2010.
- [PAG09] Jorge Pérez, Marcelo Arenas, and Claudio Gutierrez. Semantics and Complexity of SPARQL. *ACM Transactions on Database Systems (TODS)*, 34(3):1–45, 2009.
- [Pet19] Alex Petrov. *Database Internals*. O'Reilly Media, Inc., 2019. ISBN: 978-1-4920-4034-7.
- [PR19] Orestis Polychroniou and Kenneth A Ross. Towards Practical Vectorized Analytical Query Engines. In *Proceedings of the 15th International Workshop on Data Management on New Hardware*, pages 1–7, 2019.
- [PR20] Orestis Polychroniou and Kenneth A Ross. VIP: A SIMD Vectorized Analytical Query Engine. *The VLDB Journal*, 29(6):1243–1261, 2020.
- [PYC⁺18] Samira Pouyanfar, Yimin Yang, Shu-Ching Chen, Mei-Ling Shyu, and S. S. Iyengar. Multimedia Big Data Analytics: A Survey. *ACM Computing Surveys (CSUR)*, 51(1), January 2018. issn: 0360-0300. doi: 10.1145/3150226.
- [Pri08] Dan Pritchett. BASE: An ACID Alternative. *ACM Queue*, 6(3):48–55, 2008.

- [PBB⁺10] Piotr Przymus, Aleksandra Boniewicz, Marta Burzańska, and Krzysztof Stencel. Recursive Query Facilities in Relational Databases: A Survey. In *Database Theory and Application, Bio-Science and Bio-Technology*, pages 89–99. Springer, 2010.
- [PTP46] Edward M Purcell, Henry Cutler Torrey, and Robert V Pound. Resonance Absorption by Nuclear Magnetic Moments in a Solid. *Physical Review*, 69(1-2):37, 1946.
- [QL01] Schuyler Quackenbush and Adam Lindsay. Overview of MPEG-7 Audio. *IEEE Transactions on Circuits and Systems for Video Technology*, 11(6):725–729, 2001. doi: 10.1109/76.927430.
- [QFB⁺19] Thorsten Quandt, Lena Frischlich, Svenja Boberg, and Tim Schatto-Eckrodt. Fake News. *The International Encyclopedia of Journalism Studies*:1–6, 2019.
- [RKH⁺21] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning Transferable Visual Models from Natural Language Supervision. In *Proceedings of the 2021 International Conference on Machine Learning*, pages 8748–8763, 2021.
- [RM97] Davood Rafiei and Alberto Mendelzon. Similarity-Based Queries for Time Series Data. In *Proceedings of the 1997 ACM SIGMOD International Conference on Management of Data*, pages 13–25, 1997.
- [RPK⁺19] Hanna Ragnarsdóttir, Þórhildur Þorleiksdóttir, Omar Shahbaz Khan, Björn Þór Jónsson, Gylfi Þór Guðmundsson, Jan Zahálka, Stevan Rudinac, Laurent Amsaleg, and Marcel Worring. Exquisitor: Breaking the Interaction Barrier for Exploration of 100 Million Images. In *Proceedings of the 27th ACM International Conference on Multimedia*, MM ’19, pages 1029–1031, New York, NY, USA. Association for Computing Machinery, 2019. ISBN: 978-1-4503-6889-6. doi: 10.1145/3343031.3350580.
- [RDR⁺98] Raghu Ramakrsihnan, Donko Donjerkovic, Arvind Ranganathan, Kevin S. Beyer, and Muralidhar Krishnaprasad. SRQL: Sorted Relational Query Language. In *Proceedings of the 10th International Conference on Scientific and Statistical Database Management*, pages 84–95, 1998. doi: 10.1109/SSDM.1998.688114.

- [RDG⁺16] Joseph Redmon, Santosh Divvala, Ross Girshick, and Ali Farhadi. You Only Look Once: Unified, Real-time Object Detection. In *Proceedings of the 2016 IEEE Conference on Computer Vision and Pattern Recognition*, pages 779–788, 2016.
- [Rei81] Raymond Reiter. On Closed World Data Bases. In *Readings in Artificial Intelligence*, pages 119–140. Elsevier, 1981.
- [RDJ12] George Ritzer, Paul Dean, and Nathan Jurgenson. The Coming of Age of the Prosumer. *American Behavioral Scientist*, 56(4):379–398, 2012.
- [RJ10] George Ritzer and Nathan Jurgenson. Production, Consumption, Prosumption: The Nature of Capitalism in the Age of the Digital Prosumer. *Journal of Consumer Culture*, 10(1):13–36, 2010.
- [Rob81] John T. Robinson. The K-D-B-tree: A Search Structure for Large Multidimensional Dynamic Indexes. In *Proceedings of the 1981 ACM SIGMOD International Conference on Management of Data*, SIGMOD ‘81, pages 10–18, New York, NY, USA. Association for Computing Machinery, 1981. ISBN: 0-89791-040-0. doi: 10.1145/582318.582321.
- [RBM22] Viktor Rosenfeld, Sebastian Breß, and Volker Markl. Query Processing on Heterogeneous CPU/GPU Systems. *ACM Computing Surveys (CSUR)*, 55(1):1–38, 2022.
- [Ros18] Luca Rossetto. *Multi-Modal Video Retrieval*. PhD thesis, University of Basel, Switzerland, September 2018.
- [RAPG⁺19] Luca Rossetto, Mahnaz Amiri Parian, Ralph Gasser, Ivan Giangreco, Silvan Heller, and Heiko Schuldt. Deep Learning-based Concept Detection in vitrivr. In *Proceedings of the 25th International Conference on Multimedia Modelling*, MMM 2019, pages 616–621, Cham. Springer International Publishing, 2019. ISBN: 978-3-030-05716-9.
- [RBG⁺21] Luca Rossetto, Matthias Baumgartner, Ralph Gasser, Lucien Heitz, Ruijie Wang, and Abraham Bernstein. Exploring Graph-Querying Approaches in LifeGraph. In *Proceedings of the 4th Annual Workshop on Lifelog Search Challenge*, LSC ’21, pages 7–10, New York, NY, USA. Association for Computing Machinery, 2021. ISBN: 978-1-4503-8533-6. doi: 10.1145/3463948.3469068.

- [RGL⁺20] Luca Rossetto, Ralph Gasser, Jakub Lokoc, Werner Bailer, Klaus Schoeffmann, Bernd Muenzer, Tomas Soucek, Phuong Anh Nguyen, Paolo Bolettieri, Andreas Leibetseder, et al. Interactive Video Retrieval in the Age of Deep Learning - Detailed Evaluation of VBS 2019. *IEEE Transactions on Multimedia*. TOMM, 2020.
- [RGS19] Luca Rossetto, Ralph Gasser, and Heiko Schuldt. Query by Semantic Sketch. *CoRR*, abs/1909.12526, 2019.
- [RGT⁺16] Luca Rossetto, Ivan Giangreco, Claudiu Tanase, and Heiko Schuldt. Vitrivr: A Flexible Retrieval Stack Supporting Multiple Query Modes for Searching in Multimedia Collections. In *Proceedings of the 24th ACM International Conference on Multimedia*, pages 1183–1186, 2016.
- [RSB21] Luca Rossetto, Klaus Schoeffmann, and Abraham Bernstein. Insights on the V3C2 Dataset. *CoRR*, abs/2105.01475, 2021.
- [RHM97] Yong Rui, Thomas S Huang, and Sharad Mehrotra. Content-Based Image Retrieval with Relevance Feedback in MARS. In *Proceedings of International Conference on Image Processing*, volume 2, pages 815–818, 1997.
- [SJ19] Ayodeji Olalekan Salau and Shruti Jain. Feature Extraction: A Survey of the Types, Techniques, Applications. In *Proceedings of the 2019 International Conference on Signal Processing and Communication*, pages 158–164, 2019. doi: 10.1109/ICSC45622.2019.8938371.
- [SFW83] Gerard Salton, Edward A. Fox, and Harry Wu. Extended Boolean Information Retrieval. *Communications of the ACM*, 26(11):1022–1036, 1983.
- [SWY75] Gerard Salton, Anita Wong, and Chung-Shu Yang. A Vector Space Model for Automatic Indexing. *Communications of the ACM*, 18(11):613–620, 1975.
- [SAPG⁺20] Loris Sauter, Mahnaz Amiri Parian, Ralph Gasser, Silvan Heller, Luca Rossetto, and Heiko Schuldt. Combining Boolean and Multimedia Retrieval in vitrivr for Large-scale Video Search. In *Proceedings of the 26th International Conference on Multimedia Modeling*, volume 11962 of *Lecture Notes in Computer Science*, pages 760–765. Springer International Publishing, 2020. ISBN: 978-3-030-37734-2. doi: 10.1007/978-3-030-37734-2_66.

- [Sch19] Klaus Schoeffmann. Video Browser Showdown 2012-2019: A Review. In *2019 International Conference on Content-based Multimedia Indexing*, pages 1–4, 2019.
- [SJD13] Felix Martin Schuhknecht, Alekh Jindal, and Jens Dittrich. The Uncracked Pieces in Database Cracking. *Proceedings of the VLDB Endowment*, 7(2):97–108, 2013.
- [SMM99] Eugenio Di Sciascio, G. Mingolla, and Marina Mongiello. Content-Based Image Retrieval over the Web using Query by Sketch and Relevance Feedback. In *Proceedings of the 3rd International Conference on Advances in Visual Information Systems*, pages 123–130, 1999.
- [SR12] Russell Sears and Raghu Ramakrishnan. bLSM: A General Purpose Log Structured Merge Tree. In *Proceedings of the 2012 ACM SIGMOD International Conference on Management of Data*, pages 217–228, 2012.
- [SHS⁺17] Daniel Seebacher, Johannes Häußler, Manuel Stein, Halldor Janetzko, Tobias Schreck, and Daniel A Keim. Visual Analytics and Similarity Search: Concepts and Challenges for Effective Retrieval Considering Users, Tasks, and Data. In *Proceedings of the 10th International Conference on Similarity Search and Applications*, volume 10609 of *Lecture Notes in Computer Science*, pages 324–332, Cham. Springer International Publishing, 2017. doi: 10.1007/978-3-319-68474-1_23.
- [SAC⁺79] Patricia G. Selinger, Morton M. Astrahan, Donald D. Chamberlin, Raymond A. Lorie, and Thomas G. Price. Access Path Selection in a Relational Database Management System. In Philip A. Bernstein, editor, *Proceedings of the 1979 ACM SIGMOD International Conference on Management of Data*, pages 23–34. ACM, 1979. doi: 10.1145/582095.582099.
- [SR06] Uri Shaft and Raghu Ramakrishnan. Theory of Nearest Neighbors Indexability. *ACM Transactions on Database Systems*, 31(3):814–838, September 2006. ISSN: 0362-5915. doi: 10.1145/1166074.1166077.
- [SHS⁺08] Jie Shao, Zi Huang, Heng Tao Shen, Xiaofang Zhou, Ee-Peng Lim, and Yijun Li. Batch Nearest Neighbor Search for Video Re-

- trieval. *IEEE Transactions on Multimedia*, 10(3):409–420, 2008. doi: 10.1109/TMM.2008.917339.
- [SXJ21] Yijie Shen, Jin Xiong, and Dejun Jiang. Using Vectorized Execution to Improve SQL Query Performance on Spark. In *Proceedings of the 50th International Conference on Parallel Processing*, ICPP 2021, New York, NY, USA. Association for Computing Machinery, 2021. ISBN: 978-1-4503-9068-2. doi: 10.1145/3472456.3472495.
- [SOV⁺21] Larissa C. Shimomura, Rafael Seidi Oyamada, Marcos R. Vieira, and Daniel S. Kaster. A Survey on Graph-based Methods for Similarity Searches in Metric Spaces. *Information Systems*, 95:101507, 2021.
- [SWD⁺21] Jihye Shin, Alexandra Waldau, Aaron Duane, and Björn Þór Jónsson. PhotoCube at the Lifelog Search Challenge 2021. In *Proceedings of the 4th Annual on Lifelog Search Challenge*, pages 59–63. 2021.
- [SJH⁺05] Rut Sigurðardóttir, Björn Þór Jónsson, Hlynur Hauksson, and Laurent Amsaleg. The Quality vs. Time trade-off for Approximate Image Descriptor Search. In *Proceedings of the 21st International Conference on Data Engineering Workshops (ICDEW'05)*, pages 1175–1175, 2005. doi: 10.1109/ICDE.2005.294.
- [SAA⁺10] Yasin N. Silva, Ahmed M. Aly, Walid G. Aref, and Per-Ake Larson. SimDB: A Similarity-aware Database System. In *Proceedings of the 2010 ACM SIGMOD International Conference on Management of Data*, pages 1243–1246, 2010. ISBN: 978-1-4503-0032-2. doi: 10.1145/1807167.1807330.
- [SWA⁺21] Harsha Vardhan Simhadri, George Williams, Martin Aumüller, Matthijs Douze, Artem Babenko, Dmitry Baranchuk, Qi Chen, Lucas Hosseini, Ravishankar Krishnaswamy, Gopal Srinivasa, et al. Results of the NeurIPS’21 Challenge on Billion-scale Approximate Nearest Neighbor Search. In *Proceedings of NeurIPS 2021 Competitions and Demonstrations Track*, volume 176 of *Proceedings of Machine Learning Research*, pages 177–189, 2021.
- [SZ03] Sivic and Zisserman. Video Google: A Text Retrieval Approach to Object Matching in Videos. In *Proceedings of the 9th IEEE Inter-*

- national Conference on Computer Vision*, 1470–1477 vol.2, 2003. doi: 10.1109/ICCV.2003.1238663.
- [SB11] Tomáš Skopal and Benjamin Bustos. On Nonmetric Similarity Search Problems in Complex Domains. *ACM Computing Surveys (CSUR)*, 43(4):1–50, 2011.
- [SWS⁺00] Arnold W. M. Smeulders, Marcel Worring, Simone Santini, Amarnath Gupta, and Ramesh C. Jain. Content-Based Image Retrieval at the End of the Early Years. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 22(12):1349–1380, 2000. doi: 10.1109/34.895972.
- [SAR⁺14] Mohamed A. Soliman, Lyublena Antova, Venkatesh Raghavan, Amr El-Helw, Zhongxian Gu, Entong Shen, George C. Caragea, Carlos Garcia-Alvarado, Foyzur Rahman, Michalis Petropoulos, Florian Waas, Sivaramakrishnan Narayanan, Konstantinos Krikellas, and Rhonda Baldwin. Orca: A Modular Query Optimizer Architecture for Big Data. In *Proceedings of the 2014 ACM SIGMOD International Conference on Management of Data*, SIGMOD ’14, pages 337–348, New York, NY, USA. Association for Computing Machinery, 2014. ISBN: 978-1-4503-2376-5. doi: 10.1145/2588555.2595637.
- [SZB11] Juliusz Sompolski, Marcin Zukowski, and Peter Boncz. Vectorization vs. Compilation in Query Execution. In *Proceedings of the 7th International Workshop on Data Management on New Hardware*, pages 33–40, 2011.
- [SML19] Tomáš Souček, Jaroslav Moravec, and Jakub Lokoč. TransNet: A Deep Network for Fast Detection of Common Shot Transitions. *CoRR*, abs/1906.03363, 2019.
- [SGH⁺22] Florian Spiess, Ralph Gasser, Silvan Heller, Mahnaz Parian-Scherb, Luca Rossetto, Loris Sauter, and Heiko Schuldt. Multi-modal Video Retrieval in Virtual Reality with vitrivr-vr. In *Proceedings of the 28th International Conference on Multimedia Modeling*, volume 13142 of *Lecture Notes in Computer Science*, pages 499–504. Springer International Publishing, 2022. ISBN: 978-3-030-98355-0. doi: 10.1007/978-3-030-98355-0_45.

- [SGH⁺21a] Florian Spiess, Ralph Gasser, Silvan Heller, Luca Rossetto, Loris Sauter, and Heiko Schuldt. Competitive Interactive Video Retrieval in Virtual Reality with vitrivr-VR. In *Proceedings of the 27th International Conference on Multimedia Modeling*, volume 12573 of *Lecture Notes in Computer Science*, pages 441–447. Springer International Publishing, 2021. ISBN: 978-3-030-67835-7. doi: 10.1007/978-3-030-67835-7_4.
- [SGH⁺21b] Florian Spiess, Ralph Gasser, Silvan Heller, Luca Rossetto, Loris Sauter, Milan van Zanten, and Heiko Schuldt. Exploring Intuitive Lifelog Retrieval and Interaction Modes in Virtual Reality with vitrivr-VR. In *Proceedings of the 4th Annual on Lifelog Search Challenge*, pages 17–22, 2021.
- [Spi09] David I. Spivak. Simplicial Databases. *CoRR*, abs/0904.2012, 2009.
- [SRA⁺01] Ioana Stanoi, Mirek Riedewald, Divyakant Agrawal, and Amr El Abbadi. Discovery of Influence Sets in Frequently Updated Databases. In *Proceedings of the 27th International Conference on Very Large Databases*, volume 2001, pages 99–108, 2001.
- [Sto81] Michael Stonebraker. Operating System Support for Database Management. *Communications of the ACM*, 24(7):412–418, 1981.
- [SBP⁺11] Michael Stonebraker, Paul Brown, Alex Poliakov, and Suchi Raman. The Architecture of SciDB. In *Proceedings of the 2011 International Conference on Scientific and Statistical Database Management*, pages 1–16, 2011.
- [SBZ⁺13] Michael Stonebraker, Paul Brown, Donghui Zhang, and Jacek Becla. SciDB: A Database Management System for Applications with Complex Analytics. *Computing in Science & Engineering*, 15(3):54–62, 2013. doi: 10.1109/MCSE.2013.19.
- [Swa89] Arun Swami. Optimization of Large Join Queries: Combining Heuristics and Combinatorial Techniques. In *Proceedings of the 1989 ACM SIGMOD International Conference on Management of Data*, pages 367–376, 1989.
- [SLJ⁺15] Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott Reed, Dragomir Anguelov, Dumitru Erhan, Vincent Vanhoucke, and Andrew Rabinovich. Going Deeper with Convolutions. In

- Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 1–9, 2015.
- [Tsa07] Chih-Fong Tsai. A Review of Image Retrieval Methods for Digital Cultural Heritage Resources. *Online Information Review*, 2007.
- [TYB16] Yi-Hsuan Tsai, Ming-Hsuan Yang, and Michael J Black. Video Segmentation via Object Flow. In *Proceedings of the 2016 IEEE Conference on Computer Vision and Pattern Recognition*, pages 3899–3908, 2016.
- [TSF⁺12] Petros Tsialiamanis, Lefteris Sidiropoulos, Irini Fundulaki, Vasilis Christophides, and Peter Boncz. Heuristics-Based Query Optimisation for SPARQL. In *Proceedings of the 15th International Conference on Extending Database Technology*, pages 324–335, 2012.
- [Uma83] Motohide Umano. Retrieval from Fuzzy Database by Fuzzy Relational Algebra. *IFAC Proceedings Volumes*, 16(13):1–6, 1983.
- [Via19] Gregory Vial. Understanding Digital Transformation: A Review and a Research Agenda. *The Journal of Strategic Information Systems*, 28(2):118–144, 2019.
- [VHS⁺20] Marco Vogt, Nils Hansen, Jan Schönholz, David Lengweiler, Isabel Geissmann, Sebastian Philipp, Alexander Stiemer, and Heiko Schuldt. Polypheny-DB: Towards Bridging the Gap Between Polystores and HTAP Systems. In *Heterogeneous Data Management, Polystores, and Analytics for Healthcare*, volume 12633 of *Lecture Notes in Computer Science*, pages 25–36. Springer, 2020.
- [VLG⁺21] Marco Vogt, David Lengweiler, Isabel Geissmann, Nils Hansen, Marc Hennemann, Cédric Mendelin, Sebastian Philipp, and Heiko Schuldt. Polystore Systems and DBMSs: Love Marriage or Marriage of Convenience? In *Heterogeneous Data Management, Polystores, and Analytics for Healthcare*, volume 12921 of *Lecture Notes in Computer Science*, pages 65–69. Springer, 2021.
- [VSS18] Marco Vogt, Alexander Stiemer, and Heiko Schuldt. Polypheny-DB: Towards a Distributed and Self-Adaptive Polystore. In *Proceedings of the 2018 IEEE International Conference on Big Data*, pages 3364–3373, 2018.

- [VHC⁺19] Stefanos Vrochidis, Benoit Huet, Edward Y Chang, and Ioannis Kompatsiaris. *Big Data Analytics for Large-scale Multimedia Search*. John Wiley & Sons, 2019.
- [Vu19] Tin Vu. Deep Query Optimization. In *Proceedings of the 2019 International Conference on Management of Data*, SIGMOD '19, pages 1856–1858, New York, NY, USA. Association for Computing Machinery, 2019. ISBN: 978-1-4503-5643-5. doi: 10.1145/3299869.3300104.
- [WYG⁺21] Jianguo Wang, Xiaomeng Yi, Rentong Guo, Hai Jin, Peng Xu, Shengjun Li, Xiangyu Wang, Xiangzhou Guo, Chengming Li, Xiaohai Xu, et al. Milvus: A Purpose-Built Vector Data Management System. In *Proceedings of the 2021 International Conference on Management of Data*, pages 2614–2627, 2021.
- [WZs⁺18] Jingdong Wang, Ting Zhang, jingkuan song, Nicu Sebe, and Heng Tao Shen. A Survey on Learning to Hash. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 40(4):769–790, 2018. doi: 10.1109/TPAMI.2017.2699960.
- [WOK⁺98] Yasuhiko Watanabe, Yoshihiro Okada, Kengo Kaneji, and Yoshi-taka Sakamoto. Multimedia Database System for TV Newscasts and Newspapers. In *International Conference on Advanced Multimedia Content Processing*, pages 208–220, 1998.
- [WMZ10] William Webber, Alistair Moffat, and Justin Zobel. A Similarity Measure for Indefinite Rankings. *ACM Transactions on Information Systems*, 28(4), November 2010. ISSN: 1046-8188. doi: 10.1145/1852102.1852106.
- [WB00] Roger Weber and Klemens Böhm. Trading Quality for Time with Nearest-Neighbor Search. In *International Conference on Extending Database Technology*, pages 21–35, 2000.
- [WSB98] Roger Weber, Hans-Jörg Schek, and Stephen Blott. A Quantitative Analysis and Performance Study for Similarity-Search Methods in High-Dimensional Spaces. In *Proceedings of the 1998 International Conference on Very Large Databases*, volume 98, pages 194–205, New York City, NY, USA. Morgan Kaufmann, 1998.

- [WLK⁺10] Kyu-Young Whang, Jae-Gil Lee, Min-Soo Kim, Min-Jae Lee, Ki-Hoon Lee, Wook-Shin Han, and Jun-Sung Kim. Tightly-Coupled Spatial Database Features in the Odysseus/OpenGIS DBMS for High-Performance. *GeoInformatica*, 14(4):425–446, 2010.
- [WLL⁺15] Kyu-Young Whang, Jae-Gil Lee, Min-Jae Lee, Wook-Shin Han, Min-Soo Kim, and Jun-Sung Kim. DB-IR Integration using Tight-Coupling in the Odysseus DBMS. *World Wide Web-internet and Web Information Systems*, 18(3):491–520, 2015.
- [WZW85] S. K. M. Wong, Wojciech Ziarko, and Patrick C. N. Wong. Generalized Vector Spaces Model in Information Retrieval. In *Proceedings of the 8th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, SIGIR ’85, pages 18–25, New York, NY, USA. Association for Computing Machinery, 1985. ISBN: 0-89791-159-8. doi: 10 . 1145 / 253495 . 253506.
- [WCZ⁺13] Wentao Wu, Yun Chi, Shenghuo Zhu, Junichi Tatenuma, Hakan Hacigümüs, and Jeffrey F. Naughton. Predicting Query Execution Time: Are Optimizer Cost Models Really Unusable. In *Proceedings of the 29th IEEE International Conference on Data Engineering*, pages 1081–1092, 2013.
- [XTZ18] Donna Xu, Ivor W. Tsang, and Ying Zhang. Online Product Quantization. *IEEE Transactions on Knowledge and Data Engineering*, 30(11):2185–2198, 2018. doi: 10 . 1109 / TKDE . 2018 . 2817526.
- [YPM⁺19] Fan Yang, Shiva K. Pentyala, Sina Mohseni, Mengnan Du, Hao Yuan, Rhema Linder, Eric D. Ragan, Shuiwang Ji, and Xia (Ben) Hu. XFake: Explainable Fake News Detector with Visualizations. In *Proceedings of the 2019 World Wide Web Conference*, WWW 2019, pages 3600–3604, New York, NY, USA. Association for Computing Machinery, 2019. ISBN: 978-1-4503-6674-8. doi: 10 . 1145 / 3308558 . 3314119.
- [YLF⁺20] Wen Yang, Tao Li, Gai Fang, and Hong Wei. PASE: PostgreSQL Ultra-High-Dimensional Approximate Nearest Neighbor Search Extension. In *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data*, pages 2241–2253, 2020.

- [YLZ07] Yubin Yang, Hui Lin, and Yao Zhang. Content-based 3-D Model Retrieval: A Survey. *IEEE Transactions on Systems, Man, and Cybernetics, Part C (Applications and Reviews)*, 37(6):1081–1098, 2007.
- [YLK10] Bin Yao, Feifei Li, and Piyush Kumar. K Nearest Neighbor Queries and kNN-joins in Large Relational Databases (almost) for Free. In *2010 IEEE 26th International Conference on Data Engineering*, pages 4–15, 2010.
- [YHM15] Shaoyi Yin, Abdelkader Hameurlain, and Franck Morvan. Robust Query Optimization Methods with Respect to Estimation Errors: A Survey. *Sigmod Record*, 44(3):25–36, December 2015. ISSN: 0163-5808. doi: 10.1145/2854006.2854012.
- [YCL⁺19] Lin Yue, Weitong Chen, Xue Li, Wanli Zuo, and Minghao Yin. A Survey of Sentiment Analysis in Social Media. *Knowledge and Information Systems*, 60(2):617–663, 2019.
- [Zad96] Lotfi A. Zadeh. Fuzzy Sets. In *Fuzzy Sets, Fuzzy Logic, and Fuzzy Systems: Selected Papers by Lotfi A Zadeh*, pages 394–432. World Scientific, 1996.
- [ZW14] Jan Zahálka and Marcel Worring. Towards Interactive, Intelligent, and Integrated Multimedia Analytics. In *Proceedings of the 2014 IEEE Conference on Visual Analytics Science and Technology (VAST)*, pages 3–12, 2014. doi: 10.1109/VAST.2014.7042476.
- [ZCB⁺18] Savvas Zannettou, Tristan Caulfield, Jeremy Blackburn, Emiliano De Cristofaro, Michael Sirivianos, Gianluca Stringhini, and Guillermo Suarez-Tangil. On the Origins of Memes by Means of Fringe Web Communities. In *Proceedings of the Internet Measurement Conference 2018*, pages 188–202, 2018.
- [ZCL⁺10] Daniel Zeng, Hsinchun Chen, Robert Lusch, and Shu-Hsing Li. Social media Analytics and Intelligence. *IEEE Intelligent Systems*, 25(6):13–16, 2010. doi: 10.1109/MIS.2010.151.
- [ZAD⁺06] Pavel Zezula, Giuseppe Amato, Vlastislav Dohnal, and Michal Batko. *Similarity Search The Metric Space Approach*. Springer Science + Business Media, Inc., 2006. ISBN: 978-0-387-29146-8.

- [ZHC⁺06] Zhen Zhang, Seung-won Hwang, Kevin Chen-Chuan Chang, Min Wang, Christian A Lang, and Yuan-chi Chang. Boolean+ ranking: Querying a Database by k-constrained Optimization. In *Proceedings of the 2006 ACM SIGMOD International Conference on Management of Data*, pages 359–370, 2006.
- [ZWN22] Wanlei Zhao, Hui Wang, and Chong-Wah Ngo. Approximate k-NN Graph Construction: A Generic Online Approach. *IEEE Transactions on Multimedia*, 24:1909–1921, 2022.
- [ZTL20] Weijie Zhao, Shulong Tan, and Ping Li. SONG: Approximate Nearest Neighbor Search on GPU. In *Proceedings of the 36th IEEE International Conference on Data Engineering*, pages 1033–1044, 2020.
- [ZZ20] Xinyi Zhou and Reza Zafarani. A Survey of Fake News: Fundamental Theories, Detection Methods, and Opportunities. *ACM Computing Surveys (CSUR)*, 53(5):1–40, 2020.
- [ZCW⁺15] Wenwu Zhu, Peng Cui, Zhi Wang, and Gang Hua. Multimedia Big Data Computing. *IEEE Multimedia*, 22(3):96–105, 2015.
- [Zih21] Gabriel Zihlmann. *Magnetic Resonance Fingerprinting Reconstruction Using Methods from Multimedia Retrieval*. PhD thesis, University of Basel, Basel, Switzerland, 2021.



*My heart has joined the Thousand, for my
friend stopped running today.*

— Richard Adams, Watership Down

There is grandeur in this view of life, with its several powers, having been originally breathed into a few forms or into one; and that, whilst this planet has gone cycling on according to the fixed law of gravity, from so simple a beginning endless forms most beautiful and most wonderful have been, and are being, evolved.

— Charles Darwin, *The Origin of Species*

